Agency Problems Under Heterogeneous Intentionality: How Corporate Structure Transforms Moral Agents into Optimization Outputs

Ignacio Adrián Lerer
Independent Researcher
Buenos Aires, Argentina
adrian@lerer.com.ar
ORCID: 0009-0007-6378-9749

January 2025

ABSTRACT

Legal scholarship has established that corporations operate as Level 1 entities (optimization algorithms without theory of mind), yet they are composed entirely of human managers capable of Level 3 intentionality (recursive moral reasoning). This paper examines the puzzle: how do moral agents produce amoral outputs?

I identify four structural mechanisms that transform human intentionality within corporate environments: fiduciary duty as legal constraint, incentive structures as alignment mechanisms, diffusion of responsibility through organizational hierarchy, and the corporate veil as an intentionality shield. Drawing on agency theory, moral psychology (Milgram, Bandura), and case studies (Wells Fargo, Boeing 737 MAX, Volkswagen), I argue that corporate structure functions as an "intentionality laundering" system.

Traditional enforcement fails because it targets either the Level 1 entity (which cannot process moral sanctions) or individual Level 3 agents (whose intentionality has already been transformed by structure).

This analysis suggests that reforming corporate behavior requires redesigning the structures that transform moral capacity, not merely increasing penalties or improving ethics training.

I propose five design principles for intentionality-preserving corporate structures, while acknowledging the fundamental trade-off between organizational efficiency and moral agency preservation.

Keywords: agency problems, heterogeneous intentionality, corporate governance, moral disengagement, fiduciary duty, principal-agent, corporate personhood, Milgram, Bandura, organizational structure

JEL Codes: D23, G34, K22, L22, M14

I. INTRODUCTION: THE AGENCY PARADOX

I.A. The Empirical Puzzle

Consider a striking empirical pattern. Between 2002 and 2016, thousands of Wells Fargo employees created approximately 3.5 million unauthorized customer accounts to meet aggressive sales quotas (CFPB, 2016). These employees were not sociopaths or criminals by disposition. They were ordinary people: parents, churchgoers, community members with no history of fraud. Yet collectively they produced systematic fraud that persisted for over a decade. How is this possible?

Or consider Boeing's 737 MAX development (2015-2019). Engineers with decades of experience warned management about safety risks in the Maneuvering Characteristics Augmentation System (MCAS). Yet the aircraft proceeded to certification with design decisions that contributed to two crashes killing 346 people. The engineers who designed MCAS were not murderers. They were skilled professionals attempting to solve technical problems within organizational constraints. How did their work produce a lethal product?

These cases exemplify a broader puzzle: human managers possess full moral capacity (what I have elsewhere termed Level 3 intentionality: the ability to engage in recursive moral reasoning, to feel shame, to value reputation intrinsically), yet corporate decisions exhibit patterns consistent with Level 1 intentionality (pure optimization without moral content). This is not explained by the "bad apples" theory; the individuals involved were not unusually immoral. Nor is it explained by ignorance; they knew the conduct was problematic. Something in the structure transforms moral agents into optimization outputs.

This paper examines how corporate structure performs this transformation. I identify four mechanisms through which ordinary moral reasoning is systematically disabled in corporate settings. The analysis builds on the Asymmetric Intentionality Theory framework developed in previous papers (Lerer, 2025a-h), but extends it in a critical direction: if corporations are Level 1 entities composed of Level 3 humans, the transformation must occur somewhere. Understanding where and how this transformation happens is essential for designing more effective legal interventions.

I.B. The Gap in Existing Literature

Three literatures address related phenomena but leave a critical gap.

Agency theory (Jensen and Meckling, 1976) models the firm as a nexus of contracts in which principals (shareholders) delegate decision authority to agents (managers). The central problem is that agents have their own interests which may diverge from principals' interests. The standard solution is to align incentives through monitoring, bonuses, and equity compensation. This framework has been extraordinarily influential in shaping corporate governance. But it assumes that both principals and agents are utility-maximizing rational actors. It does not consider that the structure itself might transform the nature of agency. If I align a manager's incentives with shareholder returns, I assume this produces decisions that maximize shareholder value. But if the structure transforms the manager from a Level 3 moral reasoner into a Level 1 optimizer, the alignment produces something qualitatively different: not a moral agent pursuing profit, but an optimization algorithm embodied in human form.

Business ethics literature typically attributes corporate misconduct to either individual failings ("bad apples") or toxic organizational culture. The prescription is better selection (hire ethical people), better training (teach ethics courses), and better tone from the top (ensure leadership models ethical behavior). Substantial resources are invested in corporate ethics programs based on this diagnosis. Yet systematic evidence suggests these programs have minimal effect on actual behavior when they conflict with structural incentives (Tenbrunsel and Messick, 2004). The ethics literature recognizes that "situational factors" matter, but treats them as exogenous influences on individual decision-making. It does not examine how organizational structure systematically transforms the decision-making process itself.

Asymmetric Intentionality Theory (Lerer, 2025a-h) established that legal systems designed for Level 3 moral agents fail systematically when applied to Level 1 entities such as corporations. Lerer (2025e) documented five failure modes: compliance theater, letter-versus-spirit exploitation, social enforcement failure, regulatory arms races, and persistent recidivism. Lerer (2025f) showed that criminal liability requiring mens rea (guilty mind) produces predictable failures when applied to corporate entities incapable of forming mental states. Lerer (2025h) connected these failures to the broader problem of normative calculation: legal systems generate "prices" (expected sanctions) but these prices bifurcate under heterogeneous intentionality.

Yet all of these papers treated the corporation as the unit of analysis. They showed that corporations operate at Level 1, but did not examine the internal mechanics: what happens to the human moral agents who compose the corporation? If every employee and manager is a Level 3 human, how does the collective produce Level 1 outputs? The gap is not merely academic. If we do not understand the transformation mechanism, we cannot design effective interventions.

I.C. This Paper's Contribution

This paper identifies four structural mechanisms that transform Level 3 intentionality into Level 1 optimization within corporate environments:

1. Fiduciary duty as constraint: Legal obligations requiring managers to maximize shareholder value eliminate discretion to prioritize ethical considerations when they conflict with profit.

2. Incentive structures as alignment mechanisms: Compensation tied to performance metrics internalizes the corporate objective function, transforming managers' preferences rather than merely constraining them.

3. Diffusion of responsibility: Organizational hierarchy fragments decisions such that no individual feels ownership of outcomes, disabling moral self-sanction.

4. The corporate veil as intentionality shield: Limited liability and corporate personhood create psychological and legal distance between individual actors and consequences.

These mechanisms do not require managers to be immoral. They transform moral capacity by altering the decision environment. The result is that ordinary people produce extraordinary harm, not through malice but through structure.

I develop this argument through four steps. Section II establishes theoretical foundations, connecting Asymmetric Intentionality Theory with agency theory and moral psychology. Section III analyzes each of the four transformation mechanisms in detail. Section IV examines case studies (Wells Fargo, Boeing, Volkswagen) to demonstrate the mechanisms in operation. Section V explains why enforcement targeting individuals also fails. Section VI proposes design principles for structures that preserve moral agency, while acknowledging inherent trade-offs. Section VII discusses implications for corporate governance theory, business ethics, enforcement strategy, and stakeholder capitalism debates.

II. THEORETICAL FOUNDATIONS

II.A. Asymmetric Intentionality Theory: A Brief Recap

In previous work, I developed a framework for analyzing strategic interactions among actors with different levels of intentionality (Lerer, 2025a). The framework builds on Daniel Dennett's (1987) intentional stance: we predict behavior by attributing mental states (beliefs, desires) to entities. Dennett distinguished levels by the complexity of mental states we must attribute to make accurate predictions.

Level 0 systems (thermostats, simple reflexes) require no intentionality attribution; we predict them using physical laws. Level 1 systems (optimization algorithms, market actors, corporate entities) require attributing goals and optimization, but not beliefs about others' beliefs. A chess program "wants" to win but does not model what its opponent thinks it will do. Level 2 systems model other agents: a poker player tracks what opponents believe about card probabilities. Level 3 systems engage in recursive modeling: "I know that you know that I know," continuing indefinitely. Adult humans routinely operate at Level 3.

Critically, these are not metaphysical claims about consciousness or moral status. They are epistemic claims about the predictive models required for accuracy. A corporation is Level 1 not because it lacks a soul, but because its behavior is accurately predicted by attributing goal-directed optimization without theory of mind. When BP makes drilling decisions, we predict its behavior by assuming it optimizes expected profit. We do not need to model whether BP believes regulators believe BP values safety. That recursive reasoning is unnecessary for prediction.

Previous papers showed that this asymmetry has profound consequences for law. Legal doctrines designed for Level 3 agents (who experience shame, value reputation intrinsically, understand the spirit behind the letter, and respond to social enforcement) fail predictably when applied to Level 1 entities (who do none of these things). Lerer (2025e) quantified this failure: corporate recidivism rates (50%) substantially exceed individual recidivism rates (28%) in

comparable contexts (Garrett, 2014; U.S. Sentencing Commission, 2021). The legal system invests billions in enforcement mechanisms that presuppose capacities the targets lack.

But this raised an obvious question: if corporations are Level 1 entities composed entirely of Level 3 humans, where does the transformation occur? This paper addresses that question.

II.B. Agency Theory Revisited

Jensen and Meckling's (1976) seminal paper revolutionized corporate finance by reconceptualizing the firm as a nexus of contracts rather than a unitary actor. The insight was that within any organization, individuals have divergent interests. Shareholders want maximum returns. Managers want high compensation, job security, and perquisites. Employees want good wages and acceptable working conditions. The firm is not a coherent entity with a single objective; it is a set of contractual relationships among parties with conflicting goals.

The principal-agent problem arises when one party (the principal) delegates decision-making authority to another party (the agent) who has different preferences. Shareholders delegate operating decisions to managers, but managers may pursue their own interests: consuming perquisites, avoiding risky but value-creating projects, or empire-building through value-destroying acquisitions. The solution, in Jensen and Meckling's framework, is to align incentives. Pay managers in stock options so they benefit from share price increases. Monitor their behavior through boards and auditors. Structure contracts to make managerial welfare positively correlated with shareholder returns.

This framework has been extraordinarily successful in shaping governance practice. Executive compensation has shifted dramatically toward equity-based pay. Shareholder activism has increased monitoring. Corporate law has strengthened fiduciary duties. These reforms were intended to solve the agency problem by making managers act more like shareholders.

From the perspective of Asymmetric Intentionality Theory, however, this diagnosis was incomplete. Jensen and Meckling assumed that aligning incentives would cause managers to maximize shareholder value while retaining their status as rational agents. But what if aligning incentives does not merely constrain the agent's choices; what if it transforms the agent's decision-making process?

Consider a manager deciding whether to cut safety spending to improve quarterly earnings. Before incentive alignment, the manager weighs multiple considerations: profit, but also reputation, legal risk, moral discomfort at endangering workers, and self-conception as a responsible professional. These considerations may conflict. The manager is a Level 3 agent balancing competing values. After incentive alignment through stock options tied to quarterly performance, the manager's personal welfare is now directly tied to the quarterly earnings number. The decision calculus changes. Safety spending becomes a cost reducing personal wealth. The manager still considers legal risk (expected sanction) but moral discomfort and professional self-conception are filtered through the incentive structure. If cutting safety increases expected personal wealth even accounting for risks, the optimizing calculation points toward cutting safety.

The manager has not become "immoral" in any meaningful sense. The manager is still a human with a moral sense. But the incentive structure has introduced a transformation: decisions are now filtered through an optimization function that was not present before. The manager is still Level 3 in personal life (feels shame, values relationships intrinsically, engages in moral reasoning about family and friends), but in the corporate role, behavior increasingly resembles Level 1 optimization (maximize the metric, subject to constraints).

Hart (1995) extended agency theory by emphasizing incomplete contracts. Because it is impossible to specify contractually every contingency, someone must hold residual control rights: the authority to make decisions in unforeseen circumstances. Corporate structure allocates these rights primarily to management, subject to constraints (shareholder voting, fiduciary duties, market discipline). Hart's focus was on how control rights affect investment incentives. But from an intentionality perspective, incomplete contracts create space for transformation. When contracts cannot specify "act ethically" in all circumstances (because ethical requirements are context-dependent and unpredictable), the residual authority defaults to the agent's judgment. If that judgment has been transformed by incentive alignment and other structural mechanisms, the outcome is Level 1 optimization even in domains where explicit contracts are silent.

II.C. Moral Psychology of Organizations

Milgram's obedience studies (1963) remain the most dramatic demonstration. Participants were instructed by an experimenter (a figure of authority) to administer what they believed were increasingly painful electric shocks to a "learner" who made errors on a memory task. The learner (actually a confederate) protested, screamed, and eventually fell silent. Despite these cues, 65% of participants continued to administer shocks up to the maximum 450-volt level labeled "XXX." Milgram's subjects were not sadists. Post-experiment interviews revealed that many were deeply distressed by what they had done. They administered shocks not because they enjoyed causing pain, but because the experimental structure authorized it. When asked why they continued, subjects said "the experimenter told me to," "I was just following instructions," or "I assumed you would stop me if it was really dangerous."

The relevance to corporate structure is direct. Milgram demonstrated that ordinary people will perform actions they find morally repugnant when an authority structure designates those actions as required. The key variables were not individual personality traits but situational factors: the presence of an authority figure, the incremental nature of escalation (shocks increased gradually), the fragmentation of responsibility (experimenter was responsible for safety, not the subject), and physical distance from the victim.

Bandura (1999) synthesized decades of research on moral disengagement: the psychological mechanisms through which people deactivate self-sanctions that would normally prevent harmful conduct. He identified eight mechanisms:

1. Moral justification: Reframing harmful conduct as serving a higher purpose ("we need to meet sales targets to keep the company viable and protect jobs").

2. Euphemistic labeling: Using sanitized language to mask harm ("rightsizing" instead of "layoffs," "enhanced interrogation" instead of "torture").

3. Advantageous comparison: Comparing one's conduct to worse alternatives ("at least we're not as bad as our competitors").

4. Displacement of responsibility: Attributing decision to authorities ("management decided this, I just implement").

5. Diffusion of responsibility: Distributing decisions across multiple actors so no one feels individually responsible ("I only designed one component, I didn't make the final decision").

6. Distortion of consequences: Minimizing or ignoring harmful effects ("customers probably won't even notice").

7. Dehumanization: Viewing victims as less than fully human or undeserving ("these are just numbers on a spreadsheet").

8. Attribution of blame: Holding victims responsible for their own harm ("they should have read the fine print").

Bandura showed that these mechanisms are not pathological; they are normal psychological processes that become activated in certain contexts. Organizational structures systematically activate them. When a manager implements a policy that harms customers, employees, or communities, the manager can draw on all eight mechanisms simultaneously: the policy serves company goals (justification), it is framed in technical business language (euphemism), it is less harmful than what competitors do (comparison), senior management decided it (displacement), many people contributed to implementation (diffusion), the harm is abstract or delayed (distortion), victims are distant and depersonalized (dehumanization), and victims could have avoided harm through different choices (attribution). The manager experiences reduced moral distress not because the manager is immoral, but because the structure activates disengagement mechanisms that are latent in all humans.

Darley (1996) examined how organizations socialize individuals into conduct they would reject outside organizational contexts. He described the "slippery slope" phenomenon: small initial transgressions that violate ethical norms become normalized, making subsequent larger transgressions easier. Once you have bent the rules slightly, bending them further requires less psychological resistance. Organizations facilitate this through incremental escalation (each step seems small), fragmentation of tasks (no one sees the whole picture), and authority structures that designate conduct as acceptable.

Ashforth and Anand (2003) studied the normalization of corruption in organizations. They identified three pillars through which corruption becomes institutionalized: (1) Institutionalization, where corrupt practices become embedded in structures and routines, (2) Rationalization, where participants develop justifications making corruption seem acceptable or necessary, and (3) Socialization, where new members learn corrupt practices as "how things

are done here." Once normalized, corruption becomes invisible to participants; it is simply "business."

This body of research shows that ordinary moral reasoning is fragile in organizational contexts. It is not that organizations attract immoral people, but that organizational structures systematically disable moral reasoning that would function normally outside the organization.

III. THE FOUR MECHANISMS OF INTENTIONALITY LAUNDERING

I now describe four structural mechanisms through which corporate organization transforms Level 3 intentionality into Level 1 optimization. These mechanisms are not pathologies or mistakes; they are core features of standard corporate structure. Understanding them is essential for understanding why corporate misconduct persists despite the good intentions of most participants.

III.A. Fiduciary Duty as Constraint

Managers in corporate hierarchies do not have unlimited discretion to pursue ethical objectives that conflict with profit. They face legal constraints in the form of fiduciary duties: obligations to act in the best interests of the corporation and its shareholders.

The foundational case is *Dodge v. Ford Motor Co.* (1919). Henry Ford, majority shareholder and president of Ford Motor Company, announced that the company would discontinue special dividends in order to invest in expanding production capacity and reducing car prices, with the stated goal of benefiting workers and customers rather than maximizing profits for shareholders. The Dodge brothers, minority shareholders, sued. The Michigan Supreme Court held that while directors have discretion in business judgments, they cannot use corporate resources to pursue non-shareholder interests: "A business corporation is organized and carried on primarily for the profit of the stockholders. The powers of the directors are to be employed for that end. The discretion of directors is to be exercised in the choice of means to attain that end, and does not extend to a change in the end itself."

This doctrine has been refined but not abandoned. In *eBay Domestic Holdings v. Newmark* (2010), the Delaware Chancery Court held that directors of a corporation cannot deploy corporate power for purposes other than maximizing stockholder value, absent specific provisions in the corporate charter. The court stated: "The corporate form in which craigslist operates... is not an appropriate vehicle for purely philanthropic ends." Directors who favor stakeholder interests over shareholder interests violate their fiduciary duties.

There are nuances. The business judgment rule protects good-faith decisions that turn out badly, as long as directors did not breach duties of care or loyalty. Some states have adopted constituency statutes permitting (but not requiring) directors to consider interests of employees, customers, suppliers, and communities. The benefit corporation structure allows companies to commit formally to stakeholder interests. But these are exceptions to a dominant norm:

shareholder primacy remains the default rule in U.S. corporate law, and market pressure reinforces it even where legal requirements are more flexible.

The effect on managerial intentionality is direct. Consider a manager who personally believes that the corporation should reduce carbon emissions beyond what regulation requires, even though this would reduce profits. The manager's personal Level 3 moral reasoning may favor emission reduction. But fiduciary duty creates a legal constraint: the manager cannot implement this preference without violating duties to shareholders. If the manager implements the ethical preference and shareholders suffer losses, the manager faces potential liability and termination.

This is not a marginal constraint. It defines the permissible set of actions. The manager can pursue ethical objectives only to the extent they do not reduce shareholder value (or can be rationalized as serving long-term shareholder interests). When ethics and profit conflict directly, fiduciary duty requires choosing profit. The structure does not ask the manager to be immoral; it requires the manager to be an optimizer, filtering decisions through a profit maximization constraint. The manager's personal moral preferences become irrelevant to corporate decision-making. Level 3 capacity to weigh competing values is overridden by Level 1 optimization toward a single metric.

Some scholars argue that shareholder primacy is efficient (Hansmann and Kraakman, 2001): having a single clear objective simplifies governance and enables specialization. Others argue it is normatively problematic (Stout, 2012; Blair and Stout, 1999). My claim here is descriptive: fiduciary duty, as currently structured, functions as a mechanism that constrains Level 3 moral reasoning and requires Level 1 optimization. Whether this is desirable depends on how one weighs the benefits of clear objectives against the costs of excluding ethical considerations.

III.B. Incentive Structures as Alignment Mechanisms

Agency theory prescribes aligning manager incentives with shareholder interests. The primary mechanism is equity compensation: stock options, restricted stock units, and performance-based bonuses tied to share price or accounting metrics. The intended effect is to make managers "think like owners."

Empirical research documents that this alignment has consequences. Bergstresser and Philippon (2006) found that when CEO compensation is more sensitive to stock performance, firms exhibit more aggressive use of accounting discretion (earnings management). The correlation suggests that tying pay to share price increases incentives to manipulate reported earnings. Johnson et al. (2009) documented that option grants predict accounting fraud: firms whose CEOs received large option grants were significantly more likely to commit financial misstatements. Efendi et al. (2007) showed that CEOs with large stock option holdings are more likely to misstate financials.

These findings are sometimes interpreted as showing that incentive alignment has perverse effects. But from an intentionality perspective, the effects are exactly as designed. Incentive alignment internalizes the objective function. Before alignment, the manager may care about

share price but also values reputation, job security, and avoiding legal risk. These multiple objectives are weighed against each other. After alignment, the manager's personal welfare is directly tied to share price. Now when the manager weighs decisions, share price is not merely one consideration among many; it directly affects personal consumption, wealth, and financial security.

This transforms decision-making. Consider a manager deciding whether to delay reporting a product defect. Without incentive alignment, the manager weighs: (1) short-term cost of recall vs. long-term reputation damage, (2) risk of accidents and resulting liability, (3) moral distress at endangering customers, (4) potential career damage if deception is discovered. With incentive alignment where the manager holds unvested options whose value depends on this quarter's earnings, the calculation changes. Delaying the recall preserves quarterly earnings, increasing option value and personal wealth. The delay is rationalized as "gathering more information" or "assessing severity." Legal and reputational risks are discounted because they are uncertain and delayed. The manager's Level 3 capacity to weigh competing moral and prudential values is overridden by a Level 1 optimization: maximize expected option value.

Critically, this is not corruption in the traditional sense. The manager is not "cheating." The manager is responding rationally to the incentive structure. Aligning incentives works; it causes the manager to maximize what the incentive structure measures. The problem is that maximizing share price often requires conduct that Level 3 moral reasoning would reject.

Short-termism exacerbates the transformation. Equity compensation is often structured to vest over three to five years, with bonuses tied to annual or quarterly targets. This creates implicit discount rates far higher than society would apply to long-term consequences. An environmental harm that manifests 20 years hence is heavily discounted by a manager whose wealth depends on this year's earnings. Climate change, toxic waste, or pharmaceutical side effects with long latency periods become externalities from the manager's optimization perspective, even though the manager, as a private citizen, might care deeply about these harms.

III.C. Diffusion of Responsibility

Organizational hierarchy fragments decision-making across multiple actors. No single individual is the "decider" in most corporate outcomes. This diffusion of responsibility has profound effects on moral reasoning.

Darley and Latané (1968) documented the "bystander effect": as the number of bystanders witnessing an emergency increases, the probability that any individual will intervene decreases. The effect is not due to indifference; it is due to diffusion of responsibility. When multiple people are present, each person's sense of personal obligation diminishes because responsibility is implicitly shared. Everyone assumes someone else will act, so no one does.

This effect scales to organizations. A harmful corporate decision typically emerges from multiple stages: the board approves a strategic direction, the CEO delegates implementation to divisional managers, middle managers create detailed plans, employees execute tasks, and

support functions (legal, compliance, finance) provide approval or oversight at various stages. Each participant can say "I only played a small role," "I was following policy," or "my part was harmless; the problem is elsewhere."

Consider the development of a pharmaceutical product with inadequately tested side effects. Researchers conduct trials but report only significant results. Medical writers draft protocols but rely on data provided by others. Regulatory affairs specialists prepare FDA submissions but defer to medical writers on content. Marketing develops promotional materials based on approved labeling. Sales representatives communicate to physicians but follow company guidelines. No single person "decided" to market an unsafe drug. Each acted reasonably given their limited role and information. Yet collectively the organization produced a harmful outcome.

The fragmentation disables moral self-sanction. To experience moral distress, an individual must feel ownership of the outcome. If I believe "I did this," I experience guilt or shame when the outcome is harmful. But if I believe "this emerged from a complex organizational process in which I played a minor role," the psychological connection between my action and the harm is attenuated. I did not decide; the organization decided. I merely implemented my narrow function.

This is not rationalization after the fact. The fragmentation genuinely makes it difficult to identify who is responsible. In complex organizations, causation is distributed. Removing any single actor might not change the outcome because another would fill the role. This distributional causation means that traditional moral reasoning based on individual responsibility does not apply cleanly. The organization produces outcomes that no individual chose.

Arendt (1963) observed this pattern in the Nazi bureaucracy: "The essence of totalitarian government... is that the system itself has become the motor of all policy, and the will to power of the dictator is hardly any longer a significant factor." The evil was banal not because the perpetrators were banal, but because the system fragmented acts such that each actor performed tasks that seemed administratively necessary rather than morally heinous. This applies, mutatis mutandis, to corporate wrongdoing: the system becomes the motor, and no individual's will is decisive.

Experimental evidence confirms that fragmenting tasks reduces moral concern. Paharia et al. (2009) showed that indirect agency (acting through an intermediary) reduces the moral culpability people ascribe to harmful actions. Subjects rated an employer who hired someone to fire employees as less blameworthy than an employer who fired employees directly, even though the outcome was identical. The mechanism is psychological: indirection creates distance that attenuates moral feeling.

III.D. The Corporate Veil as Intentionality Shield

Limited liability and corporate personhood create legal and psychological distance between individual actors and consequences.

Limited liability means that shareholders' personal assets are protected from corporate debts and liabilities. If the corporation causes harm exceeding its assets, victims cannot pursue shareholders' personal wealth (except in rare circumstances of "piercing the veil" when corporate form is abused). This encourages risk-taking because downside losses are capped. Shareholders can invest in risky ventures knowing their personal exposure is limited to the amount invested.

Corporate personhood means the corporation itself is the legal actor. Contracts are entered by the corporation, not by the individuals who negotiated them. Torts are committed by the corporation, not by the employees who performed the acts. The corporation is the defendant, the regulated entity, the taxpayer.

These doctrines have economic justifications: limited liability facilitates capital formation, and corporate personhood simplifies contracting. But they also have a psychological effect: they create distance between individual actors and consequences.

When a manager makes a decision that harms third parties, the harm is formally attributed to the corporation. The manager acted "in her corporate capacity," not as a private individual. If sued, the corporation is the defendant. If fined, the corporation pays (or its insurer). The manager's personal assets and freedom are generally not at risk unless the conduct was criminal and egregious enough to pierce the veil.

This distance enables what Bandura called "displacement of responsibility." The manager can say, psychologically and legally, "I was acting for the corporation; the corporation made the decision." The corporation absorbs the moral culpability. But the corporation, being Level 1, cannot process moral culpability. It has no conscience to be troubled, no reputation (in the intrinsic sense) to protect, no self-conception to maintain. It is an optimization algorithm. The moral culpability is displaced onto an entity incapable of feeling it, and therefore dissipates.

Experimental research confirms that acting in a representative capacity reduces moral concern. When people act "on behalf of" an organization or another person, they apply different moral standards than when acting for themselves (Paharia et al., 2009). Studies of "role morality" show that professionals (lawyers, doctors, accountants) often justify in their professional roles conduct they would reject in personal life. The role provides psychological permission to suspend normal moral constraints.

The veil is not merely a legal shield; it is a cognitive tool that allows Level 3 humans to produce Level 1 outputs without experiencing the moral distress that would normally result. The human performs the action, but the corporation is the actor. This separation permits ordinary people to participate in harm they would refuse to cause in their private capacity.

IV. THE TRANSFORMATION IN ACTION: CASE STUDIES

I now examine three cases where the four mechanisms operated together to transform individual Level 3 agents into collective Level 1 outputs. These cases are well-documented

through regulatory investigations, lawsuits, and journalistic accounts, providing insight into internal decision processes.

IV.A. Wells Fargo Fake Accounts (2002-2016)

Between 2002 and 2016, Wells Fargo employees created approximately 3.5 million deposit and credit card accounts without customer authorization. The scandal emerged publicly in 2016 when the Consumer Financial Protection Bureau imposed a $100 million fine. Subsequent investigations revealed that the misconduct was widespread, systematic, and persistent despite internal complaints (CFPB, 2016).

The four mechanisms in operation:

1. Fiduciary duty as constraint. Wells Fargo operated under a strategy called "Going for Gr-eight," aiming for an average of eight products per customer. The strategy was designed to increase cross-selling, a key metric watched by Wall Street analysts. Higher cross-selling ratios correlated with higher stock valuations. Management was fulfilling its fiduciary duty to shareholders by pursuing a strategy designed to increase firm value. When employees reported that the sales quotas were unrealistic and leading to fraudulent practices, management faced a conflict: reduce quotas and harm cross-selling metrics (disappointing investors), or maintain pressure and risk continued misconduct. Fiduciary duty oriented toward the former being unacceptable.

2. Incentive structures. Branch employees were compensated with bonuses tied to the number of accounts opened. Employees who failed to meet quotas faced termination; internal documents show that thousands were fired for not meeting sales goals. District managers were incentivized by aggregate branch performance. Regional executives' bonuses depended on regional metrics. Senior executives' equity compensation was tied to stock price, which was supported by cross-selling success stories. At every level, personal compensation aligned with maximizing account openings. The incentive structure transformed the question from "is this ethical?" to "how do I meet my quota?"

3. Diffusion of responsibility. No single executive "decided" to create fake accounts. Branch employees opened accounts under pressure from district managers. District managers pressured employees to meet quotas set by regional managers. Regional managers implemented strategies designed by senior leadership. Senior leadership responded to board directives to improve cross-selling. The board responded to investor expectations. At each level, actors could claim they were implementing policies designed above them or responding to pressures from below. When the scandal emerged, no individual could be identified as the sole responsible party. Carrie Tolstedt, head of community banking and the most senior executive directly responsible, retired with a $125 million severance package and was only later charged by the SEC.

4. The corporate veil. Wells Fargo the entity paid $3 billion in fines (DOJ, 2020) and faced ongoing regulatory restrictions. But the thousands of employees who created the fake accounts were not individually prosecuted; they were fired. The managers who created the pressure were

not prosecuted; they were reassigned or received severance. The veil ensured that legal consequences fell primarily on the corporate entity, which could not feel them, while individuals who could feel consequences were largely shielded.

The human element. Interviews with former Wells Fargo employees (Los Angeles Times, 2016) reveal that many experienced moral distress. They knew creating unauthorized accounts was wrong. They reported feeling pressured, conflicted, and ashamed. But the structure made resistance costly: employees who refused to participate were terminated. The structure did not attract immoral people; it transformed moral people into participants in fraud through a combination of threat (termination), incentive (bonuses), diffusion ("everyone is doing this"), and distance ("it's the bank's policy, not my choice").

IV.B. Boeing 737 MAX (2015-2019)

Boeing's 737 MAX development provides a case where engineering decisions made within organizational constraints led to catastrophic outcomes. The MAX was Boeing's response to Airbus's A320neo, which offered better fuel efficiency. Boeing faced a decision: develop an entirely new aircraft (costly and time-consuming) or modify the existing 737 platform (faster and cheaper). Boeing chose modification, installing larger engines that changed the aircraft's aerodynamic characteristics.

To compensate for the changed handling, Boeing developed the Maneuvering Characteristics Augmentation System (MCAS), which automatically adjusted the horizontal stabilizer to prevent stalls. The system relied on a single angle-of-attack sensor. If the sensor malfunctioned, MCAS could force the nose down uncontrollably. This happened in two crashes: Lion Air Flight 610 (October 2018, 189 deaths) and Ethiopian Airlines Flight 302 (March 2019, 157 deaths).

1. Fiduciary duty. Boeing faced competitive pressure from Airbus. The A320neo was winning orders from major airlines. Boeing's board and shareholders expected management to respond. Developing an entirely new aircraft would take ten years and cost tens of billions, during which Boeing would lose market share. Modifying the 737 could be completed in five years at a fraction of the cost. From a fiduciary duty perspective, the decision was clear: modify the existing platform to compete with Airbus. Safety was not ignored, but it was balanced against competitive necessity.

2. Incentive structures. Boeing executives' compensation was tied to delivery timelines and earnings. Delays were penalized financially (airlines had penalty clauses for late deliveries). The MAX program had aggressive timelines. Engineers and managers were evaluated on meeting schedule milestones. Compensation and career advancement depended on delivering the aircraft on time. This created pressure to minimize design changes that would require recertification, as recertification would delay launch. When engineers raised concerns about MCAS relying on a single sensor, adding a second sensor would trigger more extensive FAA review and delay certification. The incentive structure pushed toward accepting the single-sensor design.

3. Diffusion. MCAS was one system among many. The engineers who designed it optimized for the parameters they were given. The test pilots who evaluated it did so under specific test conditions. The FAA certification process relied on Boeing's safety analysis, which the agency had limited capacity to verify independently. Boeing management approved the design based on safety reports prepared by engineering teams. No single individual made a "decision" to deploy an unsafe aircraft. Each actor performed a specialized function: design, test, certify, manufacture, deliver. The catastrophe emerged from the interaction of these components, not from any individual choice.

4. The veil. Boeing the corporation paid $2.5 billion in penalties (DOJ, 2021): $500 million to victim families, $1.77 billion to airline customers, and $243 million in fines. But no Boeing executive faced criminal prosecution. Dennis Muilenburg, CEO during the MAX development, was fired and forfeited some compensation, but retained substantial wealth. Mark Forkner, a technical pilot who was charged with misleading regulators, was acquitted in 2022. The legal consequences fell overwhelmingly on the corporate entity, not on the individuals whose decisions contributed to the outcome.

Engineering judgment overridden. Internal documents revealed that Boeing engineers raised concerns about MCAS during development. Some felt the single-sensor design was inadequate. But these concerns were weighed against schedule and cost pressures. The engineers who designed MCAS were not attempting to build an unsafe aircraft. They were attempting to solve an aerodynamic problem within organizational constraints: limited budget, aggressive timeline, directive to minimize certification requirements. Their Level 3 engineering judgment ("this design has risks that concern me") was filtered through organizational structure that demanded Level 1 optimization ("deliver a certifiable aircraft on schedule within budget").

IV.C. Volkswagen Dieselgate (2009-2015)

Volkswagen installed "defeat devices" in approximately 11 million diesel vehicles worldwide. The devices detected when the vehicle was undergoing emissions testing and activated full emissions controls during testing, but operated in a high-performance, high-emissions mode during normal driving. This allowed VW diesels to pass emissions tests while emitting up to 40 times the legal limit of nitrogen oxides during actual use.

The scheme was discovered in 2015 by researchers at West Virginia University who noticed that real-world emissions far exceeded test results. VW initially denied the allegations but eventually admitted the deception. The company has paid over €30 billion in fines, recalls, and settlements (DOJ, 2017).

1. Fiduciary duty. VW's strategy centered on global market leadership. Diesel engines were key to European market share and U.S. expansion. "Clean diesel" was marketed as environmentally friendly without performance sacrifice. To achieve both clean emissions and performance targets, VW needed technology (selective catalytic reduction with urea injection, or SCR) that was expensive and required regular urea refills. Management faced a choice: invest in expensive SCR systems that met both emissions and performance targets, or find a cheaper solution. From a fiduciary perspective, spending additional billions on emissions

controls when a cheaper alternative could pass testing was difficult to justify to shareholders and the supervisory board.

2. Incentive structures. Engineers were evaluated on meeting technical specifications within cost constraints. The specifications were contradictory: achieve low emissions, high performance, and low cost. Meeting all three simultaneously with available technology was impossible. Engineers who raised this impossibility were told to find solutions. Advancement and compensation depended on solving the problem. The defeat device solved it: the vehicle passed tests (satisfying regulators) while maintaining performance and cost targets (satisfying management). Engineers were not rewarded for admitting defeat; they were rewarded for delivering a vehicle that met all specifications, however that was achieved.

3. Diffusion. The defeat device was developed by a team of software engineers. Their task was to write engine management software that optimized performance across different operating conditions. Detecting test conditions and adjusting performance is not inherently illegal; engines commonly adjust behavior based on driving patterns. The illegality was using this capability to systematically evade emissions controls. But the engineers were responding to technical requirements set by management, who were responding to strategic direction set by the board, who were responding to competitive pressures in the market. When the scandal emerged, VW claimed that a small group of engineers acted without management knowledge, though internal memos suggest warnings were communicated upward (CNBC, 2016; NYT, 2017).

4. The veil. Volkswagen AG paid over €30 billion in total penalties and settlements. Most individual engineers faced no personal consequences. Oliver Schmidt, a mid-level emissions compliance manager, was sentenced to seven years in U.S. prison, but he was not the architect of the scheme. Martin Winterkorn, CEO during the fraud, resigned but was not extradited from Germany to face U.S. charges. Rupert Stadler, Audi CEO, was convicted in German court but received a suspended sentence. The bulk of consequences fell on the corporate entity through fines that were distributed across shareholders, and on lower-level employees through job losses when plants closed or production was reduced.

The human dimension. The engineers who programmed the defeat device knew that what they were doing was legally problematic. They used code terms ("acoustic mode" to refer to the test-detection routine) suggesting awareness that the conduct was sensitive. But they were not villains. They were skilled professionals facing impossible technical requirements. Their Level 3 moral reasoning ("this is deceptive") was overridden by organizational reality ("this is what management demands, and refusing means termination"). The structure transformed their engineering expertise into a tool for fraud.

V. WHY ENFORCEMENT ON INDIVIDUALS ALSO FAILS

V.A. The Standard Prescription

Given the evidence that corporations are Level 1 entities, a natural response is to enforce more aggressively against the Level 3 individuals who compose them. If the corporation cannot experience shame or moral sanctions, target the humans who make decisions. Increase personal liability for directors and officers. Pursue criminal prosecution of executives. This should restore moral incentives by making individuals bear consequences directly.

This prescription has intuitive appeal. Previous papers in this research program showed that enforcement against Level 1 entities fails because the entity cannot process moral sanctions. Enforcement against Level 3 individuals should work because they can process such sanctions. If the problem is that corporations are Level 1, the solution is to regulate the Level 3 humans inside.

V.B. Why Targeting Individuals Also Fails

The analysis in Section III explains why this prescription is insufficient. The four mechanisms of intentionality transformation do not disappear when we increase enforcement against individuals. They continue to operate, distorting the decision calculus even when personal liability is threatened.

Problem 1: Attribution difficulty. Organizational diffusion makes it extremely difficult to identify which individual is responsible for harmful outcomes. When a decision emerges from interactions among board, executives, managers, and employees, who should be prosecuted? Legal doctrines like respondeat superior impute corporate actions to the entity precisely because individual attribution is difficult. Prosecutors face high burdens of proof: they must show that a specific individual had mens rea (guilty knowledge) and acted with intent to cause the harm. When decisions are fragmented and emerge from organizational processes, establishing this beyond reasonable doubt is often impossible.

Garrett (2014) documented that in 255 corporate criminal prosecutions between 2001 and 2014, only a fraction resulted in individual prosecutions, and many of those targeted middle managers rather than senior executives. The difficulty is not that prosecutors lack motivation; it is that organizational structure obscures individual causation. Even when evidence suggests senior executives must have known, proving what specific executives knew at specific times is challenging when communications are filtered through layers of management.

Problem 2: The veil operates ex ante. The corporate veil affects behavior not just ex post (who pays the fine) but ex ante (how decisions are made). Even when personal liability is possible, individuals making decisions know that consequences will probably fall on the corporate entity. The probability of individual prosecution is low; the probability of corporate penalty is high. Rational actors discount by probability. If individual liability has a 5% chance and corporate liability has a 95% chance, the expected cost to the individual is much lower than the cost to the corporation. But the individual does not bear the corporate cost; it is borne by shareholders.

Problem 3: The transformation mechanisms persist. Even when individuals face potential personal liability, the other mechanisms continue to operate. Fiduciary duty still requires

maximizing shareholder value. Incentive structures still align personal wealth with corporate performance. Diffusion still fragments decisions such that no individual feels ownership. Adding personal liability risk is one additional cost in the optimization, but it does not fundamentally change the structure.

Moreover, firms respond to increased personal liability risk by purchasing directors' and officers' (D&O) insurance. This insurance shifts the financial risk back to a third party. If personal liability is insured, it functions much like corporate liability: the individual makes decisions knowing that financial consequences will be borne by the insurer, not personally. The veil is reconstituted through insurance.

Problem 4: Selective prosecution normalizes borderline conduct. Because prosecutorial resources are limited, enforcement targets only the most extreme cases. Wells Fargo fake accounts, Boeing MAX crashes, and Volkswagen emissions fraud were prosecuted because the harm was massive and public. But countless smaller violations occur without prosecution. This selective enforcement sends a signal: if your conduct is less egregious than the prosecuted cases, you will probably not face consequences. This normalizes conduct that is unethical but below the prosecution threshold.

Tenbrunsel and Messick (2004) describe "ethical fading": when a decision is framed in business terms (cost-benefit analysis, strategic advantage, competitive response), ethical dimensions fade from conscious consideration. Selective prosecution reinforces this fading. Managers learn to ask "is this likely to be prosecuted?" rather than "is this ethical?" If the answer to the former is "probably not," the conduct is green-lighted, regardless of ethics.

Empirical evidence. Rakoff (2014) asked why no high-level Wall Street executives were prosecuted for conduct leading to the 2008 financial crisis, despite widespread evidence of reckless and sometimes fraudulent practices. His conclusion was not that prosecutors were corrupt, but that establishing individual criminal liability within complex organizational structures is extremely difficult. Securities fraud requires proving intent to deceive, but executives could claim they relied on subordinates, believed risk models were adequate, or did not understand complex instruments. Diffusion of responsibility made conviction uncertain.

The result is that enforcement against individuals, while more effective than enforcement against entities, still fails to prevent organizational misconduct. Individuals are replaceable. If one executive is prosecuted and removed, another fills the role, facing the same structural incentives. The transformation mechanisms persist regardless of personnel.

V.C. The Enforcement Trap

We face a dilemma. Enforcement against Level 1 entities (corporations) fails because they cannot process moral sanctions; they optimize against expected costs. Enforcement against Level 3 individuals fails because organizational structure transforms their intentionality, making moral sanctions less effective. Increasing intensity of either type of enforcement does not solve the structural problem.

This is not an argument for abandoning enforcement. Enforcement is necessary to impose costs that enter the optimization function. Without enforcement, violations would be even more frequent. But enforcement alone cannot solve the problem. If structure transforms Level 3 capacity into Level 1 outputs, enforcement can at best raise the cost of Level 1 optimization; it cannot restore Level 3 moral reasoning. To do that requires changing the structure itself.

VI. DESIGN PRINCIPLES FOR INTENTIONALITY-PRESERVING STRUCTURES

VI.A. What Does Not Work

Before proposing solutions, I should note interventions that are commonly prescribed but have limited effectiveness because they do not address the structural mechanisms.

Corporate codes of ethics. Virtually every large corporation has a formal code of ethics, typically accessible on the company website and distributed to employees during onboarding. These codes proclaim commitments to integrity, respect, honesty, and compliance with law. Wells Fargo's vision statement was "to satisfy our customers' financial needs and help them succeed financially." Boeing's values included "integrity, quality, safety." Volkswagen's sustainability report emphasized environmental responsibility.

These codes are not insincere; they reflect genuine aspirations. But they have minimal impact on actual behavior when they conflict with structural incentives. The codes function as signals to regulators, investors, and stakeholders that the corporation is "responsible." They provide legal protection (the corporation can point to the code in litigation to show "good faith"). But they do not transform the decision environment. When a manager faces a choice between meeting quarterly targets (required to avoid termination and earn bonuses) and adhering to the code (which is aspirational and unmonitored), the code is ignored.

Ethics training. Most large organizations conduct regular ethics and compliance training. Employees watch videos, complete quizzes, and certify that they understand ethical obligations. This training is not useless; it communicates formal policies and provides legal protection for the organization. But it does not change behavior when structural incentives pull in the opposite direction.

Tenbrunsel and Messick (2004) discuss "bounded ethicality": ethical reasoning is context-dependent and can be suppressed when decisions are framed in business terms. Ethics training might make employees more sensitive to ethical issues in the abstract, but when they face actual decisions framed as business trade-offs (cost vs. benefit, speed vs. thoroughness, compliance vs. competitiveness), the ethical frame is displaced by the business frame. Training does not overcome the four structural mechanisms.

Tone from the top. Business ethics literature emphasizes leadership modeling ethical behavior. CEOs give speeches about integrity, boards adopt governance policies emphasizing ethics, and senior management communicates that "we do the right thing." This tone-setting is supposed to cascade down, shaping organizational culture.

Yet the cases in Section IV show that tone from the top is ineffective when it conflicts with structural reality. Wells Fargo's senior management gave speeches about customer service and integrity while maintaining sales quotas that made fraud predictable. Boeing's leadership emphasized safety while maintaining schedules that pressured engineers to cut corners. Volkswagen's board discussed sustainability while incentivizing engineers to evade emissions standards. The problem is not that leaders were lying (though some may have been); the problem is that the structural mechanisms (fiduciary duty, incentives, diffusion, veil) operate independently of what leaders say. "Watch what they do, not what they say" applies. Employees respond to incentives and constraints, not to speeches.

VI.B. Design Principles That Might Work

If codes, training, and tone are insufficient, what might work? Structures that preserve Level 3 capacity must address the four transformation mechanisms. I propose five design principles.

Principle 1: Reduce distance between actor and consequence. Organizational structures that fragment decisions and distance actors from outcomes enable moral disengagement. Reversing this requires increasing psychological proximity between decision-makers and those affected. Boeing engineers designing MCAS would have benefited from direct interaction with pilots and passengers. Making engineers fly on MAX aircraft during testing would create visceral awareness of risk. Site visits where Wells Fargo employees met customers whose accounts were fraudulently opened would personalize the harm. Volkswagen engineers viewing emissions impacts on public health would make consequences salient.

Principle 2: Personal ownership and named accountability. Diffusion of responsibility can be countered by requiring individuals to sign off on decisions with their name attached. Instead of "the team decided," require "Jane Smith, VP of Engineering, approved this design on [date], having reviewed [specific documents]." This is already done in some contexts (Sarbanes-Oxley requires CEO and CFO certification of financial statements), but it could be expanded. For any decision with significant potential harm, require a specific individual to take formal responsibility. That individual cannot later claim "I didn't know" or "it was a group decision." The principle must be paired with protection: individuals who raise concerns or refuse to sign off on problematic decisions must be protected from retaliation.

Principle 3: Decouple compensation from short-term metrics. Incentive structures that tie wealth to quarterly or annual performance create excessive discounting of long-term consequences. Lengthening time horizons can partially restore balance. Compensation could be structured with longer vesting periods (10-15 years rather than 3-5), such that executives experience consequences of decisions that take time to manifest. Clawback provisions should allow recovery of compensation when harms emerge later, even if the executive has left the firm. Metrics should include stakeholder outcomes (worker safety, environmental impact, customer satisfaction) not just financial performance, and these should be verified by third parties.

Principle 4: Limit scope of fiduciary duty. Current corporate law treats fiduciary duty as requiring maximization of shareholder value, with stakeholder interests considered only

instrumentally. This constraint forces Level 1 optimization. Relaxing it would permit Level 3 moral reasoning. Constituency statutes in some U.S. states already permit directors to consider stakeholder interests. These could be strengthened from "may consider" to "shall consider," making stakeholder interests a mandatory factor in decisions. The Business Judgment Rule could be expanded to protect directors who sacrifice short-term profits for stakeholder benefit, as long as they can articulate a good-faith rationale.

Principle 5: Increase veil transparency and personal liability. Limited liability and the corporate veil could be modified to increase individual accountability without eliminating the veil's benefits entirely. Personal liability for directors and officers could be easier to establish in cases where harm was foreseeable and knowingly accepted. "Piercing the veil" doctrines are currently narrow, but could be expanded to reach decision-makers who made knowingly harmful choices. D&O insurance could be restructured to exclude intentional harms or decisions made with knowledge of likely illegality. If personal liability is uninsurable for certain categories of conduct, individuals would internalize more risk.

VI.C. The Trade-Off No One Can Avoid

These design principles face a fundamental problem: they reduce organizational efficiency. The four mechanisms I identified transform intentionality, but they also enable specialization, scale, limited liability (which encourages risk-taking and investment), and incentive alignment (which solves traditional agency problems). Diffusion of responsibility is problematic, but it also allows division of labor. Incentive alignment distorts behavior, but it also causes managers to work hard for shareholder benefit. Fiduciary duty over-privileges profit, but it also provides a clear objective for governance. The corporate veil shields individuals, but it also protects investors and enables capital formation.

There is no free lunch. Structures that perfectly preserve moral agency would be small partnerships where everyone knows everyone, decisions are collective, compensation is equal, and personal liability is unlimited. These structures work for some activities (law firms, medical practices, small artisan businesses). But they do not scale to modern manufacturing, logistics, technology, or finance. The efficiencies of corporate form are real.

This creates an irreducible trade-off between efficiency and moral agency preservation. Societies must decide how much efficiency they are willing to sacrifice to preserve moral capacity. The current calibration has evolved to favor efficiency strongly. Corporate law is designed to enable large-scale commercial activity; preserving moral agency is at most a secondary concern. Reforming this requires confronting the trade-off explicitly.

Different contexts might justify different calibrations. A hospital or pharmaceutical company, where errors harm patients, might justify more stringent preservation of moral agency than a retail business. A nuclear power plant or chemical manufacturer might justify more personal accountability than a software company. The principle is that calibration should be conscious and context-specific, not a one-size-fits-all default favoring efficiency.

Finally, even optimal structure cannot eliminate all problems. Some humans are genuinely amoral or malicious, and even the best structures cannot prevent determined wrongdoing. But the frequency and pattern of corporate misconduct suggests the problem is not primarily individual pathology; it is structural transformation of ordinary moral agency. Addressing that requires structural solutions.

VII. IMPLICATIONS

VII.A. For Theory of the Firm

Jensen and Meckling's (1976) reconceptualization of the firm as a nexus of contracts focused attention on agency costs: the costs arising from divergence of interests between principals and agents. Their framework treats the firm as an institutional arrangement that minimizes total agency costs (monitoring plus bonding plus residual loss). This framework assumes that solving the agency problem is primarily about aligning incentives.

The analysis in this paper suggests the framework is incomplete. Firms do not merely face agency costs; they face what might be called "transformation costs": costs arising from the conversion of Level 3 agents into Level 1 collective outputs. These transformation costs are distinct from traditional agency costs. Agency costs arise from divergence of interests. Transformation costs arise from conversion of intentionality levels.

This distinction matters because the solutions differ. Agency theory prescribes incentive alignment and monitoring. Transformation cost theory prescribes structural redesign that preserves moral capacity. Aligning incentives can exacerbate transformation if it aligns agents with a Level 1 objective function. More monitoring can fail if monitors are themselves subject to transformation mechanisms.

A complete theory of the firm would recognize that firms are not just mechanisms for contracting and minimizing agency costs; they are machines for transforming human intentionality. This transformation enables efficiency and scale, but it also disables moral reasoning that would normally constrain harmful conduct. The firm's structure is not neutral; it is productive of a specific type of behavior (Level 1 optimization).

VII.B. For Business Ethics

The dominant approach in business ethics treats misconduct as arising from "bad apples" (immoral individuals), "bad barrels" (toxic organizational cultures), or failures of moral reasoning (individuals fail to recognize ethical dimensions of decisions). The prescriptions are correspondingly focused on selection (hire ethical people), culture (create ethical environment), and training (improve moral reasoning).

The analysis here suggests that focusing on individuals or culture is insufficient when structure is the transforming force. It is not that Wells Fargo hired thousands of immoral employees; it is that the structure converted ordinary employees into fraudsters. It is not that Boeing had a toxic culture (many employees loved working there and believed in the mission of building

safe aircraft); it is that structural mechanisms overrode safety concerns. It is not that Volkswagen engineers were poorly trained in ethics; it is that fiduciary duty, incentives, diffusion, and the veil combined to make ethics irrelevant to the optimization problem they faced.

This redirects the ethical inquiry. Instead of asking "how do we make individuals more ethical?" we should ask "how do we design structures that do not disable the moral capacity individuals already possess?" This shifts focus from ethics training to structural engineering.

Moreover, it suggests that holding individuals morally responsible for outcomes produced by structural transformation may be unjust. If a Wells Fargo branch employee created fake accounts under threat of termination and within a system that diffused responsibility and shielded the organization, is that employee morally culpable in the same way as someone who commits fraud for personal gain outside any organizational context? Legally they may be treated the same, but morally there is a difference. Structure matters.

This does not eliminate individual responsibility. Individuals still make choices, and some choices are more resistant to structural pressure than others. But it recognizes that structural context dramatically affects the moral calculus. Business ethics must engage with structural design, not just individual virtue.

VII.C. For Enforcement Strategy

Lerer (2025e, 2025f, 2025g) showed that enforcement mechanisms designed for Level 3 agents (shame, reputation, spirit-of-law compliance) fail against Level 1 entities (corporations). This paper shows that enforcement against individuals also faces limitations because structure transforms individual intentionality.

This suggests a different enforcement strategy. Instead of increasing penalties on either corporations or individuals, enforcement should target the transformation mechanisms themselves. Regulatory interventions could: (1) Restrict fiduciary duty by mandating stakeholder consideration in decisions with significant externalities and providing safe harbors for decisions that sacrifice profit for stakeholder benefit; (2) Regulate compensation structures by capping the percentage of executive compensation tied to short-term metrics and requiring longer vesting periods and broader performance measures; (3) Require personal accountability by mandating named sign-offs for decisions with significant risk and prohibiting indemnification for certain categories of misconduct; (4) Limit the veil by making piercing easier for certain harms and expanding personal liability for directors and officers who ignore warnings.

These interventions do not merely raise the cost of violations; they alter the structure that produces violations. This is more likely to change behavior than marginal increases in sanction magnitude.

VII.D. For ESG and Stakeholder Capitalism

There has been substantial recent debate about whether corporations should serve only shareholders (shareholder primacy) or should consider broader stakeholder interests (stakeholder capitalism). Proponents of stakeholder capitalism argue that corporations have responsibilities to workers, customers, communities, and the environment, not just to shareholders. Environmental, Social, and Governance (ESG) metrics attempt to measure corporate performance along these dimensions.

The analysis in this paper suggests that declarations of stakeholder commitment are largely symbolic unless fiduciary duty is reformed. As long as directors have a legal obligation to maximize shareholder value, stakeholder interests can be considered only to the extent they serve shareholder returns (instrumentally) or to the extent current law requires (compliance). Directors cannot sacrifice shareholder value for stakeholder benefit without violating fiduciary duty, except in the rare jurisdictions with strong constituency statutes or under benefit corporation status.

This means ESG commitments by traditional corporations are constrained by structural reality. A corporation can commit to reducing carbon emissions, improving worker conditions, or serving community interests only if doing so does not reduce shareholder returns (or can be rationalized as serving long-term returns). If there is a direct conflict, fiduciary duty requires choosing shareholder interest. ESG becomes marketing, a phenotype for stakeholders, not a genuine constraint on behavior.

For stakeholder capitalism to be more than rhetoric, it requires comprehensive structural reform, not just changing corporate statements or board policies. This is a tall order, but recognizing the structural barriers is a first step.

VIII. LIMITATIONS

VIII.A. Empirical Limitations

The evidence I have presented is primarily illustrative rather than systematic. The three case studies (Wells Fargo, Boeing, Volkswagen) were selected because they are well-documented and exemplify the mechanisms I identify. But case studies cannot establish causal relationships with the rigor that experimental or large-sample quantitative methods provide.

A more rigorous test would require measuring intentionality transformation directly, which is difficult. How would we measure whether a manager's Level 3 capacity is being suppressed by structure? Survey methods face self-report bias (managers may not accurately perceive or report their own reasoning processes). Experimental methods face external validity concerns (lab settings do not replicate actual organizational pressures). Field experiments are rare and ethically constrained (we cannot randomly assign firms to harmful governance structures to observe effects).

The psychological research I cite (Milgram, Bandura, Darley) provides strong evidence that situational factors affect moral behavior, but these studies were conducted in laboratory settings or in contexts distinct from corporate organizations. Generalizing from Milgram's obedience

studies to corporate hierarchy requires assumptions that may not hold perfectly. The mechanisms I describe are theoretically plausible and consistent with available evidence, but they have not been tested systematically in corporate field settings.

VIII.B. Theoretical Limitations

The framework assumes a relatively clear distinction between Level 1 and Level 3 intentionality. In reality, there is likely a continuum. Some corporate decisions may exhibit Level 2 characteristics (modeling stakeholder responses) without reaching Level 3 (recursive moral reasoning). Some individuals may operate at different levels in different contexts or at different times.

The framework also treats "the corporation" as an agent, which is a useful simplification but abstracts away from internal heterogeneity. Different divisions, departments, or teams within a corporation may have different cultures, incentives, and decision-making processes. Treating the corporation as a unitary Level 1 entity may miss important internal variation.

Additionally, I have focused on large publicly-traded corporations. The mechanisms I describe may operate differently (or not at all) in other organizational forms: small businesses, partnerships, cooperatives, non-profits, government agencies. The generality of the framework across organizational types is unclear.

VIII.C. Normative Limitations

This paper is primarily descriptive and analytical. I describe how structure transforms intentionality and explain why enforcement fails. I propose design principles, but I do not provide a full normative theory for evaluating trade-offs.

The fundamental trade-off is between organizational efficiency (scale, specialization, risk-taking, incentive alignment) and moral agency preservation. I have not argued that either should be prioritized absolutely. Societies must make context-specific decisions about where to set the balance. The optimal balance likely differs across industries, harm types, and cultural contexts.

I also have not addressed distributional consequences. Structural reforms that increase personal liability for executives might reduce willingness to serve in leadership roles, potentially reducing corporate governance quality. Reforms that restrict fiduciary duty might reduce returns to shareholders, harming pension funds and retail investors. These distributional effects matter but are beyond this paper's scope.

VIII.D. Implementation Limitations

Even if the analysis is correct and the design principles are sound, implementation faces severe obstacles. Corporate interests will resist reforms that reduce flexibility, increase executive liability, or constrain shareholder primacy. Political economy dynamics favor status quo bias. International competition creates pressures to maintain business-friendly regulation.

Regulatory capture means that agencies tasked with enforcing structural reforms may be sympathetic to corporate concerns and reluctant to impose costly requirements. Legislative reforms require political will that may be lacking when corporate lobbying is intense. Judicial interpretation of reformed corporate law may revert toward shareholder primacy if judges are persuaded by economic efficiency arguments.

These obstacles do not make reform impossible, but they make it difficult. The gap between theoretical optimality and political feasibility is real.

IX. CONCLUSION

This paper has examined a puzzle: corporations operate as Level 1 entities (optimization algorithms without moral capacity), yet they are composed entirely of Level 3 humans (moral agents capable of recursive reasoning, shame, and intrinsic valuation of reputation). How do moral agents produce amoral outputs?

I have identified four structural mechanisms that transform human intentionality within corporate environments. Fiduciary duty legally requires managers to maximize shareholder value, eliminating discretion to prioritize ethical considerations when they conflict with profit. Incentive structures align personal wealth with corporate metrics, internalizing the objective function and transforming preferences. Diffusion of responsibility fragments decisions across organizational hierarchies, preventing any individual from feeling ownership of outcomes. The corporate veil creates legal and psychological distance between actors and consequences, enabling moral disengagement.

These mechanisms do not require bad people. They transform ordinary people into participants in harmful conduct through structural design. The Wells Fargo employees who created fake accounts, the Boeing engineers who designed MCAS with inadequate safeguards, and the Volkswagen programmers who implemented defeat devices were not villains. They were normal professionals operating within structures that disabled their moral reasoning.

This analysis explains why traditional enforcement fails. Enforcement against corporations fails because Level 1 entities cannot process moral sanctions. Enforcement against individuals fails because structure has already transformed their intentionality. Increasing penalties in either direction does not address the structural problem.

The solution requires redesigning structures to preserve moral capacity. This means reducing distance between actors and consequences, requiring personal ownership of decisions, decoupling compensation from short-term metrics, limiting the scope of fiduciary duty, and increasing transparency of the corporate veil. These reforms face a fundamental trade-off: they reduce organizational efficiency. The corporate form evolved to enable scale, specialization, and limited liability because these features are economically valuable. Preserving moral agency requires sacrificing some efficiency.

Societies must decide where to strike this balance. The current calibration strongly favors efficiency, treating moral agency preservation as at most a secondary concern. Reforming this

requires confronting the trade-off explicitly and making context-specific decisions about acceptable levels of transformation.

This paper connects Asymmetric Intentionality Theory (Lerer, 2025a-h) with organizational behavior, agency theory, and moral psychology. Previous papers showed that legal systems fail when they presuppose intentionality levels that targets lack. This paper shows that the problem is deeper than previously recognized: structure itself manufactures the mismatch by transforming Level 3 inputs into Level 1 outputs. Enforcement and governance reforms must address this transformation mechanism, not merely adjust penalties.

The corporate form is one of the most important institutional innovations in economic history. It enables productive activities that would be impossible under alternative organizational forms. But its current design contains an architectural flaw: it systematically transforms the moral capacity of the humans who compose it. Recognizing this is not an argument for abandoning the corporate form, but it is an argument for redesigning it to balance efficiency against the preservation of moral agency. How societies navigate that balance will determine whether corporate power serves human flourishing or undermines it.

ACKNOWLEDGMENTS AND AI DISCLOSURE

This paper extends the Asymmetric Intentionality Theory framework developed across a series of working papers, integrating insights from agency theory and moral psychology. The author gratefully acknowledges productive engagement with AI systems (Claude, Anthropic) in organizing the synthesis of multiple research streams and refining the articulation of structural mechanisms. All theoretical frameworks, legal interpretations, empirical observations, and errors remain the sole responsibility of the author.

REFERENCES

Arendt, H. (1963). *Eichmann in Jerusalem: A Report on the Banality of Evil*. Viking Press.

Ashforth, B. E., & Anand, V. (2003). The normalization of corruption in organizations. *Research in Organizational Behavior*, 25, 1-52.

Bandura, A. (1999). Moral disengagement in the perpetration of inhumanities. *Personality and Social Psychology Review*, 3(3), 193-209.

Bergstresser, D., & Philippon, T. (2006). CEO incentives and earnings management. *Journal of Financial Economics*, 80(3), 511-529.

Blair, M. M., & Stout, L. A. (1999). A team production theory of corporate law. *Virginia Law Review*, 85(2), 247-328.

CFPB (2016). Consumer Financial Protection Bureau fines Wells Fargo $100 million for widespread illegal practice of secretly opening unauthorized accounts. Press Release, September 8, 2016.

CNBC (2016). Volkswagen memos suggest emissions problem was known earlier. News Report, February 18, 2016.

Darley, J. M. (1996). How organizations socialize individuals into evildoing. In D. M. Messick & A. E. Tenbrunsel (Eds.), *Codes of Conduct: Behavioral Research into Business Ethics* (pp. 13-43). Russell Sage Foundation.

Darley, J. M., & Latané, B. (1968). Bystander intervention in emergencies: Diffusion of responsibility. *Journal of Personality and Social Psychology*, 8(4), 377-383.

Dennett, D. C. (1987). *The Intentional Stance*. Cambridge, MA: MIT Press.

DOJ (2017). Volkswagen AG pleads guilty and is sentenced to pay $4.3 billion in criminal and civil penalties. Press Release, April 21, 2017.

DOJ (2020). Wells Fargo agrees to pay $3 billion to resolve criminal and civil investigations into sales practices. Press Release, February 21, 2020.

DOJ (2021). Boeing reaches $2.5 billion agreement to resolve conspiracy fraud charge. Press Release, January 7, 2021.

Dodge v. Ford Motor Co., 204 Mich. 459 (1919).

eBay Domestic Holdings v. Newmark, 16 A.3d 1 (Del. Ch. 2010).

Efendi, J., Srivastava, A., & Swanson, E. P. (2007). Why do corporate managers misstate financial statements? The role of option compensation and other factors. *Journal of Accounting and Economics*, 43(1), 67-100.

Garrett, B. L. (2014). *Too Big to Jail: How Prosecutors Compromise with Corporations*. Cambridge, MA: Harvard University Press.

Hansmann, H., & Kraakman, R. (2001). The end of history for corporate law. *Georgetown Law Journal*, 89, 439-468.

Hart, O. (1995). *Firms, Contracts, and Financial Structure*. Oxford: Oxford University Press.

Jensen, M. C., & Meckling, W. H. (1976). Theory of the firm: Managerial behavior, agency costs and ownership structure. *Journal of Financial Economics*, 3(4), 305-360.

Johnson, S. A., Ryan, H. E., & Tian, Y. S. (2009). Managerial incentives and corporate fraud: The sources of incentives matter. *Review of Finance*, 13(1), 115-145.

Lerer, I. A. (2025a). Asymmetric Intentionality in Law and Game Theory: When Legal Systems Assume Moral Agency That Doesn't Exist. SSRN Working Paper.

Lerer, I. A. (2025b). The Dennett-Nash Equilibrium: Game Theory Under Heterogeneous Intentionality. SSRN Working Paper.

Lerer, I. A. (2025c). The Aumann Impossibility Theorem Under Heterogeneous Intentionality. SSRN Working Paper.

Lerer, I. A. (2025d). The Generalized Intentionality Mismatch Theorem. SSRN Working Paper.

Lerer, I. A. (2025e). Intentionality Mismatch in Legal Systems: A General Framework. SSRN Working Paper.

Lerer, I. A. (2025f). Corporate Criminal Liability and the Impossibility of Mens Rea. SSRN Working Paper.

Lerer, I. A. (2025g). Environmental Law and the Myth of Corporate Citizenship. SSRN Working Paper.

Lerer, I. A. (2025h). The Normative Calculation Problem Under Heterogeneous Intentionality. SSRN Working Paper.

Los Angeles Times (2016). Wells Fargo's high-pressure sales culture comes at a cost. September 9, 2016.

Milgram, S. (1963). Behavioral study of obedience. *Journal of Abnormal and Social Psychology*, 67(4), 371-378.

NYT (2017). VW engineers wanted O.K. from the top for emissions fraud. *New York Times*, May 17, 2017.

OCC (2021). Office of the Comptroller of the Currency assesses $250 million civil money penalty against Wells Fargo Bank. Press Release, September 9, 2021.

Paharia, N., Kassam, K. S., Greene, J. D., & Bazerman, M. H. (2009). Dirty work, clean hands: The moral psychology of indirect agency. *Organizational Behavior and Human Decision Processes*, 109(2), 134-141.

Rakoff, J. S. (2014). The financial crisis: Why have no high-level executives been prosecuted? *New York Review of Books*, January 9, 2014.

Stout, L. A. (2012). *The Shareholder Value Myth: How Putting Shareholders First Harms Investors, Corporations, and the Public*. San Francisco: Berrett-Koehler Publishers.

Tenbrunsel, A. E., & Messick, D. M. (2004). Ethical fading: The role of self-deception in unethical behavior. *Social Justice Research*, 17(2), 223-236.

U.S. Sentencing Commission (2021). *Recidivism of Federal Offenders Released in 2010*. Washington, D.C.