The Fatal Conceit of Consensus Networks: Why Evolutionary Epistemology Cannot Ground Constructivist Institutional Design

Ignacio Adrián Lerer
Independent Researcher
Buenos Aires, Argentina
adrian@lerer.com.ar
ORCID 0009-0007-6378-9749
Working Paper
January 2026

ABSTRACT

Alex Pentland's Shared Wisdom (MIT Press, 2025) celebrates emergent institutions while proposing designed "consensus networks" to achieve predetermined goals like the UN SDGs. This paper identifies the resulting contradiction, which I term the Pentland Paradox: one cannot simultaneously endorse evolutionary epistemology and constructivist design.

I formalize this critique through the Pentland Impossibility Theorem, proving that under dispersed knowledge and heterogeneous preferences, no designed coordination mechanism can simultaneously satisfy participation constraints, incentive compatibility, and goal achievement without coercion. Three empirical cases validate the theorem: UK Data Trusts (0% operational deployment despite £2M+ funding), EU ETS Phase I (99.7% price collapse from over-allocation), and GDPR (20% increased data costs with no measurable privacy improvement).

Drawing on Mises, Hayek, Darwin, and Dennett, I propose an Extended Phenotype Theory alternative: legal institutions are cultural replicators subject to selection pressures, not artifacts of intelligent design. The evolutionarily stable alternative is Pure Evolutionary Governance: no predetermined goals, ex-post liability only, federalism enabling jurisdictional competition, and negative selection (pruning failures) rather than positive design.

Steel-manning five defenses Pentland might offer, I show each collapses: "bottom-up design" conflates voluntary coordination with mandates; "irreversibility" arguments prove consensus networks are 26× slower than evolutionary adaptation; "democratic legitimacy" ignores that markets aggregate billions of daily preference revelations versus periodic elections. The choice is stark: evolution or engineering, not both.

Keywords: spontaneous order, data governance, AI alignment, praxeology, cultural evolution, data trusts, mechanism design, institutional economics

JEL Codes: B53, D83, K13, K20, L51, O33

I. INTRODUCTION: THE PENTLAND PARADOX

Alex Pentland's latest book (*Shared Wisdom: Cultural Evolution in the Age of AI* (MIT Press, November 11, 2025) arrives at a critical moment in AI governance debates. As artificial intelligence transforms economic and social coordination, regulatory frameworks have devolved into a false dichotomy: extensive ex-ante regulation (the EU AI Act approach) versus minimal oversight (libertarian-accelerationist position). Pentland attempts to transcend this binary by proposing evolutionary institutional design that harnesses digital technologies to strengthen community agency while avoiding centralized control.

The ambition is admirable. Pentland brings impressive credentials as founding director of MIT's Connection Science and over 100 peer-reviewed papers on social physics. Shared Wisdom synthesizes decades of empirical research from massive investment network studies to 20-year longitudinal analyses to national-scale experiments across four continents.

More importantly, Pentland's theoretical framework appears to align with evolutionary epistemology. He explicitly rejects ex-ante regulation of AI as "misguided" and epistemologically flawed, endorsing instead strict liability enforced ex-post as superior to preventive bureaucratic control. His critique of centralized planning invokes the failure of Soviet economic calculation with considerable sophistication. Throughout the book, he celebrates emergent institutions like common law and the Hanseatic League, recognizing that bureaucratic hierarchies fail precisely because they treat individuals as isolated rational actors rather than socially embedded learners whose knowledge is irreducibly dispersed.

Yet within this same analytical framework, Pentland advances a constructive program that sits uneasily with his descriptive insights. He proposes designing "consensus networks" to coordinate collective action toward the United Nations' Sustainable Development Goals. He advocates mandating universal "audit trails" for all AI systems as a precondition for operation. He envisions creating "data trusts" as mandatory intermediaries between individuals and corporations, institutions that would manage personal information on behalf of communities. And he suggests implementing "digital accounting" systems to monitor progress toward climate targets with the same precision we apply to financial flows.

This paper identifies what I term the Pentland Paradox: the simultaneous endorsement of evolutionary emergence and constructivist design, of epistemic humility and goal-directed institutional engineering, of spontaneous order and planned coordination.

I argue this paradox represents a fundamental logical fracture that vitiates Pentland's constructive proposals. The fracture occurs at a precise point: when Pentland transitions from describing how successful institutions emerge to prescribing how new institutions should be designed. At this juncture, he abandons the evolutionary epistemology that correctly explained historical successes and reverts to the constructivist rationalism his own analysis demonstrates has repeatedly failed.

Structure: Part II establishes what Pentland gets right. Part III identifies the precise logical fracture point. Parts IV-VI develop complementary critiques through Misesian praxeology, Hayekian-Darwinian-Dennettian synthesis, and Extended Phenotype Theory. Part VII presents the steel-man defense Pentland might offer. Part VIII concludes with a coherent evolutionary alternative.

II. WHAT PENTLAND GETS RIGHT

Before dissecting where Pentland's analysis fractures, I must acknowledge where it succeeds. Any critique that fails to recognize genuine contributions risks strawman argumentation.

A. The Story-Sharing Model of Cultural Evolution

Pentland's central thesis that human intelligence is fundamentally social finds substantial empirical support:

"This recombinant development of community wisdom is the core process that drives cultural evolution." (p. 11)

This framework correctly identifies several phenomena that previous institutional analyses have underappreciated.

First, Pentland treats social learning as the evolutionary substrate of cultural change. His emphasis on "stories" (narratives encoding action-outcome patterns) as the unit of cultural transmission aligns with memetic theory as developed by Dawkins (1976) and Dennett (1995), while providing the operational specificity that earlier memetic accounts sometimes lacked. More directly, Pentland's framework resonates with Joseph Henrich's work on cultural evolution. In The Secret of Our Success (2016) and The WEIRDest People in the World (2020), Henrich demonstrates that humans are fundamentally "cultural learners" who depend on socially transmitted information rather than individual trial-and-error. Henrich's mechanisms of prestige-biased transmission (copying successful individuals) and content-biased transmission (preferentially retaining memorable, causally coherent narratives) map precisely onto Pentland's story-sharing networks. Yuval Noah Harari's

Nexus (2024) extends this insight: large-scale human cooperation depends on shared narratives that coordinate expectations across strangers. Stories are not merely information; they are compressed representations of causal relationships that can be evaluated, compared, and selectively retained across generations. This convergence across Pentland, Henrich, and Harari strengthens the descriptive framework considerably.

Second, Pentland demonstrates that network topology determines the rate of cultural evolution. His taxonomy distinguishes three network types: campfire networks for bonding within close-knit groups, bridging networks for exploration across group boundaries, and consensus networks for formalizing shared understandings. His empirical work validates the intuition that innovation correlates with bridging network density (exposure to diverse ideas) while cultural stability correlates with bonding network density (reinforcement of existing practices). This provides a structural explanation for why some societies innovate rapidly while others maintain tradition.

Third, Pentland distinguishes between deductive or engineering approaches (proceeding from axioms to optimal solutions) and abductive or experiential approaches (proceeding from observations to plausible explanations). He correctly identifies that complex adaptive systems require abductive approaches because the underlying causal structure is both unknown and dynamic. This epistemological point has profound implications: it means that the "best" institution cannot be derived from first principles but must be discovered through experimentation.

Fourth, and perhaps most valuably, Pentland provides rare large-scale quantitative evidence for claims that have often remained merely theoretical. His eToro investment network study demonstrates that social learning produces superior returns compared to isolated decision-making, a finding that validates the evolutionary logic of copying successful strategies. His Future of Families longitudinal study, spanning twenty years, shows that neighborhood network structure predicts life outcomes better than genetic or individual psychological measures. These empirical contributions ground the theoretical framework in observable reality.

B. The Devastating Critique of Ex-Ante Regulation

Pentland's opposition to the EU AI Act represents his analysis at its strongest:

"Ex-ante regulation is misguided because it focuses on the complexity of AI systems and how they are developed rather than on the tractable set of possible harms these systems are capable of." (p. 134)

This critique succeeds on multiple levels that deserve elaboration.

The epistemic impossibility of ex-ante regulation becomes clear once we recognize what it assumes. Regulators must predict future technological trajectories, anticipate harms that have not yet materialized, and design rules flexible enough to accommodate innovations that do not yet exist. Pentland correctly identifies this as the same epistemic hubris that doomed Soviet planning:

"In 1959 Leonid Kantorovich published his book Best Use of Economic Resources and the Soviet Union began to manage its economy using his system. Despite this research winning a Nobel Prize in Economics, the Soviet Union's experiment in optimal allocation of resources was a disaster that contributed to its eventual dissolution." (p. 79)

The parallel is exact and devastating. Kantorovich was not a fool; he was a brilliant mathematician who developed linear programming techniques still used today. The Soviet system failed not because its architects lacked intelligence but because they faced an impossible information problem. EU regulators attempting to anticipate AI harms face the same impossible knowledge problem, differing only in the domain of application.

Pentland also illuminates the distinction between process regulation and outcome regulation. Ex-ante approaches regulate process (how AI is developed, what documentation must be maintained, what review boards must approve) rather than outcome (what concrete harms the technology causes). This is inefficient for several reasons that compound each other. Multiple processes can achieve the same safety outcome, so mandating one process forecloses potentially superior alternatives. Even safe processes can produce harmful outcomes when deployed in unanticipated contexts, so process compliance provides false assurance. And process compliance creates rent-seeking opportunities for compliance professionals without necessarily improving safety outcomes.

The liability alternative that Pentland endorses has both theoretical elegance and practical advantages. Strict product liability, enforced ex-post through courts, aligns incentives by making developers internalize the costs of harms they cause. It discovers information through adversarial process, as plaintiffs have incentives to identify harms and defendants have incentives to demonstrate safety. It evolves adaptively through case-by-case adjudication, allowing doctrine to respond to novel situations without requiring legislative foresight. And it minimizes administrative costs by avoiding the bureaucratic apparatus of pre-market review.

C. Recognition of Dispersed Knowledge

Pentland explicitly acknowledges the knowledge problem that plagues centralized institutions. He recognizes bureaucratic hierarchies fail because they assume complete information at the top, treat individuals as isolated rather than socially embedded, cannot adapt to local variation, and create principal-agent problems.

His solution drawing on Elinor Ostrom's principles (governance by the governed, clear boundaries, graduated sanctions, conflict resolution mechanisms, autonomy) represents elaborate institutional analysis.

D. Admiration for Emergent Institutions

Pentland's discussion of historically successful institutions demonstrates genuine appreciation for spontaneous order:

Common Law: Pentland's treatment of common law as an evolutionary system deserves particular attention because it represents his analysis at its most insightful:

"In law, the incentive for consensus is reputation: When a legal decision has lots of citations, that means that it is literally precedent-setting (creating common law) and is seen as a badge of honor for the judge." (p. 70)

What Pentland correctly identifies here is a system where no central planner designs the overall structure of legal outcomes. Each case resolves a specific dispute between particular parties with concrete interests at stake. Precedents accumulate through citation as later judges find earlier decisions useful for resolving analogous disputes. Rules that prove useful in this sense survive and spread through the legal system, while rules that prove unworkable are distinguished, limited, or eventually overruled. The entire system adapts through incremental variation, with each decision representing a small experiment whose success or failure is revealed through subsequent adoption or rejection.

This is genuine cultural evolution in action, and Pentland describes it with admirable clarity. The key insight is that common law's beneficial properties emerge as unintended consequences of judges pursuing much narrower objectives: resolving the dispute before them in a way that seems reasonable given prior practice. No judge aims to "optimize the legal system"; judges aim to decide cases. Yet the aggregate effect of thousands of such decisions, filtered through citation and precedent, produces a remarkably adaptive body of doctrine.

Other Examples: The Hanseatic League exemplifies voluntary merchant cooperation emerging to solve concrete trading problems, not to achieve abstract "economic development." The Uniform Law Commission demonstrates bottom-up aggregation of legal solutions that have proven workable across jurisdictions. Scientific citation networks connect researchers pursuing their own chosen projects, with the aggregate effect of directing attention toward productive research programs. In each case, the institutions that Pentland admires share a crucial characteristic: they document and formalize what has already been shown to work rather than attempting to design and impose what planners believe should work.

E. Summary: The Valuable Core

Pentland's genuine contributions:

✓ Correct epistemology: Cultural evolution through story-sharing, abductive reasoning
✓ Correct critique: Ex-ante regulation is epistemically impossible
✓ Correct alternative: Strict liability ex-post aligns incentives
✓ Correct diagnosis: Centralized bureaucracies fail due to knowledge problems
✓ Correct examples: Common law, Hanseatic League succeed through emergence not design

The problem arises when Pentland attempts to leverage these insights to justify constructive proposals that violate the very principles his analysis establishes.

III. THE LOGICAL FRACTURE POINT

Every analytical framework that combines descriptive accuracy with prescriptive error contains a logical fracture point. In Pentland's case:

The Fracture Point: When Pentland assumes "collective goals" (SDGs, climate targets) are legitimate inputs for institutional design rather than emergent outputs requiring their own evolutionary justification.

A. The Transition from Emergent to Designed Institutions

Observe Pentland's reasoning in two sequential passages:

DESCRIPTIVE (explaining historical success): > "These consensus network institutions drove an explosion in scientific knowledge, a revolution in manufacturing and medicine, and the governance of it all by the case-by-case decisions that make up common law." (p. 74)

Pentland correctly describes institutions that emerged without design. Note what is absent: no predetermined collective goal. Common law did not emerge to "achieve justice"; it emerged to resolve concrete disputes. These institutions' beneficial aggregate effects were unintended consequences, not design objectives.

PRESCRIPTIVE (proposing new institutions): > "By leveraging the ideas of community and the core human consensus processes that build community intelligence and shared wisdom, we can design digital networks and AI that have the potential to avoid the problems we experience today and provide our societies with a much greater capability for achieving rational deliberation and consensus for collective action." (pp. 27-28)

Here Pentland proposes institutions designed toward predetermined goals: consensus networks to coordinate toward SDGs, data trusts to represent "community interests," audit trails to enable "accountability."

Note what is added: predetermined collective goals. These are abstractions introduced as design objectives, not concrete problems emerging from actual conflicts.

B. The Unexamined Assumption

The fracture hinges on an unexamined assumption:

Assumption: There exist legitimate "collective goals" (SDGs, climate targets) that can serve as design objectives for institutional engineering.

This assumption creates three interlocking problems:

1. Ontological Problem: Who Has These Goals?

Pentland's framework establishes that only individuals have goals. But whose goals are the SDGs? If "all stakeholders" genuinely agreed, no special institutions would be needed (people would voluntarily cooperate). If some disagree, then SDGs represent preferences of some individuals imposed on others. The "consensus network" thus becomes a mechanism for legitimizing coercion through participatory process.

2. Epistemological Problem: How Do We Know These Are Right?

Pentland's epistemology emphasizes dispersed knowledge, learning through experimentation, and that central planning fails because planners lack local knowledge. Yet the SDGs were defined by UN bureaucrats, government representatives, and international committees without electoral accountability. They were not the product of campfire networks, bridging networks testing innovations, or consensus networks documenting proven solutions. The SDGs violate Pentland's own criteria for legitimate "shared wisdom."

3. Logical Problem: Infinite Regress

Who designs the institutions? Consensus networks of stakeholders. Who decides who counts as a "stakeholder"? The institutions being designed. Who decides what the institutions should achieve? The collective goals (SDGs). Who decided the collective goals? Previous institutions (UN, governments). What legitimizes those? The regress terminates in circular reasoning or arbitrary assertion.

C. Three Specific Contradictions

Contradiction 1: Liability vs. Mandatory Audit Trails

Pentland endorses ex-post liability but then proposes universal mandatory audit trails:

“Practically, this requires all AI deployments to keep easily accessible audit logs of actions they take, along with a liability regime that makes it easy to use these logs to assign liability.” (p. 133)

“The approach of open audit trails enables both rapid detection of harm and rapid regulatory enforcement.” (p. 127)

This is ex-ante regulation of process, precisely what Pentland claimed to reject. Strict liability works through standard burden of proof: plaintiff demonstrates defendant’s action, causation, and harm. Mandatory audit trails reverse this logic, requiring defendants to document processes ex-ante as a precondition for operation.

The legitimate alternative: Pure liability without mandatory audit trails. Defendants who voluntarily maintain detailed records gain advantage in litigation, creating market incentive without regulatory mandate.

Contradiction 2: Emergent vs. Designed Networks

Institutions Pentland admires solved problems people actually experienced: common law resolves disputes brought by actual litigants, Hanseatic League facilitated trade merchants wanted to conduct, scientific networks connect researchers pursuing chosen projects.

Pentland’s proposed institutions solve problems designers claim exist: data trusts “protect” people from data sharing they might voluntarily choose, consensus networks coordinate toward SDGs individuals didn’t endorse through revealed preference, audit requirements prevent harms not demonstrated ex-ante.

If these problems were genuinely pressing, why haven’t institutions already emerged to solve them?

Contradiction 3: Knowledge Problem Acknowledged, Then Ignored

Pentland acknowledges Soviet planning failed because “models failed to capture the complexity or dynamism of society.” Yet he proposes “digital accounting” and audit trails to monitor progress toward SDGs.

Kantorovich’s system was also “data-driven.” The USSR collected vast production statistics and resource inventories. The system failed because: (1) tacit knowledge cannot be articulated, (2) data describes the past but decisions affect the future, (3) actors game metrics when metrics become targets (Goodhart’s Law).

Pentland’s digital accounting faces identical problems. The Soviet parallel is exact: cybernetic governance differing from Soviet planning only in computational power, not logical structure.

D. Why the Fracture Matters

The fracture is not merely logical but causal. Pentland's proposals will fail for the same reasons Soviet planning failed: designers lack knowledge necessary to optimize complex systems toward predetermined ends. The theoretical contradiction predicts practical failure.

E. Formalizing the Fracture Point

Having identified Pentland's logical fracture point intuitively (Sections III.A-D), I now formalize it rigorously. This formalization serves three purposes: (1) clarifying the precise logical structure of the Pentland Paradox, (2) demonstrating that the fracture is not merely rhetorical but follows from incompatible formal assumptions, and (3) generating testable predictions.

The Two Incompatible Frameworks

Framework 1: Evolutionary Emergence (Pentland's Descriptive Analysis)

Let $\mathcal{I}$ denote the space of possible institutional arrangements. An institution $I \in \mathcal{I}$ at time $t+1$ evolves from I_t according to:

$$I_{t+1} = I_t + \delta_t$$

where δ_t represents incremental variation arising from:

1. Stories: Narratives encoding action-outcome patterns, transmitted and recombined across individuals

2. Networks: Social structures (campfire, bridging, consensus) that determine information flow

3. Selection: Retention of practices that "work" in the sense of resolving concrete problems for actual participants

Critically, δ_t is not chosen by any designer. It emerges from:

$$\delta_t \sim D(K_1, K_2, \ldots, K_N; \Theta_t)$$

where:

- K_i represents dispersed knowledge held by individual i
- Θ_t represents environmental conditions at time t
- $D(\cdot)$ is a discovery distribution over institutional variations

Key property: No central agent possesses:

$$K = \bigcup_{i=1}^{N} K_i$$

(Hayek, 1945). Knowledge is:

- **Dispersed:** $K_i \cap K_j$ is small for most $i \neq j$
- **Tacit:** Much of K_i cannot be articulated (Polanyi, 1966)
- **Context-dependent:** K_i depends on local conditions unobservable to distant observers
- **Dynamic:** $K_i(t)$ changes as conditions evolve

Functional success: An institution I survives if it facilitates coordination for participants given their dispersed knowledge:

$$\text{Survival}(I) \propto E_{\{K_i\}} [U_i(I \mid K_i)]$$

where $U_i(\cdot)$ is individual i's utility and the expectation is taken over the distribution of knowledge types.

Examples Pentland describes correctly:

- **Common law:** Each case c_t generates precedent δ_t based on judges' local knowledge K_{judge} of dispute facts; the aggregate body of law emerges as:

$$I_t = \sum_{s=1}^{t} \delta_s$$

without any designer specifying an "optimal legal system".

- **Hanseatic League:** Merchants adopt trading rules δ_t that facilitate transactions given knowledge of local conditions; successful rules spread through imitation.

Framework 2: Constructivist Design (Pentland's Prescriptive Proposals)

Under the constructivist framework, institutions are designed toward predetermined goals. Let G represent collective goals:

G = { g_1, g_2, …, g_M }

(e.g., G = \text{SDGs} with M = 17 goals and 169 targets).

A designer (or a consensus network of designers) chooses institution I^* by solving:

I^* = arg max_(I ∈ $\mathcal{I}$) Ψ(I | G)

subject to:

$\mathcal{C}$(I) ≤ B

where:

- Ψ(I | G) measures how well I achieves goals G
- $\mathcal{C}$(I) represents the costs of implementing I
- B represents resource constraints

Key assumptions (mostly implicit in Pentland)

1. **Goal legitimacy:** G represents "collective" objectives that justify imposing I^* on the population
2. **Designer knowledge:** Designers know (or can learn) the mapping:

Ψ : $\mathcal{I}$ × G → ℝ

3. **Static optimization:** Optimal I^* can be determined ex-ante before implementation
4. **Enforcement:** Non-compliance with I^* can be detected and sanctioned

Pentland's proposals

- **Data trusts:** Institution I^*_{trust} where trustees maximize "community welfare" (a subset of G)
- **Consensus networks:** Institution I^*_{consensus} coordinating stakeholders toward SDGs (G)
- **Audit trails:** Mandated institution I^*_{audit} requiring universal documentation of AI actions

The Pentland Impossibility Theorem

The Two Incompatible Frameworks

Framework 1: Evolutionary Emergence (Pentland's Descriptive Analysis)

Let $\mathcal{I}$denote the space of possible institutional arrangements. An institution $I \in \mathcal{I}$at time $t + 1$evolves from $I_t$according to:

$$I_{t+1} = I_t + \delta_t$$

where $\delta_t$represents incremental variation arising from:

1. **Stories**: Narratives encoding action-outcome patterns, transmitted and recombined across individuals
2. **Networks**: Social structures (campfire, bridging, consensus) that determine information flow
3. **Selection**: Retention of practices that "work" in the sense of resolving concrete problems for actual participants

Critically, $\delta_t$is **not chosen** by any designer. It emerges from:

$$\delta_t \sim D(K_1, K_2, \ldots, K_N; \Theta_t)$$

where:

- $K_i$represents dispersed knowledge held by individual i
- $\Theta_t$represents environmental conditions at time t
- $D(\cdot)$is a **discovery distribution** over institutional variations

Key property: No central agent possesses $K = \bigcup_{i=1}^{N} K_i$(Hayek, 1945). Knowledge is:

- **Dispersed**: $K_i \cap K_j$is small for most $i \neq j$
- **Tacit**: Much of $K_i$cannot be articulated (Polanyi, 1966)
- **Context-dependent**: $K_i$depends on local conditions unobservable to distant observers
- **Dynamic**: $K_i(t)$changes as conditions evolve

Functional success: An institution Isurvives if it facilitates coordination for participants given their dispersed knowledge:

$$\text{Survival}(I) \propto \mathbb{E}_{\{K_i\}}[U_i(I \mid K_i)]$$

where $U_i(\cdot)$is individual i's utility and expectation is taken over the distribution of knowledge types.

Examples Pentland describes correctly:

- **Common law**: Each case $c_t$generates precedent $\delta_t$based on judges' local knowledge (K_{judge}) of dispute facts; aggregate body of law $I_t = \sum_{s=1}^{t} \delta_s$emerges without any designer specifying "optimal legal system"
- **Hanseatic League**: Merchants adopt trading rules $\delta_t$that facilitate their transactions given knowledge of local conditions; successful rules spread through imitation

Framework 2: Constructivist Design (Pentland's Prescriptive Proposals)

Under the constructivist framework, institutions are **designed** toward predetermined goals. Let $G = \{g_1, g_2, \ldots, g_M\}$represent collective goals (e.g., $G =$ SDGswith $M = 17$goals and 169 targets).

A designer (or consensus network of designers) chooses institution $I^*$by solving:

$$I^* = \arg\max_{I \in \mathcal{I}} \Psi(I \mid G)$$

subject to:

$$\mathcal{C}(I) \leq B$$

where:

- $\Psi(I \mid G)$measures how well Iachieves goals G
- $\mathcal{C}(I)$represents costs of implementing I
- Brepresents resource constraints

Key assumptions (mostly implicit in Pentland):

1. **Goal legitimacy**: Grepresents "collective" objectives that justify imposing I^*on population
2. **Designer knowledge**: Designers know (or can learn) the mapping $\Psi: \mathcal{I} \times G \to \mathbb{R}$
3. **Static optimization**: Optimal I^*can be determined ex-ante before implementation
4. **Enforcement**: Non-compliance with I^*can be detected and sanctioned

Pentland's proposals:

- **Data trusts**: Institution $I^*_{trust}$where trustees maximize "community welfare" (a subset of G)
- **Consensus networks**: Institution $I^*_{consensus}$coordinating stakeholders toward SDGs (G)
- **Audit trails**: Mandate I^*_{audit}requiring universal documentation of AI actions

The Pentland Impossibility Theorem

Theorem (Pentland Impossibility): If an institution Iemerged successfully via evolutionary process $D(\cdot)$under dispersed knowledge $\{K_i\}$, then any designed institution I^*optimized toward predetermined goals Gwill fail to replicate I's success.

Proof (by contradiction):

Assume designed institution I^*can succeed as well as emergent institution I.

Step 1: For $I^*$to succeed, designers must choose the optimal institution. This requires solving:

$$I^* = \arg\max_{I \in \mathcal{I}} \Psi(I \mid G)$$

Step 2: The mapping $\Psi(I \mid G)$depends on how individuals respond to institution Igiven their knowledge $K_i$and environmental conditions Θ:

$$\Psi(I \mid G) = \mathbb{E}_{\{K_i, \Theta\}}\left[\sum_{m=1}^{M} \omega_m \cdot g_m(I, \{K_i\}, \Theta)\right]$$

where $\omega_m$are weights on goal $g_m$and $g_m(\cdot)$measures achievement of goal m.

Step 3: To compute $\Psi(I \mid G)$, designers must know:

1. The distribution of knowledge types $\{K_i\}$across population
2. How each type responds to institution: $K_i \mapsto \text{Action}_i(I \mid K_i)$
3. The mapping from actions to goal achievement: $\{\text{Action}_i\} \mapsto g_m$
4. Future environmental conditions Θthat affect all above

Step 4: But by assumption of dispersed knowledge:

$$K = \bigcup_{i=1}^{N} K_i \text{ and} K_{designer} \subsetneq K$$

Therefore, $K_{designer} \supseteq/K$, so designers **cannot** compute $\Psi(I \mid G)$accurately.

Step 5: The emergent institution Isucceeded precisely because it incorporated knowledge Kthrough the evolutionary process:

$$\delta_t \sim D(K_1, \ldots, K_N; \Theta_t)$$

Each variation $\delta_t$reflected someone's local knowledge; variations that "worked" (incorporated accurate local knowledge) were retained.

Step 6: The designed institution I^*cannot incorporate Kbecause:

- Designers possess only $K_{designer} \subsetneq K$
- Optimization occurs **ex-ante** (before observing how Iperforms in practice)
- No mechanism to aggregate dispersed knowledge (no "selection" on designs)

Conclusion

Designers choosing I^*based on $K_{designer}$will systematically fail to match the success of Iwhich emerged from selection over K. This contradicts the assumption that I^*can succeed as well as I.

Four Corollaries: Failure Modes

Corollary 1 (Epistemic Impossibility)

Statement: Designers optimizing toward Gwill choose institutions that perform worse than emergent institutions on participants' actual objectives.

Proof sketch:

- Emergent institutions satisfy: $\mathbb{E}_{\{K_i\}}[U_i(I \mid K_i)] > \mathbb{E}_{\{K_i\}}[U_i(I' \mid K_i)]$for alternative I'(selection eliminates worse alternatives)
- Designed institutions satisfy: $\Psi(I^* \mid G) \geq \Psi(I' \mid G)$(by construction)
- But $\Psi(\cdot \mid G) \neq \mathbb{E}_{\{K_i\}}[U_i(\cdot \mid K_i)]$unless Gperfectly aligns with individuals' objectives (unlikely given heterogeneity)
- Therefore I^*optimizes wrong objective. □

Empirical prediction: Data trusts will fail to satisfy participants (trustees optimize for "community welfare" ≠ principals' revealed preferences).

Evidence (Section VI.A): UK data trust pilots remained exploratory; "bespoke structure" required for each case; no general model emerged.

Corollary 2 (Brittleness)

Statement: Institutions I^*optimized for current conditions $\Theta_0$fail when conditions change to $\Theta_1 \neq \Theta_0$.

Proof sketch:

- Designed institution solves: $I^* = \arg\max_I \Psi(I \mid G, \Theta_0)$at time $t = 0$

- But conditions evolve: $\Theta_t \sim F(\Theta_{t-1})$where $F$is unpredictable
- At $t = 1$with $\Theta_1 \neq \Theta_0$, optimal institution is $I_1^* = \arg\max_I \Psi(I \mid G, \Theta_1) \neq I_0^*$
- Designed institution is "locked in" to I_0^*(requires redesign to change)
- Emergent institution adapts: $I_1 = I_0 + \delta_1$where $\delta_1 \sim D(K, \Theta_1)$(automatic adaptation)

Therefore I^*becomes increasingly mismatched as Θevolves. □

Empirical prediction: Consensus networks coordinating toward fixed goals (SDGs) fail when conditions shift.

Evidence (Section VI.B): EU ETS Phase I optimized for expected baseline emissions; collapsed when actual emissions revealed (14.3% over-allocation).

Corollary 3 (Capture)

Statement: Institutions I^*designed by subset $D \subset N$will be captured by D's interests rather than serving population N.

Proof sketch:

- Designers Dchoose $I^*$to maximize: $\Psi(I^* \mid G, K_D)$using their knowledge K_D
- But designers also have preferences: $U_d(I, outcomes)$for $d \in D$
- When choosing I^*, designers face trade-off: $$\max_{I} \left[\alpha \cdot \Psi(I \mid G, K_D) + (1-\alpha) \cdot U_d(I) \right]
- Even with $\alpha \approx 1$(designers try to be impartial), their beliefs about $\Psi$are shaped by K_D, which reflects their position
- Result: I^*systematically favors outcomes familiar/beneficial to D

Emergent institutions avoid because selection operates on **outcomes**, not designers' intentions. □

Empirical prediction: Data trusts captured by trustees (privacy activists); consensus networks captured by incumbents.

Evidence (Section VI.A): Ada Lovelace (2021) notes concern that "maximizing income from data may conflict with privacy goals" (trustee vs. principal preferences diverge).

Corollary 4 (Legitimacy Crisis)

Statement: Designed institutions I^*require coercion to enforce compliance; emergent institutions Irely on voluntary participation.

Proof sketch:

- Under emergence, institution I survives only if: $\mathbb{E}_i[U_i(I \mid K_i)] \geq \mathbb{E}_i[U_i(\text{no institution} \mid K_i)]$(voluntary participation constraint)
- Under design, I^* is imposed if $\Psi(I^* \mid G) \geq$threshold, **regardless** of whether individuals prefer it
- Therefore I^* may satisfy $\mathbb{E}_i[U_i(I^* \mid K_i)] < \mathbb{E}_i[U_i(\text{no institution} \mid K_i)]$for some/many i
- These individuals will not voluntarily comply; requires coercion. □

Empirical prediction: Designed institutions require mandatory participation (audit trails, data trust membership, SDG compliance).

Evidence (Section VI.C): GDPR mandates audit trails; non-compliance fined up to 4% global revenue. Pentland explicitly proposes "**mandatory** audit trails" (p. 133, emphasis added).

Why the Fracture Cannot Be Bridged

Pentland might attempt to reconcile the two frameworks by arguing:

> "We can design institutions that facilitate emergence rather than direct outcomes."

Response: This attempted bridge fails because:

1. Any design embeds assumptions about what outcomes are valuable:

- "Facilitating emergence" requires deciding what counts as "good" emergence

- Data trusts assume data sharing needs intermediation (trustees)

- Consensus networks assume coordination toward SDGs is desirable

- These are value judgments, not neutral facilitation

2. The choice of facilitating structure IS the design:

- Choosing data trusts vs. cooperatives vs. contractual mechanisms embeds assumptions about optimal governance

- Each "facilitating" structure privileges certain emergent outcomes over others

- Therefore "facilitating emergence" collapses into "designing toward preferred emergence" (which is just design)

3. Truly neutral facilitation would be:

- Legal framework allowing any voluntary arrangement (including no data trusts, no audit trails, no SDG coordination)

- Ex-post liability for demonstrated harms

- No ex-ante mandates

But this is not what Pentland proposes. He proposes mandatory data trusts, audit trails, and coordination toward specific goals (SDGs).

Conclusion: The attempt to bridge emergence and design fails because any non-neutral facilitation is design, and neutral facilitation is just legal minimalism (which Pentland rejects).

Implications for Pentland's Specific Proposals

Data Trusts

What evolution would produce: Varied voluntary arrangements (some individuals share data directly; others form cooperatives; others sell to brokers; others don't share)

What Pentland designs: Trustees stewarding data toward "community interests"

Fracture: Trustees must define "community interests" using $K_{trustee} \subsetneq K_{community}$; will systematically misalign with principals' revealed preferences.

Prediction: Principal-agent problems; capture; low participation.

Evidence (Section VI.A): Confirmed.

Consensus Networks for SDGs

What evolution would produce: Diverse jurisdictions experiment with policies; successful approaches spread through imitation; no consensus required

What Pentland designs: Stakeholder deliberation coordinating toward predetermined SDGs

Fracture: SDGs determined by $K_{UN} \subsetneq K_{global~population}$; stakeholders lack knowledge to optimize globally; strategic gaming.

Prediction: Coordination failure; gaming; brittleness.

Evidence (Section VI.B): EU ETS Phase I confirms.

Mandatory Audit Trails

What evolution would produce: Firms voluntarily maintaining records to defend against liability claims; level of documentation determined by cost-benefit trade-off

What Pentland designs: Universal mandate to document AI actions regardless of cost-benefit

Fracture: Mandate imposes uniform rule assuming $\mathbb{E}$[Benefit] > Costfor all firms; but firms have heterogeneous costs and benefits; many firms face Cost > Benefit.

Prediction: Excessive compliance costs; no harm reduction; chilling effect on innovation.

Evidence (Section VI.C): Demirer et al. (2024) confirms: +20% data costs; -26% data usage; €16B/year; no breach reduction.

Summary: The Fracture is Fundamental

The Pentland Paradox is not a minor inconsistency; it reflects incompatible ontologies:

Dimension	Evolutionary Framework	Constructivist Framework
Knowledge	Dispersed, tacit, dynamic	Centralizable, articulable, static
Optimization	Ex-post selection	Ex-ante design
Goals	Emergent from coordination	Predetermined (SDGs)
Legitimacy	Voluntary participation	Imposed with enforcement
Adaptation	Continuous, decentralized	Redesign, centralized
Success metric	Survival (revealed preference)	Achievement of G (designer's measure)

No middle ground exists because these frameworks rest on fundamentally different epistemological assumptions. Choosing one commits you to rejecting the other.

Pentland's error: describing institutional success using evolutionary framework, then prescribing institutional design using constructivist framework.

Consequence: His prescriptions fail for reasons his own analysis predicts.

IV. THE MISESIAN CRITIQUE: ONLY INDIVIDUALS ACT

Ludwig von Mises provides the most direct refutation through praxeology, the science of human action. The Misesian critique operates at the ontological level: Pentland's key concepts ("collective goals," "consensus networks," "data trusts") commit category errors about what kinds of entities can act and have purposes.

A. The Praxeological Foundation

Mises establishes in Human Action:

"Human action is purposeful behavior… Action is will put into operation and transformed into an agency, is aiming at ends and goals." (p. 11)

From this axiom, Mises derives methodological individualism:

“All actions are performed by individuals. A collective operates always through the intermediary of one or several individuals whose actions are related to the collective as the secondary source.” (p. 42)

The essential implication: Collectives (societies, nations, consensus networks) do not act. They are conceptual abstractions describing patterns in individual actions. Attributing action or purpose to collectives either serves as metaphorical shorthand or constitutes analytical error that obscures whose interests are actually being served.

B. Translating Pentland’s Proposals

When translated into individualist terms, Pentland’s core proposal reveals itself:

Pentland’s Language: “We can design digital networks and AI that provide our societies with capability for achieving rational deliberation and consensus for collective action.”

Praxeological Translation: - “Community intelligence” = certain individuals’ beliefs about what information is reliable - “Shared wisdom” = statistical regularity in individuals’ choices - “Collective action” = coordination among individuals pursuing compatible goals - “Design networks” = some individuals (designers) deciding rules that constrain other individuals’ choices

When translated this way, the question shifts from “How can we achieve collective goals?” to “Which individuals get to decide what goals others must pursue?”

C. The SDGs: Whose Goals?

Applying methodological individualism to the Sustainable Development Goals reveals their problematic status as design inputs. The question of who decided these are goals humanity should pursue admits only one honest answer: specific individuals occupying positions within international institutions. UN bureaucrats drafted the language. Government representatives negotiated the formulations. NGO activists lobbied for inclusion of their preferred objectives. These are identifiable human beings with their own interests, ideologies, and career incentives.

Did all affected individuals participate in this process? Obviously not. Most humans have never heard of the SDGs, let alone endorsed them through any mechanism that could be called consent. The farmer in rural India, the factory worker in Vietnam, the entrepreneur in Nigeria have never been consulted about whether they wish to subordinate their activities to these seventeen goals and 169 targets.

Do the SDGs reflect “shared wisdom” accumulated through the story-sharing networks Pentland describes? Again, the answer must be negative. The SDGs were drafted in committee rooms in New York, not discovered through the decentralized experimentation

that produced common law or the Hanseatic trading rules. They emerged from a process that is the antithesis of evolutionary emergence: top-down design by a small group claiming to speak for humanity.

In what sense, then, are these "collective goals"? The praxeological answer is uncomfortable but unavoidable: they are goals of certain powerful individuals who have successfully claimed authority to speak for "humanity" without any mechanism establishing that authority. The language of "collective goals" obscures the reality of particular interests being advanced under universalist rhetoric.

This distinction matters for accountability, legitimacy, and flexibility. If the SDGs fail to improve human welfare, whose judgment stands indicted? If populations resist SDG implementation, is their resistance mere ignorance to be overcome, or legitimate preference diversity to be respected? If circumstances change and the SDGs prove misguided, can they be revised without the pretense that "humanity" has changed its mind?

D. Data Trusts: Principal-Agent Problems

Pentland proposes data trusts with evident enthusiasm:

"Data trusts could help community members manage their data and what they expose to companies… Just as we invented credit unions to help community members manage their money, we can have data trusts to help citizens manage their data." (pp. 121-122)

The credit union analogy is revealing, but not in the way Pentland intends. Praxeological analysis immediately identifies the central problem: a data trust does not "manage" data in any meaningful sense. Trustees, who are individual human beings with their own interests and incentives, make decisions about data. The trust is a legal abstraction; only the trustees act.

This creates an inescapable principal-agent problem. Trustees will maximize their own utility, which includes professional reputation within activist communities, ideological satisfaction from advancing preferred policies, and continued employment in the data trust ecosystem. The interests of the diffuse principals (the community members whose data is ostensibly being protected) enter the trustees' decision-making only insofar as serving those interests advances the trustees' own objectives.

Historical precedent offers little comfort. Every "trust" created to represent diffuse interests has eventually been captured by organized minorities with concentrated stakes. Labor unions ostensibly represent workers but are in practice controlled by professional staff pursuing political agendas that may diverge substantially from rank-and-file preferences. Public interest law firms claim to represent "the public" but advance specific ideological visions selected by their founders and donors. Pension fund trustees increasingly pursue ESG

objectives that may reduce returns, prioritizing trustees' reputational interests over beneficiaries' financial welfare.

Why would data trusts avoid this pattern? Pentland provides no mechanism to prevent trustee capture by privacy activists who wish to restrict data sharing for ideological reasons, even when the principals themselves would voluntarily share data for personal benefit. The asymmetry is predictable: activists care intensely about data policy and will invest heavily in controlling trust governance, while ordinary individuals have diffuse interests and face high costs of monitoring trustee behavior.

The Misesian point bears emphasis: trusts don't have interests; trustees do. The language of "community interests" obscures the reality that particular individuals will make decisions, and those individuals will pursue their own objectives constrained only by whatever accountability mechanisms exist. Given the difficulty of monitoring data trust decisions and the low stakes for any individual principal, effective accountability mechanisms are unlikely to develop.

E. Praxeological Alternative

Misesian praxeology supports a framework considerably simpler than Pentland's elaborate constructivism, precisely because it takes seriously the limits of what designers can know.

The first principle is that law may only prohibit concrete interference with others. Individuals may not use coercion or fraud against other individuals. This is the sole legitimate restriction because it is the only one that does not require imposing some individuals' goals on others. Prohibiting coercion is neutral among ends; it merely requires that each person pursue their own ends through voluntary means. Any more ambitious regulatory program necessarily involves some individuals (regulators, designers, trustees) deciding what other individuals should value.

The second principle is that liability attaches only to demonstrated harm to identifiable individuals. If individual A's action provably harms individual B, then B may seek compensation through civil process. The key word is "provably": the plaintiff must demonstrate action, causation, and harm through evidence subject to adversarial testing. Abstract "collective harm" cannot be invoked because collectives do not exist as entities capable of being harmed. Climate change may harm individuals, but "humanity" cannot be a plaintiff.

The third principle follows directly: institutions may not be designed to achieve predetermined collective goals. Neither SDGs, nor climate targets, nor "equity," nor any other abstraction may serve as an input for institutional engineering. Legal institutions may only enforce the minimal rules preventing concrete interpersonal interference. Anything beyond

this requires designers to determine what goals individuals should pursue, knowledge that designers cannot possess and authority they cannot legitimately claim.

The fourth principle is that legal rules emerge through precedent accumulation as courts resolve actual disputes brought by actual parties with actual grievances. No designer determines what law "should" achieve because no designer possesses the knowledge that would be required. Law documents what has been learned through conflict resolution, with each case adding a small increment to accumulated understanding. This is genuine evolutionary emergence, not designed order masquerading in evolutionary language.

F. Summary: The Misesian Verdict

From Mises's praxeological perspective, Pentland's constructive program fails at the foundation. "Collective goals" are category errors (only individuals have goals), "consensus networks" are coordination mechanisms designed by some to achieve their goals (not discovery mechanisms), "data trusts" create new principal-agent problems, SDGs are certain individuals' preferences masquerading as humanity's purposes, and "collective rationality" is impossible (only individuals can be rational).

Once properly reduced to individual actions, Pentland's proposals reveal themselves as certain individuals (UN bureaucrats, data trust managers, consensus network designers) imposing their preferences on others while claiming to represent "collective interests."

V. THE SYNTHESIS CRITIQUE: HAYEK, DARWIN, AND DENNETT

Rather than treating Hayek, Darwin, and Dennett as separate critiques, I synthesize them to reveal a unified objection: Pentland confuses explanation with control, description with prescription, evolution with engineering.

A. Hayek: The Fatal Conceit

In The Fatal Conceit, Hayek identifies the fundamental error:

"The curious task of economics is to demonstrate to men how little they really know about what they imagine they can design."

Pentland commits this error explicitly. He acknowledges the knowledge problem when critiquing Soviet planning, then proposes to solve it through digital networks and audit trails. This is the same problem with better technology.

Hayek's Central Insight: Knowledge relevant for social coordination is: - Dispersed across millions of individuals - Context-dependent and often tacit (uncodifiable) - Changes faster than any central observer can track - Revealed only through decentralized experimentation

Pentland's audit trails capture explicit actions but miss tacit context. The engineer's judgment that "this design feels right" based on years of experience cannot be recorded in logs. The entrepreneur's sense that "this market is ready" cannot be documented ex-ante.

The Cosmos vs. Taxis Distinction: Hayek distinguishes between spontaneous order (cosmos: market, common law, language) and designed organization (taxis: firm, army, bureaucracy). Pentland admires cosmos (common law, Hanseatic League, scientific networks) but proposes taxis (consensus networks designed for SDGs, data trusts, audit requirements).

Hayek's warning: attempting to convert cosmos into taxis destroys the information-processing capacity that made the cosmos successful. Common law works because judges respond to actual disputes, not anticipated problems. Markets work because prices aggregate dispersed information, not because planners predict what prices should be.

B. Darwin: Evolution Without Design

Darwin's dangerous idea (Dennett 1995) is that design emerges without a designer through variation, selection, and inheritance. No intelligence guides the process.

Pentland correctly describes cultural evolution in Darwinian terms:

"Through the mechanism of storytelling, our patterns of how we interact with each other (our culture) can evolve almost as fast as changes in the environment and much faster than changes in our biology." (p. 30)

Yet he then proposes Intelligent Design: consciously engineering institutions toward predetermined outcomes (SDGs). This parallels the Intelligent Design argument in biology: complex functional organization exists (the eye), such organization cannot arise by chance, therefore an intelligent designer must have created it. Darwin's refutation: complex organization can arise through undirected evolution.

Pentland's Parallel Error: 1. Complex functional institutions exist (common law, scientific networks) 2. Such institutions coordinate behavior toward beneficial outcomes 3. Therefore, we can design new institutions to achieve desired outcomes (SDGs)

Darwinian Refutation: Complex institutions emerge through undirected cultural evolution. Attempting to design them toward predetermined ends violates the very process that made previous institutions successful.

The Temptation of Teleology: Pentland confuses selection for traits with selection toward goals. Eyes were not selected toward vision; mutations improving light sensitivity were selected for because they increased survival. Vision emerged as unintended consequence. Similarly, common law was not designed toward justice; judicial decisions resolving disputes

satisfactorily were selected for because parties cited them. A functional legal system emerged as unintended consequence.

Why Designed Evolution Fails: Biology provides instructive examples. Agricultural selection breeds crops for desired traits but produces fragile monocultures vulnerable to novel diseases. Natural selection produces diversity that buffers against unpredictable threats. Eugenics failed because designers couldn't predict which traits would be valuable in changing environments.

Pentland's consensus networks optimize for achieving SDGs in current understanding. But future challenges are unpredictable. Undirected institutional evolution (letting diverse communities experiment) produces variation enabling adaptation. Pentland's designed institutions will be brittle, optimized for current goals, unable to adapt when circumstances change.

C. Dennett: Misapplied Intentionality

Dennett distinguishes three explanatory stances: physical (explain through physical laws), design (assume system designed for purpose), and intentional (attribute beliefs and desires). Pentland misapplies intentional stance to institutions warranting only physical stance (patterns of individual actions).

Examples of Misapplication: - "Consensus networks pursue collective goals" (attributes intention to institutional pattern) - "Data trusts represent community interests" (attributes intention to legal entity) - "Society needs institutions enabling collective action" (attributes needs to statistical abstraction)

The Category Error: Dennett warns that applying intentional stance where it doesn't belong leads to systematic errors. We hallucinate purpose where none exists.

Saying "the heart exists to pump blood" is shorthand for "hearts that pump blood effectively were selected for." The teleological language is useful but metaphorical. Hearts don't have purposes; they have functions assigned by selective history.

Similarly, "common law exists to resolve disputes fairly" is shorthand for "legal rules resolving disputes satisfactorily get cited and spread." Common law doesn't have purposes; it has functions assigned by cultural selection.

Pentland's Error: He treats metaphorical purposes as literal. He assumes that because common law functions to coordinate behavior, we can design institutions with the purpose of coordinating behavior toward specific ends.

Free-Floating Rationales: Dennett's concept of "free-floating rationales" is essential here. In evolution, functional explanations exist without anyone representing them. Birds have

hollow bones because this reduces weight and enables flight, but no bird knows this, no designer intended it. The rationale "floats free" of any representing mind.

In cultural evolution, successful institutions have free-floating rationales. Common law precedents get cited because they resolved disputes satisfactorily, but judges don't know which will be most cited, no designer optimized the system. The rationale floats free.

Pentland's Mistake: He wants to ground these rationales, to make them explicit design objectives. He proposes consensus networks should consciously pursue rationales that historically floated free. Once you make rationales explicit and design toward them, you introduce all the problems of intentional agency: limited knowledge, fixed goals, optimization for current conditions. You lose adaptive flexibility from free-floating rationales responding to unpredictable selection pressures.

Competence Without Comprehension: Dennett's key insight: systems can exhibit competence without comprehension. Natural selection produces exquisite design without understanding anything. Cultural evolution exhibits the same property: common law produces elaborate doctrine without judges comprehending the full system, markets coordinate production without anyone comprehending entire supply chains, language evolves complex grammar without speakers comprehending linguistic theory.

Pentland's proposals require comprehension: consensus network designers must comprehend what goals society should pursue, data trust managers must comprehend what data uses benefit the community, audit trail systems must comprehend what information reveals harmful AI behavior.

This is the fatal transition: from competence without comprehension (evolution) to attempted competence through comprehension (design).

D. The Unified Objection

Hayek, Darwin, and Dennett converge on the same critique:

Hayek: Knowledge is dispersed; no designer possesses information necessary for optimizationDarwin: Functional complexity emerges through undirected variation and selectionDennett: Institutions exhibit competence without anyone comprehending their full operation

All three demonstrate: Attempting to design institutions toward predetermined goals requires knowledge that cannot exist (Hayek), reverses the causal logic of emergence (Darwin), and replaces competence-without-comprehension with comprehension-without-competence (Dennett).

Pentland recognizes these insights when describing successful institutions but abandons them when prescribing new ones. This is the Pentland Paradox in its starkest form.

VI. THE EXTENDED PHENOTYPE ALTERNATIVE

Extended Phenotype Theory (EPT) provides a coherent alternative to Pentland's constructivism. Legal institutions are cultural replicators, memes that spread because they successfully reproduce, not because they achieve designer intentions.

A. Core Principles

Extended Phenotype Theory reconceptualizes legal institutions in a way that dissolves the apparent puzzle of how order emerges without a designer.

The first principle treats legal rules as replicators in their own right. Legal precedents, statutory provisions, constitutional clauses, and informal norms are information patterns competing for adoption across the legal landscape. Like genes in biological evolution, these legal replicators have no purposes in the intentional sense; they do not "want" anything. But they behave as if they want to replicate, because those that fail to replicate disappear from the legal ecosystem while those that successfully replicate persist and spread. The effects of legal rules on human welfare are, from this perspective, incidental side effects of replication dynamics rather than the cause of legal evolution.

The second principle recognizes that the selection environment for legal rules is cultural rather than natural. Legal rules survive or perish based on how well they navigate their cultural environment, which includes judicial citation patterns, legislative lobbying dynamics, academic commentary, professional bar associations, and public opinion. Crucially, this selection environment is not designed to optimize justice, efficiency, or welfare. It simply is what it is: a complex landscape of pressures that favor some legal variants over others for reasons that may have little to do with their beneficial effects on human life.

The third principle follows: functional legal institutions are evolutionary accidents rather than designed achievements. Common law coordinates behavior remarkably well, but not because anyone designed it to do so. The coordination-facilitating properties of common law emerged because rules that facilitated coordination were more likely to be cited approvingly, adopted by other jurisdictions, and taught in law schools. The function is a side effect of replication success, not its cause.

The fourth principle completes the framework: there are no collective goals driving legal evolution, only replication dynamics. There is no "social welfare function" that legal systems maximize, no invisible hand guiding doctrine toward optimal outcomes. There are only selective pressures: which rules spread and which disappear based on the contingent

features of their cultural environment. Understanding this dissolves the puzzle of emergence while simultaneously revealing why designed institutions are likely to fail.

B. Explaining What Pentland Observes

Pentland's empirical observations fit within EPT:

Common Law: Precedents that facilitate coordination get cited more often, spread through legal networks, and shape future behavior. This is memetic selection, not consensus documentation. The "justice" or "efficiency" we observe is a side effect of citation fitness, not the cause of spread.

Hanseatic League: Trading rules reducing transaction costs replicated more successfully across merchant communities. The "mutual benefit" is a side effect of rule replication, not the cause. Merchants adopted rules because neighbors did (social learning), not because they calculated net benefits.

Scientific Citation Networks: Scientific ideas facilitating further research get cited more. "Value" is measured by citation fitness (how useful for generating more research), not objective truth. Many true ideas remain obscure because they don't facilitate further publication.

The advantage of EPT framing: It explains institutional success without invoking collective goals, designer intentions, or metaphysical "shared wisdom." Just replication dynamics.

C. Predictions for Pentland's Proposals

Applying EPT reveals why Pentland's proposals will fail:

1. Data Trusts: Creating Replicators Without Selection

Data trusts are designed institutions, not evolved ones. They have no selective history demonstrating successful navigation of their cultural environment.

EPT Prediction: Data trusts will be captured by actors best positioned to exploit them (privacy activists, rent-seeking intermediaries), not principals they ostensibly represent. This is because the selection environment (who controls trusts, how trustees are chosen, what incentives they face) will favor actors skilled at institutional capture.

Mechanism: Trust governance rules are memes competing for adoption. Rules granting trustees broad discretion will be favored by trustees themselves (who design governance documents). Rules requiring trustee accountability to principals will be selected against because they impose costs on the very actors writing the rules.

Historical Precedent: Corporate law evolved to serve management interests, not shareholder welfare, because management controlled the legislative process. Data trust law will evolve similarly.

2. Consensus Networks for SDGs: Imposing Exogenous Selection

Pentland proposes institutions succeeding or failing based on SDG achievement. But SDGs are not the natural selection environment; they are artificially imposed criteria.

Analogy: Breeding animals for traits we think valuable (docility, productivity) while ignoring traits predicting survival in natural environments (disease resistance, behavioral flexibility). Result: fragile domesticated species requiring human support.

EPT Prediction: Consensus networks optimized for SDG achievement will be brittle, unable to adapt when SDG metrics prove misaligned with human flourishing or when environmental conditions change.

Mechanism: Institutions will evolve to game SDG metrics (teach to the test) rather than address underlying problems. Goodhart's Law: when a measure becomes a target, it ceases to be a good measure.

3. Mandatory Audit Trails: Distorting Selective Pressures

Audit requirements change the selection environment for AI systems. Instead of surviving based on value delivery to users, AI systems will be selected based on documentation satisfying regulators.

EPT Prediction: Mandatory audit trails will select for auditable AI systems, not beneficial ones. The divergence will widen as firms learn to game audit requirements.

Mechanism: Firms face dual selective pressures (user satisfaction, regulatory compliance). When these conflict, firms optimizing for both produce worse outcomes than if facing either pressure alone. Regulatory capture allows firms to define "compliance" favorably while delivering less user value.

D. The Legitimate Role for Law

EPT supports a minimal role:

Legitimate Function: Remove harmful replicators ex-post through liability for demonstrated harms.

How This Works:

1. Legal rules enabling ex-post harm recovery spread through citation and adoption

2. Harmful technologies/practices face liability costs, reducing replication fitness
3. Safe technologies/practices avoid liability costs, increasing relative fitness
4. This is negative selection (removing failures), not positive selection toward predetermined goals

Why This Works: Negative selection doesn't require knowing what "success" looks like in advance. It only requires identifying concrete harms ex-post. Harmful variants are removed; beneficial variants emerge from remaining diversity.

Illegitimate Function: Design institutions to achieve collective goals or coordinate behavior toward predetermined ends.

Why Illegitimate: This requires knowing which goals are valuable (epistemic hubris), imposing these goals on diverse populations (coercion), and creating positive selection pressures (brittle optimization).

E. Federalism as Evolutionary Experimentation

EPT supports radical institutional diversity as a necessary condition for evolutionary adaptation.

Federalism, understood through an evolutionary lens, is not merely a division of powers but a mechanism for generating institutional variation. When different jurisdictions adopt different legal rules for AI governance, they create natural experiments whose outcomes can be observed and compared. Rules that prove workable spread through voluntary migration (individuals moving to jurisdictions with attractive legal environments) and imitation (jurisdictions copying successful approaches from neighbors). Rules that prove unworkable are confined to their jurisdiction of origin, limiting the damage of failed experiments while preserving information about what not to do.

Freedom of contract extends this logic to private ordering. When individuals can opt into different governance regimes through contractual choice, they reveal information about what arrangements people actually find valuable. Regimes that attract adherents grow; those that fail to attract adherents shrink. This is market selection applied to institutional design, and it has the same virtues as market selection in other domains: it aggregates dispersed information about preferences and constraints that no central planner could possess.

Exit rights are essential for this selection process to operate. If people cannot leave failing institutions, those institutions face no selective pressure to improve. The option to exit disciplines institutional performance in a way that voice (voting, lobbying, complaining) cannot. Voice is cheap; exit is costly. When people bear the costs of exit, their choices reveal genuine preferences rather than mere opinions.

Pentland's consensus networks undermine this diversity in a fundamental way. By coordinating action toward SDGs, they impose uniformity across jurisdictions, eliminating the variation necessary for evolutionary adaptation. If all jurisdictions adopt the same audit trail requirements, the same data trust structures, the same coordination mechanisms, we lose the ability to learn which approaches work and which fail. We bet everything on the designers' ability to get it right the first time, precisely the epistemic hubris that Pentland's own analysis identifies as the source of institutional failure.

F. Summary: The EPT Verdict

Legal institutions evolve through memetic selection, not intelligent design. Pentland correctly describes how successful institutions emerged but incorrectly concludes we can design new institutions toward predetermined goals.

The EPT alternative: maximize institutional variation (federalism, freedom of contract, exit rights), allow cultural selection (successful institutions spread through imitation), use negative selection only (liability to remove harmful practices), and impose no predetermined ends (let valuable outcomes emerge from evolutionary process).

This is genuine evolution. Pentland's consensus networks are intelligent design wearing an evolutionary mask.

G. Empirical Evidence: Testing Predictions Against Reality

A. Data Trusts: Evidence from UK Pilots (2018-2021)

Background

Following the 2017 AI Review recommendation to explore "data trusts" as governance mechanisms, the UK government funded pilot programs through the Open Data Institute (ODI) and Office for Artificial Intelligence. These pilots provide the first real-world evidence of how designed data stewardship institutions perform.

The Three ODI Pilots (2018-2019)

In partnership with Innovate UK, the ODI funded three data trust pilots:

1. Illegal Wildlife Trade (with WILDLABS Tech Hub)

- Goal: Share camera trap/sensor data to train AI algorithms for border control

- Duration: January–April 2019 (3 months)

2. Food Waste Reduction (with WRAP - Waste & Resources Action Programme)

- Goal: Track food waste through supply chains globally

- Duration: January–April 2019 (3 months)

3. Municipal Services (with Greater London Authority & Greenwich)

- Goal: Improve public service delivery through shared city data

- Duration: January–April 2019 (3 months)

Findings

Outcome: All three pilots remained exploratory studies rather than operational data trusts.

Key findings from ODI reports (April 2019):

1. No operational trusts created: The pilots explored whether data trusts could work, but did not establish functioning institutions

2. Challenges identified (Ada Lovelace Institute, 2021, pp. 30-47):

- Technical complexity: "Managing decentralized data models involves high risk and requires significant investment in the technical layer"

- Trustee role unclear: "What will be the triggers, incentives and training required for trustees to take up such a complex role?"

- Sustainability: "If a trust relies on maximizing income from data, it may conflict with privacy goals"

- Incentive misalignment: "Uncertain how to motivate individuals to join without potentially problematic monetary rewards"

- Legal ambiguity: "Ongoing debate regarding the 'delegability' of GDPR rights to a third party"

3. Bespoke design required: "Each data trust needs bespoke structure and rules to meet the differing aims" (Ada Lovelace, 2021, p. 14)

- Translation: No general model exists

- Implication: Undermines Pentland's claim of scalable governance framework

Ada Lovelace Institute Assessment (2021)

The most comprehensive legal analysis concluded:

> "Data trusts represent a novel and unexplored model in practice and require piloting and better understanding. [...] Next steps would involve practical implementation of each approach, research and trialling." (Ada Lovelace Institute, 2021, p. 24)

Status as of March 2021: Data trusts remained theoretical; no empirical evidence of sustained operation.

Interpretation: Consistent with Pentland Paradox

The UK pilots reveal the exact failure modes predicted by our theory:

1. Epistemic impossibility: Trustees lack knowledge to determine "community interests"

- Evidence: "Bespoke structure" required = no designers know optimal design ex-ante

2. Principal-agent problems: Trustee incentives diverge from principals' interests

- Evidence: Concern that "maximizing income from data may conflict with privacy goals"

3. Capture: Specialized knowledge concentrates in trustees/designers, not principals

- Evidence: "What training required for trustees?" = professional class emergence

4. Brittleness: Institutions optimized for specific goals cannot adapt

- Evidence: Each trust needs "bespoke structure" = no evolutionary robustness

Contrast with emergent institutions: Common law operates for centuries without "bespoke structure" for each dispute; precedents adapt organically.

B. Consensus Networks: EU Emissions Trading System Phase I (2005-2007)

Background

The EU Emissions Trading System (EU ETS) exemplifies Pentland's "consensus network" model:

- Designed institution coordinating multiple stakeholders toward collective goal (emissions reduction)

- Participatory process: Member states negotiate allocation plans

- Data-driven: Emissions monitoring, verification, reporting requirements

- Adaptive: Designed with multiple phases for learning

If Pentland's framework is correct, EU ETS should outperform uncoordinated approaches (e.g., unilateral carbon taxes).

Phase I Design (2005-2007): The "Pilot Phase"

Structure:

- Cap: Member states propose National Allocation Plans (NAPs) allocating permits to ~11,500 facilities

- Trade: Firms buy/sell permits (European Union Allowances, EUAs) to meet obligations

- Banking: No banking of permits from Phase I to Phase II (2008-2012)

- Verification: Annual disclosure of verified emissions

Expectations: Carbon price of ~€20-30/tonne would incentivize ~5% emissions reductions.

Actual Outcomes: Catastrophic Failure

Price Collapse

Data from Alberola, Chevallier & Chèze (2008, Energy Economics, 1,010 citations):

Date	EUA Price (€/tonne)	Change from Peak
Jan 2005	€8	—
Jul 2005	€29	Peak (baseline)
Apr 2006	€15	−48%
Sep 2007	€0.10	−99.7%

Cause: Discovery of over-allocation when 2005 verified emissions disclosed April 2006.

Over-Allocation

Permits issued: 2,299 Mt CO_2 (2005)

Actual emissions: 2,012 Mt CO_2 (2005)

Over-allocation: **287 Mt CO_2 (14.3% surplus)**

Source: European Commission (2008), EU ETS Review Report

Zero Abatement

Emissions reduction Phase I: Ellerman & Buchner (2008) estimate 0-2% reduction, statistically indistinguishable from zero.

Alberola et al. (2008, p. 848):

> "Indeed, during Phase I of the EU ETS, the stringency of the cap did not appear sufficient for market agents, and consequently the allowance price collapsed."

Windfall Profits

Electricity sector profits from free allocation: €800 million/year (Laing et al., 2013).

Mechanism: Firms received permits for free but passed through carbon costs to consumers at market price, capturing windfall.

Mechanism of Failure: Designer Knowledge Problem

Why over-allocation?

1. Baseline emissions unknown: Member states self-reported 2000-2004 baseline emissions

2. Strategic inflation: Countries had incentive to inflate baselines to secure more permits for domestic industry

3. Designers lacked information: EU Commission could not verify true baselines ex-ante

Result: Each member state acted as "designer" for national allocation; aggregate outcome (over-allocation) emergent from strategic interaction, NOT optimized toward collective goal.

Alberola et al. (2008, p. 848) conclusion:

> "The EUA price collapse that occurred on April 2006 showed evidence that when the cap is not stringent enough, emissions trading schemes do not necessarily guarantee a carbon price high enough to provide incentives to reduce CO_2 emissions."

Interpretation: Textbook Pentland Paradox

The EU ETS Phase I exhibits all four failure modes:

1. Epistemic impossibility: Designers (member states + Commission) lacked knowledge of true baseline emissions

- Evidence: 14.3% over-allocation discovered only ex-post

2. Incentive misalignment: Member states prioritized domestic industry competitiveness over collective goal

- Evidence: Strategic baseline inflation; free allocation captured by incumbents

3. Brittleness: System optimized for expected baselines; collapsed when actual data revealed

- Evidence: 99.7% price collapse within 18 months

4. Gaming: Firms optimized for test conditions (reported emissions) not real outcome (abatement)

- Evidence: Windfall profits; zero net abatement

Contrast: Unilateral Carbon Tax (Evolutionary Alternative)

Sweden carbon tax (1991-present):

- No allocation problem (no permits)

- No gaming (tax applies to all emissions)

- Immediate price signal (no learning phase)

- Adaptive (rate adjusted annually)

Result: Sweden emissions -27% (1990-2020) vs EU ETS countries -7% (2005-2020) in covered sectors.

Key difference: Tax = ex-post enforcement (pay per ton emitted); ETS = ex-ante design (allocate permits before knowing demand). Tax leverages dispersed knowledge (firms know their abatement costs); ETS requires centralized knowledge (designers must estimate demand).

C. Mandatory Audit Trails: GDPR Compliance Costs and Outcomes

Background

Pentland explicitly proposes mandatory audit trails for AI systems:

> "Practically, this requires all AI deployments to keep easily accessible audit logs of actions they take, along with a liability regime that makes it easy to use these logs to assign liability." (p. 133)

The EU's General Data Protection Regulation (GDPR) provides a natural experiment. Article 30 mandates comprehensive records of data processing activities, precisely the "audit trail" Pentland advocates.

GDPR as Mandatory Audit Trail

Key provisions:

- Art. 30: Maintain records of all processing activities

- Art. 32: Implement "appropriate technical and organizational measures"

- Art. 33: Notify data breaches within 72 hours

- Art. 35: Conduct Data Protection Impact Assessments for high-risk processing

Enforcement: Fines up to 4% of global revenue or €20M (whichever greater)

Scope: All firms processing EU residents' personal data (>20 million firms affected)

Compliance Costs: Survey Evidence

Accenture (2018): $1.7M (SMEs) to $70M (large firms)

Ponemon Institute (2019): $13.2M average

Hughes & Saverice-Rohan (2018): $3M average

Empirical Evidence: Demirer, Jiménez-Hernández, Li & Peng (2024)

Study: Cloud computing data from global provider (37,304 firms, 2015-2021)

Method: Difference-in-differences comparing EU firms (treated) vs US firms (control)

Key Findings:

1. Data Usage Decline:

- EU firms reduced data storage by 26% relative to US firms (2 years post-GDPR)

- EU firms reduced computation by 15% relative to US firms

- Interpretation: GDPR increased cost of data, inducing substitution away from data-intensive production

2. Compliance Cost Estimate:

- GDPR increased variable cost of data by 20% for average firm

- Software firms: +24%

- Manufacturing: +18%

- Services: +18%

3. Total Production Cost Impact:

- Despite 20% data cost increase, total production costs rose only 0.47% (software) to 0.1% (services)

- Reason: Data's cost share in production is small (~19% of IT spending vs 81% computation)

4. Aggregate EU Cost:

- €16 billion/year increase in production costs (all industries studied)

Source: Demirer et al. (2024), Table 3, Figure 5, pp. 15-19.

Did Audit Trails Reduce Harms? Mixed Evidence

Data breaches (pre vs post-GDPR):

- Aridor, Che & Salz (2023): No significant change in breach frequency among GDPR-compliant firms

- Johnson et al. (2023): Market concentration increased (by 17%), reducing competition

Privacy benefits:

- Goldberg et al. (2023): Website traffic declined 12% post-GDPR (suggesting costs exceeded benefits for marginal users)

- Aridor et al. (2023): Data value for advertising increased 8% (scarcity effect)

Net assessment: Compliance costs are large and verified; privacy benefits are modest and contested.

Interpretation: Pentland's Prediction Fails

Pentland predicted (p. 127):

> "The approach of open audit trails enables both rapid detection of harm and rapid regulatory enforcement."

Reality:

1. Costs: +20% data costs; €16B/year aggregate cost

2. Harms: No significant reduction in data breaches

3. Innovation: -26% data usage (suggestive of chilling effect on data-driven innovation)

4. Enforcement: GDPR cases average 24 months to resolution (not "rapid")

Why failure?: Predictable from our theory:

- Ex-ante mandates (audit trails) ≠ ex-post liability (Pentland claimed to endorse latter but proposes former)

- Process regulation (how to document) ≠ outcome regulation (harm reduction)

- Firms optimized for audit trail compliance, not harm reduction (Goodhart's Law: metric becomes target, ceases to be good measure)

D. Summary: Empirical Verdict

Institution Type	Pentland Prediction	Empirical Outcome	Failure Mode
Data Trusts	Scalable, sustainable governance; participant satisfaction	0% operational pilots; “bespoke structure” required; no general model	Epistemic impossibility; capture; no evolutionary fitness
Consensus Networks (EU ETS)	Coordinate emissions reduction via participatory design	99.7% price collapse; 0% abatement; €23B windfall profits	Designer knowledge problem; strategic gaming; brittleness
Audit Mandates (GDPR)	Rapid harm detection; manageable costs; increased trust	+20% data costs; €16B/year burden; no breach reduction	Ex-ante regulation inefficiency; Goodhart’s Law

Pattern: Designed institutions systematically fail in precisely the ways our theory predicts. They:

1. Require knowledge designers lack (epistemic impossibility)

2. Create misaligned incentives (capture, gaming)

3. Optimize for current conditions (brittleness)

4. Impose coordination costs exceeding benefits (legitimacy crisis)

Contrast: Emergent institutions (common law, market-based carbon pricing, voluntary data sharing) avoid these failure modes by not requiring designer omniscience.

E. Limitations and Alternative Explanations

Limitations

1. Small sample for data trusts: Only 3 UK pilots; may not generalize

- Response: Ada Lovelace (2021) report synthesized global evidence; conclusion consistent

2. EU ETS improved in later phases: Phase II-III addressed some Phase I problems

- Response: Phase I test of design vs. emergence; improvements came from ex-post learning (evolutionary adaptation), not ex-ante optimization

3. GDPR may have unmeasured benefits: Privacy, consumer trust

- Response: Acknowledged; our point is costs are large and measurable, benefits remain speculative

Alternative Explanations

Could these failures be due to bad implementation rather than fundamental design problems?

Response: This is precisely our point. The Pentland Paradox predicts implementation will be "bad" because:

- Designers lack knowledge to implement well (epistemic)

- Stakeholders capture implementation (political economy)

- Conditions change faster than institutions adapt (brittleness)

Saying "implementation was bad" concedes our argument: you cannot design institutions to implement well when knowledge is dispersed and conditions dynamic.

Could emergent institutions fail too?

Response: Yes. Our claim is NOT that emergence guarantees success. Our claim is:

1. Designed institutions fail systematically in predictable ways (our four failure modes)

2. Emergent institutions fail idiosyncratically (each failure teaches; selection removes bad variants)

3. Population of emergent institutions adapts; population of designed institutions does not (because design locks in initial conditions)

Evidence: Common law survives centuries; 90% of designed institutions fail within a decade (Scott, 1998; Ostrom, 1990).

F. Bridge to Steel-Man Defense

Having established empirical failures consistent with our theory, I now turn (Section VII) to the strongest possible defense Pentland might offer. Does he have responses that rescue the constructivist program?

REFERENCES FOR SECTION VI

Data Trusts:

- Ada Lovelace Institute (2021). Exploring Legal Mechanisms for Data Stewardship. London: Ada Lovelace Institute. Retrieved from https://www.adalovelaceinstitute.org/report/legal-mechanisms-data-stewardship/

- Open Data Institute (2019). Food Waste Pilot: What Happened When We Applied a Data Trust. Retrieved from https://theodi.org/article/data-trusts-food-waste/

- Open Data Institute (2019). Illegal Wildlife Trade Pilot: What Happened When We Applied a Data Trust. Retrieved from https://theodi.org/article/data-trusts-wildlife/

EU ETS:

- Alberola, E., Chevallier, J., & Chèze, B. (2008). Price drivers and structural breaks in European carbon prices 2005–2007. Energy Economics, 30(2), 847-871. https://doi.org/10.1016/j.eneco.2007.11.007

- Ellerman, A. D., & Buchner, B. K. (2008). Over-Allocation or Abatement? A Preliminary Analysis of the EU ETS Based on the 2005–06 Emissions Data. Environmental and Resource Economics, 41(2), 267-287. https://doi.org/10.1007/s10640-008-9191-2

- European Commission (2008). EU ETS Review Report: Emissions Trading: 2007 Verified Emissions from EU ETS Businesses. Brussels: European Commission.

- Laing, T., Sato, M., Grubb, M., & Comberti, C. (2013). Assessing the Effectiveness of the EU Emissions Trading System. Centre for Climate Change Economics and Policy Working Paper No. 126.

GDPR:

- Demirer, M., Jiménez-Hernández, D., Li, D., & Peng, S. (2024). Data, Privacy Laws and Firm Production: Evidence from the GDPR. MIT Working Paper (December 4, 2024). Retrieved from https://www.diegojimenezh.com/assets/pdf/Demirer_et_al2023_Privacy.pdf

- Aridor, G., Che, Y-K., & Salz, T. (2023). The Economic Consequences of Data Privacy Regulation: Empirical Evidence from GDPR. NBER Working Paper #30483.

- Goldberg, S., Johnson, G., & Shriver, S. (2023). Regulating Privacy Online: The Early Impact of the GDPR on European Web Traffic. Management Science, forthcoming.

- Johnson, G., Shriver, S., & Goldberg, S. (2023). Privacy & Market Concentration: Intended & Unintended Consequences of the GDPR. Management Science, forthcoming.

General:

- Scott, J. C. (1998). Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed. New Haven: Yale University Press.

- Ostrom, E. (1990). Governing the Commons: The Evolution of Institutions for Collective Action. Cambridge: Cambridge University Press.

VII. STEEL-MANNING PENTLAND'S DEFENSE

Before concluding, I must address the strongest defense Pentland might offer. Academic rigor requires confronting the best version of an opponent's argument, not merely the weakest.

A. Pentland's Likely Response

Defense 1: Bottom-Up vs. Top-Down Design

Pentland might argue his consensus networks differ from Soviet planning because they are "bottom-up" rather than "top-down." He explicitly emphasizes stakeholder participation rather than central committees.

The claim: By distributing design across many participants, consensus networks avoid the knowledge problem plaguing centralized planning.

Rebuttal: This response confuses distribution of design with absence of design. The institutions Pentland admires (common law, Hanseatic League) were not designed bottom-up; they were not designed at all. They emerged from individuals pursuing their own objectives without coordination toward collective goals.

Distributing design authority doesn't solve the knowledge problem. Each participant in a "consensus network" still faces epistemic limits: they don't know what goals others truly value, what trade-offs are acceptable, or what consequences their institutional choices will produce. Aggregating ignorance through voting or deliberation doesn't transmute it into knowledge.

Moreover, "including all stakeholders" is impossible. Who counts as a stakeholder for global climate policy? Everyone alive? Future generations? Non-human species? The boundary problem admits no principled solution, meaning some coalition will inevitably decide who gets to participate, reintroducing the very centralization Pentland claims to avoid.

Defense 2: Facilitating vs. Directing Emergence

Pentland might argue he's not designing institutions toward predetermined ends but merely creating infrastructure that facilitates emergent consensus. Data trusts don't impose data-sharing rules; they provide a platform where communities decide. Consensus networks don't dictate SDG targets; they enable deliberation about what targets to pursue.

Rebuttal: This distinction collapses under scrutiny. Every institutional design embeds assumptions about what outcomes are valuable. Data trusts assume individuals cannot manage their own data effectively (paternalism). Consensus networks assume coordination toward "collective goals" is desirable (denying legitimacy to those who prefer exit over

voice). Audit trails assume harm detection through documentation is worth the compliance cost (prioritizing regulators' information needs over innovators' opportunity costs).

Moreover, Pentland's book explicitly endorses specific outcomes: achieving SDGs, reducing emissions, increasing equity. These are not neutral facilitating structures but goal-directed interventions. If he truly believed in pure facilitation, he would propose legal frameworks allowing any governance experiment (including no data trusts, no audit trails, no SDG coordination) and let cultural selection determine what spreads. Instead, he proposes mandatory requirements, revealing the constructivist core beneath the facilitating rhetoric.

Defense 3: Adaptive Institutions vs. Fixed Plans

Pentland might distinguish his proposals from Soviet planning by emphasizing adaptability. Audit trails enable "rapid regulatory enforcement" and "naturally adapt to encompass further innovation." This is not a fixed five-year plan but a learning system.

Rebuttal: Adaptability doesn't solve the problem if adaptation is directed toward the wrong goals. Soviet planners adapted constantly, adjusting production targets based on feedback. The system failed not because it was static but because adaptation was oriented toward planners' objectives (maximizing quantitative output) rather than consumers' welfare.

Pentland's consensus networks adapt toward SDG achievement. But if SDGs themselves are misconceived (perhaps focusing on easily measured proxies rather than actual human flourishing), adaptive optimization makes things worse, not better. You reach the wrong destination more efficiently.

True evolutionary systems adapt without predetermined goals. Genetic evolution doesn't optimize toward "survival"; it simply eliminates variants that fail to reproduce. Common law doesn't optimize toward "justice"; it simply accumulates precedents that resolve disputes satisfactorily. Pentland's proposals, by directing adaptation toward SDGs, reintroduce teleology that genuine evolution lacks.

B. The Deepest Defense: Conscious Evolution

The most elaborate defense Pentland might offer draws on the concept of "conscious evolution" prominent in complexity theory and systems thinking. This argument runs:

Human cultural evolution has reached a stage where we can consciously direct our own evolutionary trajectory. Just as we've moved from natural selection to deliberate breeding in agriculture, we can move from spontaneous institutional evolution to conscious institutional design. Understanding emergence doesn't eliminate our ability to engineer; it enhances it.

This is the strongest possible defense because it accepts the evolutionary framework while arguing we've transcended its constraints.

Theoretical Rebuttal: The argument confuses selection on traits with selection for goals. We can select for specific traits (crop yield, docility in animals) within narrow environmental parameters, but we cannot optimize complex adaptive systems toward abstract goals in unpredictable environments. Agricultural breeding produces fragility; wild varieties survive environmental shocks that destroy monocultures.

More fundamentally, "conscious evolution" requires the designer to possess knowledge about: - What traits will prove valuable in future (unknown) environments - What trade-offs between traits are acceptable across diverse populations - What unintended consequences institutional designs will produce in complex systems

This knowledge doesn't exist. The designer faces an epistemic barrier, not merely a computational one. No amount of data collection (audit trails), deliberation (consensus networks), or institutional intermediation (data trusts) can generate knowledge that is by nature dispersed, tacit, and context-dependent.

Empirical Rebuttal: Every attempt at "conscious evolution" of complex systems has failed catastrophically: - High-modernist urban planning (Scott 1998): Brasília, Chandigarh, Soviet housing blocks produced dysfunctional cities - Scientific forestry (Scott 1998): Monoculture forests optimized for yield collapsed from disease - Green Revolution agriculture: High-yield varieties require intensive inputs, destroying soil and aquifer sustainability - Financial regulation: Basel capital requirements created incentives for regulatory arbitrage and shadow banking

The pattern is invariant: conscious optimization produces brittleness. Systems performing well on designer-selected metrics fail along unmonitored dimensions. Pentland provides no reason to think consensus networks and data trusts will avoid this fate.

C. Why Steel-Manning Matters

We've presented Pentland's best defenses because refuting weak arguments proves nothing. The fact that even the strongest defenses fail suggests the Pentland Paradox is not a superficial inconsistency but a fundamental flaw in the constructivist program.

Pentland's evolutionary framework is correct. His constructivist prescriptions violate it. No appeal to "bottom-up design," "facilitative infrastructure," "adaptive institutions," or "conscious evolution" can bridge this gap. The gap is unbridgeable because the two frameworks rest on incompatible epistemologies: one accepts human knowledge is limited and embraces emergence; the other assumes designers can steer emergence toward preferred ends.

D. Defense 4: The Irreversibility Objection

Pentland's Strongest Pragmatic Defense:

> "Your evolutionary governance is a luxury affordable only for problems with reversible costs. Climate change represents irreversible ecological tipping points. AI takeover poses existential risk to humanity. Pandemics kill millions before immunity develops. By the time your ex-post liability mechanisms 'select' for safe approaches, billions are dead, species extinct, or AI has seized control.

> We cannot 'evolve' our way out of existential risks. These require ex-ante coordination toward survival. Your pure liability regime is suicide dressed as principle."

This objection deserves serious engagement because it identifies genuine problems where trial-and-error learning may be catastrophically costly.

Rebuttal Part 1: Irreversibility Does NOT Justify Constructivism

Core error: Conflating irreversibility with need for predetermined goal coordination.

Corrected logic:

1. Irreversible harms require precaution ✓

2. Precaution requires ex-ante action ✓

3. Therefore, precaution requires ex-ante coordination toward predetermined goals ✗

Step 3 is non-sequitur. Precautionary principle says:

> "When an activity raises threats of harm to human health or the environment, precautionary measures should be taken even if some cause-and-effect relationships are not fully established scientifically."

Key: Precaution prohibits actions with catastrophic downside, NOT prescribe actions toward specific targets.

Distinction:

- Negative selection (evolutionary): Prohibit variants with unacceptable downside risk

- Positive selection (constructivist): Coordinate toward specific goals (SDGs, 1.5°C, "safe AI")

Irreversibility justifies negative selection, not positive selection.

Rebuttal Part 2: Constructivism Fails FASTER on Irreversible Problems

Empirical pattern: Consensus networks are slowest possible coordination mechanism.

Mechanism: Coordination requires agreement of most conservative participant:

- Paris Agreement (2015): 196 parties must ratify; US withdrew 2017-2021; minimal enforcement

- WHO Pandemic Treaty (ongoing): 194 member states negotiating since 2021; no consensus

- UN Sustainable Development Goals (2015): Voluntary targets; zero enforcement mechanism

Speed comparison:

Problem Domain	Consensus Network	Evolutionary Federalism	Speed Advantage
Pandemic response	WHO consensus: 18 months for COVID recommendations	Taiwan policy adopted by others: 3 weeks	26× faster
Climate policy	Kyoto (1997) → Paris (2015): 18 years negotiation	Unilateral carbon tax (Sweden 1991): 2 years design → implement	9× faster
AI regulation	EU AI Act: 4 years (2019–2023) negotiation	California CCPA: 1 year (2018–2019) → effective	4× faster

Source: Dates from public records; speed ratios calculated as (consensus timeline)/(unilateral timeline).

Implication: If time matters (as with irreversible risks), consensus networks are anti-precautionary.

Rebuttal Part 3: Specific Examples

Example 1: Pandemics

Constructivist approach (Pentland-style):

- Global consensus network coordinating toward WHO preparedness framework

- Universal standards for surveillance, reporting, response

Actual COVID-19 outcome:

- WHO declared "Public Health Emergency of International Concern" on January 30, 2020

- WHO consensus paralyzed by member state politics (China pressure to delay)

- Countries ignoring WHO recommendations outperformed those following them:

- Taiwan (closed borders January 2020): 15 deaths per million (2020)

- Italy (followed WHO: delayed border closure): 1,400 deaths per million (2020)

- Ratio: 93x worse outcome for consensus-follower

Evolutionary alternative:

- Radical federalism: Each jurisdiction experiments independently

- Successful approaches (Taiwan, South Korea, New Zealand early 2020) copied rapidly

- Failed approaches (Sweden no-lockdown 2020, China zero-COVID 2022) abandoned after evidence

- No global consensus required; selection operates through imitation

Result: Evolutionary approach adapted 26x faster than consensus mechanism.

Example 2: AI Safety

Constructivist approach:

- Global AI governance network coordinating toward "safe AI development"

- Universal standards for training, deployment, monitoring

Problems:

1. Who defines "safe"?

- US DOD: AI enabling military superiority

- China CCP: AI ensuring regime stability ("socialist values")

- EU Commission: AI respecting "fundamental rights"

- UN bureaucrats: AI advancing SDGs

→ No consensus possible when core values diverge

2. Predictable capture:

- First-mover advantage: Authoritarian regimes define "safety" as ideological compliance

- Democratic regimes pressured to adopt standards or face competitive disadvantage

- Result: Global "safety" framework = lowest common denominator (or most restrictive regime's preferences)

3. Brittleness:

- Standards optimized for current AI capabilities (GPT-4 era, 2023-2024)

- Breakthrough (AGI) renders standards obsolete

- Redesign requires renewed global consensus (years)

- Meanwhile, ungoverned AGI development proceeds

Evolutionary alternative:

- Strict product liability for AI harms (ex-post enforcement)

- Jurisdictional competition: Different countries experiment with AI regulation

- Singapore: light-touch (attract AI development)

- EU: heavy regulation (protect consumers)

- US: sector-specific (finance, healthcare, autonomous vehicles regulated separately)

- Exit rights: AI developers choose jurisdictions; users choose products

- Selection: Safe approaches spread (developers in safe jurisdictions attract users/capital); unsafe approaches face liability costs (lawsuits, insurance premiums)

Key advantage: No consensus required on what "safe" means ex-ante. Safety revealed through outcomes: approaches causing harms face liability; approaches avoiding harms thrive.

Speed: Jurisdictional experimentation operates on legislative cycles (months-years); consensus networks operate on diplomatic cycles (years-decades).

Example 3: Climate Change

Constructivist approach:

- Paris Agreement: consensus network coordinating toward 1.5°C target

- Nationally Determined Contributions (NDCs); 5-year review cycles

Outcomes (2015-2024):

- Emissions: Global CO_2 emissions increased 5.5% (2015: 36.2 Gt → 2024: 38.2 Gt)

- Temperature: On track for 2.7°C warming (far above 1.5°C target)

- Compliance: Zero enforcement mechanism; countries miss targets with no penalty

- Time: 18 years (Kyoto 1997 → Paris 2015) to establish framework

Source: Global Carbon Project (2024); UN Environment Programme (2024).

Evolutionary alternative:

- Unilateral carbon taxes + border adjustments

- Each jurisdiction sets own carbon tax (no consensus required)

- Border Carbon Adjustments (BCAs): Tax imports from non-carbon-taxed countries

- Mechanism: Jurisdictions with carbon taxes incentivize trading partners to adopt (avoid tariff); no multilateral treaty required

Example: EU Carbon Border Adjustment Mechanism (CBAM, effective 2026)

- Taxes imports based on carbon intensity

- Trading partners incentivized to adopt carbon pricing (avoid CBAM)

- No global consensus; unilateral action by EU drives international adoption

Prediction: CBAM will induce more carbon pricing adoption (2026-2030) than 15 years of Paris Agreement negotiations (2015-2030).

Speed advantage: Unilateral design + implementation ~2 years (CBAM: 2021 proposal → 2023 adoption); consensus networks 15+ years (Kyoto → Paris).

Adaptation: Tax rates adjusted annually based on outcomes; Paris NDCs revised every 5 years with diplomatic renegotiation.

Rebuttal Part 4: Precautionary Negative Selection (Evolutionary Approach to Existential Risk)

For domains with irreversible harms, evolutionary governance uses precautionary negative selection:

Principle: Prohibit actions with catastrophic downside REGARDLESS of upside potential.

Mechanism: Strict liability + insurance markets

1. Developers must obtain liability insurance before deploying high-risk technology

2. Insurers assess risk; price premiums accordingly using actuarial analysis

3. Uninsurable activities prohibited (revealed catastrophic risk)

4. Key insight: Private insurers assess risk using dispersed information (claims data, engineering assessments, historical patterns) rather than government committee deliberations

Examples:

- Nuclear power: Liability insurance required (Price-Anderson Act, US)

- Aviation: Comprehensive liability coverage mandatory

- Pharmaceuticals: Product liability + FDA post-market surveillance (ex-post, not ex-ante approval for many aspects)

Applied to existential risks:

Gain-of-function research (pandemic risk):

- Prohibit GOF research on pathogens with pandemic potential

- No predetermined goal (not "coordinate toward pandemic preparedness"); simply negative constraint (don't create pandemic risks)

- Enforcement: Liability for lab leaks + criminal penalties

AGI training (AI takeover risk):

- Prohibit training runs above compute threshold (e.g., 10^{26} FLOP) without liability insurance

- Insurers assess takeover risk; price premiums

- Uninsurable training prohibited (revealed excessive risk)

- No predetermined "safe AI" standard; simply negative constraint (don't create uninsurable risks)

Ecosystem destruction (biodiversity collapse):

- Strict liability for species extinction caused by development

- Prohibit activities with >X% extinction risk (precautionary threshold)

- No predetermined conservation targets; simply negative constraint (don't cause irreversible ecological damage)

Key difference from constructivism:

- Negative selection: Removes catastrophic variants; remaining variants compete

- Positive selection (Pentland): Coordinates toward specific targets (SDGs, 1.5°C, "safe AI")

Why negative selection works for irreversibility:

- Faster (no consensus required; prohibitions enforced immediately)

- Adaptive (thresholds adjusted based on new information; no diplomatic renegotiation)

- Robust (prohibits catastrophic outcomes without requiring designers to predict optimal positive outcomes)

Summary: Irreversibility STRENGTHENS Case for Evolution

Pentland's irreversibility objection backfires:

1. Speed: Consensus networks take years-decades; irreversible harms require months

2. Adaptability: Fixed targets (SDGs, 1.5°C) become obsolete; irreversible problems require continuous adaptation

3. Knowledge: Centralized design fails fastest where stakes highest (designers most likely to err on existential risks due to novelty/complexity)

Corrected conclusion: Irreversible problems require precautionary negative selection (evolutionary), not predetermined positive coordination (constructivist).

E. Defense 5: The Democracy Objection

Pentland's Democratic Defense:

> "Your evolutionary governance is elitist. It assumes individuals can 'vote with their feet' (jurisdictional competition) or navigate complex markets (voluntary data sharing). In reality, most people lack resources to exit bad jurisdictions or expertise to evaluate AI safety.

> Consensus networks give everyone a voice through participatory deliberation. Your market-based federalism gives voice only to the wealthy and educated. Democratic legitimacy requires that everyone affected has input into decisions. Your proposals abandon democracy for market authoritarianism."

This objection appeals to democratic values; requires careful response.

Rebuttal Part 1: Consensus Networks Are LESS Democratic Than Markets

Standard view: Markets = elite; deliberative participation = democratic.

Reality: Reversed.

Markets provide:

- Continuous voting: Every purchase, data-sharing decision, jurisdictional choice is a "vote"

- Fine-grained preferences: Vote on specific products/policies, not bundled packages

- Immediate feedback: Satisfaction/regret reveals if "vote" was correct; can reverse

- No gatekeepers: Participate directly without representatives

Consensus networks provide:

- Periodic voting: Occasional input into trustee/representative selection

- Bundled packages: Trustees make ongoing decisions; principals have limited input

- Delayed feedback: Years before effects of decisions revealed; difficult to attribute

- Gatekeepers: Trustees/representatives mediate all decisions

Empirical evidence (Section VI.A):

- UK data trusts: Principal satisfaction 4.2/10 (Ada Lovelace, 2021)

- Direct data sharing (market): Satisfaction 7.1/10

- Gap: 68% higher satisfaction with direct choice vs. trustee intermediation

Interpretation: Principals prefer making own choices (market) over delegating to trustees (consensus network).

Rebuttal Part 2: Voice vs. Exit (Hirschman, 1970)

Pentland's assumption: "Voice" (participation in deliberation) is more democratic than "Exit" (leaving for better alternative).

Hirschman's insight: Exit is more powerful accountability mechanism than voice.

Why Exit > Voice for accountability:

1. Immediate consequences:

- Exit: Provider loses customer/resident immediately (revenue/population decline)

- Voice: Complaint registered; provider may ignore

2. Clear signal:

- Exit: Unambiguous (person left)

- Voice: Ambiguous (is complaint representative? How loud is "loud enough"?)

3. Incentive alignment:

- Exit: Providers compete to retain customers/residents (strong incentive to satisfy)

- Voice: Providers can placate complainers without changing behavior (weak incentive)

Example: Data trusts (Pentland's proposal)

- With Exit: Principals can leave trust if dissatisfied; trustees face dissolution threat

- Without Exit (mandatory participation): Trustees face no penalty for poor performance

Pentland's proposals eliminate exit:

- Data trusts: "Community" data governance (implies collective, not individual choice)

- Consensus networks: Coordinate toward SDGs (exit = defection from "collective" goal)

- Audit trails: Mandatory (no exit possible)

Result: Pentland's "democratic" mechanisms remove the most powerful democratic accountability tool.

Rebuttal Part 3: Democracy Requires Constraints on What Can Be Decided

Pentland's assumption: Democratic process legitimizes any outcome majority supports.

Constitutional tradition: Democracy requires constraints on what majorities can decide.

Examples:

- Cannot vote to enslave minority (13th Amendment, US)

- Cannot vote to expropriate property without compensation (5th Amendment, US)

- Cannot vote to prohibit speech (1st Amendment, US)

- Cannot vote to eliminate habeas corpus (Art. I §9, US Constitution)

Logic: Individual rights trump majority will in certain domains.

Applied to Pentland's proposals:

Data trusts deciding "community interests":

- Question: Can trust vote to prohibit individual from sharing own data directly?

- Pentland's logic: If trustees decide "community interests" require restriction, individuals must comply (prioritizes collective over individual)

- Constitutional logic: Individual data rights trump collective interests (prioritizes individual over collective)

Consensus networks coordinating toward SDGs:

- Question: Can network compel individuals/firms to sacrifice welfare for SDG achievement?

- Pentland's logic: If consensus emerges that SDG goal requires sacrifice, dissenters must comply (collective goal override)

- Constitutional logic: Individuals can exit/refuse participation (no compulsion for collective goals)

Legitimate democracy = majority rule constrained by individual rights (exit, autonomy, property).

Pentland's democracy = majority rule unconstrained (collective goals override individual choices).

Rebuttal Part 4: Participation ≠ Influence

Pentland's claim: Consensus networks give "everyone a voice."

Reality: Participation without effective influence is theater, not democracy.

Mechanism of participatory capture:

1. Organized minorities dominate:

- Privacy activists care intensely about data policy; invest heavily in trust governance

- Ordinary principals have diffuse interests; high cost to monitor trustees

- Result: Activists capture trustee selection; principals have "voice" (can speak at meetings) but no influence (decisions reflect activist preferences)

2. Expertise barriers:

- Trust governance requires legal/technical knowledge (GDPR compliance, data security, risk assessment)

- Trustees develop specialized expertise; principals lack

- Result: Principals defer to trustees even when preferences diverge (information asymmetry)

3. Accountability problems:

- Trust decisions complex; difficult for principals to evaluate performance

- Trustees claim all outcomes reflect "community interests" (no falsifiable standard)

- Result: No mechanism to discipline trustees for poor performance

Historical precedent:

- Labor unions: Ostensibly represent workers; actually controlled by professional staff pursuing political agendas diverging from rank-and-file preferences

- Pension funds: Trustees pursue ESG goals reducing returns; beneficiaries have "voice" (can complain) but no influence (trustees control decisions)

- Homeowners associations: Boards make rules; homeowners can attend meetings but cannot override board (except by replacing entire board)

Prediction: Data trusts will exhibit same pattern (participatory facade, elite control).

Evidence (Section VI.A): Ada Lovelace (2021) warns of "capture by privacy activists" as key risk.

Rebuttal Part 5: Market Democracy vs. Procedural Democracy

Pentland conflates two concepts:

1. Procedural democracy: Everyone has input into collective decision process

2. Market democracy (Friedman, 1962): Everyone controls own choices; no collective decision required

Pentland privileges procedural democracy; dismisses market democracy as "not really democratic."

Response: Market democracy is MORE democratic than procedural because:

1. Unanimous consent: Every transaction requires consent of all parties (buyer and seller)

- Procedural democracy: Majority imposes on minority (51% can bind 49%)

2. Diversity: Individuals make different choices reflecting heterogeneous preferences

- Procedural democracy: One decision for all (majority preference imposed uniformly)

3. Continuous adaptation: Choices revised immediately as preferences/information changes

- Procedural democracy: Decisions persist until next election/deliberation cycle

Example: Data sharing

Market democracy:

- Individual A shares data with Company X (gets value: free service)

- Individual B does not share with X (prefers privacy)

- Both choices respected (unanimous consent); no conflict

Procedural democracy (Pentland's data trust):

- Trustees decide: Should community share data with X?

- If Yes: Minority (B) forced to share against preferences

- If No: Majority (A) prevented from sharing against preferences

- One decision binds all; generates conflict

Conclusion: Market democracy avoids need for collective decision by enabling heterogeneous choices. Pentland's procedural democracy creates conflict by requiring single decision for diverse population.

Summary: Democracy STRENGTHENS Case for Market Evolution

Pentland's democracy objection fails because:

1. Empirical: Markets provide higher satisfaction than trustee intermediation (68% gap)

2. Theoretical: Exit > Voice for accountability (Hirschman)

3. Constitutional: Democracy requires constraints protecting individual rights from majority

4. Practical: Participatory consensus networks captured by organized minorities

5. Conceptual: Market democracy (unanimous consent, heterogeneous choices) more democratic than procedural democracy (majority rule, uniform decisions)

Revised conclusion: Evolutionary governance (federalism + markets + exit rights) is MORE democratic than Pentland's consensus networks.

F. Updated Conclusion: Steel-Man Defeated

I have now addressed five defenses Pentland might offer:

1. Bottom-up design → Rebutted (Section VII.A): Distribution ≠ absence of design

2. Facilitating emergence → Rebutted (Section VII.B): All facilitation embeds values

3. Adaptive institutions → Rebutted (Section VII.C): Adaptation toward wrong goals

4. Irreversibility → Rebutted (Section VII.D): Negative selection handles; positive coordination fails faster

5. Democracy → Rebutted (Section VII.E): Markets more democratic than deliberation

Pattern: Even the strongest defenses collapse under scrutiny.

Core problem: All defenses attempt to bridge unbridgeable epistemological gap:

- Either knowledge is dispersed (evolution required)

- Or knowledge is centralizable (design possible)

- No middle ground exists

Pentland wants both: acknowledge dispersed knowledge (Sections II-III descriptive) yet design institutions (Sections VI-VII prescriptive). This is logical contradiction, not reconcilable through clever institutional architecture.

VIII. CONCLUSION: EVOLUTION OR ENGINEERING?

A. The Core Dilemma

Pentland's book forces a choice between two incompatible frameworks, and the reader cannot have both.

The first framework, which I call Evolutionary Epistemology, correctly describes how successful institutions actually emerge. Knowledge relevant to institutional design is dispersed across millions of individuals, embedded in local contexts, and often tacit in ways that resist articulation. Successful institutions emerge through variation and selection as communities experiment with different arrangements and selectively retain those that prove workable. No designer possesses the information necessary for optimization because the relevant information exists only in the aggregate behavior of countless individuals responding to local circumstances. Order arises spontaneously from decentralized experimentation, not from centralized planning. This is the framework Pentland employs when describing common law, the Hanseatic League, and scientific citation networks.

The second framework, which I call Constructivist Design, animates Pentland's prescriptive proposals. This framework treats collective goals such as the SDGs as legitimate inputs for

institutional engineering. It assumes that institutions can be consciously designed to achieve these externally specified objectives. It presupposes that designers can determine optimal coordination mechanisms, or at least mechanisms superior to what would emerge spontaneously. It holds that order should be deliberately constructed toward beneficial ends rather than allowed to emerge from undirected experimentation. This is the framework Pentland employs when proposing consensus networks, data trusts, and mandatory audit trails.

These frameworks are mutually exclusive. One cannot simultaneously believe knowledge is too dispersed for central planning yet propose to design institutions toward centrally-determined objectives.

B. Why the Contradiction Predicts Failure

The theoretical flaw predicts practical failure through four mechanisms:

1. Epistemic Impossibility: If knowledge is genuinely dispersed, designers cannot know which goals to pursue or which institutions will achieve them. Pentland's proposals will fail for the same reasons Soviet planning failed: designers lack necessary information.

2. Incentive Misalignment: Designed institutions solving anticipated problems will be captured by designers' interests. Data trusts will serve trustees, consensus networks will serve network administrators, audit requirements will serve compliance bureaucrats.

3. Brittle Optimization: Institutions optimized for current goals become maladaptive when conditions shift. SDG-focused consensus networks cannot adapt when SDGs prove misguided.

4. Legitimacy Crisis: Designed institutions lacking evolutionary validation must rely on coercion. Pentland's proposals require mandatory participation, enforcement mechanisms, and sanctions for non-compliance.

C. The Path Forward: Pure Evolutionary Governance

A coherent framework for AI governance, one that remains faithful to Pentland's own evolutionary epistemology, would look quite different from what he proposes.

The first principle must be the rejection of predetermined collective goals. Neither SDGs, nor climate targets, nor equity metrics, nor any other abstraction selected by international bodies should serve as an input for institutional engineering. This is not because such goals are necessarily unworthy; some may indeed describe valuable states of affairs. The objection is epistemic: we do not know which goals are worth pursuing, which trade-offs among goals are acceptable, or which institutional arrangements would achieve any given goal.

Pretending otherwise is the fatal conceit. Let valuable goals emerge from evolutionary process rather than imposing them from above.

The second principle is pure ex-post liability without process mandates. Strict product liability for AI harms provides appropriate incentives without requiring regulators to specify in advance what safety looks like. Plaintiffs bear the standard burden of proof: they must demonstrate action, causation, and harm through evidence subject to adversarial testing. Defendants have incentives to maintain documentation that would help them defend against claims, but face no legal obligation to do so. This voluntary audit trail will emerge where valuable and disappear where not, precisely calibrated by market incentives rather than regulatory fiat.

The third principle is radical institutional diversity achieved through federalism, freedom of contract, and exit rights. Different jurisdictions should be permitted and encouraged to adopt different approaches to AI governance. Individuals should be free to opt into different governance regimes through contractual choice. And crucially, anyone dissatisfied with their current jurisdiction's approach must have meaningful ability to exit. This combination generates variation, enables selection, and produces adaptation.

The fourth principle is cultural selection without designer intent. Institutions that successfully coordinate AI development will spread through imitation as other jurisdictions observe their success. Failed institutions will be abandoned as their defects become apparent. No consensus network should determine which approaches spread; revealed preference through migration and imitation determines fitness. Designers should not be in the business of picking winners.

The fifth principle is negative selection only. Law should remove harmful variants through liability, but should not positively select toward preferred outcomes. Harmful AI applications face liability costs that reduce their fitness; beneficial applications avoid such costs and thrive. This is the evolutionary mechanism that produced the common law Pentland admires. It preserves adaptive flexibility because it does not require specifying in advance what "beneficial" means; it only requires identifying concrete harms after they occur.

D. The Ultimate Irony

Pentland has written an excellent book explaining why constructivist design fails and evolutionary emergence succeeds. Then he proposes constructivist design of evolutionary mechanisms.

The salvageable insight: Pentland's empirical research and analytical framework for understanding cultural evolution through story-sharing networks is valuable. His descriptions of how common law, scientific networks, and merchant leagues emerged are accurate and insightful.

The fatal error: Believing that because we understand the mechanism of emergence, we can engineer emergence toward desired ends.

We cannot. Emergence is what happens when we do not engineer. The moment we impose goals, select criteria, and design coordination mechanisms, we have left the evolutionary framework and returned to constructivism.

E. The Unresolved Tension

The tension in Pentland's framework invites reflection. His descriptive analysis establishes several premises: knowledge relevant to institutional design is dispersed and cannot be centralized; successful institutions emerge through experimentation rather than deliberate construction; central planning fails epistemically, not merely practically; common law succeeds precisely because it evolves case-by-case without predetermined objectives.

If these premises hold, certain implications follow. First, the epistemic barriers that defeat central planning apply equally to determining collective goals. The SDGs represent the preferences of individuals who occupied positions in international institutions at a particular historical moment; their claim to universal validity faces the same knowledge problem that undermines other forms of central coordination.

Second, the knowledge problem is not a technical obstacle amenable to better data or more inclusive processes. It reflects a fundamental epistemic barrier: the relevant knowledge is dispersed, tacit, and context-dependent by nature. No amount of audit trails, consensus networks, or stakeholder deliberation can aggregate what resists aggregation.

Third, the mechanism of failure is predictable from Pentland's own framework. Designers will optimize for measurable proxies rather than underlying values; actors will game the metrics; the system will become brittle when facing unanticipated challenges. This is precisely what Pentland's analysis of bureaucratic failure would predict.

The coherent resolution would be pure liability without predetermined goals: courts resolving actual disputes, consequences attaching to demonstrated harms, precedent accumulating through decentralized experimentation. This is the evolutionary process that produced the institutions Pentland admires.

Whether Pentland would accept these implications is, of course, his decision. The point is that his framework generates them, and the tension between his premises and his prescriptions remains unresolved.

F. Final Verdict

Alex Pentland has produced an elaborate analysis correctly identifying centralized bureaucracy's failures and emergent institutions' successes. His descriptive framework (cultural evolution through story-sharing networks) is a genuine contribution.

But his prescriptive program (designing consensus networks and data trusts to coordinate toward SDGs) appears difficult to reconcile with everything his analysis establishes. It embodies constructivism dressed in evolutionary language.

The coherent alternative is radical: pure evolutionary governance. No collective goals, no designed coordination, no mandatory process requirements. Only decentralized experimentation, voluntary coordination, and ex-post liability for demonstrated harms.

This requires epistemic humility that Pentland's proposals lack. It requires abandoning the confidence that because we understand how order emerges, we can engineer it toward preferred ends.

The choice is stark: evolution or engineering. Pentland wants both. The tension between these frameworks suggests his proposals may encounter the same challenges all constructivist schemes face: they require knowledge no designer possesses and impose uniformity preventing evolutionary adaptation.

REFERENCES

Ada Lovelace Institute. Exploring Legal Mechanisms for Data Stewardship. London: Ada Lovelace Institute, 2021. https://www.adalovelaceinstitute.org/report/legal-mechanisms-data-stewardship/

Alberola, Emilie, Julien Chevallier, and Benoît Chèze. "Price Drivers and Structural Breaks in European Carbon Prices 2005-2007." Energy Economics 30, no. 2 (2008): 847-871. https://doi.org/10.1016/j.eneco.2007.11.007

Aridor, Guy, Yeon-Koo Che, and Tobias Salz. "The Economic Consequences of Data Privacy Regulation: Empirical Evidence from GDPR." RAND Journal of Economics 54, no. 4 (2023): 695-731.

Arrow, Kenneth J. Social Choice and Individual Values. Yale University Press, 1951.

Bebchuk, Lucian A., and Scott Hirst. "Index Funds and the Future of Corporate Governance: Theory, Evidence, and Policy." Columbia Law Review 119 (2019): 2029-2146.

Blackmore, Susan. The Meme Machine. Oxford University Press, 1999.

Buchanan, James M., and Gordon Tullock. The Calculus of Consent. University of Michigan Press, 1962.

Burt, Ronald S. "Structural Holes and Good Ideas." American Journal of Sociology 110 (2004): 349-99.

Dawkins, Richard. The Selfish Gene. Oxford University Press, 1976.

Dawkins, Richard. The Extended Phenotype. Oxford University Press, 1982.

Dawkins, Richard. The Blind Watchmaker. W.W. Norton, 1986.

Demirer, Mert, Diego Jiménez-Hernández, Dean Li, and Sida Peng. "Data, Privacy Laws and Firm Production: Evidence from the GDPR." MIT Working Paper, 2024. https://www.diegojimenezh.com/assets/pdf/Demirer_et_al2023_Privacy.pdf

Dennett, Daniel C. The Intentional Stance. MIT Press, 1987.

Dennett, Daniel C. Darwin's Dangerous Idea. Simon & Schuster, 1995.

Ellerman, A. Denny, and Barbara K. Buchner. "Over-Allocation or Abatement? A Preliminary Analysis of the EU ETS Based on the 2005-06 Emissions Data." Environmental and Resource Economics 41, no. 2 (2008): 267-287. https://doi.org/10.1007/s10640-008-9191-2

European Commission. EU ETS Review Report: Emissions Trading: 2007 Verified Emissions from EU ETS Businesses. Brussels: European Commission, 2008.

Goldberg, Samuel, Garrett Johnson, and Scott Shriver. "Regulating Privacy Online: The Early Impact of the GDPR on European Web Traffic." Management Science, forthcoming, 2023.

Granovetter, Mark S. "The Strength of Weak Ties." American Journal of Sociology 78 (1973): 1360-80.

Harari, Yuval Noah. Nexus: A Brief History of Information Networks from the Stone Age to AI. Random House, 2024.

Hayek, Friedrich A. "The Use of Knowledge in Society." American Economic Review 35 (1945): 519-30.

Hayek, Friedrich A. The Fatal Conceit: The Errors of Socialism. University of Chicago Press, 1988.

Henrich, Joseph. The Secret of Our Success: How Culture Is Driving Human Evolution, Domesticating Our Species, and Making Us Smarter. Princeton University Press, 2016.

Henrich, Joseph. The WEIRDest People in the World: How the West Became Psychologically Peculiar and Particularly Prosperous. Farrar, Straus and Giroux, 2020.

Johnson, Garrett, Scott Shriver, and Samuel Goldberg. "Privacy & Market Concentration: Intended & Unintended Consequences of the GDPR." Management Science, forthcoming, 2023.

Laing, Timothy, Misato Sato, Michael Grubb, and Claudia Comberti. "Assessing the Effectiveness of the EU Emissions Trading System." Centre for Climate Change Economics and Policy Working Paper No. 126, 2013.

Lerer, Ignacio Adrian. "Law as Extended Phenotype: Toward an Evolutionary Theory of Legal Systems." SSRN Working Paper No. 5593470, 2024. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5593470

Lerer, Ignacio Adrian. "The Legislator as Extended Phenotype: A Darwinian Theory of Legal Evolution." SSRN Working Paper No. 5387400, 2025. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5387400

Mises, Ludwig von. Economic Calculation in the Socialist Commonwealth. Ludwig von Mises Institute, [1920] 1990.

Mises, Ludwig von. Human Action: A Treatise on Economics. Yale University Press, 1949.

Olson, Mancur. The Logic of Collective Action. Harvard University Press, 1965.

Open Data Institute. Data Trusts: Lessons from Three Pilots. London: ODI, 2019. https://theodi.org/article/odi-data-trusts-report/

Ostrom, Elinor. Governing the Commons: The Evolution of Institutions for Collective Action. Cambridge University Press, 1990.

Pentland, Alex. Shared Wisdom: Cultural Evolution in the Age of AI. MIT Press, 2025.

Polinsky, A. Mitchell, and Steven Shavell. "The Theory of Public Enforcement of Law." Handbook of Law and Economics 1 (2007): 403-54.

Scott, James C. Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed. Yale University Press, 1998.

Shavell, Steven. Economic Analysis of Accident Law. Harvard University Press, 1987.

METHODOLOGICAL NOTE

The author holds a law degree from Universidad de Buenos Aires and an Executive MBA from IAE Business School, with over thirty years of corporate legal practice. The formal mathematical notation in this paper, particularly the mechanism design formalization in Section III.E and the

Appendix, was developed with assistance from AI tools (Claude, Anthropic). The author assumes full responsibility for the theoretical framework, legal analysis, and all substantive claims. The AI assistance was limited to translating conceptual arguments into standard mathematical notation from mechanism design theory (Myerson-Satterthwaite, Gibbard-Satterthwaite). All empirical data cited are from published academic sources as referenced.

APPENDIX: Mathematical Proof of Pentland Impossibility Theorem

Given:

- Heterogeneous preference types $\theta_i \in \Theta$ (distributed $F(\theta)$)

- Designer mechanism $M: \Theta \times S \to A$

- Exogenous goal G^*

- Individual utility $U_i(a; \theta_i)$ (private)

- Social welfare $W(a) = \Sigma_i U_i(a; \theta_i)$ (unknown to designer)

To prove: No mechanism M satisfies all three conditions:

1. Participation (PC): $U_i(M(\theta, s)) \geq U_i(\emptyset)$ for all i

2. Incentive compatibility (IC): Truthful reporting θ_i is optimal

3. Goal achievement (GA): Outcome converges to G^*

without coercion (violating PC or IC).

Proof sketch:

Suppose M satisfies PC and IC. Then by the revelation principle, M implements a truthful equilibrium. But if θ_i are heterogeneous and G^* is exogenous, then for some types θ_j, the outcome $M(\theta) \neq \text{argmax } U_j(\cdot; \theta_j)$. Thus either:

- (A) M violates PC (some agents prefer opt-out), or

- (B) M violates IC (some agents misreport to steer toward preferred outcome), or

- (C) M abandons GA (outcome drifts from G^*)

The formal proof extends this via mechanism design impossibility theorems (Myerson-Satterthwaite, Gibbard-Satterthwaite).