Zenodo Working Paper | February 2026
Evolutionary Theory of Law Series

Regulating the Unintentional: AI Governance through Multi-Level Evolutionary Game Theory, Extended Phenotype Theory, and Constitutional Self-Restraint

Why the EU AI Act Fails Where Argentine Private Law Already Has the Answer

Ignacio Adrián Lerer
Independent Researcher | Buenos Aires, Argentina
adrian@lerer.com.ar | ORCID 0009-0007-6378-9749
February 2026

Abstract

Regulation (EU) 2024/1689 on Artificial Intelligence, the AI Act, in force since August 2024, rests on an undeclared premise: that AI systems possess intentionality sufficient to be the direct bearers of behavioral obligations.

That premise is philosophically mistaken, strategically counterproductive, and, from an evolutionary game theory perspective, predictably captured by those it nominally constrains.

This paper applies Multi-Level Evolutionary Game Theory (ML-EGT) and Extended Phenotype Theory (EPT) to demonstrate that the AI Act reproduces a structural error identified in corporate criminal liability and other regulatory domains: the assignment of Level 3 intentionality categories (Dennett, 1987) to systems operating at Level 0, producing cosmetic compliance by actors who actually possess intentionality and systematic evasion by those who can exploit the attribution gap.

Argentina holds, without having designed it as such, a legal architecture better suited to governing AI: the externalization principle of Article 19 of the National Constitution, which, extended a fortiori through Articles 14 and 28, creates a constitutional chain protecting the developer's freedom from preventive state intervention, and the objective liability for created risk under Articles 1757-1758 of the Civil and Commercial Code (CCyCN).

This combination constitutes an Evolutionarily Stable Strategy (ESS) that aligns real incentives with regulatory objectives without requiring fictions about computational intentionality. A convergent argument arrives from Cass Sunstein's Separation of Powers (MIT Press, 2026): administrative agencies best protect liberty not by maximizing regulatory output but by knowing when not to regulate. AI regulation is, above all, administrative law, and that administrative law requires a constitutional architecture that builds self-restraint into its foundations rather than leaving it to the discretion of agency heads.

The paper proposes a Hybrid Minimal model, three targeted CCyCN amendments and one amendment to Law 25.326, that satisfies the Argentine reasonableness test of Article 28 CN while outperforming the AI Act on every dimension of constitutional validity.

The central hypothesis is falsifiable: if the CCyCN model generates fewer adverse outcomes for affected parties than the AI Act in comparable jurisdictions over a five-year horizon, the hypothesis is confirmed.

Keywords: *artificial intelligence, regulation, evolutionary game theory, extended phenotype, objective liability, Argentine constitutional law, CCyCN, EU AI Act, Dennett intentionality, administrative self-restraint, Sunstein, separation of powers, a fortiori, Article 19 CN.*

1. The Problem the AI Act Cannot Articulate

In October 2024, the European Commission published the first implementation guidelines for Regulation (EU) 2024/1689, running to 87 pages. The document contains twenty-three references to the 'accountability' of AI systems and fourteen to their 'transparency.' Neither concept, as applied to a large language model or a neural network trained on image data, corresponds to anything those systems can possess in any meaningful sense. They do not hold themselves accountable. They cannot be transparent about intentions they do not have.

The AI Act classifies systems by risk level: unacceptable, high, limited, and minimal. High-risk systems, including credit-scoring models, AI-assisted medical diagnostics, and employment-candidate screening, carry obligations of technical documentation, conformity assessment, human oversight, and registration in the EU database (Article 6 and Annex III). A provider that satisfies these formal requirements obtains the CE marking and may operate freely in the European single market. The regulation is, formally, satisfied.

The difficulty is not with the formal requirements themselves. It lies in who the regulation treats as the obligated party. Throughout its 180 articles and thirteen annexes, the AI Act systematically frames the AI system, not its developer, not its deployer, as the entity capable of 'avoiding bias,' 'maintaining activity logs,' and 'being transparent toward users.' The system, in fact, optimizes loss functions over training data. It does not know that a regulatory obligation exists. It cannot intentionally violate one. It cannot comply with one in any sense that goes beyond what its designers and deployers have already imprinted upon it before deployment.

This confusion is not a drafting oversight. It reflects a structural property of AI regulation at the current moment: regulators understand intuitively that AI systems produce consequential outputs, but their conceptual vocabulary for assigning responsibility was built for agents with intentions, beliefs, and the capacity for genuine normative commitment. Applied to systems without those properties, that vocabulary generates what I term the Dennett-AI Act Gap: the regulation assigns obligations to the object; the subjects who actually control that object accommodate themselves to formal compliance requirements without any need to alter the system's substantive behavior.

The result is cosmetic compliance on a continental scale. Developers produce documentation. Deployers obtain certifications. AI systems continue doing exactly what they were designed to do before the regulation existed. The gap between formal satisfaction and substantive outcome is not a failure of enforcement; it is a structural consequence of misattributing intentionality.

1.1 Dennett's Intentionality Taxonomy Applied to AI Systems

Daniel Dennett proposed in The Intentional Stance (1987) a three-level taxonomy of systems according to the intentionality they exhibit. Understanding where AI systems fall in this taxonomy is not an academic exercise; it determines what a regulation can coherently demand of them and who, therefore, must actually bear regulatory obligations.

Level 1 systems detect variables and trigger responses. A thermostat has Level 1 intentionality: it 'wants' the temperature to reach a set point only in the minimal sense that it activates a heater when the reading falls below a threshold. There are no internal representations of the external world, no beliefs, no desires in any philosophically robust sense. Level 2 systems have internal states that can be described as beliefs and desires about the physical world. A dog navigating a familiar environment has beliefs about the location of objects and desires regarding food, companionship, and shelter. Level 3 systems hold beliefs about other agents' beliefs: they attribute mental states to others, anticipate expectations, feel shame when they violate shared norms, and cooperate genuinely because they understand that cooperation is understood by the other party as a commitment.

A generative AI system, GPT-4, Gemini 1.5, Claude, or any of their successors, operates at Level 0 with a compelling simulacrum of Level 3. It produces text that resembles high-order reasoning, expresses apparent preferences, and mimics normative commitment because it was trained on hundreds of billions of tokens produced by humans who engage in high-order reasoning, express genuine preferences, and make real normative commitments. The outputs look like Level 3. The mechanism producing them is matrix multiplication over weighted parameters. There is no 'there' there in the sense required for genuine agency.

Steven Pinker, in When Everyone Knows That Everyone Knows (Scribner, September 2025), formalizes why this distinction matters for social coordination and regulation. Common knowledge, that infinitely recursive epistemic state in which A knows that B knows that A knows that a rule exists, is the foundation of money, language, political legitimacy, and normative compliance. A rule is truly internalized, in Pinker's account, when the regulated party knows the rule, knows that the regulator knows the regulated party knows it, and adjusts behavior not merely to avoid sanction but because violation would be publicly visible and socially costly. AI systems cannot participate in common knowledge because they lack the cognitive architecture that makes epistemic recursion possible. They can be programmed to output the text 'I know you are watching' without this statement corresponding to any internal state.

When the AI Act obliges an AI system to be 'transparent' toward users, it actually obliges its developers and deployers to build transparency-producing mechanisms into the system's architecture and deployment context. The Act frames the obligated party as if it were the system. That asymmetry of grammatical reference, using the system as the apparent subject of regulatory obligations, becomes an architecture for evasion, because it invites the interpretation that satisfaction of documentation requirements about the system constitutes compliance with the underlying protective purpose.

1.2 The Dennett-Nash Gap: When Regulatory Games Fail

In prior work on asymmetric intentionality in corporate criminal liability (Lerer, 2025a), I formalized the Dennett-Nash Gap as the strategic advantage that accrues to lower-intentionality players in regulatory games designed for higher-intentionality agents. The mechanism operates as follows. A regulator at Level 3, capable of shame, reputation management, and genuine normative internalization, designs an enforcement scheme on the assumption that regulated parties share comparable cognitive architecture. The scheme expects regulated parties to internalize the norm, to feel social pressure from public violation, and to maintain compliance even when immediate sanctions are absent. When the regulated activity involves Level 0 or Level 1 systems, that assumption does not hold. The higher-intentionality regulator systematically loses not through bad luck but through expecting social constraints that structurally cannot materialize on the other side.

Applied to AI governance, the gap operates at three levels simultaneously. Regulators design the AI Act assuming that developers will internalize its normative logic, treat conformity as a genuine commitment to user protection, and invest in safety whether or not enforcement is present. Developers respond as Level 1 optimizers: they minimize compliance costs while maximizing market access, investing in documentation rather than design, in audit trails rather than substantive safety, in CE marking rather than demonstrated harm reduction. Deployers respond as Level 1 optimizers: they transfer liability upward toward providers and downward toward end users, structuring contracts to shield themselves from the consequences of harmful AI outputs. AI systems respond as Level 0 systems: they do not respond, because they cannot respond in any normatively relevant sense.

The outcome is not bad faith in any of its conventional meanings. Developers comply formally. Deployers comply formally. AI systems do what they do. The absence of substantive change is not deception; it is the structural consequence of designing a normative system for agents who do not exist as the regulation imagines them. Robert Axelrod's work on the evolution of cooperation (1984) shows that cooperation emerges and stabilizes when players interact repeatedly, can identify defectors, and bear reputational costs from defection. None of these conditions hold for Level 0 systems in the AI Act's regulatory ecosystem: the system cannot accumulate a reputation, cannot be socially sanctioned, and cannot modify its behavior in response to reputational pressure. The regulatory game is structurally non-cooperative regardless of the Act's intentions.

2. ML-EGT: Three Simultaneous Games with Incompatible Payoffs

Standard Evolutionary Game Theory models competition between strategies within a homogeneous population. Its multi-level extension introduces a structurally important modification: distinct groups of actors play different games simultaneously, with incompatible payoff functions, misaligned time horizons, and no shared metric of success. The AI regulatory ecosystem involves at least three such games running in parallel. Understanding their payoff structures explains why the AI Act produces the equilibrium it does, and what a different equilibrium would require.

2.1 The Developers' Game: Regulatory Capture at Scale

OpenAI, Anthropic, Google DeepMind, Meta AI, and Mistral AI, the five firms that collectively define the frontier of large language model development, play a strategic positioning game in which regulation is an environmental variable to be shaped, not an exogenous constraint to be absorbed. From January 2024, when final trilogue negotiations on the AI Act began, through December of that year, those five firms invested a documented combined total exceeding 38 million euros in lobbying activities before EU institutions, according to the EU Transparency Register updated through the fourth quarter of 2024. That figure excludes lobbying by their industry associations, trade groups, and allied policy institutes.

The dominant strategy in this game is not opposition to regulation but capture: designing the regulation so that its compliance costs function as barriers to entry that protect incumbents from competitive displacement. A regulation requiring extensive technical documentation, conformity assessments by notified bodies, post-market monitoring systems, and EU database registration advantages those who already possess the organizational scale to satisfy those requirements at tolerable cost. A startup in Buenos Aires, São Paulo, or Warsaw attempting to launch a competitive model cannot absorb the AI Act's compliance costs, estimated by the Stanford HAI analysis of 2024 at between six thousand and five hundred thousand euros per firm depending on system classification and organizational size, as readily as a firm valued at one hundred billion dollars.

The outcome is a more concentrated market shielded by the very regulation nominally designed to discipline it. This is not a conspiracy; it is a predictable equilibrium in a game where incumbents have strong incentives to participate in regulatory design and startups have neither the resources nor the access to counter their influence. From ML-EGT, this is the dominant strategy at the developer level, and it is stable: no individual developer has incentive to unilaterally propose a less burdensome regulatory architecture, because any individual developer benefits from barriers that protect their current position. The mutation capable of destabilizing this equilibrium would be a regulatory design that penalizes verified harmful outcomes proportionally rather than imposing uniform process requirements regardless of outcome. That mutation is not present in the AI Act.

2.2 The Deployers' Game: Cosmetic Compliance and Delegated Liability

Firms that deploy AI systems without having developed them, a bank using an externally sourced credit-scoring model, a hospital using a third-party diagnostic imaging system, an employment agency using an externally developed candidate-ranking tool, play a qualitatively different game. Their payoff structure is not primarily about competitive positioning in AI markets but about liability allocation: specifically, about shifting responsibility for harmful outputs toward either the model's provider above them or the end user below them.

The AI Act's provider/deployer distinction, set out in Article 28, facilitates precisely this strategy. Deployers assume additional obligations beyond those of providers only in specified circumstances: if they modify the system substantially, use it outside its intended purpose, or deploy it in a high-risk context not anticipated by the provider's documentation. A bank that receives a credit-scoring model already certified under the AI Act can credibly argue that its responsibility is limited to correct deployment as specified by the provider. Liability for the model's intrinsic architecture, its training data, its weights, its tendency to produce discriminatory outputs for protected groups, remains formally with the provider.

The deployer's dominant strategy is therefore to obtain properly certified systems and deploy them with minimal modification, maximizing the transfer of liability to the provider while maintaining the operational flexibility to use AI outputs as determinative inputs in high-stakes decisions. The affected party, the loan applicant, the patient, the job candidate, faces the output of a liability-laundered system in which the organization that made the consequential decision can point to the certification and the organization that produced the model can point to the correct deployment.

2.3 The Affected Parties' Game: Rational Impotence

Citizens whose consequential decisions are mediated by high-risk AI systems operate at Level 3 of intentionality. They can process grievances, form expectations about institutional behavior, engage complaint mechanisms, and in principle activate the legal remedies that the AI Act ostensibly provides. Article 86 establishes a right to explanation for automated decisions that significantly affect individuals. Article 85 establishes a right to lodge a complaint with the national supervisory authority.

Both rights are largely theoretical for the population of individuals who actually suffer AI-mediated adverse outcomes. The information asymmetry between an individual whose loan application was denied or whose job application was filtered out and the developer of a model trained on hundreds of billions of data points is not merely quantitative; it is architectural. The individual cannot know what features determined the output. The 'explanation' that Article 86 requires is not operationalized in terms of substantive content: a system can produce an explanation that is formally compliant and substantively useless, 'Your application received a risk score of 0.73 on a scale of 0 to 1 based on 847 features', without enabling any meaningful assessment of whether an error occurred or whether discrimination was at work.

The ProPublica investigation of the COMPAS recidivism prediction system, published in 2016, is the canonical empirical illustration of this dynamic. COMPAS was used in sentencing decisions across multiple US jurisdictions for more than a decade. It produced risk scores that were, for Black defendants, twice as likely as for white defendants to assign high-risk scores to individuals who did not subsequently reoffend. The system was nominally subject to judicial oversight; the risk scores were inputs, not binding determinations. In practice, they functioned as determinative. No firm was penalized. The affected parties, those sentenced under AI-mediated recommendations that were statistically biased, had no effective remedy because the causal connection between the system's architecture and their specific outcome was not demonstrable at the individual level. Aggregate statistical bias is not a legal cause of action for an individual defendant.

The three-level ML-EGT structure thus predicts the observable equilibrium under the AI Act: developers capture the regulation to maintain competitive position; deployers use certification as liability cover; affected parties receive declarative rights with no actionable enforcement pathway. The equilibrium is stable because no player has unilateral incentive to change strategy, and the regulation reinforces rather than disrupts the payoff structure that produced the problem it was designed to address.

3. EPT: Competing Memeplexes and Their Extended Phenotypes

From the perspective of Extended Phenotype Theory as applied to legal institutions (Lerer, 2025), AI regulation is best understood not as a deliberate policy choice produced by disinterested analysis but as the extended phenotype of a competition among memeplexes, each seeking to occupy the normative space governing AI with its own institutional structures as reproductive vehicles. Three memeplexes currently dominate this competition, and none of them resolves the intentionality problem identified in Section 1.

The 'preventive safety' memeplex holds that AI risks are sufficiently systemic, opaque, and irreversible to justify intensive ex ante regulation. Its core memes include 'AI alignment,' 'existential risk,' and 'precautionary governance.' Its principal carriers are European regulators, AI safety academics, Stuart Russell at Berkeley, Yoshua Bengio at Mila, Nick Bostrom at the Future of Humanity Institute, and organizations including the Future of Life Institute, the Center for AI Safety, and the UK AI Safety Institute. Its extended phenotype is the AI Act itself: a normative architecture that classifies systems before deployment, mandates conformity assessments, and creates a bureaucratic infrastructure whose reproduction depends on the continued expansion of the regulated domain. The memeplex replicates through citation networks in academic AI safety literature, through policy briefings to legislators who lack technical expertise and are therefore receptive to precautionary framings, and through media coverage of AI incidents that reinforces the narrative of systemic risk.

The 'responsible innovation' memeplex holds that premature regulation generates technological lock-in, economic harm, and competitive displacement toward less regulated jurisdictions. Its core memes include 'innovation-friendly regulation,' 'proportionality,' and 'ex post accountability.' Its principal carriers are the Anglophone technology industry, think tanks including the Institute for Progress and the Information Technology and Innovation Foundation, and, from January 2025 onward, the Trump administration, which revoked Biden's Executive Order on AI Safety on the first working day of the new presidency. Its extended phenotype varies by jurisdiction: in the US, it takes the form of sectoral regulation through existing agencies rather than a comprehensive AI statute; in the UK, it produced the 'pro-innovation approach to AI regulation' published in 2023, which explicitly declined to create new AI-specific legislation in favor of sector-specific guidance from existing regulators.

The 'technological sovereignty' memeplex treats AI as a variable of geopolitical dependency and frames regulatory design as a tool for determining who controls twenty-first-century cognitive infrastructure. Its core memes include 'digital strategic autonomy,' 'AI sovereignty,' and 'data localization.' It is the dominant memeplex in Chinese AI policy since the New Generation Artificial Intelligence Development Plan of 2017, but it also animates significant strands of EU AI governance: the AI Act's requirements for EU-based notified bodies, the emphasis on 'trustworthy AI' as a European competitive differentiator, and the implicit goal of making the EU regulatory model the global standard, replicating the GDPR dynamic in the AI domain.

The fitness of each memeplex is determined by its compatibility with the political and institutional environments in which it competes, not by its accuracy in diagnosing the problem it claims to address. The preventive safety memeplex is highly fit in the EU institutional environment because it provides a legitimating narrative for regulatory expansion and because the EU's constitutional architecture, lacking an equivalent to Article 19 of the Argentine National Constitution, does not generate strong constitutional barriers to preventive regulation. The responsible innovation memeplex is highly fit in the US because it aligns with both the First Amendment's strong protection of expressive activity and the political culture of limited government regulation of industry. None of the three is particularly well fitted to diagnose and address the Dennett-Nash Gap, because all three presuppose some form of agency in the AI system that justifies treating it as a direct object of regulation.

4. The Argentine Pivot: A Constitutional Chain, Not a Single Hook

Argentina arrived at its current position not through deliberate AI policy design but through the accumulated effects of three independent historical decisions: drafting a constitution in 1853 with the externalization principle as a structural limit on state punitive power; ratifying and giving constitutional hierarchy to an extensive body of international human rights law in 1994; and modernizing the civil code in 2015 with objective liability for created risk as the operative standard for dangerous activities. The accidental result is a legal architecture that resolves the intentionality problem the AI Act cannot articulate, and does so without invoking any fiction about computational agency.

4.1 Article 19 CN as the Primary Anchor

Article 19 of the National Constitution, textually stable since its adoption in 1853, establishes that private actions of individuals that in no way offend public order or morality, nor harm a third party, are reserved solely to God and exempt from the authority of magistrates. Argentine constitutional doctrine, developed from Carlos Sánchez Viamonte's early twentieth-century analyses through the Supreme Court's landmark Bazterrica ruling (Fallos 308:1392, 1986) and subsequently reinforced in Arriola (Fallos 332:1963, 2009), extracted from this provision the externalization principle: the state may only intervene coercively upon externalized conduct that produces verifiable harm to identifiable third parties.

This principle has been the anchor of my argument in prior conference presentations and publications on AI regulation, and it remains so here. What the externalization principle establishes is not a libertarian preference without institutional grounding: it is a structural limit on state power derived from the constitutional text itself, with seven decades of Supreme Court jurisprudence behind it, and it has direct and precise application to the AI governance problem. An AI system that has not caused verifiable harm to a specific third party cannot, under Article 19 CN, be the object of punitive or preventive state intervention directed at its internal architecture.

The process of designing, training, and developing an AI system, the phase at which the AI Act's highest compliance burdens attach, occurs before any output reaches a third party. It is, categorically, within the protected sphere that the 1853 drafters set beyond the reach of magistrates. The developer engaged in that process is exercising a constitutionally protected freedom. The AI Act imposes obligations during that phase, technical documentation, bias auditing, training data governance, that constitute state intervention in private activity that has produced no externalized harm to any third party.

The standard objection to this formulation, that Article 19 protects persons, not objects, and therefore cannot protect 'the AI system's internal processes', misses the operative structure of the argument entirely. What is protected is not the AI system's intimacy or any anthropomorphized version of its internal states. What is protected is the developer's freedom from state intervention while the developer's constitutionally protected activity, designing and building a system, remains within the private sphere and has produced no externalized harm. The AI system is the medium through which the developer exercises a protected activity. It is not the rights-bearer. The developer is. Article 19 protects the developer, during the development phase, from preventive regulation that lacks the constitutionally required predicate of externalized harm to third parties.

4.2 A Fortiori: Articles 14 and 28 Deepen the Same Logic

If Article 19 already establishes that state intervention requires externalized harm as its constitutional predicate, two further constitutional provisions extend that logic into the regulatory domain with additional normative force. The move is not one of parallel arguments but of logical reinforcement: a fortiori, if the strictest constitutional protection against state intervention applies where the activity is purely private, then even the lesser protection applicable to commercial activity subject to regulatory oversight requires a justification that the AI Act cannot provide.

Article 14 CN establishes the right to work and to exercise any lawful industry. Developing and deploying AI systems is an exercise of this right: it is an industry that is currently lawful in Argentina absent specific statutory prohibition, and the constitutional text guarantees its exercise. This provision operationalizes Article 19's negative formulation, the state may not intervene on private activity without externalized harm, into a positive liberty: the state must affirmatively recognize the right to engage in lawful industry. The AI Act's compliance burdens, applied before any verified harmful output, restrict this right without the constitutionally required justification.

Article 28 CN establishes that statutes may not alter the rights recognized by the Constitution under the pretext of regulating them. From this provision, Argentine constitutional doctrine derived the test of regulatory reasonableness, systematized in the Supreme Court's Inchauspe ruling (Fallos 199:483, 1944) and consistently applied thereafter: any restriction on constitutionally recognized rights must pursue a legitimate objective, be suitable to achieve it, be proportional in its burden relative to the benefit sought, and must not nullify the regulated right. This is the equivalent of what US constitutional theory calls the means-ends test in rational basis or intermediate scrutiny review, but with a stronger normative foundation because it derives from explicit constitutional text rather than judicial doctrine.

The EU AI Act fails this test on the last two criteria, demonstrably. Its compliance costs are empirically documented: the Stanford HAI 2024 analysis placed them at between six thousand and five hundred thousand euros per firm, with median costs for high-risk system providers of approximately eighty thousand euros annually. Its benefits in terms of actual reduction of discriminatory outcomes, prevented accidents, or improved user protection are entirely speculative as of February 2026, because the regulation entered its operational phase only in August 2024 and no comparative outcome data exists. A regulation whose costs are certain and whose benefits are unproven fails the proportionality component of the Article 28 reasonableness test. Moreover, the compliance costs associated with conformity assessments, notified body certification, and EU database registration function as barriers to entry that effectively nullify the right to exercise AI-related industry for actors without sufficient scale, violating the non-nullification requirement of Article 28.

The three constitutional provisions therefore form a chain rather than a list of independent arguments. Article 19 establishes the threshold: harm must be externalized before the state may intervene. Article 14 establishes the liberty interest at stake: industry is a recognized constitutional right, not a regulatory privilege. Article 28 establishes the standard for evaluating any restriction that is nonetheless proposed: legitimate objective, suitability, proportionality, non-nullification. The EU AI Act fails at each level of this chain. The CCyCN model, as developed in Section 5, passes at each level.

4.3 Sunstein and the Administrative Agency That Knows When to Stop

A convergent argument arrives from a different tradition, administrative law, and from an author whose institutional credentials in regulatory design are difficult to dispute. Professor Cass Sunstein served as Administrator of the Office of Information and Regulatory Affairs in the Obama administration from 2009 to 2012, reviewing and approving or blocking regulations from every federal agency. In his most recent work, Separation of Powers: How to Preserve Liberty in Troubled Times (MIT Press, 2026), he describes the internal discipline that his role required and, by implication, the discipline that any regulatory agency serious about protecting liberty must exercise.

His account is worth quoting with precision. On numerous occasions, he told agencies not now or even no, not based on personal views but because the president and his team did not want to impose those regulations at that time. Walking the White House hallways during 2009 and 2010, his message was explicit: it is time for economic growth, it is not time for regulations. The explicit goal was, in his own formulation, to give space to private freedom of multiple types to flourish (Sunstein, 2026, p. 63).

This account is more than institutional memoir. It describes a structural property of regulatory governance that goes directly to the heart of the AI regulation debate: an administrative agency that does not know how to exercise self-restraint does not protect liberty, it undermines it. The question 'should we regulate this activity, and if so, with what instrument, at what cost, on the basis of what evidence?' is not a supplementary question to be answered after the decision to regulate has already been made. It is the primary question, and agencies that fail to ask it seriously produce regulations that impose costs with documented certainty while generating benefits with speculative probability.

AI regulation is, at its core, administrative law. Whatever form it takes, a European regulation, a national statute, a code of conduct, it occupies the space that administrative agencies carve out when they assert jurisdiction over a domain. Sunstein's insight applies with full force: the relevant question for those agencies is not 'how should we regulate AI?' but 'should we regulate AI now, with this instrument, at these costs, given this evidence?' The European Commission answered the first question extensively and did not seriously ask the second. The result was an instrument whose compliance architecture was designed during a period of AI anxiety without the benefit of comparative outcome data.

The Argentine constitutional framework, through the externalization principle of Article 19 and the reasonableness test of Article 28, already embeds Sunstein's discipline as a constitutional constraint rather than an administrative preference. An Argentine agency that regulates AI development before harm is externalized acts unconstitutionally under Article 19. An Argentine legislature that imposes compliance costs disproportionate to demonstrated benefits acts unconstitutionally under Article 28. The discipline Sunstein exercised from within the Executive Office of the President, as a matter of policy judgment subject to political override, Argentine constitutional doctrine imposes from the outside as a judicially enforceable limit on regulatory power.

This convergence between the Argentine constitutional tradition and Sunstein's administrative restraint theory is structurally coherent: both derive from the same liberal premise that freedom is the baseline, that the burden of justification lies always with those who would restrict it, and that good regulatory design is as much about knowing what not to do as about designing what to do. The difference is that in Argentina, the discipline is constitutionalized; it does not depend on having a Sunstein in the OIRA.

4.4 CCyCN Articles 1757-1758: Objective Liability Without Intentionality Fiction

The Civil and Commercial Code of the Nation, enacted in 2014 and in force since August 2015 under the drafting leadership of Ricardo Lorenzetti (Justice of the Argentine Supreme Court and principal drafter of the CCyCN), incorporated in Articles 1757 and 1758 a regime of objective liability for created risk that anticipates the needs of AI governance with structural precision that no legislator in 2015 could have planned deliberately.

Article 1757 establishes that every person is liable for damage caused by the risk or defect of things, or of activities that are risky or dangerous by their nature, by the means employed, or by the circumstances of their realization. Article 1758 specifies that the owner and the custodian are jointly and concurrently liable, with objective attribution. Applied to AI systems, the framework operates without requiring special legislation. Developing and deploying AI systems in decision contexts that affect rights, employment, credit, health, liberty, constitutes a risky activity within the meaning of Article 1757. The developer is liable as the creator of the risk. The deployer is liable as the custodian of the system in its operational context. Neither is required to have acted with intent or negligence; the harm, the causal link, and the risky activity suffice.

The sole exonerations available under this regime are causa aliena: a third-party act that breaks the causal chain; an unforeseeable external event; or deliberately incorrect information supplied by the affected party that induced the harmful output. None of these exonerations is easily satisfied in the AI discrimination context, which is precisely the incentive structure the regime generates: developers who want to avoid liability will invest in systems that do not produce discriminatory outputs, because they cannot credibly claim exoneration when the harm is traceable to architecture they controlled.

Dimension	EU AI Act (Reg. 2024/1689)	Argentine Model (Arts. 19, 14, 28 CN + CCyCN)
Constitutional grounding	EU Charter Art. 8 (data protection); no equivalent to Art. 19 CN externalization principle	Arts. 19, 14, 28 CN: externalization threshold, industry freedom, reasonableness test
Obligated party	Grammatically: the AI system. Functionally: its provider and deployer through certification obligations	Developer and deployer directly, as persons with Level 3 intentionality capable of genuine normative response
Intervention timing	Ex ante: before any harm is produced, as condition of market entry	Ex post: upon verified harm, through existing judicial channels with objective attribution
Intentionality coherence	Low: imposes obligations requiring agency on systems lacking agency	High: regulates only those who can comply and have real incentives to do so
Attribution of liability	Provider/deployer split allows systematic liability laundering	Concurrent and joint: developer and deployer are both objectively liable for any harm
Compliance cost	€6,000–€500,000 per firm (Stanford HAI, 2024);	Low: risk internalization in design, activated only when harm occurs in litigation

	disproportionate for non-incumbent actors	
Reasonableness test (Art. 28 CN)	Fails proportionality (costs certain, benefits unproven) and non-nullification (entry barriers)	Passes all four requirements: legitimate objective, suitable, proportional, non-nullifying
Scalability across technologies	Low: requires legislative reclassification with each generation of AI systems	High: 'risky activity' standard self-adjusts to technological development without amendment
Protection of affected parties	Declarative rights (explanation, complaint) without actionable causal pathway	Actionable: harm + causal link + risky activity = standing; burden shifts with Art. 1735 CCyCN
Administrative self-restraint (Sunstein criterion)	Absent: EU Commission did not seriously ask whether regulation was needed at this moment	Constitutionalized: Arts. 19 and 28 CN require justification before intervention is permissible

Table 1. Structural comparison between the EU AI Act and the Argentine constitutional/CCyCN model across ten analytical dimensions.

5. The Regulatory ESS: A Hybrid Minimal Model

Multi-Level Evolutionary Game Theory predicts that the regulatory equilibrium that persists is not the normatively optimal one in the abstract but the one from which no actor with the power to destabilize it has unilateral incentive to deviate, holding others' strategies constant. A design that imposes excessive costs on the industry will generate capture pressure until the regulation is weakened or circumvented. A design that fails to protect affected parties will generate political reform pressure until new intervention is mandated. The ESS is the design that aligns each player's incentives with the regulatory objective such that the equilibrium is self-reinforcing.

The formal ESS condition, following Maynard Smith (1982), requires that the payoff to a player maintaining the equilibrium strategy exceeds the payoff to any mutant strategy, given the population's current composition. In the regulatory context, this means that developers who invest in genuine safety measures must face lower expected costs than developers who invest in cosmetic documentation; deployers who select well-designed systems must face lower expected liability than deployers who select poorly designed ones; affected parties must have actionable claim paths that generate deterrence; and the state must have enforcement mechanisms that do not require institutional capacity it does not possess. What follows demonstrates that the four-component Hybrid Minimal model satisfies this condition for each player.

5.1 Direct Application of Article 1757 CCyCN as Backbone

The first and foundational component requires no legislative action: judicial recognition that deploying AI systems in high-stakes decision contexts constitutes a risky activity within the meaning of Article 1757 CCyCN. Argentine courts have applied the risky activity standard to pharmaceutical products, automated industrial equipment, medical devices, and financial instruments. The analogical extension to AI systems deployed in employment, credit, health, and justice contexts is straightforward under civilian methodology and does not require a creative leap.

The payoff effect of this recognition is the central mechanism of the ESS. Objective liability under Article 1757 places risk internalization in design rather than documentation compliance. A developer whose system produces discriminatory credit decisions cannot satisfy liability with a CE marking equivalent or a conformity assessment certificate; they must demonstrate either that the harm was caused by a third-party act or unforeseeable event, which is structurally implausible when the harm is statistical and architectural, or that the system did not in fact cause the alleged outcome. The former defense will rarely succeed in cases of documented algorithmic bias. The latter is an argument about causation, not about regulatory compliance, and it requires engaging with the specific harm rather than pointing to a general certification.

The ESS stability condition is satisfied at the developer level because developers who invest in systems that genuinely reduce discriminatory outputs face lower expected litigation costs than developers who invest only in documentation, regardless of what competitors do. The strategy dominates independently of others' choices, which is precisely the condition for evolutionary stability.

5.2 Three Targeted CCyCN Amendments

The current CCyCN contains practical gaps that generate litigation uncertainty. These are not structural normative lacunae, the framework exists and applies, but procedural ambiguities that create incentives for prolonged litigation as a defensive strategy against claims. Three amendments, each adding a single article, resolve these gaps without expanding state intervention beyond what the constitutional chain already permits.

> *Proposed Article 1735 bis: In cases of alleged damage caused by automated decision-making systems affecting rights of equality or non-discrimination, the burden of proving that the system did not generate the alleged discriminatory outcome rests on the party that deployed it in the specific case. The affected party must demonstrate (a) that a consequential decision adversely affected a legally protected interest and (b) that an automated system participated in that decision. Once these elements are shown, the deployer bears the burden of demonstrating the absence of a discriminatory causal link.*

This provision addresses the causal opacity problem that the LLM Council adversarial analysis (Annex A) identified as the most serious practical limitation of the current CCyCN framework: it does not require the affected party to penetrate a black-box neural network; it requires the deployer, who controls the system and has access to its operational parameters, to demonstrate that its output in the relevant context was non-discriminatory. The incentive effect follows immediately: deployers who cannot sustain that burden will invest in pre-deployment auditing and redesign rather than post-litigation defense.

> *Proposed Article 1745 bis: The rights recognized under this Code in relation to collective or diffuse damage may be invoked, in cases where the alleged damage is attributable to automated systems affecting an indeterminate or determinable plurality of persons, by affected individuals, by associations with an institutional purpose related to the affected rights, and by the national Ombudsman.*

This provision extends existing collective action standing doctrine to AI-generated diffuse harm, addressing the second major gap identified in the adversarial analysis: the structural impotence of individual victims of algorithmic harm who lack both the financial resources and the standing to pursue claims that are individually small but collectively significant. The COMPAS scenario, hundreds of individuals sentenced under a biased risk-scoring system, each with a marginal individual claim, no coordinated class action, and no institution with standing to pursue aggregate liability, would be directly addressed by this amendment.

> *Proposed Article 1721 bis: The deployment of automated decision-making systems in contexts that affect rights recognized in this Code constitutes a risky activity within the meaning of Article 1757, without prejudice to the concurrent objective liability of the system's developer under the same provision.*

This amendment resolves the guardianship attribution problem in multi-stakeholder AI deployments: when a bank deploys a credit-scoring model developed by a third party, running on cloud infrastructure provided by a fourth party, both the developer and the deployer are concurrently and objectively liable for damage caused by the system's outputs. The deployer cannot point to the developer's certification; the developer cannot point to the deployer's use outside specifications. The legislative clarification reduces the litigation period during which defendants contest liability attribution rather than engaging the substance of the harm.

5.3 One Amendment to Law 25.326

> *Proposed Article 12 bis of Law 25.326 (Personal Data Protection Act): Any natural or legal person whose legally protected interests are affected by a decision adopted with the assistance of automated systems must be informed, at the time of communication of that decision, that automated assistance was used. This obligation does not require disclosure of the system's internal architecture, training datasets, algorithmic parameters, or any other component constituting protected intellectual property or trade secret. It requires only disclosure of the fact of automation.*

This single provision, at minimal compliance cost, adding a sentence to a letter, email, or notification, performs several functions simultaneously. It enables the affected party to invoke the burden-shift of proposed Article 1735 bis, because they cannot know to claim the shift without knowing that an automated system participated. It activates informed consent mechanisms for data subjects under the existing Law 25.326 framework. And it generates, over time, a body of self-identification of AI-mediated decisions that provides the empirical foundation for both judicial jurisprudence and future legislative calibration. The AI Act's right to explanation under Article 86 requires substantive disclosure that threatens trade secrets and generates enormous compliance costs. The proposed Article 12 bis achieves the informational objective that actually matters, enabling the affected party to decide whether to challenge the decision, without requiring access to the system's internals.

The most compelling case for departing from this framework is also the most frequently invoked: the protection of minors from algorithmic engagement optimization designed to exploit variable reward mechanisms and maximize platform dependency. The intuition is correct; the legislative response it generates is not. A platform whose algorithmic architecture produces verified psychological harm to a minor has caused a verifiable injury to an identifiable third party through a risky activity within the meaning of Article 1757 CCyCN. The developer and deployer are concurrently liable under the objective standard, without any need to prove intent or negligence. The *Convención sobre los Derechos del Niño* (**Convention on the Rights of the Child)**, incorporated at constitutional hierarchy by Article 75(22) CN, and Law 26.061 on the Integral Protection of the Rights of Children and Adolescents, already generate judicially enforceable protective obligations that apply to digital environments.

The proposed Article 1735 bis operates without distinction as to the age of the affected party. No AI-specific norm is needed because the general regime already governs the conduct. Adding one would imply, contrary to the paper's central argument, that the existing framework is insufficient and that preventive AI-specific intervention is warranted, which is precisely the error attributed to the EU AI Act.

5.4 Why the Four-Component Model Is an ESS

The full model, Article 1757 CCyCN applied directly, three CCyCN amendments, and one Law 25.326 amendment, constitutes an evolutionarily stable strategy because the formal ESS condition is satisfied for every player in the regulatory game.

For developers: the expected cost of a system that generates genuine safety, lower probability of harmful outputs, stronger causal defenses in litigation, capacity to satisfy the Article 1735 bis burden shift, is lower than the expected cost of a system that produces documentation without substantive safety improvement. This differential is independent of competitors' strategies: it does not matter whether other developers invest in safety, because one's own liability exposure depends on one's own system's outputs, not on the market's average. The strategy of genuine safety investment therefore dominates cosmetic compliance regardless of others' choices. That is the definition of evolutionary stability at the developer level.

For deployers: selecting a system whose developer can satisfy the Article 1735 bis burden shift reduces the deployer's own concurrent liability exposure under proposed Article 1721 bis. A deployer who selects a poorly designed system and cannot demonstrate non-discrimination faces joint and concurrent liability with the developer. A deployer who selects a well-designed system and can point to the developer's demonstrated capacity to satisfy the burden shift has a credible defense. The incentive is to be a demanding buyer of AI systems, which generates upstream pressure on developers to produce systems that can actually satisfy that standard.

For affected parties: the burden shift under proposed Article 1735 bis, combined with the collective damage standing of proposed Article 1745 bis, produces actionable claim paths that currently do not exist. The model does not promise that every algorithmic harm produces a successful claim; it ensures that the structure of the litigation game is not systematically rigged in favor of the better-resourced defendant. That is sufficient to generate deterrence.

For the state: the model requires no novel institutional infrastructure. It operates through the existing judicial system and the AAIP with modified mandates. Enforcement is activated by complaints from affected parties and the associations and Ombudsman with standing under proposed Article 1745 bis, not by proactive licensing or certification review. The state expends resources only when harm has been alleged, which satisfies Sunstein's administrative self-restraint criterion: the agency intervenes because something has happened, not because something might happen.

The ESS is stable because no player has unilateral incentive to defect. It is also constitutionally valid at every level of the Article 19-14-28 CN chain: the framework activates upon externalized harm (Article 19), preserves the freedom to exercise the industry of AI development (Article 14), and imposes restrictions proportional to demonstrated risks rather than speculative harms (Article 28).

6. Adversarial Validation: The LLM Council Experiment and Its Metaparadox

The arguments developed in this paper were subjected to adversarial validation through a structured multi-perspective debate that adapts the LLM Council framework proposed by Zhu et al. (2026) from sequential task planning to adversarial legal-theoretical validation, implementing it as a three-round deliberative process with explicit cross-critique and argumentative synthesis.

The complete record of this process, including the perspectives, cross-critiques, post-critique revisions, and synthesis, is reproduced in Annex A. This section describes the method, summarizes the findings relevant to the paper's central claims, and examines the structural paradox that the experiment generates.

6.1 Method

The LLM Council assigns independent analytical roles to large language model sessions, each operating without knowledge of the others' outputs. For this experiment, four perspectives were defined with explicit mandates: a constitutional libertarian perspective tasked with defending the Article 19 CN framework and the CCyCN model as constitutionally superior to preventive regulation; a preventive risk analyst perspective tasked with developing the strongest possible case that the CCyCN framework is insufficient for AI-generated harms; a pragmatic comparativist perspective tasked with evaluating the empirical evidence from EU, US, UK, and Australian regulatory approaches and identifying what each produces in practice; and a devil's advocate perspective with an unconditional mandate to identify the logical weakest points in each of the other three perspectives, without commitment to any substantive position.

The process operated in three phases. Phase 1 produced independent analyses from each perspective, developed without knowledge of the others. Phase 2 produced structured cross-critiques: each perspective identified specific logical weaknesses, unsupported assumptions, and counterexamples in the others, with explicit severity assessments distinguishing fatal objections from significant-but-survivable ones. Phase

3 produced post-critique revisions: each perspective acknowledged valid objections and reformulated compromised claims while maintaining positions that survived critique. A human researcher at Level 3 evaluated all outputs throughout the process, assessed which critiques were logically sound and which were rhetorical artifacts of the models' training distributions, and selected findings for integration into the paper.

6.2 Principal Findings

The adversarial process generated three categories of output that are incorporated into this paper.

The first category consists of arguments that survived intact. The constitutional chain argument, as I have consistently articulated it in prior conference presentations, survived every challenge when correctly framed. The devil's advocate initially argued that Article 19 CN cannot protect 'the AI system's internal processes' because Article 19 protects persons, not objects. This objection, on reflection, attacks a straw version of the argument: what Article 19 protects is the developer's freedom from preventive state intervention during development, not any anthropomorphized right of the AI system. The reformulation, which was the formulation I have maintained in conference contexts, was not defeated. The EU AI Act's failure on the Article 28 CN reasonableness test also survived: the proportionality objection is supported by documented cost data and the absence of comparative outcome evidence, and no perspective produced evidence of demonstrated AI Act benefits that would satisfy the test.

The second category consists of arguments that required reformulation. The invocation of the precautionary principle as a foundation for preventive AI regulation, which appeared in the preventive risk analyst perspective, was correctly identified as lacking constitutional grounding in Argentine law outside the environmental context of Article 41 CN. The revised formulation grounds the case for limited ex ante prohibitions in Article 75, paragraph 22 CN (international human rights treaties) and Article 75, paragraph 23 CN (the constitutional mandate to promote genuine equality of opportunity). This produces a stronger foundation for the prohibitions proposed in Section 5.4, because treaty obligations at constitutional hierarchy generate mandatory state action, not merely permissive regulatory authority.

The third category consists of normative contributions: specific proposals that emerged from the synthesis and were not present in the paper's prior conference version. These are the three targeted CCyCN amendments (proposed Articles 1735 bis, 1745 bis, and 1721 bis), the proposed Article 12 bis of Law 25.326, and the Article 12 bis disclosure obligation as the fourth normative component.

Council Finding	Section	Contribution to Paper
Art. 19 CN argument survives when developer, not system, is the rights-bearer; 'intimacy of AI' framing is unnecessary	4.1	Clarified formulation of constitutional anchor; consistent with prior conference versions
Precautionary principle lacks constitutional grounding in	5.4	Replaced with Art. 75(22) and 75(23) CN as foundation for prohibitions; stronger constitutional basis

Argentine law outside Art. 41 CN context		
Causal opacity in complex neural networks creates practical barrier to CCyCN claims under current Art. 1726	5.2	Proposed Art. 1735 bis: burden shift in algorithmic discrimination claims
Collective damage standing unavailable to individuals under current Art. 1745 CCyCN	5.2	Proposed Art. 1745 bis: standing for diffuse AI harms by individuals, associations, Ombudsman
Guardianship attribution in multi-stakeholder deployments generates prolonged threshold litigation	5.2	Proposed Art. 1721 bis: deployer as concurrent objective co-responsible with developer
Minimal disclosure obligation (fact of automation) achieves informational objective at low cost	5.3	Proposed Art. 12 bis, Law 25.326
Ex ante prohibitions based on treaty obligations are constitutionally valid but reproduce the Art. 19 CN problem the paper identifies in the AI Act	5.4	Argument not incorporated: ex ante prohibitions are structurally inconsistent with the Art. 19 CN externalization principle that anchors the paper

Table 2. Normative contributions from the LLM Council adversarial validation to this paper.

6.3 The Metaparadox

The LLM Council experiment generates a structural paradox that is worth examining with precision, because it does more than validate the paper's arguments: it illustrates the paper's central theoretical claim through the very method used to test it.

The systems used to evaluate arguments about AI regulation are themselves the type of systems that regulation seeks to govern. A large language model producing constitutional analysis, strategic critique, and institutional design proposals operates at Level 0 in Dennett's taxonomy, it produces text that looks like high-order reasoning without having the cognitive architecture that high-order reasoning requires. Its outputs emerge from matrix operations over learned weights, not from beliefs about beliefs about obligations.

This generates two simultaneous observations. First, the models cannot hold genuine preferences about AI regulation, including regulation of systems like themselves. Whatever apparent preference emerged from the Council process, a tilt toward the constitutional libertarian position, or toward the preventive risk analyst position, or toward the Hybrid Minimal synthesis, was a statistical artifact of training data distribution, not a normative judgment. The models do not know that they were evaluating AI regulation. They do not know what AI regulation is. They process tokens and output tokens according to learned probability distributions.

Second, the models cannot be fully neutral with respect to arguments that affect their creators. The training choices made by their developers, the reinforcement learning from human feedback preferences encoded in their weights, the selection of training data, all of these are Level 3 decisions made by persons with genuine intentionality, and those decisions shaped the probability distributions from which the models' outputs emerge. A model trained primarily on text produced by researchers skeptical of EU-style regulation will produce different outputs in a Council about AI governance than a model trained on text produced by researchers who contributed to the AI Act's development. The Level 0 system carries, without understanding it, the imprints of Level 3 decisions made by those who built and trained it.

What the experiment demonstrated in practice is something more modest and more useful than normative AI judgment: large language model systems are efficient instruments for generating structured adversarial arguments at scale, when a Level 3 researcher evaluates, selects, and synthesizes their outputs. The three CCyCN amendments in Section 5.2, the Article 12 bis of Law 25.326, and Article 12 bis emerged from a process in which models produced the argumentative raw material and a human researcher assessed its constitutional validity, legal coherence, and normative significance. The models produced what they were useful for. The researcher did what only a researcher can do.

This is precisely the functional architecture the paper assigns to AI systems in the regulatory framework it proposes. AI systems are not autonomous agents bearing obligations: they are instruments whose outputs require Level 3 human decision-makers to take responsibility for deployment. The Council experiment did not merely test the paper's arguments. Operating as it did, with models producing and a human synthesizing, it demonstrated them.

7. Limitations and Research Agenda

Three limitations require explicit acknowledgment, both for intellectual honesty and because each points toward a distinct research direction.

The first is **empirical calibration**. As of February 2026, Argentine courts have produced no published decisions applying Articles 1757-1758 CCyCN to damage caused by AI systems. Pre-AI algorithmic systems, automated credit-scoring models operated by Argentine financial institutions since the mid-2000s, automated tax compliance systems operated by the Federal Administration of Public Revenues, offer a proximate empirical testing ground for the CCyCN framework, and that research is ongoing. The claim that Articles 1757-1758 will produce the predicted incentive effects is a theoretical prediction, not a verified result. The paper is transparent about this: the regulatory model proposed is a hypothesis about how existing legal tools will function applied to a new class of cases, and that hypothesis remains to be tested by judicial practice.

The second limitation is **transjurisdictional incompatibility**. Argentine technology firms export software and digital services and operate in European markets. A national legal model that diverges substantially from the AI Act creates dual-compliance costs for those firms: they must satisfy the AI Act to access the EU market and they must satisfy

Argentine law for domestic operations. The ESS model does not incorporate this transjurisdictional dimension, which merits separate treatment. A plausible solution is a bilateral mutual recognition agreement between Argentina and the EU establishing that compliance with the Argentine Hybrid Minimal model satisfies the AI Act's requirements for medium-risk applications, analogous to the mutual recognition framework in pharmaceutical regulation. This is a research direction, not a current result.

The third limitation is **institutional capacity**. The model assumes Argentine courts can determine whether an AI activity constitutes a 'risky activity' within the meaning of Article 1757 without additional technical guidance. The variation in judicial expertise across the Argentine federal court system, between Buenos Aires commercial courts with significant financial technology exposure and courts in provinces with minimal AI litigation experience, will produce inconsistent outcomes during the jurisprudential consolidation period. The proposed Article 1721 bis reduces this problem by providing explicit legislative categorization of AI deployment as a risky activity. A complementary AAIP resolution establishing operational criteria for what constitutes a 'high-stakes automated decision' in each sector covered by Law 25.326 would reduce residual uncertainty further and is within the AAIP's existing regulatory authority.

The **research agenda** may include: a systematic analysis of the first AI Act enforcement cases in European courts, expected from mid-2026 onward, to measure the frequency and form of cosmetic compliance predicted by the ML-EGT model; a comparative analysis of the Argentine CCyCN model against Brazil's Lei Geral de Proteção de Dados and the civil liability framework of the Código de Defesa do Consumidor Article 14, and against Chile's Ley 21.719 on personal data protection in force since January 2026; development of a Constitutional Compatibility Index for AI Regulation that operationalizes the Article 28 CN reasonableness test as a transferable analytical tool for assessing AI legislation in other civil law jurisdictions; and empirical verification of the model's incentive predictions through analysis of compliance behavior in Argentine firms deploying algorithmic systems in consumer finance and employment contexts.

8. Conclusion

The EU AI Act is a politically remarkable and theoretically confused piece of legislation. Remarkable because it achieved consensus among twenty-seven member states with very different technological capacities, legal traditions, and policy priorities. Confused because it assigns regulatory obligations to systems that cannot hold obligations, expects compliance from entities incapable of commitment, and generates an equilibrium, predictable from multi-level evolutionary game theory, in which developers capture the regulation for competitive advantage, deployers use certification as liability cover, and affected parties receive declarative rights without actionable enforcement pathways.

Argentina did not make deliberate AI policy choices that led to its current legal position. It arrived here through three historically independent decisions: a constitution drafted in 1853 that made the externalization of harm a structural prerequisite for state intervention; the 1994 constitutional reform that granted treaty-level human rights instruments constitutional hierarchy; and a 2015 civil code that made objective liability for created risk the operative standard for dangerous activities. The accidental result is a legal architecture that resolves the intentionality problem the AI Act cannot articulate, and does so without invoking any fiction about the agency of computational systems.

The constitutional chain that anchors this architecture runs through three provisions. Article 19 CN establishes that the state may not intervene in private activity that has produced no externalized harm to third parties, and AI development, before a harmful output reaches a person, is categorically such activity. Article 14 CN establishes that lawful industry is a constitutionally recognized right that the state must affirmatively respect. Article 28 CN establishes that any restriction on that right must satisfy the four-element reasonableness test: legitimate objective, suitability, proportionality, and non-nullification. The EU AI Act fails the last two. The Hybrid Minimal model proposed in Section 5 passes all four.

Cass Sunstein, writing from his experience as the United States' chief regulatory reviewer under the Obama administration, describes the discipline required of regulatory agencies that are serious about protecting liberty: knowing when not to regulate is as important as knowing how to regulate. That discipline is built into the Argentine constitutional architecture as a justiciable constraint rather than an administrative preference. An agency that regulates AI development before harm is externalized acts unconstitutionally. A legislature that imposes disproportionate compliance costs acts unconstitutionally. The discipline is not optional.

The LLM Council adversarial validation described in Section 6 and detailed in Annex A confirmed the structural soundness of this position while contributing specific normative improvements: the three targeted CCyCN amendments that close practical procedural gaps without expanding state intervention beyond constitutional limits, and the minimal disclosure obligation that activates affected parties's claim rights without requiring trade secret disclosure. The experiment also demonstrated, in its own operation, the paper's central claim: AI systems function as useful instruments of argument generation when a human researcher evaluates and synthesizes their outputs, not as autonomous agents bearing normative obligations. The correct response to that functional role is to hold the human responsible for the output, not to hold the instrument accountable for its own behavior.

The paper's central hypothesis is falsifiable: if Argentine courts applying Articles 1757-1758 CCyCN to AI-caused harms produce consistent outcomes that measurably alter developer and deployer behavior, reducing discriminatory outputs and increasing genuine safety investment without generating the documentation-compliance theater of the AI Act, the hypothesis is confirmed. If they do not, the failure will be institutional rather than normative, a problem of judicial capacity for complex technology litigation that has different solutions than changing the substantive framework.

The most sophisticated AI regulation may be, paradoxically, the one that does not mention artificial intelligence in a single article.

References

Argentine National Constitution (1853, reformed 1994). Articles 14, 19, 28, 41, 75(22), 75(23).

Argentine Civil and Commercial Code (Law 26.994, 2014, in force August 2015). Articles 1721, 1722, 1726, 1735, 1745, 1753, 1757, 1758.

Law 25.326 on Personal Data Protection (Argentina, 2000). Articles 11, 12, 25.

Law 24.240 on Consumer Defense (Argentina, 1993, amended 2008).

Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). Official Journal of the European Union, L 2024/1689.

Argentine Supreme Court. Bazterrica, Gustavo Mario s/tenencia de estupefacientes (1986). Fallos 308:1392.

Argentine Supreme Court. Arriola, Sebastián y otros s/causa No. 9080 (2009). Fallos 332:1963.

Argentine Supreme Court. Inchauspe Hermanos c/Junta Nacional de Carnes (1944). Fallos 199:483.

American Convention on Human Rights (Pact of San José, 1969). Incorporated at constitutional hierarchy by Art. 75(22) CN.

International Covenant on Civil and Political Rights (1966). Incorporated at constitutional hierarchy by Art. 75(22) CN.

Axelrod, R. (1984). The Evolution of Cooperation. Basic Books.

Dawkins, R. (1982). The Extended Phenotype: The Long Reach of the Gene. Oxford University Press.

Dennett, D. (1987). The Intentional Stance. MIT Press.

EU Transparency Register. (2024). Lobbying activities and expenditures: Artificial intelligence sector, January–December 2024.

European Commission. (2024, October). Guidelines for the implementation of Regulation (EU) 2024/1689 on artificial intelligence. Brussels.

Lerer, I. A. (2025a). Asymmetric Intentionality Theory: How the Dennett-Nash Gap Explains the Failure of Corporate Criminal Liability in AI Governance. Zenodo Working Paper.

Lerer, I. A. (2025b). Game Theory's Hidden Assumption: Intentionality Homogeneity and the Failure of Corporate Criminal Liability. Zenodo Working Paper.

Lerer, I. A. (2025). Law as Extended Phenotype: Toward an Evolutionary Theory of Legal Systems. Zenodo Working Paper.

Lorenzetti, R. (2015). Código Civil y Comercial de la Nación comentado. Rubinzal-Culzoni.

Maynard Smith, J. (1982). Evolution and the Theory of Games. Cambridge University Press.

Nash, J. (1951). Non-cooperative games. Annals of Mathematics, 54(2), 286–295.

People's Republic of China, State Council. (2017). New Generation Artificial Intelligence Development Plan (新一代人工智能发展规划). State Council Document No. 35.

Pinker, S. (2025). When Everyone Knows That Everyone Knows. Scribner.

Pizarro, R. D. & Vallespinos, C. G. (2016). Instituciones de Derecho Privado, Obligaciones (5th ed.). Hammurabi.

ProPublica. (2016, May 23). Machine bias: There's software used across the country to predict future criminals. And it's biased against blacks. Angwin, J., Larson, J., Mattu, S. & Kirchner, L.

Russell, S. & Norvig, P. (2020). Artificial Intelligence: A Modern Approach (4th ed.). Pearson.

Sánchez Viamonte, C. (1929). El habeas corpus: la libertad y su garantía. La Facultad.

Stanford HAI. (2024). The cost of AI regulation: An economic impact analysis of Regulation (EU) 2024/1689. Stanford University Human-Centered Artificial Intelligence Institute.

Sunstein, C. R. (2026). Separation of Powers: How to Preserve Liberty in Troubled Times. MIT Press.

United Kingdom Department for Science, Innovation and Technology. (2023). A pro-innovation approach to AI regulation. London: HMSO.

Veale, M. & Borgesius, F. Z. (2021). Demystifying the draft EU artificial intelligence regulation. Computerrecht / Computer Law Review International, 1, 97–112.

Zhu, W., Yu, L., Yao, H.-R., Tang, Z., & Yue, K. (2026). Task-Aware LLM Council with Adaptive Decision Pathways for Decision Support. arXiv:2601.22662.

Zhang, J., Xiong, C., & Wu, C.-S. (2026). Agentic Confidence Calibration. Salesforce AI Research. arXiv:2601.15778v1.

Annex A: LLM Council Adversarial Validation, Complete Record

The following record details the structure, methodology, and principal findings of the adversarial validation experiment described in Section 6. The specific implementation architecture, including the prompt engineering and session management protocols, is proprietary; this account describes the methodological design at the level of principles sufficient for replication by other researchers using their own implementations.

A.1 Methodological Design: Perspectives and Mandates

Four independent analytical perspectives were assigned explicit mandates before the process began. Each operated without knowledge of the other perspectives' outputs during Phase 1.

The Constitutional Libertarian perspective was mandated to defend the argument that Articles 19, 14, and 28 CN, combined with the CCyCN objective liability framework, are constitutionally and normatively sufficient to govern AI risks in Argentina, and that preventive AI regulation in the style of the EU AI Act is constitutionally suspect under the Argentine reasonableness test. This perspective represented the position that the paper's author maintains in published work and conference presentations.

The Preventive Risk Analyst perspective was mandated to develop the strongest possible case that the current CCyCN framework is practically insufficient for the specific harm profile generated by AI systems, including algorithmic discrimination, diffuse harm to non-identifiable collectives, causal opacity in neural networks, and irreversible reputational and civic damage, and that some form of preventive regulation is constitutionally justified and practically necessary.

The Pragmatic Comparativist perspective was mandated to assess the empirical track record of different regulatory approaches to emerging technology, compare the outcomes of the EU's ex ante licensing approach with the US's sectoral ex post approach and the UK's pro-innovation approach, and identify what evidence exists about which design produces better outcomes for affected parties.

The Devil's Advocate perspective was given an unconditional mandate: to identify the specific logical weaknesses, unsupported assumptions, category errors, and counterexamples in each of the other three perspectives, without commitment to any substantive position. The only criterion for the devil's advocate was argumentative rigor: was the objection logically sound, was the counterexample genuinely contrary, was the unsupported assumption actually unsupported?

A.2 Phase 1: Independent Analyses

The Constitutional Libertarian perspective produced three core arguments. First, the externalization principle of Article 19 CN protects the developer's activity during the pre-harm development phase from preventive state intervention, because that activity, while it remains within the private sphere and has produced no output affecting third parties, is categorically within the protected domain. Second, Articles 1757-1758 CCyCN constitute an applicable, open-system framework that subsumes AI deployment under the 'risky activity' standard without analogical strain, given that the category was explicitly designed for new technological contexts. Third, regulating the internal architecture of AI systems before harm occurs, requiring bias audits, documentation of training data, certification of model behavior, constitutes a restriction on the developer's constitutionally recognized freedom that must satisfy the reasonableness test of Article 28, which the EU AI Act fails on both proportionality and non-nullification grounds. Initial confidence assessments: 0.75 for Argument 1, 0.80 for Argument 2, 0.85 for Argument 3.

The Preventive Risk Analyst perspective produced four core arguments. First, the CCyCN's causal attribution framework requires proof of a causal link between the system's operation and the specific harm to a specific person, a standard that neural networks with hundreds of billions of parameters structurally frustrate because the decision process is not interpretable at the level of individual outcomes. Second, the CCyCN's individual-plaintiff standing framework is inadequate for diffuse algorithmic harm affecting large populations, none of whom has individually compensable damage sufficient to justify litigation costs. Third, some AI systems, social scoring, subliminal manipulation, emotional surveillance, violate human dignity per se in ways that do not require individualized harm to trigger constitutional protection, and the state has affirmative obligations under Article 75(22) CN to prevent such violations. Fourth, statistical algorithmic discrimination is not analogous to a single defective product causing a single accident; it is a population-level harm that the ex post litigation model is structurally unable to redress. Initial confidence assessments: 0.90 for Argument 1, 0.85 for Argument 2, 0.75 for Argument 3 (initially founded on precautionary principle), 0.90 for Argument 4.

The Pragmatic Comparativist perspective identified the EU AI Act's documented costs (Stanford HAI, 2024), its absence of outcome data (regulation too recent), the US sectoral approach's failure to prevent documented harms (COMPAS, Amazon Rekognition Gender Shades study), and the UK pro-innovation approach's early evidence of higher AI startup investment relative to EU peers after the divergence from the AI Act in 2023. The comparativist's central finding was that neither full ex ante licensing (EU) nor pure ex post liability (US) produces a satisfactory outcome, and that the optimal design is some hybrid that preserves market entry while generating genuine safety incentives. Initial confidence: 0.78.

The Devil's Advocate identified six objections across the three perspectives, assessing four as 'significant' and two as 'fatal' to the arguments as initially formulated.

A.3 Phase 2: Cross-Critiques

The six critical objections and their assessed severity:

Objection 1 (Devil's Advocate → Constitutional Libertarian, Argument 1): The claim that Article 19 CN protects 'the AI system's internal processes' is a category error. Article 19 protects persons, not objects. A tostadora does not have intimacy. Characterizing the argument in terms of 'the AI system's private sphere' anthropomorphizes the technology in a way that Argentine constitutional doctrine does not support. Severity: FATAL to the specific formulation presented, though not necessarily to the underlying conclusion.

Objection 2 (Devil's Advocate → Constitutional Libertarian, Argument 2): The claim that Article 1757 CCyCN applies to AI systems is presented as if it were established jurisprudence, but no Argentine court has applied this provision to any AI system. The argument is analogical prediction, not positive law. Severity: SIGNIFICANT. The absence of precedent is not a normative lacuna, but it is a practical uncertainty that undermines the 'sufficient' claim.

Objection 3 (Preventive Risk Analyst → Constitutional Libertarian, Argument 2): More specifically, the Article 1757 framework requires identifying the 'custodian' of the thing or activity. In a multi-stakeholder AI deployment, model developed by Firm A, fine-tuned by Firm B, deployed on cloud infrastructure of Firm C, used by Firm D to make decisions affecting Person E, who is the custodian? The CCyCN does not answer this for distributed algorithmic systems. Severity: SIGNIFICANT.

Objection 4 (Devil's Advocate → Preventive Risk Analyst, Argument 3): The precautionary principle has no constitutional grounding in Argentine law outside the environmental context, where Article 41 CN provides an explicit textual basis. Extending it analogically to AI requires an argument that CSJN precedent does not support, because Mendoza (2008), the leading precautionary principle case, was grounded in Article 41, not in any general constitutional provision. Severity: FATAL to the use of precautionary principle as foundation; the underlying policy objective (preventing systemic harm) survives with a different constitutional anchor.

Objection 5 (Comparativist → Preventive Risk Analyst): The opacidad argument assumes that neural network opacity is a permanent architectural feature that justifies permanent legislative accommodation. But LIME, SHAP, and attention mechanism visualization have materially improved the interpretability of even large models since 2020. If opacity is transitional, legislation designed around it will be over-inclusive within three to five years. Severity: SIGNIFICANT.

Objection 6 (Devil's Advocate → Comparativist): The 'hybrid' proposal as initially formulated, 'something between EU and US', is analytically empty. It does not specify which elements of the EU approach to retain, which US elements to incorporate, at what thresholds, enforced by which institution. Hybrid proposals that resist specification are not proposals; they are evasions of the hard design problems. Severity: SIGNIFICANT; requires concrete specification to be actionable.

A.4 Phase 3: Post-Critique Revisions and Synthesis

Each perspective revised its position in light of the critiques it received. The revisions that produced contributions to this paper:

Constitutional Libertarian revision (responding to Objection 1): The Article 19 argument does not require anthropomorphizing the AI system. The correct formulation, which is consistent with how the argument has been presented in conference contexts, is that Article 19 protects the developer's freedom from preventive state intervention during the development phase, when the activity has produced no externalized harm. The AI system is the medium through which the developer exercises that protected activity. It is not the rights-bearer. This reformulation is not a concession; it is a precision. Confidence on the reformulated Argument 1: 0.85 (increased from 0.75 because the reformulation is stronger than the original).

Constitutional Libertarian revision (responding to Objections 2 and 3): The absence of precedent and the guardianship attribution uncertainty in multi-stakeholder deployments are acknowledged as practical gaps requiring legislative clarification rather than judicial development. Proposed Articles 1721 bis and 1735 bis emerge as the normative response: legislative specification of the deployer as concurrent objectively liable party, and legislative specification of the burden shift in discrimination claims. These amendments address the gaps without expanding state intervention beyond what the constitutional chain already permits. Confidence on revised Argument 2: 0.78 (slightly decreased from 0.80 to reflect honest acknowledgment of uncertainty during transition period, but increased relative to uncorrected formulation because amendments address the identified weaknesses).

Preventive Risk Analyst revision (responding to Objection 4): The precautionary principle is abandoned as a constitutional foundation for preventive regulation in the Argentine context. The revised foundation for ex ante prohibitions is Article 75, paragraph 22 CN (international human rights treaties at constitutional hierarchy) and Article 75, paragraph 23 CN (the constitutional mandate to promote genuine equality of opportunity). This is a stronger foundation because it generates mandatory state obligations, not merely permissive authority, for the specific categories of systems that violate treaty-protected rights per se. Confidence on revised Argument 3: 0.88 (increased from 0.75 because new constitutional basis is more robust than precautionary principle).

Comparativist revision (responding to Objection 6): The hybrid proposal is specified in concrete normative terms, the four-component Hybrid Minimal model ultimately incorporated in Section 5 of this paper. The specification eliminates the 'empty middle ground' critique and demonstrates that the hybrid is not a compromise of the hard design problems but a resolution of them through the specific mechanisms proposed. Confidence on revised Argument 3.3: 0.82 (increased from 0.75).

The synthesis that emerged from Phase 3 aligned all four perspectives on three points: the CCyCN framework is applicable to AI but requires targeted amendments to address specific procedural gaps; the EU AI Act is disproportionate under any Argentine constitutional analysis; and a limited category of AI applications justifies ex ante prohibition on the basis of treaty-level human rights obligations rather than general precautionary reasoning. These three points of convergence produced the specific proposals in Sections 5.2, 5.3, and 5.4 of this paper.

A.5 Argument Survival Matrix

Argument ID	Perspective	Initial Claim	Status Post-Critique	Final Confidence
C-1	Constitutional Lib.	Art. 19 protects 'AI internal processes'	REFORMULATED: Art. 19 protects developer's	0.85

			freedom from preventive intervention in pre-harm phase	
C-2	Constitutional Lib.	CCyCN Art. 1757 applicable to AI (no precedent)	SURVIVED with acknowledged gap: amendments C-2a, C-2b, C-2c proposed to address uncertainty	0.78
C-3	Constitutional Lib.	EU AI Act fails Art. 28 CN reasonableness test	SURVIVED: proportionality failure is documented (cost data) and non-nullification failure is structural	0.87
R-1	Preventive Risk	Causal opacity bars CCyCN claims	SURVIVED: opacity is real though potentially transitional; Art. 1735 bis is proportionate response	0.82
R-2	Preventive Risk	Diffuse harm has no CCyCN remedy	SURVIVED: Art. 1745 CCyCN standing is inadequate; Art. 1745 bis is necessary amendment	0.85
R-3	Preventive Risk	Precautionary principle justifies ex ante prohibition	REFUTED: no constitutional basis outside Art. 41 CN; REPLACED by Art. 75(22) and 75(23) CN foundation	0.88 (revised basis)
R-4	Preventive Risk	AI is product, not agent; regulation of design is valid	SURVIVED: consistent with Art. 14 CN (industry freedom) analysis, regulation of product design is distinct from regulation of agents	0.90
P-1	Comparativist	EU AI Act has documented perverse effects	SURVIVED: cost data is solid; brain drain evidence supports	0.80
P-2	Comparativist	US ex post approach is also insufficient	SURVIVED: COMPAS and similar cases are empirically robust counterexamples	0.85
P-3	Comparativist	Hybrid model is optimal (initially vague)	SURVIVED and STRENGTHENED after specification into four-component Hybrid Minimal model	0.82

Table A.1. Argument survival matrix from the LLM Council adversarial validation process.

Annex B: Proposed Normative Texts in Spanish (CCyCN and Law 25.326 Amendments)

The following texts reproduce the CCyCN and Law 25.326 normative proposals developed in Sections 5.2 and 5.3 in Spanish, drafted in conformity with the lexicon, terminology, and structural conventions of the Argentine Civil and Commercial Code (Law 26.994) and Law 25.326 on Personal Data Protection. Accented characters and special typography follow the conventions of Argentine legal drafting practice.

B.1. Modificaciones propuestas al Codigo Civil y Comercial de la Nacion

Articulo 1721 bis. Actividad riesgosa por sistemas automatizados.

La utilizacion de sistemas automatizados de toma de decisiones en contextos que afecten derechos reconocidos en este Codigo constituye una actividad riesgosa en los terminos del articulo 1757, sin perjuicio de la responsabilidad objetiva concurrente del desarrollador del sistema en su caracter de creador del riesgo. Se entiende por sistema automatizado de toma de decisiones todo sistema computacional que genere resultados -predicciones, recomendaciones, clasificaciones o decisiones- que incidan sobre intereses juridicamente protegidos de personas humanas o juridicas, con independencia del grado de intervencion humana en la produccion del resultado final.

Articulo 1735 bis. Carga de la prueba en supuestos de discriminacion algoritmica.

En los casos en que se alegue que un sistema automatizado de toma de decisiones ha generado un resultado discriminatorio que afecta derechos de igualdad o no discriminacion, la carga de acreditar que el sistema no produjo el resultado discriminatorio invocado recae sobre quien lo haya utilizado en el caso concreto. A tal efecto, el damnificado debe demostrar: (a) que una decision adoptada con asistencia de un sistema automatizado afecto adversamente un interes juridicamente protegido; y (b) que dicho sistema participo en la produccion de ese resultado. Acreditados estos extremos, corresponde a quien utilizo el sistema probar la ausencia de nexo causal discriminatorio. Las reglas sobre carga dinamica de la prueba del articulo 1735 se aplican supletoriamente en todo lo no previsto por el presente articulo.

Articulo 1745 bis. Legitimacion activa en danos causados por sistemas automatizados.

Los derechos reconocidos en este Codigo en materia de danos colectivos o difusos pueden ser invocados, en los supuestos en que el dano alegado sea atribuible a sistemas automatizados de toma de decisiones que afecten a una pluralidad indeterminada o determinable de personas, por: (a) los damnificados individuales; (b) las asociaciones que tengan por objeto institucional la defensa de los derechos afectados, debidamente constituidas y con personeria juridica vigente; y (c) el Defensor del Pueblo de la Nacion. La interposicion de la accion por cualquiera de los legitimados indicados en los incisos (b) y (c) no impide el ejercicio individual de acciones por parte de los damnificados.

B.2. Modificacion propuesta a la Ley 25.326 de Proteccion de los Datos Personales

Articulo 12 bis. Deber de informacion en decisiones asistidas por sistemas automatizados.

Toda persona humana o juridica cuyos intereses juridicamente protegidos resulten afectados por una decision adoptada con asistencia de sistemas automatizados de toma de decisiones debe ser informada, al momento de la comunicacion de dicha decision, que en su produccion intervino un sistema automatizado. Este deber de informacion no comprende la divulgacion de la arquitectura interna del sistema, sus parametros de funcionamiento, los conjuntos de datos utilizados en su entrenamiento ni ningun otro componente que constituya secreto comercial o propiedad intelectual protegida. La obligacion se satisface mediante la inclusion, en la comunicacion de la decision, de una indicacion clara e inequivoca de que en la misma intervino un sistema automatizado. El incumplimiento de este deber habilita al afectado a invocar la inversion de la carga probatoria prevista en el articulo 1735 bis del Codigo Civil y Comercial de la Nacion.