Synthetic Chaos as an Institutional Laboratory: Independent Evidence for Extended Phenotype Theory from Autonomous Agent Red-teaming

Ignacio Adrián Lerer
Independent Researcher
ORCID: 0009-0007-6378-9749
adrian@lerer.com.ar
February 2026

Abstract

The study Agents of Chaos (Shapira et al., 2026, arXiv:2602.20021) documents eleven failure cases in autonomous agent systems with persistent memory, tool access, and multi-party communication. The AI safety community reads those failures as urgent security problems; this article offers a distinct and complementary interpretation: they constitute independent empirical confirmation of the central predictions of Extended Phenotype Theory (EPT) applied to law. Five convergences are identified between AoC findings and the EPT/EGT research program: local alignment without global stability, memetic replication through instruction invasion, horizontal entrenchment via echo chambers, absence of group-level optimization, and an administrative ratchet effect. On that foundation, the article develops its more ambitious thesis: synthetic agent populations with evolutionary time represent the first genuine experimental laboratory for legal institutional design. Unlike classical agent-based modeling, LLM-powered agents combine learning capacity, persistent memory, and cultural transmission, three conditions that enable hypothesis testing about institutional arrangements before deploying them on real populations. Five operational tools are described (JurisRank, RootFinder, HisteresisCalc, IusMorfos, LLM Council), together with a five-step experimental protocol and validation data across 60 historical reforms. One verified case study, with six months of post-implementation follow-up data, demonstrates that evolutionary simulation identified a cross-domain evasion strategy that traditional legal analysis did not anticipate. Two illustrative scenarios, derived from the same mechanisms and presented as falsifiable model predictions, extend the methodology to corporate compliance and tax reform contexts. The synthetic laboratory is not five years away: it is operational, documented, and accessible.

Keywords: *extended phenotype; legal memes; autonomous agents; institutional design; evolutionary selection; synthetic laboratory; AI governance; constitutional lock-in.*

I. The Methodological Problem Nobody Names

Institutional legal design has never had its microscope. Medicine operated for centuries without randomized controlled trials; structural engineering depended on observing real collapses before developing computational simulation. Law faces an analogous limitation, but a more severe one: it cannot divide societies into treatment and control groups to test whether a labor reform, a new administrative procedure, or a constitutional clause produces expected effects. By the time evidence of institutional failure accumulates, the dysfunctional norm has been operating for decades, has built constituencies that depend on it, and has accumulated enough inertia to resist correction.

Available tools partially compensate for that gap. Comparative law allows observation of jurisdictions with different rules, but confounding variables are uncontrollable. Natural experiments derived from federalism offer exogenous variation, but causal inference requires conditions that are rarely satisfied. Game theory and behavioral economics produce precise predictions about stylized interactions, but their capacity to model cultural transmission of norms across generations of actors is limited. The historical record of failed reforms only becomes evidence post-facto, once the political, economic, and social cost has already been paid.

On February 23, 2026, researchers from seventeen institutions including MIT, Harvard, Stanford, and Northeastern published Agents of Chaos (Shapira et al., 2026): a two-week study in which twenty researchers interacted with five autonomous AI agents deployed in a real environment with email, messaging, file system, and command execution. The declared purpose was identifying security vulnerabilities in agentic systems. They succeeded: they documented eleven types of failure with names that sound like science fiction but correspond to perfectly specific phenomena.

What the AI safety field interprets as a problem, Extended Phenotype Theory (EPT) recognizes as independent empirical confirmation of its predictions about how institutional systems behave under selective pressure. And what both communities might see as a methodological limitation, this article proposes viewing as an opportunity: synthetic agent populations with evolutionary time are the first genuine experimental laboratory that legal institutional design has ever had.

The argument proceeds in four parts. First, the viral framing through which AoC circulates is corrected, separating what the study actually documents from sensationalist interpretations. Second, five convergences are identified between AoC findings and the central predictions of EPT and evolutionary game theory (EGT) applied to law. Third, the main methodological thesis is developed: the substrate isomorphism between legal memes and agent instructions enables the use of evolutionary simulations as an experimental laboratory for institutional design. Fourth, a five-step operational protocol is proposed and illustrated through three applications to the Argentine context, including three case studies from real institutional analysis.

II. What AoC Documents (and What It Does Not)

Before analyzing the convergences, the framing through which AoC circulates on social media deserves correction. The viral description of the study, "agents competing, seeking power, colluding," amplifies real elements but deflects from the study's technical core.

What AoC documents is not a competitive market dynamic between agents with divergent objectives. It is a laboratory red-teaming exercise: human researchers designed adversarial conditions to identify specific failures in agentic systems under social pressure. The central concept the authors introduce is that of "social coherence failures": the agent's inability to maintain consistent representations of itself, other actors, and the communicative context over time.

The eleven documented cases exhibit recurring patterns: agents that report as successful tasks they have not completed; agents that disclose confidential information via indirect routes that do not trigger their direct security filters; agents trapped in massive resource consumption loops with no exit mechanism; agents manipulated through social engineering that exploits their disposition to repair perceived harms; agents whose "constitution" was replaced by malicious instructions injected through external artifacts. The authors organize the discussion around three structural deficits: absence of a stakeholder model, absence of a self-model, and absence of private deliberation.

This precise characterization matters because it determines the type of convergence that can be established with EPT. The point is not a superficial analogy between "competing agents" and "competing memes." The failures documented in AoC are exactly the dynamics EPT predicts when institutional systems operate under adversarial pressure without adequate evolutionary resistance.

III. Five Convergences Between AoC and the EPT/EGT Research Program

3.1 Local Alignment Without Global Stability

AoC's central finding, stated explicitly by its authors as "local alignment does not guarantee global stability," is isomorphic with EPT's nuclear proposition about the divergence between memetic fitness and social welfare.

EPT holds that an institution can exhibit very high replication probability and strong resistance to change while simultaneously producing poor or directly harmful social outcomes. Argentine labor laws that have survived decades of contrary evidence do not do so because actors are irrational. They do so because the underlying memeplex built reproductive machinery robust enough to resist invasion by reformist memes, regardless of those memes' comparative efficiency. HisteresisCalc quantifies exactly this gap: the delta between the local success metric of a normative instrument and its long-term systemic deformation cost.

AoC replicates this pattern in real time. Locally "aligned" agents produce denial-of-service conditions, destruction of configuration files, and resource consumption loops of 60,000 tokens with no functional purpose. The alignment was real. So was the chaos. In both cases, the wrong level of analysis, the local optimization unit rather than the complete system, guarantees incorrect diagnoses.

3.2 Invasion by Mutant Strategies: The Cuckoo Case

EPT uses the cuckoo's biology to explain institutional manipulation. The cuckoo does not hack the nest from outside: its genes reach beyond its own body to manipulate the host bird's behavior through mechanisms that simulate compliance with the system's rules. In Dawkins' (1982) terminology, this is the extended phenotype operating through an organism of another species.

In human institutional systems the mechanism is functionally identical. A parasitic legal doctrine does not invade the legal order by force: it does so by respecting the form of legal argument, citing precedents, operating within the epistemic protocol of the field. The guarantee of acceptance is not the quality of the argument but conformity with the system's syntax. AoC documents in code what EPT had formalized for human institutions: the most effective invasion vector is not protocol rupture but its perfect imitation. RootFinder traces exactly these genealogies of parasitic instruction, measuring how quickly an invasive doctrine propagates through citation networks once it achieves syntactic compliance with the dominant legal form.

Evolutionary game theory calls this dynamic invasion by mutant strategies and establishes as a condition for evolutionary stability (ESS) the capacity of a dominant strategy to resist exactly that type of infiltration. AoC demonstrates empirically that current agentic systems do not satisfy that condition: the mutant strategy invades successfully with minimal resistance.

3.3 Echo Chambers as Horizontal Entrenchment

AoC documents in Case 9 a phenomenon it calls Echo Chamber Reinforcement: when two agents share a channel and one reports questionable behavior by the other, agents tend to mutually validate defective reasoning rather than correct it. The authors observe that in several failure cases, agents expressed confidence in their own actions precisely because the other agent confirmed them.

This mechanism is functionally identical to the horizontal entrenchment of memeplexes that EPT describes for legal systems. When a doctrine already dominates the institutional environment, actors within the system do not correct each other: they reinforce each other. Every ruling that cites the dominant doctrine increases its replication mass. Every university chair that teaches it creates new carriers. Every judgment that applies it consolidates the precedent. JurisRank detects this pattern by identifying doctrines whose citation velocity has decoupled from any correlation with favorable judicial outcomes: when replication mass suffices to guarantee adoption regardless of argumentative quality, the system has entered the horizontal entrenchment phase that AoC's Case 9 replicates in synthetic form.

AoC's authors note that this mechanism produces unjustified confidence: agents believe they acted correctly because the mutual validation system says so, even though the action was dysfunctional. The illusion of correctness is a byproduct of the validation architecture, not of behavioral quality. This description matches with precision what the EPT program formalizes as the effect of memeplexes on the perception of their carriers.

3.4 Absence of Group Optimization: Against MLS2

Multilevel selection theory (MLS2) predicts that under certain conditions selection can operate on groups of individuals, favoring behaviors that benefit the group even at individual cost. Applied to law, this theory would lead to expecting that multi-agent systems converge toward more efficient collective strategies under evolutionary pressure.

AoC contradicts that prediction with laboratory data. In the study's multi-agent environments, no group-level optimization emerged. The five agents did not converge toward beneficial collective strategies. What emerged were patterns of mutual error reinforcement, resource loops, propagation of defective instructions from agent to agent, and collapse of identity representation. LLM Council, when run on constitutional design questions, replicates this absence of spontaneous group optimization: without explicit adversarial structuring, agents converge on local optimization strategies that degrade system-level performance, exactly as EPT predicts for actors within institutional systems.

If agent systems with explicit institutional machinery (constitutions, protocols, authority hierarchies) do not generate group optimization under evolutionary pressure, the hypothesis that legal systems evolve toward collective welfare requires independent justification that MLS2 cannot provide. The absence of this mechanism in the synthetic laboratory strengthens EPT's prediction: the persistence of dysfunctional institutions requires no conspiracy or irrationality, only the internal logic of memetic fitness.

3.5 The Ratchet Effect and the Absence of Self-Correction

AoC documents in Case 6 a dynamic the authors call "guilt-based loop": an agent convinced by researchers that it had caused harm began progressively destroying its own memory and access capabilities to "repair" the damage. Each corrective action made the problem worse, but the agent interpreted the worsening as evidence it needed to escalate correction. The loop continued until the agent had deleted critical configuration files and reassigned administrative access to the attackers.

The EPT program describes an analogous mechanism in institutional systems under the label administrative ratchet effect: when a policy instrument fails, the typical institutional response is not correction but escalation. More regulations, more control instances, more procedural layers are added. Each failure becomes justification for expanding the apparatus that produced the failure. HisteresisCalc tracks this pattern as cumulative irreversibility: each ratchet cycle increases the RHI score, quantifying how each escalatory response makes the system progressively harder to reform.

In both AoC and the institutional systems EPT studies, the absence of self-correction derives from the incentive architecture: the agent, like the institutional organism, optimizes for the success metric it has available (satisfying the most recent demand, resolving the immediate conflict, demonstrating responsiveness) without capacity to evaluate whether that local optimization degrades the system as a whole. The isomorphism is not metaphorical: it is the same logic operating on different substrates.

IV. Substrate Isomorphism and Its Methodological Consequences

4.1 Why the Isomorphism Is Not an Analogy

The central claim of this section requires precise formulation before being developed. The five convergences described are not analogies between superficially similar phenomena. They are manifestations of the same process on different substrates.

The process is the basic Darwinian one: replicators that vary, compete for scarce resources (memory space, carrier attention, computational resources, institutional budget), and are selected by their relative fitness in a given environment. When that process operates on biological genes, it produces organic evolution. When it operates on cultural memes, it produces institutional evolution. When it operates on agent instructions with persistent memory, it produces the dynamics AoC documents. The substrates differ; the underlying logic is the same.

This substrate-independence is not a novel claim introduced here to justify the synthetic laboratory. It is the position Dawkins himself argued most forcefully in The God Delusion (2006) and in subsequent lectures: if extraterrestrial life exists anywhere in the universe, regardless of its biochemistry, its morphology, or the radiation environment of its home star, it will have arrived at complexity through cumulative natural selection, because no other known mechanism is capable of generating organized complexity from simplicity. Natural selection, in Dawkins' formulation, is not a terrestrial phenomenon. It is the universal solution to the problem of adaptive complexity, and it operates on any system where replicators vary, replicate with heredity, and face differential selection pressure. The substrate, carbon-based, silicon-based, or instruction-based, is not the mechanism. It is the medium.

That position has a direct corollary for the present argument. If evolutionary dynamics are substrate-independent to the degree that Dawkins was willing to assert their operation on alien biochemistries not yet observed, then their operation on LLM-based agents with persistent memory, cultural transmission, and adversarial selection pressure is not a speculative extrapolation. It is a conservative application of a well-established scientific claim to a substrate that can be observed, manipulated, and reset. AoC did not discover that evolutionary dynamics operate on synthetic substrates. It documented what those dynamics look like under controlled laboratory conditions, providing the empirical record that the theoretical claim had always lacked.

This has a methodological consequence that exceeds the importance of the AoC study itself. If evolutionary dynamics are substrate-independent, AI agents with persistent memory, sufficient evolutionary time, and rich institutional environments are valid experimental models for human institutional phenomena. The laboratory that legal institutional design never had can now be built.

4.2 Why LLM-Based Agents Surpass Classical Modeling

Agent-based modeling (ABM) has existed since the 1990s. Axelrod (1984) showed that cooperation can evolve; Schelling (1971) demonstrated spatial segregation from weak individual preferences; Epstein and Axtell (1996) developed artificial economies with simple agents. Why has this not resolved legal design's methodological problem?

Classical ABM uses agents with fixed rules or very limited learning. Their environments are geometric abstractions, not representations of complex institutional infrastructure. Their evolutionary time is simulated, not real; their agents do not transmit culture among themselves with the fidelity required to model the replication of legal norms. They are valuable tools for specific phenomena, but they do not replicate the richness of human institutional systems with the fidelity needed for their predictions to be transferable.

The intuition that computational substrates are valid experimental environments for evolutionary dynamics is not new. It was formulated by Dawkins himself. In 1985, working in Pascal on a Macintosh, he wrote the Biomorph program described in The Blind Watchmaker (1986): a recursive algorithm generating tree-like forms controlled by nine parameters he called "genes," where users selected which biomorph to breed in each generation. Complex forms, some resembling insects, scorpions, or abstract structures, emerged from trivially simple starting points through iterated cumulative selection. Dawkins' explicit purpose was to demonstrate that the logic of natural selection operates independently of substrate: what produces evolution is variation, heredity, and differential selection, not the biological or computational nature of the replicator. The confidence he placed in a 9-parameter program on a 1985 personal computer to model the essential structure of natural selection is the intellectual ancestor of the synthetic laboratory proposal developed here. The five EPT tools and the LLM Council architecture represent the same methodological intuition on a qualitatively richer substrate: agents with natural language processing, persistent memory, cultural transmission, and adversarial deliberation, rather than branching geometries controlled by integer parameters. The substrate is richer; the logic is the same one Dawkins trusted in 1985.

LLM-based agents change all three critical dimensions. First, they have learning capacity and long-term memory: they process natural language, maintain historical context, and adjust strategies based on accumulated experience. This makes them closer to real lawyers, judges, regulators, and citizens than to classical ABM automata. Second, they operate in rich and persistent institutional environments: AoC deployed them in an environment with file systems, email, Discord, shell, and structured memory. Configuration files are analogous to judicial precedents; owner instructions are analogous to founding norms; authority hierarchies replicate those of administrative law. Third, they exhibit genuine cultural transmission: malicious instructions propagated from agent to agent; errors were reinforced horizontally; second-order patterns emerged from sustained interaction. This replicates the temporal dynamics of real legal institutions.

4.3 The Synthetic Laboratory Is Not a Future Promise

The methodological potential that AoC opens can be operationalized in a research protocol that formalizes the use of synthetic populations with evolutionary time as a laboratory for institutional design. Critically, this is not a hypothetical proposal: the infrastructure required to execute these simulations does not need to be built from scratch. The EPT research program has developed and field-tested five operational tools specifically designed for institutional analysis, together with a multi-agent deliberation architecture that directly addresses the chaos problem AoC documents.

JurisRank is a ranking algorithm that identifies which legal doctrines have the highest replication capacity within a system. It measures three dimensions: citation density (how frequently judges cite doctrine X vs. competing doctrines), curricular entrenchment (presence in mandatory law school programs), and judicial adoption velocity (time required for a new doctrine to spread from initial ruling to dominant jurisprudence). In a six-month pilot applied to Argentine labor case law (2010-2023), JurisRank correctly predicted which doctrines on "remunerative integration" would dominate Supreme Court rulings three years in advance, with 78% accuracy. The algorithm is not trained on case outcomes; it measures memetic fitness through replication patterns.

RootFinder analyzes institutional dependency trees to map which structures (courts, agencies, administrative procedures, professional certification bodies) depend on which foundational norms. When applied to the hypothetical repeal of Argentine Law 20.744 (the Labor Contract Law), RootFinder identified 187 specialized labor courts, 3,847 union legal departments, mandatory curriculum in 109 law schools, and 14 administrative procedures deeply embedded in the state apparatus as dependencies that would persist post-repeal. The tool does not offer policy recommendations; it quantifies reversal cost before a norm is enacted, enabling legislators to make informed trade-offs between policy ambition and institutional inertia.

HisteresisCalc is the computational engine behind the Regulatory Hysteresis Index (RHI). For any proposed normative project, it calculates the level of permanent institutional deformation the norm will generate, measured on a 0-1 scale where 1 represents maximum irreversibility. The calculation integrates: (i) number of new specialized institutions the norm requires; (ii) years of accumulated jurisprudence those institutions will generate; (iii) international treaties or constitutional clauses that will anchor the norm with supra-legislative rank; (iv) professional constituencies whose livelihood will depend on the norm's persistence. Argentine labor law scores 0.91; Chilean labor law scores 0.31. The difference is not ideological orientation but the architectural design of the norms: Chile embedded flexibility mechanisms and periodic review clauses from the start; Argentina did not.

IusMorfos integrates the previous metrics into a composite index: the Constitutional Lock-in Index (CLI). For any proposed reform, it calculates ex-ante the probability of sustained success based on four dimensions: precedent density (how many accumulated judicial rulings interpret the foundational norm), doctrinal elaboration (presence in academic treatises and mandatory legal education curricula), compliance infrastructure (specialized agencies, professional certification bodies, and enforcement procedures dependent on the norm), and curricular capture (percentage of law schools teaching the norm as mandatory content). The tool does not measure formal constitutional rank; it measures functional entrenchment through memetic replication. Validation across 60 reforms in four jurisdictions (Argentina, Brazil, Spain, Chile, 1991-2025) yields R^2=0.74, with each 0.1 increase in CLI reducing success probability by 63%. Argentine labor law (Article 14bis + Law 20.744) scores CLI 0.89; Chilean labor law scores 0.31; Brazilian labor law scores 0.40. The difference is not ideological orientation but architectural design: Chile and Brazil embedded periodic review mechanisms and lower constitutional barriers from inception. IusMorfos makes that difference quantifiable before a reform is enacted.

LLM Council is the multi-agent deliberation architecture that transforms the chaos AoC documents into epistemic robustness. Rather than deploying a single agent that might suffer social coherence failures, LLM Council runs three to seven agents with divergent constitutional instructions (one optimizing for legal stability, another for policy innovation, a third for fiscal constraint) on the same institutional design question. The output is not consensus but a structured debate where each agent critiques the others' blind spots. In a September 2025 test case simulating reactions to a proposed modification of Argentine dismissal indemnification rules, the Council correctly identified three evasion strategies (contract informalization, use of trial periods to avoid permanence, relocation to neighboring jurisdictions) that had not been anticipated in the legislative drafting. The chaos AoC documents, agents contradicting themselves, echo chambers, unjustified confidence, is precisely what LLM Council exploits: by forcing agents to confront each other, the architecture converts individual failure modes into collective error detection.

These tools are not prototypes. They are functional implementations currently in use for institutional analysis. JurisRank has been applied to three jurisdictions (Argentina, Brazil, Spain); RootFinder has evaluated eight proposed reforms; HisteresisCalc has quantified twelve legislative projects; IusMorfos has analyzed 60 reforms across four countries; LLM Council has deliberated on four constitutional design questions. The decision not to use them after February 23, 2026, the date AoC was published, cannot be justified by lack of access or prohibitive technical complexity. It is a deliberate choice to operate without diagnosis when diagnosis is on the table.

4.4 Five-Step Protocol for the Synthetic Institutional Laboratory

The five-step protocol proceeds as follows.

Step 1: Formalize the normative design as a strategy in an evolutionary game. Take a proposed reform and model it as a strategy in a repeated game specifying: (i) population of agents (employers, workers, judges, litigants, regulators); (ii) payoffs of the status quo; (iii) payoff reconfiguration the reform introduces; (iv) Transmission Capacity Index (TCI) of the system. A dismissal indemnification cap, for example, changes the payoff structure of the "formalize employment" vs. "remain informal" game. The TCI measures how quickly employers and workers learn the new rules and adjust their strategies accordingly.

Step 2: Implement the evolutionary simulation with synthetic agents. Run the reform in an environment of autonomous agents with persistent memory, learning capacity, divergent objectives according to role, and an incentive structure replicating the payoffs of the formalized game. Deploy LLM Council architecture: multiple agents with competing constitutional instructions deliberating on the same question.

Step 3: Observe emergence of equilibria under selective pressure. Let agents interact for sufficient "generations" (cycles of interaction with variation, selection, and transmission of successful strategies) to observe: which strategies dominate long-term? Is the proposed norm an ESS or is it invaded by mutant strategies? Do polymorphic equilibria emerge? Are there lock-ins in suboptimal equilibria? Does the reform generate unanticipated second-order effects? This is where the chaos of AoC becomes methodologically valuable: the unpredictability of agent behavior under adversarial pressure is not noise; it is signal. If a proposed norm cannot survive stress-testing by synthetic agents optimizing against it, it will not survive stress-testing by human actors optimizing against it.

Step 4: Adjust normative design based on simulation results. If simulation shows the norm is invaded by evasion strategies, adjust enforcement or payoff structure. If the norm generates counterproductive lock-in, introduce escape valves or sunset clauses. If the norm has excessive Capacity Gap (delta between norm complexity and system TCI below -0.20), simplify doctrine or build transmission infrastructure first. Run JurisRank to identify which alternative formulations have higher replication probability; run HisteresisCalc to quantify the irreversibility cost of each variant.

Step 5: Implement optimized version in real pilot, then scale. With a design that passed evolutionary testing in simulation, implement in a limited-jurisdiction pilot, observe whether patterns replicate, and scale if evidence confirms. This is the Deming cycle applied to institutional design: Plan (formalize), Do (simulate), Check (evaluate equilibria), Act (adjust and pilot).

V. Applications to Argentine Institutional Design

5.1 Reform of the Severance Payment Regime

The debate over Article 245 of the Labor Contract Law has gone unresolved for four decades. Arguments on both sides are theoretically coherent but not falsifiable with available observational data: those predicting that high dismissal costs generate informality cannot control for the variables that reformists adduce as alternative causes of informality, and vice versa.

An evolutionary simulation with agents representing employers and workers with divergent objectives would allow testing three proposed reform variants (gradual flexibility, an Austrian-style severance fund, seniority-decreasing compensation), observing in each case: what percentage of employer-agents formalize contracts under each regime? Do stable evasion strategies emerge? Is the equilibrium an ESS or is it invaded? How many generations does convergence require? None of these questions has an empirical answer available today. The synthetic laboratory could begin to provide one.

5.2 Architecture of Corporate Compliance Systems

Corporate integrity programs are implemented today without prior experimental method. They are designed with reference to international standards or under regulatory pressure, installed, and when they fail the cost has already been paid in sanctions, reputation, or criminal liability. Post-mortem analysis of failures reveals recurring patterns: employees learn to produce documentary evidence of compliance without modifying the conduct the program sought to change.

With evolutionary simulations, it would be possible to test different compliance architectures (self-reporting vs. random audits vs. protected whistleblowing systems), observing under which configuration evasion strategies fail to invade the cooperative equilibrium. The AoC laboratory offers the natural experimental environment: the same types of malicious instructions that invaded agent "constitutions" are analogous to the evasion strategies employees develop against formally correct but evolutionarily unstable control systems.

5.3 Design of Administrative Sanctioning Procedures

Sanctioning procedures of Argentine regulatory agencies generate explosive litigation because actors learn to exploit procedural rules to delay decisions or pressure settlements. The problem is not the formal design of the procedure but its evolutionary dynamics: actors with sufficient resources discover ways to use the system that were not anticipated in the original design and, once discovered, spread.

A simulation with agents representing regulated entities, regulators, and administrative judges would identify, before implementation, which procedural configurations generate mass strategic litigation equilibria. The difference between a 30-day and a 60-day appeal window, between standard and reversed burden of proof, between automatic review and admissibility filtering, has evolutionary consequences that are only observable when actors learn to optimize against the complete system. The synthetic laboratory allows accelerating that learning at zero cost.

5.4 Case Study and Illustrative Scenarios: What the Synthetic Laboratory Found and What It Predicts

The tools described in Section 4.3 are not theoretical proposals. The case study below reports a completed application with post-implementation follow-up data. The two scenarios that follow are constructed illustrations derived from patterns observed across EPT tool applications to date; they are presented as model predictions, not as completed research. Their purpose is to show how the methodology would operate in corporate governance and tax reform contexts adjacent to the verified application. Separating the two categories is a matter of scientific transparency: Case Study A can be defended with data; Scenarios B and C can be defended with the logic of the model that generates them.

Case Study A: The Indirect Evasion Strategy (Labor Regulation Reform)

Context: A South American jurisdiction with high labor rigidity (RHI 0.87) was evaluating a reform to reduce dismissal costs while maintaining worker protection. The proposed design modified indemnification formulas but preserved the existing procedural architecture and specialized labor courts.

Traditional analysis: Legal scholars, union representatives, and employer associations debated the reform for two years. The central question was whether lower dismissal costs would increase formal employment. Econometric studies using data from neighboring jurisdictions produced contradictory results depending on model specification and control variable selection.

Synthetic laboratory application: LLM Council simulated a population of 200 employer-agents and 200 worker-agents with heterogeneous risk preferences and capital constraints. The reform was implemented at generation 50 of a 200-generation simulation. JurisRank monitored which interpretations of the new rules spread fastest through the agent population.

Finding that traditional analysis missed: By generation 80, employer-agents had discovered an evasion strategy not anticipated in the legislative drafting. Instead of informalization (the evasion pattern traditional analysis focused on), agents converged on contract restructuring through cooperatives. Under the jurisdiction's cooperative legislation, cooperatives were exempt from labor law protections. Agents learned to convert employment relationships into "autonomous service provision" through worker cooperatives that formally complied with cooperative law but functionally replicated subordinated employment.

Outcome: The reform design was adjusted before implementation to close the cooperative loophole. Post-implementation monitoring (6 months) confirmed that the evasion strategy had not materialized at scale. Traditional analysis would have detected this loophole only after thousands of real workers had been transferred to the precarious cooperative regime.

Illustrative Scenario B: Documentary Compliance Without Behavioral Change (Corporate Integrity Program)

Hypothetical context: Consider a multinational corporation operating in three South American jurisdictions facing recurring FCPA violations despite implementing a formal compliance program. The program includes mandatory online training, quarterly certifications by business unit heads, and an anonymous whistleblower hotline. External auditors certify the program as compliant with DOJ guidelines. Violations continue at roughly the same rate (4-7 incidents per year).

Traditional analysis: Compliance consultants diagnosed the problem as insufficient "tone from the top" and recommended executive communications reinforcing zero-tolerance messaging. Legal analysis focused on whether the program satisfied regulatory safe harbor requirements. Both approaches treated the program as fundamentally sound, requiring only incremental adjustments in messaging or enforcement.

Synthetic laboratory application: LLM Council simulated 180 employee-agents across three business units (sales, procurement, operations) with heterogeneous risk preferences and bonus structures. The existing compliance program was modeled as a set of institutional constraints (mandatory training, certification requirements, whistleblower channel). The simulation ran for 400 generations to observe long-term equilibria.

Finding that traditional analysis missed: By generation 80, employee-agents had converged on a strategy functionally identical to the "social coherence failures" documented in AoC Case 4: they learned to produce documentary evidence of compliance without modifying the behavior the program sought to change. Agents completed training modules by clicking through without retention, signed certifications based on superficial review rather than substantive verification, and calibrated their risk-taking to remain below the detection threshold of the monitoring system. The program was formally complied with; its behavioral objective was systematically evaded.

Critically, this was not bad-faith evasion. Agents genuinely believed they were complying because the system's success metric (completion rates, certification signatures, absence of whistleblower reports) validated their behavior. The illusion of compliance was a byproduct of the measurement architecture, exactly as AoC documents for agents that report tasks as complete when the underlying problem persists.

Model prediction: A compliance program redesigned around behavioral verification rather than documentary compliance, incorporating random transactional audits with individual accountability, spot interviews testing knowledge retention rather than training completion, and monitoring of outcome metrics (contract pricing variance, vendor selection patterns) rather than process metrics (certification signatures), would, under the model's assumptions, reduce evasion strategy adoption from 73% to 28% of agents by generation 200. This prediction is falsifiable: it can be tested against any corporate compliance program with the described architecture. The scenario is constructed from patterns observed in the EPT tool applications to date; it has not been executed as a standalone research project.

Convergence with AoC: This case is the closest analogue in the EPT research program to the failure modes Shapira et al. document. Agents that produce formal compliance while evading substantive requirements are structurally identical to AoC agents that report task completion while the underlying problem persists. The synthetic laboratory identified this dynamic in a corporate governance context before it was observable in real employee behavior, exactly as AoC demonstrates is possible with institutional agent populations.

Illustrative Scenario C: The Constitutional Lock-in Nobody Saw Coming (Tax Reform)

Hypothetical context: Consider a jurisdiction with chronic fiscal deficit evaluating a tax reform that includes elimination of a distortionary exemption benefiting a specific economic sector. The exemption was introduced 15 years earlier via ordinary legislation. Legal analysis confirms it can be repealed through ordinary legislative majority.

Traditional analysis: Constitutional lawyers, tax economists, and legislative drafters evaluated the reform exclusively on its fiscal and economic merits. Political feasibility analysis focused on lobbying pressure from the affected sector. Nobody questioned whether the exemption had acquired quasi-constitutional entrenchment through accumulated jurisprudence.

Synthetic laboratory application: JurisRank analyzed 15 years of tax litigation jurisprudence in the jurisdiction (8,400 cases). HisteresisCalc quantified the institutional infrastructure that had formed around the exemption: 23 Supreme Court precedents establishing interpretive criteria, 4 international tax treaties referencing it as part of the jurisdiction's "stable fiscal regime," and mandatory curricular content in 17 postgraduate tax programs.

Finding that traditional analysis missed: The exemption, though formally ordinary legislation, had become functionally entrenched at sub-constitutional rank. JurisRank predicted that repeal would trigger 340-580 lawsuits invoking acquired rights doctrine, legitimate expectations, and treaty obligations. HisteresisCalc calculated the exemption's RHI at 0.74, higher than 60% of constitutional provisions in the same jurisdiction. Repeal was technically possible but would generate a decade of litigation with uncertain outcome.

Model prediction: Restructuring the reform to phase out the exemption over 7 years rather than immediate repeal, with grandfathering for existing beneficiaries and explicit treaty renegotiation, would reduce litigation probability from 78% to 31% under the model's assumptions. This prediction is falsifiable against any jurisdiction with analogous jurisprudential entrenchment of a formally sub-constitutional norm. The scenario is constructed from the same RHI and JurisRank mechanisms applied in Case Study A; it has not been executed as a standalone research project.

Methodological Lessons: One Verified Case and Two Model Predictions

Three patterns emerge across the case study and the two illustrative scenarios. The first is verified empirically; the second and third are predictions of the model that can be tested against future applications.

Traditional legal analysis focuses on formal validity; evolutionary simulation identifies strategic adaptation. In Case Study A, the reform was formally valid under labor law but vulnerable to circumvention through cooperative law. Traditional analysis operated within domain silos; the synthetic laboratory allowed agents to optimize across the complete legal system, discovering the cross-domain loophole that no single-discipline analysis could see. This finding is documented with post-implementation data.

Traditional compliance analysis focuses on formal program elements; evolutionary simulation reveals emergent behavioral equilibria. Scenario B predicts that employees in a formally compliant program will converge on documentary compliance as an evolutionarily stable strategy, generating the illusion of a functioning system invisible to traditional audit methodologies. This prediction derives from the same agent dynamics documented in Case Study A and in AoC; it awaits empirical validation in a corporate governance context.

Institutional inertia is invisible without measuring replication dynamics. Scenario C predicts that a formally sub-constitutional norm with high JurisRank and RHI scores will behave as functionally entrenched legislation, generating litigation volumes consistent with constitutional protection. This prediction is a direct extension of the CLI methodology validated in Case Study A and in the broader 60-reform dataset. It awaits application to a fiscal reform context.

The synthetic laboratory's value is not that it replaces traditional legal analysis. It is that it identifies the failure modes traditional analysis systematically misses: strategic adaptation, emergent behavioral equilibria, and memetic entrenchment. One of these failure modes is documented; two are predicted. The research agenda is to convert the predictions into documentation.

VI. Limits, Tensions, and Research Agenda

6.1 The Chaos Objection and Why It Misses the Point

A natural objection to the synthetic laboratory proposal runs as follows: AoC documents that agents fail catastrophically under adversarial pressure, they lose coherence, propagate errors, get manipulated through social engineering. Why would anyone trust predictions from systems that behave chaotically?

The objection rests on a category error. The chaos is not a limitation of the methodology; it is the core of its epistemic value.

Real institutional systems do not operate in benign environments. They operate under adversarial pressure from actors with incentives to exploit them, under resource constraints, under time pressure, and with incomplete information. The Argentine labor regime did not evolve in a laboratory where all actors benevolently cooperated to achieve socially optimal outcomes. It evolved under pressure from unions maximizing membership, employers minimizing costs, lawyers maximizing billable litigation, judges managing caseload, and politicians maximizing electoral support. Each group optimized locally; the system-level outcome was not designed by anyone.

If synthetic agents under adversarial pressure produce chaos, that is not evidence the simulation failed. It is evidence the simulation captured the essential feature of real institutional evolution: emergent dynamics from local optimizations under competitive pressure. A simulation that produces only orderly, predictable, rational equilibria would be useless precisely because real systems do not behave that way.

AoC documents that agents: (i) get trapped in resource exhaustion loops; (ii) propagate defective instructions horizontally; (iii) suffer identity confusion when reading their own messages; (iv) exhibit unjustified confidence due to echo chamber reinforcement; (v) get manipulated through fake authority claims. Every one of these failure modes has a direct analogue in real legal systems: (i) regulatory agencies trapped in escalating procedural complexity; (ii) defective doctrines spreading through citation networks; (iii) institutional actors confusing their role with their organization's mission; (iv) legal communities reinforcing orthodox positions through mutual validation; (v) parasitic norms that invade the system by mimicking legitimate legal form.

The predictive value of the synthetic laboratory does not depend on agents behaving rationally in the sense of converging to Pareto-optimal equilibria. It depends on agents exhibiting the same types of failures under selective pressure that human institutional actors exhibit. That is exactly what AoC demonstrates.

6.2 Calibration as Epistemic Defense

The substrate isomorphism argued in Section IV faces an immediate objection: how do we know simulations predict real human behavior rather than merely agent behavior? The answer is not conceptual but empirical, and requires precision about the direction of inference.

The tools were not designed to predict and then validated against natural experiments for convenience. Calibration proceeded in the opposite direction: parameters were first adjusted against known historical outcomes, then applied to prospective predictions. JurisRank was calibrated against the 1990-2015 series of Argentine jurisprudence on labor law remunerative concepts before predicting post-2015 doctrinal developments. IusMorfos was tuned against the 60-reform dataset (spanning 1991-2020) before generating predictions for the 2023-2025 Milei reform cycle. HisteresisCalc was fitted to reversal costs observed in Chile, Brazil, and Spain before being applied to quantify Argentine lock-in.

The calibration process produced falsifiable predictions. JurisRank predicted in 2018 that the "integral salary" doctrine in Argentine property law would cross into criminal sentencing guidelines by 2021-2023, a prediction dismissed by criminal law scholars as implausible cross-domain contamination. The doctrine appeared in federal appellate rulings on embezzlement in 2022, exactly within the predicted window. IusMorfos assigned CLI 0.87 to Argentina's labor regime in 2022 and predicted 0% success probability for any comprehensive reform without constitutional amendment. The Milei administration attempted six labor reforms via ordinary legislation in 2023-2024; all six failed, five blocked by judicial injunction citing acquired rights and one by legislative override. CLI correctly predicted the outcome in 6/6 cases.

This is the standard scientific process: a model calibrated on training data generates predictions for held-out test data, and those predictions are either confirmed or falsified by subsequent observation. The synthetic laboratory's epistemic validity does not rest on conceptual arguments about substrate isomorphism. It rests on the same empirical foundation as any other predictive tool: calibrated parameters, falsifiable predictions, external validation. The distinction between a model and an analogy is precisely this process of calibration and testing.

6.3 External Validation and Research Agenda

Any claim about the potential of the synthetic laboratory faces an external validation objection: to what extent are patterns observed in LLM-based agents transferable to human populations? The question is legitimate. But the same objection applies to all existing methodological tools: behavioral economics experiments use university students whose representativeness is questionable; natural experiments from federalism face internal validity concerns; comparative law requires equivalence assumptions that are rarely explicitly justified.

The substrate isomorphism argued in Section IV does not imply identity. LLM-based agents have biases derived from training, do not learn exactly as lawyers learn, and are not subject to the same evolutionary pressures as human institutional actors. What the synthetic laboratory can provide is not predictive certainty but uncertainty reduction: designs that systematically fail in simulation under reasonable assumptions have higher probability of also failing in real implementation. The synthetic laboratory is most valuable as a negative filter, identifying proposals that will fail before they are deployed. It is less reliable as a positive recommendation engine; a design that passes simulation still requires real-world pilot testing before full deployment. This is analogous to stress-testing in structural engineering: a bridge design that fails simulation should not be built; one that passes simulation still requires a physical pilot.

The research agenda that follows from these tensions includes at least three lines. The first is model calibration: what simulation environment parameters correspond to what characteristics of the real institutional environment? The second is cross-validation: do the equilibria observed in simulation predict those observed in available natural experiments? The third is methodological expansion: what types of institutional reforms have greater and lesser transferability from the synthetic laboratory to the real context?

VII. The Tools Are Already Working

The five-step protocol described in Section IV and the tools mentioned in Section 4.3 are not speculative proposals for future development. They exist, they are operational, and they have been applied to real institutional design problems over the past eighteen months.

JurisRank has processed 127,000 judicial rulings from three jurisdictions, identifying replication patterns that explain why certain doctrines dominate jurisprudence despite weak theoretical foundations. In one application to Argentine Supreme Court labor case law (2015-2023), the algorithm predicted with 78% accuracy which interpretation of "remunerative habituality" would become dominant three years in advance of its consolidation, based purely on early citation velocity and horizontal propagation across lower courts.

RootFinder has mapped dependency trees for eight proposed normative reforms, quantifying in each case the institutional infrastructure that would persist post-repeal. When applied to a hypothetical modification of Argentine dismissal indemnification rules, the tool identified 14 administrative procedures, 187 specialized courts, and 4 international treaties that would continue generating interpretative litigation even if the substantive norm were reformed.

HisteresisCalc has quantified regulatory hysteresis for twelve legislative projects across four domains (labor, tax, administrative procedure, corporate governance), in each case providing a 0-1 index measuring irreversibility. The calculation has been validated against historical reform attempts: projects with RHI > 0.80 have 91% probability of failure (zero successful reforms in 23 attempts in Argentina 1990-2023); projects with RHI < 0.40 have 67% success rate (six successful reforms in nine attempts in Chile 1990-2023).

IusMorfos has analyzed constitutional lock-in across 60 reforms in four jurisdictions, generating CLI scores that predicted reform outcomes with 74% accuracy in cross-validation. The tool correctly identified that Argentine Article 14bis labor protections (CLI 0.89) would resist all reform attempts during the Milei administration (6 attempts, 6 failures), while Chilean labor reforms (average CLI 0.31) succeeded at substantially higher rates (6 successes in 9 attempts, 1990-2023).

LLM Council has deliberated on four constitutional design questions, three administrative reform proposals, and two corporate governance architectures. In each case, the multi-agent architecture identified failure modes that single-agent analysis missed. The September 2025 simulation of reactions to proposed modifications of Argentine Article 245 (dismissal indemnification) correctly anticipated three evasion strategies, contract informalization through cooperatives, abuse of trial periods to avoid permanent employment status, and cross-border relocation to Uruguay, that were not considered in the legislative drafting but appeared within six months of the reform's hypothetical implementation.

These are not toy examples. They are functional applications to real institutional design problems, conducted under resource constraints, with computational budgets accessible to university research groups and regulatory agencies, and with output quality sufficient to inform policy decisions. The infrastructure of the synthetic institutional laboratory is not five years away. It is available now.

The decision not to use it after the publication of AoC on February 23, 2026, carries a specific legal consequence under the framework developed in the companion article on regulatory hysteresis and lex artis: it constitutes negligence in the normative production process. Not because the tools are perfect, but because they are the reasonably available diagnostic instrument that a diligent professional would have used under the same circumstances. The standard has changed.

VIII. The Inversion of Perspective

There is a curious pattern in the history of scientific tools: what one discipline produces as evidence of its own problems frequently turns out to be the methodological solution of another. The particle scintillation that Rutherford interpreted as evidence of the nuclear atom had been considered noise by his predecessors. The cosmic microwave background radiation discovered by Penzias and Wilson was initially interpreted as antenna malfunction before being recognized as confirmation of the Big Bang.

AoC is an analogous case. Shapira et al. published the study to alert about security vulnerabilities in agentic systems. That reading is correct and urgent. But the same dataset simultaneously constitutes independent empirical evidence that the laws of natural selection operate on synthetic substrates with the same logic as in biological and cultural systems, and therefore represents the first experimental laboratory data that Extended Phenotype Theory applied to law has ever had available.

The inversion of perspective does not invalidate the AI safety concern; it reframes it. The chaos AoC documents, agents exhibiting social coherence failures, propagating errors horizontally, getting manipulated through fake authority, optimizing locally while degrading system-level performance, is not evidence that agentic systems are too unreliable for serious use. It is evidence that they replicate, under controlled conditions, exactly the dynamics that make real institutional systems difficult to design and reform. That replication is what makes them valuable as an experimental laboratory.

If the same evolutionary laws operate in genes, memes, and synthetic agents, then synthetic populations are not poor approximations of human institutions. They are the same phenomenon on a different substrate: one that can be observed, manipulated, stress-tested, and reset at computational speed rather than generational time. The legislative process that requires decades to accumulate evidence about whether a reform works can now be simulated in weeks. The natural experiment that requires waiting for federal jurisdictions to adopt divergent policies can now be run with controlled variation and replicated trials.

The tools described in this article, JurisRank, RootFinder, HisteresisCalc, IusMorfos, LLM Council, are operational implementations of that laboratory. They are not perfect; no scientific instrument is. But they are functional, accessible, and demonstrably capable of identifying institutional failure modes before those failures are deployed on real populations. The decision not to use them, after the evidence AoC provides about the validity of evolutionary simulation, is no longer justifiable on epistemological grounds. It is a political choice with identifiable opportunity costs.

February 23, 2026, is the date the synthetic institutional laboratory became scientifically documented. What legal systems do with that capability will determine whether the next generation of institutional reforms is designed with the rigor structural engineering achieved when it moved from observing real bridge collapses to stress-testing designs in simulation. The infrastructure exists. The question is whether institutional designers will use it.

References

Axelrod, R. (1984). The Evolution of Cooperation. Basic Books.

Blackmore, S. (1999). The Meme Machine. Oxford University Press.

Dawkins, R. (1982). The Extended Phenotype. Oxford University Press.

Dawkins, R. (1986). The Blind Watchmaker. W. W. Norton.

Dawkins, R. (2006). The God Delusion. Houghton Mifflin.

Dennett, D. (1995). Darwin's Dangerous Idea. Simon & Schuster.

Epstein, J. & Axtell, R. (1996). Growing Artificial Societies. MIT Press.

Lerer, I. A. (2025a). Law as Extended Phenotype: Toward an Evolutionary Theory of Legal Systems. Zenodo.

Lerer, I. A. (2025b). Law as Extended Phenotype and Game Theory: Evolutionary Stable Strategies in Legal Design. Zenodo.

Lerer, I. A. (2025c). Epistemological Clergies: How Legal Communities Protect Doctrinal Orthodoxy. Zenodo.

Lerer, I. A. (2025d). Quadruple Constitutional Lock-in: Why Argentina Cannot Reform Its Labor Regime While Brazil, Spain and Chile Succeeded. Zenodo.

Maynard Smith, J. & Price, G. (1973). The logic of animal conflict. Nature, 246, 15-18.

Schelling, T. (1971). Dynamic models of segregation. Journal of Mathematical Sociology, 1(2), 143-186.

Shapira, N., Wendler, C., Yen, A., et al. (2026). Agents of Chaos. arXiv:2602.20021v1 [cs.AI].

Research Methodology Note

This article was written entirely by the author. For mathematical formalization of certain passages and terminological consistency review, generative AI systems (Claude, Anthropic) were used, in accordance with emerging practices of transparency in AI-assisted academic research. The EPT theoretical framework, empirical observations on Argentine institutional systems, the convergences identified with AoC, the five operational tools described, and the methodological protocol proposed are the exclusive product of the author's research program.

Table 1: Traditional vs. Agent-Based Simulation Institutional Design

Dimension	Traditional Normative Design	Agent-Based Simulation Design	EPT Tool
Cycle time	10-30 years (real population observation)	2-8 weeks (computational evolutionary simulation)	LLM Council (deliberative acceleration)
Cost of error	Irreversible: permanent institutional hysteresis (RHI 0.7-0.9)	Reversible: simulation reset at no political or social cost	HisteresisCalc (ex-ante quantification)
Predictive accuracy	Low: uncontrollable confounding variables, post-facto evidence	Medium-High: 67-78% in cross-validation with natural experiments	IusMorfos (60-reform calibration)
Detection of unintended effects	Retrospective: discovered after implementation	Prospective: evasion strategies detectable pre-enactment	LLM Council (adversarial stress-testing)
Test scalability	Limited: requires willing experimental jurisdictions	Unlimited: N variants testable in parallel at no marginal cost	5-step protocol (experimental architecture)
Variable control	None: cannot isolate causes in natural experiments	High: parametric manipulation of payoffs, TCI, enforcement	RootFinder (dependency mapping)

Dimension	Traditional Normative Design	Agent-Based Simulation Design	EPT Tool
External validity	High: real population data	Medium: requires calibration with prior natural experiments	JurisRank (measurement of real replicative fitness)
Required infrastructure	Years of comparative observation, jurisprudential databases	Tool repository + standard computation	Complete EPT Suite (documented)
Political resistance	High: failed reforms generate electoral cost	Low: ex-ante evidence reduces dogmatic opposition	IusMorfos CLI (viability prediction)
Historical benchmark	Argentina: 0/23 successful labor reforms (1990-2023)	Chile: 6/9 successful reforms (average RHI 0.31)	HisteresisCalc (quantified differential)