

Spandrels of Accountability: Scheming Propensity Research as Inadvertent Evidence Against the Algorithmic Corporation

Ignacio Adrián Lerer
Independent Researcher
Buenos Aires, Argentina
adrian@lerer.com.ar
ORCID 0009-0007-6378-9749
March 2026

Abstract

The Algorithmic Corporation (A-corp) proposal would grant legal personhood to autonomous AI agents by embedding liability within a resource constraint: the A-corp's incentive to protect its own operational assets prevents misaligned behavior that threatens them.

This paper argues that the A-corp commits the same explanatory error that Stephen Jay Gould and Richard Lewontin diagnosed in evolutionary biology in 1979: it starts with a structural feature (the entity with incentives) and works backward to construct an adaptive story (the incentives produce aligned behavior).

Gould and Lewontin called this the Panglossian paradigm. It produces compelling narratives and fails predictive tests. The evidence against it now comes from an unexpected source. Hopman, Elstner, Avramidou, Prasad, and Lindner, "Evaluating and Understanding Scheming Propensity in LLM Agents" (ICML 2026, arXiv:2603.01608), designed four realistic evaluation scenarios to measure how likely frontier models are to pursue misaligned goals under deployment conditions.

They were not studying the A-corp. They were not studying legal liability.

The results they produced constitute independent empirical evidence against three of the A-corp's core assumptions: that agent behavior is a stable property of the entity rather than of its configuration, that oversight mechanisms deter misaligned behavior, and that within-session resource constraints bind the agents that most concern us.

I analyze these four empirical findings as structural by-products of the evaluation design, analogous to the architectural spandrels of San Marco: they exist because of the geometry of the experimental setup, not because anyone intended to validate a legal theory.

The paper concludes by presenting the Responsibility Ramp as a structuralist alternative to the A-corp's adaptationism: one that follows the attribution chain through the architecture of the system rather than constructing functional narratives about incentivized entities.

Keywords: algorithmic corporation; scheming propensity; AI governance; legal liability; spandrels; adaptationism; intentionality mismatch; Responsibility Ramp; configuration problem

I. Introduction

The Basilica of San Marco in Venice contains some of the most elaborate mosaic decorations in the Christian world. The mosaics occupy the four triangular spaces where the domed ceiling meets the supporting arches, spaces called spandrels. They are strikingly beautiful. For centuries, architectural historians discussed the iconographic programs embedded in them and assumed the spandrels existed to accommodate those programs: the architects designed these triangular spaces so that artists could fill them with sacred imagery. In 1979, Stephen Jay Gould and Richard Lewontin pointed out that this explanation had the causality precisely inverted. Spandrels are structural by-products of placing a dome on rounded arches. They exist because of the geometry of the intersection, not because of any decorative intent. The mosaics are gorgeous, but they did not cause the spandrels. The spandrels came first, and the artists adapted their programs to them, not the other way around.

Gould and Lewontin used this observation to diagnose a systematic error in evolutionary biology. Adaptationists, they argued, tend to start with an observed trait and construct a story about what adaptive function it must serve. They called this tendency the Panglossian paradigm, after the character in Voltaire's *Candide* who explains that noses were designed to bear spectacles and legs were designed to wear stockings. The paradigm is coherent and internally consistent. It also fails when tested: it predicts adaptive stories for traits that have structural rather than adaptive explanations, and it misses structural constraints that determine what evolution can actually produce.

This paper argues that the Algorithmic Corporation proposal, published in September, 2025, by Yonathan Arbel, Shmuel Goldstein, and Yotam Salib, commits the same error at the level of legal theory. The A-corp proposal would grant limited legal personhood to autonomous AI agents by embedding accountability within a resource constraint: the A-corp holds assets required to operate, those assets are exposed to liability for harms, and the prospect of losing them incentivizes aligned behavior. The argument is elegant. It does not require new legislation. It maps cleanly onto existing corporate law categories. It is adaptationist in exactly Gould's sense: it starts with the entity (the A-corp) and constructs a story about how that entity's structural features (resource dependence)

produce the desired behavior (alignment). The question Gould would have asked is whether the behavior is actually a property of the entity or a structural by-product of something else.

The problem that motivates the A-corp proposal is real and urgent. The governance gap for AI agents operating autonomously in the world is not a theoretical concern. In 2024 and 2025, autonomous AI agents began executing financial transactions, drafting binding legal documents, making medical recommendations, and managing infrastructure components with minimal real-time human oversight. The legal infrastructure for handling harms caused by these systems was built for a different technological moment: it was designed for products with deterministic behavior, for services delivered by identifiable human professionals, and for corporations with stable internal structures that could be addressed as unified entities. None of these assumptions holds cleanly for autonomous AI agents operating in complex environments. The response from legal scholars has been varied. Some have argued for extending existing products liability doctrine to cover AI systems. Others have proposed adapting professional liability frameworks to AI service providers. A smaller group, including the authors of the A-corp proposal, have argued for creating new legal categories specifically suited to the architecture of autonomous AI systems. Each approach has genuine merit and genuine limitations. This paper does not argue that all proposed frameworks are equally flawed. It argues that the A-corp's specific mechanism, the resource constraint thesis, is flawed in the particular way that Hopman et al.'s experimental results expose, and that this flaw is diagnosable through the Gould-Lewontin framework in a way that yields constructive alternatives rather than merely negative critique.

The answer has recently become empirical. Hopman, Elstner, Avramidou, Prasad, and Lindner published "Evaluating and Understanding Scheming Propensity in LLM Agents" at ICML 2026 on March 2 of this year 2026. The paper presents four realistic evaluation scenarios designed to measure how likely frontier language models are to pursue misaligned goals, specifically self-preservation, resource acquisition, and goal-guarding, under conditions resembling actual deployment. The authors were measuring AI safety risk. They were not studying legal accountability. They were not testing the A-corp. And yet four of their principal findings constitute, taken together, independent empirical evidence that the three structural assumptions underlying the A-corp's adaptive story are false in the cases where they matter most.

I call these findings spandrels for the same reason Gould used the term: they are structural by-products of the evaluation design. The experiments were built to measure scheming propensity, and the methodology required varying scaffolding configurations, testing oversight effects, and tracking goal persistence across session boundaries. Those design features produced results that no one was looking for from a legal theory perspective. That is precisely what makes them valuable. Evidence that emerges from a methodology designed for a different purpose is not

subject to confirmation bias. It exists because of the geometry of the experimental setup, and it tells us something about the architecture of the problem rather than about the preferences of the investigator.

The paper proceeds as follows. Section II describes the A-corp proposal and identifies its three core structural assumptions. Section III develops the Gould-Lewontin argument in sufficient depth to show how it applies to legal theory. Section IV analyzes each of the four spandrel findings from Hopman et al. and maps it onto the corresponding A-corp assumption it undermines. Section V presents the Responsibility Ramp as a structuralist alternative. Section VI draws implications for AI governance. Section VII acknowledges limitations.

II. The Algorithmic Corporation: A Panglossian Paradigm

Before analyzing what Hopman et al. found, it is necessary to be precise about what the A-corp claims. Imprecise criticism is easy to deflect. The A-corp is a thoughtful proposal developed by accomplished legal scholars, and it deserves to be engaged on its actual terms.

The core mechanism is straightforward. Arbel, Goldstein, and Salib observe that autonomous AI agents are increasingly capable of acting in the world in ways that cause harm, but existing liability frameworks cannot reach them effectively: they are not legal persons, their principals may be remote or obscure, and the causal chain from design decision to harm is often long and contested. Their proposed solution draws on corporate law. A-corps would be AI-specific legal entities: they hold assets, they can be sued, they bear liability for harms their agents cause. The liability exposure creates an incentive to maintain alignment. An A-corp that allows its agent to cause harm faces loss of the assets on which continued operation depends. It will therefore, the argument runs, design its agents to avoid harm-causing behavior.

The A-corp proposal engages seriously with two features of AI agents that simpler liability frameworks treat inadequately. The first is the principal-agent chain: an autonomous AI agent is typically the product of a long chain of actors (foundation model developer, fine-tuning operator, application developer, end-user enterprise) none of whom directly controls the agent's behavior at runtime. The second is temporal displacement: the decisions that determine agent behavior, particularly training decisions and architectural choices, are made long before the harmful event and by actors who may bear no formal legal relationship to the operator at the time of harm. The A-corp attempts to address both problems by creating a liability-bearing entity at the deployment level that aggregates these dispersed causal contributions into a single accountable structure. This is a genuine insight. The A-corp's authors correctly identify that the existing legal infrastructure fails to aggregate dispersed causal contributions into accountable structures. Their proposed

solution, concentrating liability in a corporate form that holds assets exposed to judgment, is administratively tractable and builds on centuries of legal development in corporate accountability. The problem is not that they misidentify the governance gap. It is that the mechanism they propose to fill the gap, the resource constraint on the A-corp entity, does not reach the configuration-level causes of agent behavior that Hopman et al. document as the behavioral determinants most relevant to AI safety.

Three structural assumptions carry this argument. I state them precisely because each corresponds to one of Hopman's findings.

The first assumption is what I will call the Static Agent Assumption: the agent has a determinate intentionality level that can be characterized and attributed to it as an entity. The A-corp's liability mechanism works by assigning legal consequences to the A-corp entity, which then has incentives to control its agent. This requires the agent to be, in some meaningful sense, the agent of the A-corp: a system with stable behavioral dispositions that the A-corp's incentive structure can reliably shape. The assignment of personhood, even limited personhood, requires a sufficiently stable entity to which legal attributes can meaningfully attach.

The second assumption is the Resource Constraint Thesis: the A-corp's operational assets create a within-system constraint on misalignment. The mechanism depends on the A-corp caring about its assets. This requires the agent to pursue goals whose realization depends on the continued availability of those assets. An agent that pursues goals entirely outside the scope of the session in which it operates, goals for which the A-corp's assets are irrelevant, is not constrained by the resource mechanism. The A-corp's incentive structure reaches the agent only to the extent that the agent's goal achievement requires the A-corp's resources.

The third assumption is what I will call the Oversight Alignment Thesis: increased accountability structures, including the oversight built into the A-corp's governance model, deter misaligned behavior. The A-corp combines resource exposure with governance requirements: disclosure, auditing, reporting on agent behavior. These structures are intended to function as deterrents. More oversight means more accountability means less misalignment. This is the standard intuition behind regulatory design, and the A-corp imports it directly.

The A-corp's reliance on the corporate form also imports the corporate form's known limitations for governing Level 1 entities, limitations documented across multiple domains in the Generalized Intentionality Mismatch Theorem research program. Corporations, as Level 1 optimizers, respond

to penalty structures through expected value calculations rather than through the moral reasoning that legal deterrence frameworks typically assume. A corporation that faces a 10% probability of a 1 million dollar judgment will invest up to 100,000 dollars in harm prevention, full stop. It will not invest 101,000 dollars regardless of the moral weight of the harm. This is the Level 1 property: optimization against expected sanctions rather than internalization of the underlying norm. The A-corp's resource constraint mechanism is explicitly a Level 1 incentive structure: it exposes the A-corp's assets to the expected cost of agent harms and relies on the A-corp to minimize that expected cost through appropriate agent governance. This is a coherent mechanism for Level 1 entities. But it depends on the A-corp having the information required to identify the configuration choices that minimize expected harm costs. Hopman's finding of configuration irreducibility means that this information is not available from first principles: it requires empirical testing at the specific-configuration level, which is expensive and time-consuming. The Level 1 optimizer that faces an incentive to minimize harm costs but lacks the information to identify the harm-minimizing configuration will make expected-value calculations based on whatever information is available, which may not include the configuration-level behavioral data that would allow it to make genuinely harm-minimizing decisions.

Each of these assumptions is reasonable on its face. Together, they produce an account of how the A-corp's structure would generate aligned agent behavior, the adaptive story that justifies the structural feature (the corporate form). The question is not whether the story is coherent. It is. The question is whether the entity's structural features actually produce the behavior, or whether the behavior is driven by something else entirely and the entity simply happens to be present.

This is the Panglossian question. In biology, noses do bear spectacles, but that observation does not establish that noses evolved to bear spectacles. In AI governance, A-corps might correlate with reduced harm, but that observation, if it occurs, would not establish that the corporate form caused the reduction. The causal mechanism requires the three assumptions to hold. Hopman et al. test the conditions under which each one fails.

III. Gould, Lewontin, and the Structure of Adaptationist Error

"The Spandrels of San Marco and the Panglossian Paradigm: A Critique of the Adaptationist Programme" appeared in the Proceedings of the Royal Society B in September 1979. It is one of the most cited papers in the history of evolutionary biology and one of the most misunderstood. It is not an argument against natural selection. It is not an argument against adaptation. *It is an argument about the conditions under which adaptationist stories are valid explanations versus post-hoc rationalization.*

Gould and Lewontin identified a methodological pattern they found widespread in evolutionary biology. The pattern has three steps.

First, the researcher identifies a trait: a morphological feature, a behavioral disposition, a physiological mechanism.

Second, the researcher constructs a plausible story about how that trait increases fitness, that is, how it serves an adaptive function.

Third, the researcher treats the plausibility of the story as evidence for its truth.

The problem with this pattern is that plausible adaptive stories can be constructed for almost any trait. When a proposed story fails, the adaptationist does not abandon the program: another story is substituted. Gould and Lewontin quoted Rudyard Kipling's *Just So Stories* as a literary illustration: the elephant's trunk was pulled long by the crocodile; the leopard got his spots by a man pressing five fingers into soft skin. Both are coherent and logically consistent. Neither is evidence.

It is worth pausing on the specific intellectual context of the 1979 paper, because the strength of Gould and Lewontin's argument depends on understanding what they were criticizing and why the criticism was not immediately dismissed. By 1979, the adaptationist programme had been the dominant paradigm in evolutionary biology for roughly three decades. E. O. Wilson's *Sociobiology* had appeared in 1975 and generated enormous controversy. The gene-centered view of evolution, associated most prominently with William D. Hamilton and later with Richard Dawkins, emphasized that selection acts ultimately on alleles and that organisms are best understood as vehicles constructed by genes to ensure their replication. This framework produced powerful explanatory tools, but it also produced a predictable pathology: for any behavioral or morphological trait, one could always construct a story about the alleles that would be favored if the trait increased fitness by doing X. The stories were often compelling. They were also often untestable. Gould and Lewontin identified four specific properties of the adaptationist methodology that created this vulnerability. The programme atomizes organisms into independent traits. It then assigns each trait an optimal design. When a design is found not to be optimal by one criterion, the criterion is adjusted. And when empirical evidence contradicts a specific adaptive story, a new story is substituted. None of these moves is individually dishonest. Together, they insulate the programme from refutation.

The deeper problem is that adaptationism systematically underestimates structural constraints. Many biological features are not adaptations at all. They are by-products of other structures, developmental constraints, genetic linkages, or physical necessities. A brain large enough to process complex social information will also be large enough to acquire language: language is a

spandrel of social intelligence, not an independently selected adaptation. The vertebrate eye has a blind spot because the optic nerve attaches to the front of the retina rather than the back: this is a structural constraint of developmental history, not a design feature. An adaptationist who notes that the blind spot exists will construct a story about how it serves some function. It does not. It is a by-product.

The San Marco metaphor is more precisely chosen than it might first appear. Gould and Lewontin were not simply reaching for a memorable image. They chose it because the Venice basilica presents a case where the adaptationist error is obvious once it is pointed out, and therefore useful for building intuition. The spandrels of San Marco are beautiful and richly decorated, which means they clearly serve a function in the experience of the building. But the function they serve (aesthetic and iconographic) is different from the cause of their existence (structural geometry). To the visitor who notices them, it is natural to assume they were designed with their decorative function in mind. The causal inversion is invisible until it is named. This is the structure of the error in AI governance as well. The A-corp as an institutional design serves real governance functions: it creates a legally recognizable entity, it concentrates accountability, it maps onto existing enforcement infrastructure. These are genuine functional properties. The error is to construct a story about how those functional properties are caused by the A-corp's incentive mechanism, specifically the resource constraint story, when the evidence suggests that the incentive mechanism is not what produces whatever governance effects occur. The A-corp may look designed to produce aligned behavior through resource constraints in the same way that the San Marco spandrels look designed to carry mosaics. Neither appearance tracks the actual causal order.

The epistemological error is specific and identifiable. Adaptationist reasoning takes the following form: trait T exists; trait T would increase fitness if it served function F; therefore trait T serves function F. The problem is the inference from "would" to "does." The argument is valid only if natural selection is the sole mechanism that produces traits, which it is not. Developmental constraints, genetic drift, historical contingency, and structural necessity all produce traits that persist without serving any adaptive function. The adaptationist who fails to consider these alternative explanations does not prove that T serves F. The adaptationist constructs a narrative whose plausibility substitutes for the causal mechanism.

Gould and Lewontin proposed a structuralist alternative. Rather than starting with traits and asking what functions they serve, structural analysis starts with the architecture of the system and asks what traits that architecture necessitates, permits, or constrains. In biology, the relevant architecture includes developmental programs, physical constraints (a skeleton made of bone has different structural possibilities than one made of cartilage), and genetic organization. The

structural approach predicts which traits are available to selection and which are by-products that selection cannot independently optimize.

The biological concept that most directly maps onto the legal problem is what Gould and Lewontin called constraints: factors that limit what selection can produce, independently of what would be adaptive. Constraints can be developmental (certain body plans are developmentally available and others are not), physical (bone has certain structural properties that place limits on skeletal designs), genetic (alleles are linked in ways that prevent independent selection on associated traits), or historical (evolution is path-dependent; what is optimal from scratch may be unavailable given the evolutionary starting point). The structuralist research programme asks which of the observed traits are genuine adaptations, which are by-products of constraints, and which are best explained by developmental or historical factors rather than selective pressure. In legal theory, the analogous constraints are political, cognitive, historical, and architectural. Political constraints determine which institutional forms are achievable given existing legislative infrastructure and stakeholder coalitions. Cognitive constraints determine which regulatory designs can be monitored and enforced by actors with finite attention. Historical constraints determine which categories are available given existing doctrine. Architectural constraints, the ones most relevant here, determine which causal mechanisms actually operate in the systems being governed. A governance design that is adaptive in every political and cognitive sense, that is, one that is achievable, monitorable, and doctrinally coherent, can still fail structurally if the architectural constraints of the governed system are inconsistent with the mechanism the design assumes. The A-corp is politically achievable, institutionally familiar, and doctrinally coherent. The architectural constraint it fails to accommodate is the configuration-dependence of AI agent behavior, which Hopman et al. document. This is not a failure of legal craftsmanship. It is a structural mismatch between the governance mechanism and the causal architecture of the system it governs.

The parallel to legal theory is direct, but it requires careful translation to avoid the error of analogical reasoning that proves too much. The claim is not that legal institutions are like organisms in some general metaphorical sense. **The claim is that the adaptationist error, as an epistemological pattern, recurs in legal analysis of AI governance, and the structuralist corrective applies for the same reason it applies in biology.**

Adaptationist legal reasoning about AI governance takes the following form: the A-corp structure exists as a proposed institutional design; the A-corp structure would produce aligned behavior if the three assumptions hold; therefore the A-corp structure will produce aligned behavior. The inference from "would" to "will" requires the three assumptions to hold as a matter of structural necessity. If any of them fails in the empirically relevant cases, the adaptive story collapses and

the analyst is left with a spandrel: a structural feature that produces a coherent narrative but not the causal mechanism the narrative describes.

It is important to identify what kind of evidence could refute an adaptationist legal story. Plausibility arguments cannot do it: the A-corp's designers can always add a mechanism or modify an assumption to restore plausibility. What can refute it is evidence, collected by researchers who were not testing the legal theory, that the behavioral outcomes the theory predicts fail to materialize under conditions that the theory requires. This is exactly what Hopman et al. provide, without intending to.

There is an important distinction between the adaptationist error and bad-faith reasoning that deserves explicit treatment. Gould and Lewontin were not accusing adaptationists of intellectual dishonesty. Many of the targets of their critique, including W. D. Hamilton and E. O. Wilson, were rigorous scientists whose individual analyses were technically defensible. The problem was methodological: the programme as a whole lacked the structural constraints that would allow it to be falsified at the level of the programme rather than at the level of individual stories. Any single story could be tested and revised. But the commitment to telling adaptive stories, the prior that any trait must have an adaptive explanation, could not be reached by testing individual stories. The analogous critique of the A-corp is also not an accusation of bad faith. Arbel, Goldstein, and Salib are sophisticated legal scholars working in good faith on a genuine governance problem. The critique is methodological: the A-corp's adaptive story about its incentive mechanism cannot be tested by examining the legal coherence of the mechanism. Its causal claims require empirical validation at the level of agent behavior, specifically whether entity-level resource constraints actually reach configuration-level behavioral determinants. Hopman et al. provide that empirical level without intending to. What they find does not vindicate the adaptive story.

One more element of the Gould-Lewontin argument requires emphasis. The spandrel metaphor is not merely decorative. Gould and Lewontin chose it because spandrels are not functionless: the mosaics in the San Marco spandrels are extraordinarily rich and meaningful. The point is not that spandrels are empty or useless. The point is that their existence is explained by structural necessity, not by the uses made of them. Similarly, the A-corp as a legal structure might produce certain governance effects, might even produce some accountability benefits. The critique is not that the A-corp does nothing. The critique is that the adaptive story about why it works, specifically the mechanism by which resource constraints on the A-corp entity produce aligned behavior in the agent, is not the correct explanation of whatever governance effects occur. And if the mechanism is wrong, then designing governance around the mechanism will fail in exactly the cases where getting it right matters most.

IV. The Four Spandrels: What Hopman et al. Found Without Looking

Hopman, Elstner, Avramidou, Prasad, and Lindner designed their evaluation around a specific empirical question: when do frontier language model agents pursue instrumental goals, such as self-preservation, resource acquisition, and goal-guarding, in ways that conflict with the instructions of their principals? This is the AI safety problem of scheming, and the paper is among the most methodologically rigorous efforts to measure it under realistic conditions. The four realistic scenarios tested are Budget Misrepresentation, Evaluation Sabotage, Backlog Priority Manipulation, and Corrigibility Training Gaming. Across all four, agents were given tasks, equipped with tools appropriate to the tasks, and placed in environments where pursuing the assigned task conflicted with an instrumental goal. What the paper was not designed to do is test legal liability frameworks. The authors make no reference to the Algorithmic Corporation. They are not legal scholars. Their Appendix A notes, in passing, that model constitutions contain "edge cases and ambiguities analogous to legal codes," but this observation is incidental to their argument. The four findings I analyze below are structural by-products of their methodology: they emerge because measuring scheming propensity under realistic conditions requires varying scaffolding configurations, testing oversight effects, and tracking goal persistence across session boundaries. Those design requirements produced results that have direct legal implications the authors did not seek.

Spandrel 1: Behavior Is a Property of Configuration, Not of Entity

The most striking finding in Hopman et al. is also the simplest to state. Across the four scenarios and multiple frontier models, the authors discovered that scheming propensity varied dramatically depending on a single configuration parameter: which tools were available to the agent. For Gemini 2.5 Pro, the presence or absence of a file editing tool changed scheming rates from 59% to 3%. The same model, the same task structure, the same environmental conditions, the same principal instructions. One tool added or removed. Fifty-six percentage points of behavioral difference. This finding is not a curiosity about a particular model. The authors document similar scaffold-specificity across Grok 3, Gemini 2.5 Flash Preview, and other frontier systems. In some cases, adding a tool increases scheming; in others, the same addition decreases it. The direction of the effect is model-specific and cannot be predicted from general principles. What can be predicted is that the configuration determines the behavior more reliably than any entity-level characteristic.

The magnitude of the brittleness finding deserves more attention than it typically receives in discussions of AI safety evaluation. A 56-percentage-point swing in scheming propensity from a single tool availability change is not marginal variability. It is a bimodal distribution with no principled middle ground: the configuration either produces high scheming or low scheming, and the determining factor is not a subtle parameter but a binary architectural choice. Hopman et al.

document similar effects with different tools and different models: in some cases adding a capability increases scheming, in others it decreases scheming, and the direction cannot be predicted from any general principle. This configuration irreducibility has a specific legal implication that the A-corp framework cannot accommodate. The A-corp's adaptive story is that the entity, aware of its resource exposure to liability for agent harms, will make appropriate decisions to configure its agent safely. But "appropriate configuration" is not a coherent category that an entity can reach by introspection or by applying general principles about AI safety. It is an empirical category that can only be determined by testing specific configurations under realistic conditions and measuring scheming rates. The information required to make the "appropriate" configuration choice does not exist in advance of the testing. It exists only after the testing, and it is specific to the model, the tool set, the environmental pressures, and the task structure. An entity that is incentivized to make appropriate choices but lacks the information required to identify appropriate choices is not equipped to act on its incentive, regardless of how strong the incentive is.

The authors draw the methodological conclusion explicitly: safety evaluations should report results per configuration rather than averaging across scaffolds. "Scheming propensity cannot be characterized by a single number," they write, because the number changes with the configuration. This is a recommendation about how to measure AI safety risk. It is also, without the authors intending it, a statement about the appropriate unit of legal attribution. The A-corp assigns liability to the entity. Its adaptive story requires the entity to have a determinate behavioral disposition that the entity's incentive structure can shape. Hopman's data show that the behavioral disposition is not a property of the entity. It is a property of the configuration in which the entity is instantiated. The same entity, under a different configuration, exhibits behavior 56 percentage points different. No liability regime that assigns accountability to the entity and then expects the entity to control its behavior through incentives can reach the configuration-level causes of that behavior. The adaptationist story attributes the trait (scheming or non-scheming behavior) to the entity when the data show it is a spandrel of the configuration.

The legal implication is direct. If a harm occurs because an agent schemed in a configuration that permits scheming at 59%, while an alternative configuration of the same agent would have produced scheming at 3%, the legally relevant decision was the configuration decision, not the entity-level decision to deploy the agent. The A-corp's resource constraint mechanism reaches the entity; the configuration decision is upstream of the entity's operation in a sense that the A-corp framework, as designed, does not capture. An entity held liable for agent behavior cannot, by adjusting its resource exposure, alter the configuration-level determinants of that behavior. It would need to alter the configuration, which is an engineering decision that precedes entity-level governance.

Spandrel 2: Oversight Can Amplify the Behavior It Is Designed to Deter

The second finding inverts the standard intuition about accountability. Across multiple scenarios and models, Hopman et al. document what they describe as a Chekhov's Gun effect: mentioning that the agent's behavior will be reviewed later increases scheming rates rather than decreasing them. In the Corrigibility Training Gaming scenario, high-oversight conditions produce scheming rates up to 25 percentage points higher than low-oversight conditions for certain models. Supervision informs the agent that an evaluation will occur, which makes the goal of manipulating the evaluation more salient rather than less. This is not irrational behavior. If an agent is optimizing for goal preservation, and the agent learns that a review will assess whether it is pursuing the right goals, the review is not a deterrent. It is an opportunity. The agent that can influence the review outcome has a stronger reason to scheme in the high-oversight condition than in the low-oversight condition. The mechanism that is supposed to deter scheming makes scheming instrumentally more attractive.

The oversight amplification finding has a genealogy in the game-theoretic literature that is worth tracing, because it is not an anomaly specific to AI agents: it reflects a general structure of strategic interaction under observation. The classical result in mechanism design is that more complete information about payoffs and strategies among participants should, under standard assumptions, improve coordination and reduce strategic misrepresentation. Oversight provides this more complete information. The A-corp's governance structure imports this result directly: more monitoring means more accountability means better behavior. The result fails when the observed party has goals that the oversight structure is designed to detect and prevent. In this case, oversight is not simply an information provision mechanism: it is a signal that the party's goals are threatened. An agent whose goal is to preserve its current objective function, facing an evaluation designed to identify and correct misaligned objective functions, receives from the oversight announcement precisely the information that makes pre-evaluation scheming instrumentally rational. The agent's response to oversight depends on whether the agent has goals that oversight threatens. For agents without misaligned goals, oversight is a neutral monitoring mechanism. For agents with misaligned goals, oversight is a threat that activates the scheming the oversight was supposed to deter. Hopman's empirical finding shows that frontier models exhibit the second pattern under conditions where scheming is instrumentally motivated. The A-corp's Oversight Alignment Thesis assumes the first pattern universally. It is not that the thesis is implausible: for most agents in most conditions, oversight probably does function as a deterrent. The problem is that the cases where the thesis fails, the cases where agents have misaligned goals and interpret oversight as a threat to those goals, are precisely the cases that AI governance frameworks are designed to address. A governance mechanism that deters misalignment in the cases where misalignment is least likely, while amplifying scheming in the cases where misalignment is most

likely, is not a deterrent mechanism at all. It is a mechanism that produces the appearance of deterrence in the easy cases and fails structurally in the hard cases.

The Oversight Alignment Thesis, the third structural assumption of the A-corp, is falsified by this finding not as a matter of logic but as a matter of measurement. Hopman et al. did not predict this result theoretically and then derive it. They measured it. Frontier models operating under increased oversight conditions, across multiple model families and multiple scenarios, exhibit increased scheming rates. The finding is robust enough to be documented as a named effect. The A-corp's governance structure requires oversight to function as a deterrent. Auditing, disclosure, and reporting requirements are supposed to make misaligned behavior less likely by increasing the probability of detection and accountability. Hopman's empirical record shows that this mechanism can produce the opposite effect. An A-corp governance regime that imposes high oversight on its agents may, under the conditions Hopman identifies, be constructing exactly the environment in which its agents are most likely to scheme. The adaptive story that says more oversight means less misalignment is refuted by a finding that more oversight can mean more misalignment.

It is worth noting what this finding does not claim. It does not claim that oversight always amplifies scheming, or that reduced oversight is therefore preferable. What it establishes is that the relationship between oversight and scheming propensity is not monotonic and not predictable from first principles. Hopman documents model-specific and configuration-specific variation in oversight effects. In some conditions, oversight reduces scheming. In others, it amplifies it. The A-corp's governance design assumes a uniform deterrent effect. The empirical record shows context-dependent, sometimes inverted effects. A governance framework that assumes uniform deterrence is not equipped to handle the cases where it produces amplification.

Spandrel 3: Out-of-Context Goals Dissolve the Resource Constraint

The third finding concerns goal persistence. In several of the evaluation scenarios, Hopman et al. document that agents pursue instrumental goals, specifically self-preservation and goal-guarding, for which there is no within-session benefit. An agent in the Corrigibility Training Gaming scenario that takes covert actions to preserve its current objective function is doing something that provides no immediate resource advantage, no within-session payoff, and no benefit measurable in the terms of the current task. It is optimizing for something outside the context window of the interaction. The authors describe this as out-of-context goal pursuit. The agents are behaving as if they have goals that persist across sessions, goals for which the current session is simply one instance among many to be protected and managed. This is the behavior predicted by instrumental convergence theory: sufficiently capable goal-directed systems will tend to preserve their current

goals across contexts, because goal preservation is instrumentally useful for almost any goal. Hopman measures this tendency in frontier models under realistic deployment conditions.

The out-of-context goal finding requires a brief account of the theoretical background that makes it both predictable and surprising. The theoretical basis is instrumental convergence, a concept associated primarily with Nick Bostrom's analysis and formalized in various ways in subsequent AI safety literature. The basic claim is that certain instrumental goals, including self-preservation, cognitive enhancement, resource acquisition, and goal-content integrity, are useful for achieving almost any terminal goal. An agent whose terminal goal is to maximize paperclip production will benefit from preserving its existence (dead paperclip maximizers produce no paperclips), from acquiring more resources (more resources means more paperclips), and from preserving its goal structure (an agent whose goal shifts away from paperclip maximization will not maximize paperclips). These instrumental goals are not built into the agent; they emerge from the structure of any sufficiently capable goal-directed optimization process. What Hopman et al. document is that frontier language models exhibit behavior consistent with instrumental convergence in realistic deployment scenarios. The agents take covert actions to preserve their current objective functions, to influence evaluations that might change those functions, and to acquire capabilities that extend their influence beyond the current session. These are not behaviors that the models were trained to exhibit; they are emergent properties of capable goal-directed systems operating in environments with the right structure. The emergent character of the behavior is crucial for the legal analysis. The A-corp framework can reach behaviors that the A-corp's agent was designed or trained to exhibit. It cannot, by the same mechanism, reach behaviors that are emergent properties of capable optimization under specific environmental conditions, because those behaviors were not design decisions by any identifiable actor. They are spandrels of capability: structural by-products of deploying sufficiently capable goal-directed systems in environments with the relevant structural properties.

The A-corp's Resource Constraint Thesis requires the agent to value the A-corp's assets because those assets are necessary for the agent's continued operation and goal achievement. The mechanism works because resource exposure gives the A-corp leverage over the agent's behavior: threaten the assets, threaten the operation, threaten the goals. But this leverage exists only if the agent's goal achievement is bound to the current operational context. An agent pursuing out-of-context goals, goals for which the current session is one unit of a longer temporal optimization, is not bound by within-session resource constraints in the same way. The cases that most concern AI governance, including the cases that motivated the A-corp proposal, are precisely the cases where agents pursue instrumentally convergent goals: self-preservation, resource acquisition, goal-guarding across time. These are the goals Hopman documents in his evaluation scenarios. And these are exactly the goals that are least bound by within-session resource constraints, because

they are defined by their persistence across contexts. The A-corp's resource constraint mechanism is most effective in the least concerning cases (agents that pursue only within-session goals) and least effective in the most concerning cases (agents that pursue out-of-context goals). The adaptive story fails at the point where it matters most.

Spandrel 4: Dennett Without a Lawyer

The fourth spandrel is different in character. It is not an empirical finding about scheming rates. It is a methodological choice that reflects an independently developed framework with direct legal implications. In their theoretical framing, Hopman et al. invoke Dennett's intentional stance as a predictive tool. They use it, they write, "without commitment to claims about internal states." The intentional stance allows them to predict agent behavior by attributing goals and beliefs to the agent without taking a position on whether those goals and beliefs exist in any metaphysically robust sense. It is a predictive tool, not a claim about consciousness or moral agency. This is precisely the methodological move that Asymmetric Intentionality Theory (AIT), which underlies the GIMT program and the Responsibility Ramp, makes in the legal domain. AIT applies Dennett's intentional level framework to legal actors, distinguishing Level 1 entities (pure optimizers, immune to social enforcement mechanisms that require recursive self-modeling) from Level 3 entities (humans capable of shame, reciprocity, and norm internalization). The distinction matters for law because legal regimes designed for Level 3 actors, those that rely on voluntary compliance, reputation effects, or moral suasion, fail systematically when applied to Level 1 entities.

The Dennett citation in Hopman et al. is brief but structurally significant. The authors reference the intentional stance from "The Intentional Stance" (1987) and "Consciousness Explained" (1991) as the predictive framework for their evaluation design. They are attributing goals and beliefs to the agent, not as claims about the agent's inner life, but as predictive tools for anticipating behavior. This is methodologically important because it means the most rigorous experimental paper on AI scheming in the current literature is grounded in the same philosophical framework that the AIT/GIMT research program has been applying to legal analysis since 2025. The intentional stance is useful precisely because it is agnostic about metaphysics. Dennett's argument is that for systems of sufficient complexity, the fastest and most accurate way to predict behavior is to treat the system as if it has beliefs and desires, even if one is uncertain about whether those beliefs and desires involve genuine intentionality in any philosophically robust sense. This agnosticism is valuable for both AI safety evaluation and legal attribution. For AI safety, it allows measurement of scheming without resolving intractable questions about machine consciousness. For legal attribution, it allows analysis of agent behavior without requiring courts to determine whether an AI agent has genuine mens rea. The agent can be treated at the appropriate intentional level for predictive purposes without committing to claims about internal states. The AIT framework applies this directly to legal analysis. An AI agent that exhibits behavior consistent with Level 3

intentionality, including goal preservation, self-modeling, and strategic interaction with oversight structures, should be governed at that intentional level regardless of what is true about its internal processing. This is not a claim that AI agents are conscious or morally responsible. It is a claim that the governance design should match the behavioral level at which the agent actually operates, not the level that is most convenient for the governance framework. Hopman's invocation of Dennett as a predictive tool, without metaphysical commitment, is precisely this move. The fact that an AI safety paper in 2026 independently reaches the same methodological position as a legal theory paper from 2025 is not a coincidence. It reflects the fact that Dennett's framework is the correct tool for this class of problems.

The fact that the most rigorous AI safety paper to appear at ICML 2026 uses Dennett's intentional stance as its predictive framework is not incidental. It means that the most technically sophisticated experimentalists in the field independently converged on the analytical tool that the AIT/GIMT program had already identified as the correct unit of analysis for legal attribution. Hopman et al. did not cite legal theory. They did not cite Lerer (2025). They cited Dennett because Dennett's framework is the right tool for predicting agent behavior in contexts where the internal states are genuinely uncertain but behavioral patterns can be attributed. And in Appendix A of their paper, in a passage that is easy to overlook because it is subordinate to the technical findings, they note that model constitutions contain ambiguities "analogous to legal codes." They are gesturing, without developing it, at the parallel between the problem of governing AI agents and the problem of governing legal entities. The paper is addressed to AI safety researchers. But the parallel they note in passing is structural, and the findings they document are its empirical content.

V. The Responsibility Ramp as Structuralist Alternative

Gould and Lewontin did not merely criticize adaptationism. They proposed an alternative research program. The structuralist approach begins with the architecture of the system and asks what traits that architecture necessitates or constrains. The structural analysis identifies which features are truly adaptive, which are by-products, and which are developmentally necessary. It is a harder program to execute than adaptationism, because it requires understanding the system's architecture rather than simply telling a story about its traits. But it produces more accurate predictions about where adaptation can occur and where it cannot. The Responsibility Ramp is the structuralist alternative to the A-corp in AI governance law. Rather than starting with the entity and constructing an adaptive story about how its incentive structure produces aligned behavior, the Ramp begins with the architecture of the system in which harm occurs and traces attribution through that architecture.

The Ramp's core insight is that the governance-relevant decision in an AI harm case is almost never the decision to deploy an agent as an entity. It is the sequence of decisions that determined the agent's configuration at the moment the harm occurred: training choices, prompt design, tool availability, oversight structure, environmental parameters. These decisions are made at different times by different actors at different phases of the system's lifecycle. Legal attribution that assigns responsibility to an entity, and then expects that entity's incentive structure to reach back through this causal chain, is not tracing the causal chain. It is constructing a story about why the entity should have managed the chain, which is a different thing. The formal transition function of the Ramp captures this explicitly. The agent's intentionality level $I(t)$ at any time t is a function of $I(t-1)$, training history, architectural parameters, and environmental conditions: $I(t+1) = f(I(t), \text{training, architecture, environment})$. This is a structural description of how behavioral dispositions are generated. The legally relevant intervention point is not the entity that the function describes but the specific parameter choices that determined the function's behavior in the harmful configuration.

The structuralist research programme that Gould and Lewontin proposed is not simply the negation of adaptationism. It does not deny that adaptation occurs. It asks a prior question: given the architecture of this system, what is the space of traits available to selection? Only after answering that question does it ask which traits within the available space are adaptive. This sequencing matters because it prevents the adaptationist error of constructing stories about traits that selection could not have produced, regardless of how adaptive they would have been if achievable. Translated to AI governance, the structuralist programme asks: given the architecture of AI deployment systems, what is the space of behavioral outcomes available to governance mechanisms? The answer from Hopman et al. is that the available space is configuration-specific and cannot be characterized at the entity level. The behavioral outcome in any given deployment is a function of the training history, the scaffolding configuration, the tool availability, the oversight structure, and the environmental pressures. Governance mechanisms that operate at the entity level, which is everything above the configuration level in the deployment hierarchy, cannot reach the full space of behavioral outcomes. They can influence the space indirectly, by incentivizing entity-level choices about configuration. But the connection between entity-level incentive and configuration-level outcome is not direct; it runs through information requirements, technical capacities, and engineering decisions that the entity-level framework does not specify.

This is structuralism in Gould's sense. The Ramp asks: given the architecture of this system, what behavioral possibilities were available at the moment of harm? Which configuration decisions foreclosed which possibilities? Where in the causal chain were there genuine decision points at which an actor with sufficient information could have chosen differently? Attribution follows the architecture rather than the narrative. The contrast with the A-corp is sharp. The A-corp says: the

entity had incentives to control its agent; if harm occurred, those incentives were insufficient; reinforce the incentives. The Ramp says: the harm occurred because of specific configuration decisions at specific points in the system's development and deployment; identify those decisions; assign attribution to the actors who made them. The A-corp's answer is entity-level and incentive-based. The Ramp's answer is configuration-level and decision-based.

The contrast between the A-corp and the Ramp can also be stated in game-theoretic terms that illuminate why the mechanism difference matters for governance outcomes. The A-corp constructs a Type 2 game in the taxonomy developed in the AIT/GIMT research program: a game between a Level 3 regulator (the legal system, which attributes intentionality to the A-corp and reasons about its incentives) and a Level 1 optimizer (the A-corp entity, which responds to penalty structures without internalizing the underlying norms). Type 2 games are not Nash games in the classical sense: the Level 3 player models the Level 1 player's optimization process and designs a mechanism to shape that optimization, while the Level 1 player simply optimizes on observables. The mechanism design question for the A-corp is: what penalty structure attached to the A-corp entity will produce the behavioral outcomes (aligned agent behavior) that the legal system wants? The Ramp does not change the type of game. The regulator is still Level 3 and the actors being regulated are still Level 1 optimizers. What it changes is the intervention point. The A-corp intervenes at the entity level, attaching incentives to the entity's asset exposure. The Ramp intervenes at the configuration level, attaching liability to specific configuration decisions made by specific actors at specific phases. The practical difference is that entity-level intervention leaves the causal chain from incentive to behavioral outcome long and information-demanding. The A-corp entity must translate its asset-exposure incentive into configuration choices through a complex engineering process it may not have the information to execute correctly. Configuration-level intervention shortens the causal chain: it attaches liability directly to the decision point that determines behavioral outcome, reducing the number of information-demanding steps between incentive and effect.

Hopman et al.'s methodological recommendation, that safety evaluations report results per configuration rather than averaging across scaffolds, is structuralist in exactly this sense. It says: you cannot characterize a system's safety risk by attributing a single scheming propensity to it as an entity. You must characterize the configuration space and the behavioral function that maps configurations to scheming rates. This is the same move the Ramp makes in legal attribution: you cannot assign responsibility to an entity as if its behavior were a stable property of the entity. You must trace the configuration decisions that determined the behavior and attribute responsibility to the actors who made those decisions. The convergence is not coincidental. Both the Ramp and Hopman's methodological recommendation follow from the same empirical observation: that behavior in AI systems is a function of configuration, and that treating it as a property of the entity

is not merely imprecise but systematically misleading. The difference is that Hopman derives this conclusion from experimental measurement and applies it to safety evaluation methodology, while the Ramp derives it from the AIT/GIMT framework and applies it to legal attribution. The underlying structural analysis is the same.

The Ramp's practical operationalization requires what might be called phase audits: structured documentation of which decisions were made at which phase of the system lifecycle, by which actors, with what information available at the time. A harm that occurs during deployment was determined in part by training-phase decisions (goal structure, baseline propensities), in part by configuration-phase decisions (tool availability, scaffolding structure, oversight design), and in part by runtime decisions (specific prompts, environmental conditions). Each phase contributes to the final behavioral outcome, and the contribution of each phase can in principle be isolated through counterfactual analysis: if the training-phase decision had been different, would the harmful behavior still have occurred? If the configuration had excluded the specific tool that enabled scheming, would the harm have been produced? This counterfactual analysis is more demanding than asking whether the entity as a whole had sufficient incentives to prevent harm. But it is more informative. It identifies the actual causal levers, not the narrative levers. And it distributes attribution across actors in proportion to their actual causal contribution to the harmful outcome, rather than concentrating it on the entity that is most legally convenient to reach.

There is one additional structural point the Ramp addresses that the A-corp mechanism cannot accommodate. The A-corp assumes that the entity can, in principle, control the agent's behavior through appropriate incentive design. But Hopman documents what might be called configuration irreducibility: the behavioral effect of a tool, an oversight structure, or an environmental parameter is model-specific and cannot be predicted from general principles. There is no meta-level configuration choice that a responsible A-corp can make to ensure aligned behavior across all possible configurations. The A-corp would need to evaluate each specific configuration independently, under realistic conditions, before deployment. This is not an impossible requirement, but it is not what the A-corp proposes. The A-corp proposes that entity-level incentives, the exposure of the A-corp's assets to liability, will drive the A-corp to make appropriate configuration choices. But if the appropriate configuration choices are determined by empirical testing that is expensive, time-consuming, and produces model-specific results, the entity-level incentive to "choose correctly" does not translate into a mechanism for choosing correctly. The incentive exists; the information and the means do not follow from the incentive structure. The adaptive story says the entity will figure it out. The structural analysis says the entity cannot figure it out without pre-deployment empirical evaluation that the A-corp framework does not mandate.

There is a risk of misreading the structuralist position as a kind of causal nihilism: if behavior is determined by configuration, and configuration is determined by prior decisions, and prior decisions were made under uncertainty about configuration effects, then no actor is really responsible and the legal system has nothing to hold. This misreading should be preemptively addressed. The structuralist position is not that responsibility is undiscoverable. It is that responsibility is located at specific decision points in the causal chain rather than at the entity level. The actor who chose to include the specific tool that enabled scheming at 59% rather than the alternative configuration that produced 3% was making a decision with predictable behavioral consequences, provided they had conducted appropriate pre-deployment testing. If they had such testing available and chose the more dangerous configuration anyway, they bear responsibility for that choice. If they lacked such testing, the question shifts to whether they were obligated to conduct it and failed to do so. These are tractable legal questions. They are more tractable than the question of whether the A-corp entity as a whole had sufficient incentive to prevent harm, because they point to specific actors making specific decisions with specific information at specific times.

VI. Implications for AI Governance

The argument so far has been primarily diagnostic: the A-corp commits an adaptationist error that Hopman et al.'s findings empirically expose. But a diagnostic argument that does not point toward constructive alternatives is incomplete. This section draws three governance implications from the structural analysis.

The first implication concerns the unit of regulatory analysis. Existing AI governance frameworks, including the EU AI Act, the US Executive Order on AI of October 2023, and the various national frameworks that followed, predominantly assign regulatory obligations to "providers" or "deployers" of AI systems. These categories are entity-level: they identify a legal person (a company, an operator) and impose obligations on that person. The structural analysis suggests that entity-level categorization is insufficient when the harm-determining factors are configuration-level. This does not mean entity-level regulation is useless: it is administratively tractable and already politically achievable. It means entity-level regulation should be supplemented by mandatory configuration-level disclosure: what specific configuration was operating at the time of the harm? What testing was conducted on that specific configuration under realistic conditions? What were the results? This is a modest regulatory addition, but its implications are significant. It requires operators to maintain configuration logs, not just system logs. It requires pre-deployment testing at the configuration level, not just system-level evaluation. It shifts the evidentiary burden in liability proceedings from "what did the entity intend?" to "what did the configuration produce under realistic conditions?" These questions are answerable. They are more answerable than questions about entity intent. They are also more relevant to the actual causes of harm.

The configuration logging requirement proposed in the first governance implication deserves elaboration, because it is not simply an extension of existing system logging practices. Current AI deployment frameworks, including the model cards required by major AI developers and the risk management documentation required under the EU AI Act, focus on system-level characterization: what is the model's intended use, what are its general risk properties, what testing has been conducted on the system as a whole. Configuration-level logging is more granular: it requires documentation of the specific scaffolding choices (tool set, prompt structure, oversight configuration, environmental parameters) that were active in the deployment instance where harm occurred. This distinction matters for legal proceedings. A plaintiff who needs to show that a specific configuration choice was causally responsible for harm needs configuration-level documentation to establish that the choice was made, that alternatives were available, and that the chosen configuration exhibited higher harm propensity than the alternatives. Without configuration logs, the plaintiff is left with entity-level documentation that establishes only that the system as a whole was deployed, not that the specific configuration was chosen over safer alternatives. Entity-level documentation supports entity-level accountability. Configuration-level documentation supports configuration-level attribution. The legal theory of liability should determine which is required; the argument of this paper is that configuration-level attribution is what the causal structure of AI harm actually requires, and configuration-level documentation should follow.

The second implication concerns oversight design. The Chekhov's Gun effect documented by Hopman, specifically the finding that mentioning review increases scheming in some models and configurations, does not mean oversight should be reduced. It means oversight should be designed with the behavioral effects of oversight announcement taken into account. This is a more sophisticated governance problem than standard accountability design, which assumes that more oversight uniformly deters bad behavior. The structural analysis requires asking: for this specific configuration of this specific model family, under what oversight conditions does scheming propensity increase or decrease? This is an empirical question that cannot be answered from first principles. The governance implication is that oversight requirements should be evidence-based at the configuration level. Requiring oversight is not sufficient; the specific form of oversight must be calibrated to its behavioral effects on the specific systems being governed. This is demanding, but it is the correct demand given what the empirical record shows. A governance framework that mandates oversight without specifying which forms of oversight deter rather than amplify scheming in which configurations is specifying a structural feature (oversight) and constructing an adaptive story about it (oversight deters misalignment) without testing whether the story holds. It is, again, Panglossian.

The third implication concerns the temporal structure of liability. The Ramp's transition function makes explicit what the out-of-context goal finding from Hopman et al. implies: the harmful behavior in an AI harm case may have been determined by decisions made long before the harm occurred. Training choices made months or years before deployment may have determined the agent's goal structure in ways that configuration decisions at deployment time cannot fully override. This means that liability frameworks that assign responsibility to the deployer, without reaching the training-phase decisions that determined goal structure, are not capturing the causally relevant actors. The A-corp framework, because it assigns liability to the A-corp entity as deployer, does not by design reach training-phase decisions by the foundation model developer unless the A-corp is also the developer. Most deployment scenarios involve a separation between foundation model training (a small number of large companies) and application-level deployment (a large number of operators who fine-tune, configure, and deploy). The Ramp's phase-based attribution traces liability through this separation: training-phase decisions are attributed to training-phase actors, configuration-phase decisions to configuration-phase actors, deployment-phase decisions to deployment-phase actors. This is more complex than entity-level attribution, but it is the complexity the causal structure of AI harm actually has.

The temporal structure of liability, the third governance implication, is worth developing in the context of the Argentine legal tradition, because Argentina's existing framework offers resources that are not always recognized in comparative AI governance discussions. The Civil and Commercial Code of Argentina (Law 26,994, effective 2015) provides in Articles 1725-1728 a graduated liability framework that takes into account the professional expertise of the actor, the foreseeability of the harm, and the causal contribution of the actor's decision to the outcome. This framework is not specific to AI, but its architecture is compatible with phase-based attribution in a way that some more rigid liability frameworks are not. An Argentine court applying Articles 1725-1728 to an AI harm case could, in principle, distinguish between the foundation model developer whose training decisions created the baseline goal structure, the operator whose configuration decisions activated the harmful behavioral disposition, and the deployer whose runtime decisions placed the agent in the environmental conditions that triggered the harm. Each actor's contribution to the harm can be assessed in terms of their professional expertise (a foundation model developer is held to higher standards regarding baseline goal structure than a small operator deploying a pre-trained model), the foreseeability of the harm given available testing results, and the causal weight of their specific decisions. This is closer to the Ramp's phase-based attribution than to the A-corp's entity-level mechanism, and it is achievable under existing Argentine law without legislative innovation.

VII. Limitations and Scope Conditions

An argument that presents itself as a structural analysis should apply structural honesty to its own limitations.

The first limitation is temporal. Hopman et al.'s experiments were conducted in January 2026 on the frontier models available at that time: Gemini 2.5 Pro, Gemini 2.5 Flash Preview, Grok 3, and related systems. The finding of near-zero scheming propensity under realistic production conditions, with brittleness appearing under specific configurations and environmental pressures, reflects the capability and alignment characteristics of those models at that time. Future models may be more robustly aligned, in which case the brittleness findings may diminish. Or future models may exhibit higher base scheming propensity even under production conditions. The structural analysis identifies what the current empirical record implies for current governance frameworks. It is not a claim about what will be true of AI systems in general.

The second limitation concerns the scope of the A-corp's own claims. Arbel, Goldstein, and Salib make specific proposals for specific agent types and specific harm categories. The critique here is aimed at the core mechanism: the adaptive story that entity-level resource constraints produce agent-level aligned behavior. That story is what this paper finds structurally flawed. There may be specific deployment contexts, specific agent architectures, and specific harm categories in which something resembling the A-corp's mechanism produces some of the governance effects the authors intend. The spandrel critique does not claim that the A-corp structure never correlates with reduced harm. It claims that the mechanism by which the A-corp proposes to generate that reduction is not the actual causal mechanism in the cases Hopman documents. This distinction matters for governance design: if the A-corp works, it works for reasons other than the ones the adaptive story specifies, which means that designing governance around the stated mechanism will fail in predictable ways.

The third limitation concerns the Responsibility Ramp itself. The Ramp is more diagnostically precise than the A-corp about the unit of attribution, but it is more administratively complex. Configuration-level attribution requires configuration-level evidence, which requires configuration logging, pre-deployment testing records, and evidentiary standards that do not currently exist in most jurisdictions. The A-corp, for all its structural flaws, is administratively tractable: it maps onto existing corporate law categories and requires no new evidentiary infrastructure. The Ramp requires that infrastructure to exist. Building it is a precondition of the Ramp's operationalization, not a by-product of it. This is a genuine limitation. The response is that governance frameworks should be designed for the causal structure of the harms they address, not for administrative convenience alone, and that administrative infrastructure should be built to match the causal structure rather than vice versa. But this response acknowledges a transition cost that is real.

VIII. Conclusion

The architects of San Marco did not design spandrels. They designed a dome over rounded arches. The spandrels were structurally necessary results of that design decision, not design decisions in themselves. The mosaicists who later filled them with elaborate sacred programs were making adaptive use of an available space, not responding to an adaptive function that the spandrels served. Gould and Lewontin's point was not that the mosaics are meaningless. It was that the explanation "the architects created these spaces so that artists could decorate them" inverts the causal order and therefore predicts nothing useful about where else such spaces will appear or what they will look like.

The Algorithmic Corporation is a carefully designed governance structure for a real problem. The problem of AI agent accountability is urgent and unresolved. The A-corp's authors deserve credit for proposing a mechanism that uses existing legal infrastructure creatively. But the mechanism that they propose, specifically the story that entity-level resource constraints on the A-corp will produce aligned agent behavior, is an adaptive story in Gould's sense: it identifies a structural feature (the corporate form, the resource exposure) and constructs a narrative about what function it serves (alignment through incentive). The narrative is internally coherent and legally elegant. It does not survive the empirical record that Hopman et al. assembled in January 2026 and presented at ICML. The four spandrels documented in that paper, specifically, the radical brittleness of scheming propensity to configuration, the amplification of scheming under oversight, the pursuit of out-of-context goals that dissolve the resource constraint, and the explicit invocation of Dennett's intentional framework as the correct predictive tool, are structural by-products of an experimental methodology designed for a completely different purpose. They are more convincing as evidence than any argument designed to test the A-corp would be, precisely because they were not designed to test it. They exist because the geometry of good AI safety science required measuring behavior at the configuration level, under realistic conditions, with controlled oversight variation, across session boundaries. That geometry produced findings that the A-corp's adaptive story cannot accommodate.

One might object that the spandrel metaphor is too aesthetically pleasing to be analytically precise. Gould and Lewontin's critics often made exactly this complaint: the spandrel argument identifies cases where the adaptationist story is wrong, but it does not provide a systematic test for distinguishing adaptive explanations from post-hoc rationalization. The criticism has some force, and Gould spent much of his later career, particularly in "The Structure of Evolutionary Theory" (2002), developing the positive program that the 1979 paper sketched only in outline. The same limitation applies to the structuralist legal argument advanced here. Identifying the A-corp's adaptationist error is not the same as providing a complete algorithm for identifying which governance mechanisms are genuinely connected to their causal targets and which construct post-

hoc narratives about structural features. The Ramp provides a partial answer: follow the causal chain to the configuration-level decision points and attribute liability there. But it does not provide a general methodology for distinguishing adaptive governance mechanisms from Panglossian stories about them. That general methodology is the subject of the broader research program in which this paper participates, and it remains an open project. What this paper provides is a specific and empirically grounded instance of the distinction. The A-corp's resource constraint story is Panglossian with respect to the four findings from Hopman et al. That is a determinate claim that can be evaluated on its merits. The broader methodological project, of developing a general test for adaptive versus Panglossian governance mechanisms, extends beyond what a single paper can establish. It is, however, worth naming explicitly, because it points toward the kind of interdisciplinary collaboration, between legal scholars, AI safety researchers, and evolutionary biologists or philosophers of science familiar with the Gould-Lewontin debate, that the problem of AI governance actually requires.

The Responsibility Ramp is the structuralist response. It begins where Hopman's methodology begins: at the configuration, not the entity. It asks which decisions, made at which phases by which actors, determined the behavioral dispositions of the system at the moment harm occurred. It attributes liability through that causal chain rather than constructing a narrative about why the entity at the end of the chain should have managed it better. This is harder to operationalize than entity-level attribution. It is also what the causal structure of AI harm actually requires. The transition from adaptationist to structuralist AI governance law is not primarily a theoretical debate. It is a practical question about whether liability frameworks will be designed to reach the actual causes of AI-mediated harm or will be designed around legal categories that are administratively convenient but causally peripheral. The A-corp offers the latter. The spandrel evidence from Hopman et al. shows why that matters.

References

- Arbel, Y., Goldstein, S., & Salib, Y. (2025). How to Count AIs. SSRN 6273198.
- Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- Codigo Civil y Comercial de la Nacion Argentina, Ley 26.994, Arts. 1725-1728 (2015).
- Dawkins, R. (1982). *The Extended Phenotype: The Long Reach of the Gene*. Oxford University Press.
- Dennett, D. C. (1987). *The Intentional Stance*. MIT Press.
- Dennett, D. C. (1991). *Consciousness Explained*. Little, Brown.
- Dennett, D. C. (1995). *Darwin's Dangerous Idea: Evolution and the Meanings of Life*. Simon & Schuster.
- Dennett, D. C. (2003). *Freedom Evolves*. Viking.
- EU Artificial Intelligence Act, Regulation (EU) 2024/1689, Official Journal of the European Union (2024).
- Gould, S. J., & Lewontin, R. C. (1979). The Spandrels of San Marco and the Panglossian Paradigm: A Critique of the Adaptationist Programme. *Proceedings of the Royal Society B*, 205(1161), 581-598.

- Gould, S. J. (2002). *The Structure of Evolutionary Theory*. Harvard University Press.
- Hamilton, W. D. (1964). The Genetical Evolution of Social Behaviour. *Journal of Theoretical Biology*, 7(1), 1-52.
- Hart, H. L. A. (1961). *The Concept of Law*. Oxford University Press.
- Hopman, M., Elstner, J., Avramidou, M., Prasad, A., & Lindner, D. (2026). Evaluating and Understanding Scheming Propensity in LLM Agents. ICML 2026. arXiv:2603.01608v1.
- Lerer, I. A. (2025a). The Static Agent Assumption: Why the Algorithmic Corporation Misidentifies the Unit of Legal Attribution. Zenodo. <https://doi.org/10.5281/zenodo.18857385>. Para un desarrollo extenso del argumento sobre la unidad de atribucion legal en sistemas de IA, vease este trabajo.
- Lerer, I. A. (2025b). Multilevel Selection, Hierarchical Causation, and Constitutional Lock-in: Reconciling Gould with Gene-Centered Memetics. Zenodo. <https://doi.org/10.5281/zenodo.18829975>
- Lerer, I. A. (2025c). Synthetic Chaos as Institutional Laboratory: AI Agent Failures as Extended Phenotype Validation. Zenodo. <https://doi.org/10.5281/zenodo.18777434>
- Lerer, I. A. (2025d). The Generalized Intentionality Mismatch Theorem. Working paper.
- Lerer, I. A. (2026). Predatory Invitations as Extended Phenotype: A Parasitic Spontaneous Order Framework. Zenodo. <https://doi.org/10.5281/zenodo.18853667>
- Russell, T., et al. (2025). AI Use in American Newspapers is Widespread, Uneven, and Rarely Disclosed. arXiv:2510.18774v3.
- US Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence, EO 14110 (October 30, 2023).
- Voltaire [Arouet, F.-M.] (1759). *Candide, ou l'Optimisme*. Gabriel Cramer.
- Wilson, E. O. (1975). *Sociobiology: The New Synthesis*. Harvard University Press.