

Intentionality Mismatch in Commercial Liability: A Four-Level Evolutionary Analysis of *Tapia Araya v. Starbucks* (Argentine Supreme Court, 2026)

Ignacio Adrián Lerer
Independent Researcher
Buenos Aires, Argentina
adrian@lerer.com.ar
ORCID 0009-0007-6378-9749
March 2026

Abstract

The Argentine Supreme Court unanimously reversed a lower court ruling that held Starbucks liable for an armed robbery committed by a third party inside one of its Buenos Aires locations (CSJN, March 5, 2026). The Buenos Aires Civil Appellate Chamber (Sala D) had reasoned that armed robberies are statistically foreseeable in Buenos Aires, and therefore do not constitute *caso fortuito* (force majeure) under Articles 1730-1731 of the Civil and Commercial Code. The Supreme Court rejected this reasoning on grounds of arbitrariness. Justice Lorenzetti issued a separate concurrence articulating a proportionality standard: security obligations extend only to risks preventable with reasonable means proportionate to the business activity and legally permitted under local regulation.

I analyze the case through three integrated frameworks developed in prior work: Asymmetric Intentionality Theory (AIT), Extended Phenotype Theory (EPT) applied to legal norms, and multilevel Evolutionary Game Theory (EGT). The central finding is that the appellate chamber committed an intentionality mismatch error: it imposed Level-3 cognitive obligations (foresee, anticipate, guarantee security) on a Level-1 optimizer (the corporation), which can only respond to cost signals. The predictable corporate response to such a rule is not improved security but cost transfer, market exit from vulnerable neighborhoods, and defensive litigation, the opposite of the regulatory intent. Justice Lorenzetti's proportionality standard constitutes an evolutionarily stable strategy (ESS) because it aligns regulatory demands with the actual decision architecture of the regulated entity.

The paper introduces a four-level EGT framework distinguishing individual humans (Level 1), corporate organizations (Level 2), judicial institutions (Level 3), and the liability regime ecosystem (Level 4). Traditional doctrinal analysis conflates Levels 1 and 2, treating corporations as moral agents capable of normative internalization. This conflation produces the design error that the appellate chamber exemplified and the Supreme Court corrected.

Keywords: Asymmetric Intentionality Theory, Extended Phenotype Theory, Evolutionary Game Theory, commercial liability, *caso fortuito*, consumer protection, Argentine Supreme Court, institutional design

1. Introduction

1.1 The case

A university student was working on his laptop inside a Starbucks cafe located at the intersection of Avenida Cabildo and Roosevelt in Buenos Aires, Argentina, accompanied by a classmate. An armed individual entered the establishment, displayed what appeared to be a firearm, shouted threats (“don’t do anything or I’ll shoot”), and seized the laptop. The victim filed both criminal and civil claims. The civil suit sought damages of ARS 112,449 plus interest from Starbucks Coffee Argentina SRL for breach of its security obligation under consumer protection law (Art. 5, Ley 24.240; Art. 42, Argentine Constitution) and the necessary deposit provisions of the Civil and Commercial Code (Arts. 1368-1375).

The trial court (Juzgado Nacional de Primera Instancia en lo Civil No. 78) ruled for the defendant, finding that the armed robbery constituted *caso fortuito*. On appeal, Sala D of the Buenos Aires Civil Appellate Chamber (Camara Nacional de Apelaciones en lo Civil) reversed by majority, holding Starbucks liable. The chamber reasoned that the security obligation arising from consumer protection law and the necessary deposit rules required the entrepreneur to adopt adequate security measures, to control their correct functioning, and to train personnel for effective surveillance. Armed robberies, the chamber held, are foreseeable because they “occur daily” in Buenos Aires, and the frequency with which such crimes are committed renders them neither unforeseeable nor unavoidable within the meaning of Articles 1730-1731 of the Civil and Commercial Code. The chamber did not identify what specific security measures Starbucks should have implemented to prevent an armed robbery, nor did it address CABA Law 1913, which prohibits private security personnel from carrying firearms in public-access spaces.

The dissenting member of the appellate panel took the opposite view: an armed assault by a stranger entering a commercial establishment constitutes a third-party act that cannot be automatically attributed to the defendant, and the use of a firearm made the event inevitable and irresistible, satisfying the requirements of Article 1731 of the Civil and Commercial Code.

Starbucks filed an extraordinary appeal (*recurso extraordinario*) alleging arbitrariness, which was denied, prompting a direct appeal to the Supreme Court (*recurso de queja*). On March 5, 2026, the Argentine Supreme Court (CSJN) unanimously reversed the appellate ruling, though with differing emphases.

Justices Rosatti and Rosenkrantz issued a joint majority opinion. They found that the appellate reasoning rested on a “dogmatic affirmation” that armed robbery is neither unforeseeable nor unavoidable, without explaining how Starbucks could have reasonably prevented the entry of an armed individual, particularly given the legal

prohibition on armed private security in public-access venues. They cited the precedent *D'Odorico* (Fallos: 326:107, 2003), where the Court had established that the irresistible force resulting from firearm possession makes any hypothesis of resistance inside a premises unthinkable, constituting an inevitable event within the meaning of the former Article 514 of the Civil Code, analogous to the current Article 1731.

Justice Lorenzetti issued an eleven-paragraph separate concurrence that went substantially beyond the majority's reasoning. Three features distinguish it. First, he separated abstract statistical foreseeability from concrete inevitability: that robberies "occur daily" is a descriptive statement about Buenos Aires crime rates, not a finding that each individual robbery was preventable by each proximate business. Second, he introduced a proportionality standard: the security measures demanded of a commercial establishment must be "reasonable and proportionate" to its business activity. A cafe is not a bank, and its security infrastructure should reflect its commercial nature. Third, he observed that the use of a firearm poses a concrete threat to the human persons (employees) inside the establishment, who are in the "same degree of defenselessness" as the customer victim, and whose physical risk cannot be transferred to satisfy a liability standard. Finally, Lorenzetti articulated a systemic proposition: the security obligation does not constitute a "guarantee of result"; the causation regime takes into account "normal and ordinary foreseeable consequences, exempting the responsible party from those that are inevitable or unforeseeable."

1.2 The doctrinal puzzle

Traditional tort analysis frames the dispute as a technical question: does armed robbery by a third party fall within or outside the *caso fortuito* exemption (Arts. 1730-1731, Civil and Commercial Code)? The appellate chamber constructed a syllogism with four steps. First, entrepreneurs must adopt adequate security measures (Art. 5, Ley 24.240). Second, armed robberies occur daily in Buenos Aires (factual premise). Third, therefore, armed robberies are foreseeable (inferential conclusion). Fourth, if foreseeable, the event does not constitute *caso fortuito*, and the entrepreneur is liable (legal consequence).

The Supreme Court rejected this syllogism but the two opinions diverge in their depth of engagement with the underlying problem. The majority identified a defect in the appellate reasoning (failure to specify omitted measures, contradiction with Law 1913) and applied the precedent *D'Odorico*. This is corrective in the narrow sense: it reverses an incorrect outcome. Lorenzetti identified a structural problem: the appellate chamber collapsed three distinctions that the legal system requires. The distinction between abstract foreseeability and concrete inevitability. The distinction between security measures proportionate to a specific commercial activity and security measures appropriate to a high-value target. The distinction between obligations that the legal system permits the obligor to fulfill and obligations it simultaneously prohibits.

Neither the majority nor Lorenzetti developed a further objection that deserves attention: the factual predicate of the necessary deposit regime was not satisfied. The appellate chamber founded its liability analysis on Articles 1368-1375 of the Civil and

Commercial Code, which extend the innkeeper's liability regime to "other establishments" for damages to effects "introduced" by customers. This regime presupposes that the establishment assumes custody of the customer's goods, a constructive deposit arising from the act of introduction. But the victim's laptop was not deposited with Starbucks in any meaningful sense. He was actively using it on the table at the moment of the robbery. He retained direct, physical, uninterrupted possession throughout. There was no transfer of custody (*tradición*), no acceptance by the depositary, and no assumption of custodial responsibility by the establishment. The victim himself, not the establishment, was the person in control of the object when the assailant seized it. The appellate chamber applied a custodial liability regime to a situation where custody was never transferred. This is not a minor technical objection; it goes to the factual predicate of the legal institute invoked. If there is no deposit, the liability framework built upon it lacks foundation, regardless of how one resolves the *caso fortuito* question. The Supreme Court chose to resolve the case on *caso fortuito* grounds, which was sufficient for reversal, but the absence of the deposit relationship represents an independent ground on which the appellate reasoning fails.

I propose that the structural problem Lorenzetti identified has a deeper root than doctrinal confusion: it is a mismatch between the intentionality architecture of the norm's target (a corporate optimizer that processes cost signals) and the intentionality level that the norm presupposes (a moral agent capable of cognitive foresight, normative internalization, and guarantee-level commitment to others' safety). This mismatch is not correctable by better legal reasoning within the existing doctrinal framework; it is a design error in how the legal system models corporate actors.

1.3 Research questions

This paper addresses three questions. First, diagnostic: why does the appellate chamber's rule fail at a level deeper than arbitrariness? I argue it fails because it regulates a Level-1 optimizer as if it were a Level-3 agent, producing behavioral responses opposite to those intended. Second, predictive: why is Justice Lorenzetti's proportionality standard more evolutionarily stable? I show it aligns regulatory demands with the actual decision architecture of the regulated entity, satisfying the conditions for an ESS in multilevel EGT. Third, methodological: what does a four-level evolutionary analysis reveal that bilateral doctrinal analysis misses? I demonstrate that the appellate analysis, focused exclusively on the plaintiff-defendant dyad, ignores corporate response dynamics (Level 2), judicial selection mechanisms (Level 3), and ecosystem sustainability (Level 4).

1.4 Theoretical frameworks

The analysis integrates three frameworks developed in prior publications. A summary more detailed than a passing citation but less exhaustive than a full re-derivation is necessary for self-containment; the full formal development appears in the cited works.

1.4.1 Asymmetric Intentionality Theory (AIT)

AIT, introduced in Lerer (2025a), classifies actors by their intentionality architecture, building on Dennett's (1987) intentional stance framework. The classification is not a spectrum but a typology of qualitatively different decision structures.

Level-0 agents (thermostats, simple reflexes) respond to stimuli via fixed rules and are irrelevant to legal analysis. Level-1 optimizers (corporations, algorithms, automated trading systems) maximize utility functions subject to constraints. They have goals and beliefs but do not model others' mental states. A corporation maximizes profit subject to legal, operational, and market constraints. It processes a court judgment as a change in the cost parameter, not as a moral signal. A fine is dC/dx (marginal cost increase), not social disapproval. Level-2 agents (capable of deception, basic teaching) model others' beliefs: "I believe you believe X." Level-3 agents (humans in social contexts) model others modeling others: "I believe you believe I believe X." They feel shame, reciprocate trust, internalize norms, calibrate behavior based on social feedback, and experience moral emotions (guilt, indignation, empathy).

The critical distinction for legal purposes is between Level-1 and Level-3 entities. Norms designed for Level-3 agents invoke moral imperatives (duty of care, foreseeability standards, good faith obligations) and expect internalization: the actor will comply not merely because violation is costly but because the norm is recognized as legitimate. This expectation is reasonable for human agents, who can and do internalize legal norms through socialization, education, and moral reasoning. It is unreasonable for Level-1 optimizers, which lack the cognitive architecture for norm internalization. A Level-1 optimizer "complies" with a norm only when compliance is the cost-minimizing strategy; it "violates" the norm when violation is cheaper than compliance. This is not moral failure; it is architectural limitation.

AIT generates a specific prediction about regulatory failure: when a norm imposes Level-3 cognitive demands (foresee, anticipate, care, guarantee) on a Level-1 target (corporation), the target's behavioral response will diverge from the regulator's intent. The regulator expects norm internalization (improved behavior). The optimizer delivers cost minimization (strategic adjustment). This divergence is the intentionality mismatch, and it is the central diagnostic concept applied to *Tapia*.

The empirical foundation for AIT's corporate classification draws on evidence from multiple domains. Garrett (2014) documents that deferred prosecution agreements (DPAs), designed to incentivize genuine corporate reform through Level-3 mechanisms (promise of ethical improvement, monitoring by independent compliance experts), produce minimal behavioral change in corporate entities while imposing substantial compliance costs. King and Lenox (2000) show that voluntary environmental compliance programs (the chemical industry's Responsible Care program) have near-zero effect on corporate emissions; the Level-3 appeal to environmental responsibility fails to alter Level-1 optimization behavior. Potoski and Prakash (2005) find similar patterns across ISO 14001 environmental certification. The consistent finding across these studies is that mechanisms designed for Level-3 agents (moral suasion, voluntary

commitment, reputational incentive) fail when applied to Level-1 entities, while mechanisms designed for Level-1 entities (defined fines, mandatory standards, structural remedies) succeed regardless of the entity's intentionality level.

1.4.2 Extended Phenotype Theory (EPT)

EPT, applied to legal systems in Lerer (2025b), models legal norms as cultural replicators whose fitness depends on their capacity to generate and sustain institutional support structures. Dawkins (1982) proposed that a gene's phenotypic effects extend beyond the organism's body: beaver dams, spider webs, cuckoo eggs mimicking host species. I extend this logic to legal norms.

A Supreme Court ruling is not merely text; it produces doctrinal phenotypes (citations, treatises, law review articles, future judicial applications that cite it as authority), institutional phenotypes (compliance departments redesigned to conform, insurance products priced to reflect the new rule, litigation strategies adapted to the new standard), and behavioral phenotypes (firms adjust protocols, consumers adjust expectations, lawyers advise clients based on the new precedent). The fitness of a legal rule depends on whether it generates a self-sustaining ecosystem of these phenotypic expressions.

Rules that impose impossible obligations fail to generate sustainable phenotypic support. If a rule demands that cafes prevent armed robbery, no compliance department can design a procedure that satisfies the obligation, no insurer can price the unbounded risk, and no defense lawyer can advise a client on how to comply. The rule's extended phenotype is stillborn. Rules that demand proportionate, feasible conduct generate durable phenotypic support: compliance manuals specify the required measures, insurers price the bounded risk, defense counsel documents compliance, and lower courts apply the standard consistently.

The EPT perspective explains why the Supreme Court's reversal in *Tapia* was overdetermined: the appellate rule could not generate the institutional support necessary for survival. It was pruned not because the Supreme Court disagreed on normative grounds (though it may have), but because the rule was environmentally unfit, it could not sustain itself in the institutional ecosystem it inhabited.

1.4.3 Multilevel Evolutionary Game Theory

Multilevel EGT, developed across the research program (Lerer 2025c, 2025d), separates nested levels of selection that traditional legal analysis conflates. Standard EGT (Maynard Smith 1982) asks whether a strategy is evolutionarily stable at a single level. Multilevel EGT recognizes that legal systems involve simultaneous selection at multiple levels: individual actors optimize personal welfare, organizations optimize institutional objectives, rule-producing institutions compete for doctrinal dominance, and the regime ecosystem imposes sustainability constraints on all of the above.

A strategy that is optimal at one level may be catastrophic at another. The appellate chamber's rule was arguably optimal at Level 1 (the individual victim receives

compensation) but catastrophic at Level 4 (the commercial ecosystem collapses). Bilateral analysis, which examines only Level 1, cannot detect this cross-level instability. Only multilevel analysis, which tracks the cascading effects across all four levels, reveals the full consequence of a rule choice.

The present paper's specific contribution is the application of a four-level structure to a concrete judicial ruling, demonstrating how the framework provides diagnostic capacity unavailable to standard doctrinal tools.

1.5 Contribution and structure

This is, to my knowledge, the first application of AIT to judicial doctrine in tort law. The paper demonstrates how intentionality mismatch analysis identifies a failure mode invisible to standard doctrinal tools: the appellate chamber's reasoning was not merely poorly supported (the majority's diagnosis) but structurally misconceived (it regulated the wrong type of entity with the wrong type of norm). The four-level EGT framework provides diagnostic precision absent in traditional analysis, which conflates corporate entities (Level 2) with human agents (Level 1) and fails to separate judicial selection mechanisms (Level 3) from ecosystem sustainability (Level 4).

The paper proceeds as follows. Section 2 develops the four-level analysis of *Tapia*. Section 3 diagnoses the intentionality mismatch error in detail. Section 4 analyzes evolutionary stability. Section 5 reconstructs Lorenzetti's three-part test as institutional innovation. Section 6 offers a brief comparative perspective. Section 7 derives implications for legal practice and design. Section 8 discusses limitations. Section 9 concludes.

2. Four-Level Analysis

Standard tort analysis treats *Tapia* as a bilateral dispute between a consumer and a commercial establishment. This framing captures at most one of the four levels of selection operating on the case.

2.1 Level 1: Individual humans

All individual actors in the case are Level-3 agents in AIT classification. The victim modeled the assailant's intentions ("if I resist, he will shoot"), experienced phenomenological harm (fear, loss, indignation), and sought reparation through the legal system. His classmate, who testified as a witness, corroborated the events and contributed to the evidentiary record. The Starbucks employees present during the robbery were also Level-3 agents: they experienced fear, calculated self-preservation options, and complied with the assailant's demands.

The assailant is also a Level-3 agent, but one who defects against social norms. His capacity to model victims' reactions, exhibiting the firearm to paralyze resistance, calculating the moment of entry, directing threats to prevent intervention, is precisely what makes the robbery effective. It is also what makes the event unpreventable by a Level-1 optimizer. The assailant operates with intentional sophistication (modeling the

behavior of everyone in the room); the corporation operates with parametric optimization (adjusting cost variables). The asymmetry is structural.

A point the civil liability discussion rarely makes explicit deserves emphasis: the rule governing commercial liability has zero effect on the actor who produced the harm. The assailant does not modify his behavior because Starbucks pays or does not pay an indemnity. He does not know about the civil suit. He does not care. The norm operates exclusively between two actors who did not cause the event, redistributing losses that neither generated. This observation does not invalidate tort law (loss redistribution serves legitimate purposes), but it places a constraint on what tort rules can accomplish: they cannot deter third-party criminal acts. They can only allocate losses *ex post*. The appellate chamber's reasoning implicitly assumed otherwise, treating the liability rule as if it could incentivize crime prevention by the commercial establishment. It cannot, because the actor whose behavior needs to change (the criminal) is outside the norm's causal reach.

The individual-level analysis reveals one additional datum. The victim, in his police testimony, described the event as "surprising and violent" (*sorpresivo y violento*). This characterization, recorded contemporaneously and under oath, sits uneasily with the appellate chamber's holding that such events are foreseeable and therefore preventable. The victim himself perceived the robbery as a shock, not as a predictable occurrence. The subjective experience of the Level-3 agent most directly affected by the event contradicts the factual premise of the appellate reasoning.

2.2 Level 2: The corporate organization

Starbucks Coffee Argentina SRL is a Level-1 optimizer in AIT classification. This is not a normative judgment about what corporations should be; it is a descriptive claim about how they process information and make decisions. The corporation maximizes a profit function subject to legal and operational constraints. It processes legal judgments as cost signals: a ruling imposing liability for armed robbery enters the decision calculus as an expected cost (frequency of robberies multiplied by expected judgment multiplied by probability of adverse ruling). It does not "foresee" armed robbery in the cognitive sense the appellate chamber attributed to it. It calculates expected liabilities and adjusts pricing, insurance coverage, and location decisions accordingly.

The distinction between Level-1 and Level-3 responses to the same legal rule is the analytical core of the mismatch diagnosis. Consider the appellate chamber's holding that Starbucks should "control the correct functioning of the implemented security mechanisms and the training of the personnel in whom it delegates effective surveillance." If Starbucks were a Level-3 agent, it might respond by genuinely improving its security culture: training employees to identify threats, investing in preventive measures out of moral commitment, building a workplace ethos of customer protection. If Starbucks is a Level-1 optimizer, it responds differently. It compares the marginal cost of enhanced security with the marginal reduction in expected liability. If the expected liability cost under the appellate rule is, say, a few thousand pesos per year per location (robbery frequency multiplied by judgment amount multiplied by

probability of conviction), and the cost of security upgrades exceeds that figure, the optimizer does not invest. Instead, it absorbs the expected cost through minor price increases, contracts insurance coverage, or adds the litigation cost to its annual legal budget.

Critically, the one response that might actually reduce the risk, deploying armed security capable of confronting an armed assailant, is legally unavailable. CABA Law 1913 prohibits private security personnel from carrying firearms in public-access venues. The corporation thus faces a constraint that makes effective prevention legally impossible. The rational Level-1 response to an obligation that cannot be satisfied with legal means is not heroic compliance; it is strategic adjustment of the cost structure.

Lorenzetti recognized this dynamic, though not in AIT terminology. His observation that commercial establishments “cannot be civilly liable for acts that are not feasible to be ‘avoided’ with the reasonable and proportionate resources that must be implemented” (Considerando 9) maps directly onto the Level-1 optimization constraint: liability should attach only to outcomes within the actor’s optimization frontier. Armed robbery with a firearm lies outside that frontier for a cafe.

His observation about employees deserves separate attention. “The use of a firearm entails a concrete threat to human persons whose risk also cannot be transferred to the employees of the establishment who are in the same degree of defenselessness” (Considerando 9). In AIT terms, the employees are Level-3 agents (biological organisms with self-preservation instincts, fear responses, and limited physical capacity to resist armed violence) trapped within the decision structure of a Level-1 optimizer. If the liability rule requires the establishment to prevent armed robbery, the only humans available to execute that obligation are the employees. But the employees’ Level-3 architecture (survival instinct, fear of death) overrides any Level-1 instruction to “stop the robber.” The rule creates an obligation that no human agent within the organization can fulfill, and no technological or procedural measure available to a cafe can substitute for armed intervention.

The behavioral predictions for a Level-1 optimizer under the appellate rule are four. First, security theater: installing visible but ineffective measures (additional cameras, warning signs) to satisfy judicial scrutiny without reducing the actual probability of armed robbery, since cameras do not deter masked assailants and warning signs do not deter individuals who have already decided to commit armed violence. Second, commercial relocation: exiting high-crime neighborhoods where expected liability costs exceed the commercial return of operating a location, producing service deserts in precisely the communities most in need of commercial access. Third, price increases: distributing the expected liability cost across all consumers as a risk premium embedded in product prices. Fourth, systematic litigation: contesting every robbery-related claim to reduce the probability of adverse rulings, increasing transaction costs for the judicial system and for injured consumers. None of these responses constitute “improved security,” yet all are optimal for a Level-1 optimizer facing the appellate rule.

2.3 Level 3: Judicial institutions as selection mechanisms

The appellate chamber and the Supreme Court function as producers of competing jurisprudential variants. In EPT terms, each ruling is a memetic variant that either generates institutional support (citations, doctrinal adoption, insurance product adaptation) or fails to do so and is discarded. Understanding this level requires identifying the two memplexes in collision and the selection mechanisms that determine which variant survives.

2.3.1 *The consumer protection memplex and its extended phenotype*

The consumer protection memplex in Argentina comprises Article 42 of the Constitution (right to safety, adequate information, and equitable treatment in consumer relations), Ley 24.240 (security obligation, pro-consumer presumptions, cost-free litigation for consumers under Art. 53), the expansive jurisprudence of consumer courts, and the institutional apparatus of consumer associations and the Defensoría del Pueblo. This memplex has been expanding for three decades, and its extended phenotype is substantial: specialized courts and judges, academic programs in consumer law, professional organizations dedicated to consumer defense, a body of doctrine that reinforces the expansive reading of consumer rights, and a litigation infrastructure that reduces access barriers for individual claimants.

Each component of the extended phenotype reinforces the others. Consumer law academics produce doctrine that consumer court judges cite. The judges produce rulings that the academics analyze in their publications. Consumer associations bring test cases that generate favorable precedent. The Defensoría del Pueblo intervenes in cases to support expansive interpretations. Law schools teach consumer protection as a constitutional imperative, producing graduates predisposed to expansive readings. The memplex replicates not through conspiracy but through decentralized cooperation among actors who share a genuine commitment to consumer welfare.

The chamber's ruling was a variant within this memplex: push the security obligation boundary until the *caso fortuito* exemption vanishes for any crime that occurs with statistical regularity. If adopted and replicated by other panels, this variant would have transformed Argentine commercial liability into a de facto strict liability regime for third-party criminal acts, achieved not through legislation but through incremental doctrinal expansion. The expansion mechanism is characteristic of common-law-style doctrinal development even within civil law systems: each case that extends the duty generates a precedent that the next case cites to extend it further, creating a ratchet effect that operates in one direction (more liability) with no institutional mechanism for reversal short of Supreme Court intervention or legislative amendment.

The chamber's doctrinal construction rested on a further assumption that neither the majority nor Lorenzetti examined: that the necessary deposit regime (Arts. 1368-1375 CCyCN) applied to the facts. This regime, originally designed for innkeepers who receive guests' luggage into their custody, was extended by the appellate chamber to cover a customer's laptop that the customer was actively using on a cafe table. The extension is analytically suspect. The necessary deposit regime presupposes a transfer

of custody from the depositor to the depositee: the guest leaves luggage at the hotel, the hotel assumes responsibility. In *Tapia*, the victim never relinquished control of his laptop. He was typing on it when the assailant seized it. Starbucks never assumed custodial responsibility for the device because the device was never in Starbucks's custody. The chamber constructed its liability framework on a deposit that did not exist, using a legal institute whose factual predicate was unsatisfied. This error precedes and is independent of the *caso fortuito* analysis: even if the chamber were correct that armed robbery does not constitute *caso fortuito*, the liability framework it invoked requires a custodial relationship that the facts did not establish. The entire doctrinal edifice rests on a foundation that was never laid.

2.3.2 The civil liability memplex and its extended phenotype

The civil liability memplex, by contrast, includes the *caso fortuito* exemption (Arts. 1730-1731), the proportionality requirements of general tort law, the precedent *D'Odorico* (2003), the insurance market's pricing assumptions (which presuppose bounded liability for commercial premises), and the doctrinal tradition distinguishing between risks inherent to an activity and risks imposed by external actors. This memplex's extended phenotype is also substantial: insurance products priced on the assumption that armed robbery is an excluded peril or at least a *caso fortuito*, defense bar practices built around the exemption, a doctrinal consensus that the establishment's obligation is of means not of result, and a body of precedent reaching back to the original Civil Code that treats irresistible third-party violence as exonerating.

The two memplexes coexist in normal conditions because they operate on different domains: consumer protection governs the quality and safety of the product or service itself (defective merchandise, misleading advertising, dangerous installations), while the civil liability memplex governs the allocation of losses from external events (natural disasters, criminal acts, force majeure). *Tapia* represents a case where the consumer protection memplex attempted to colonize the civil liability domain, absorbing armed robbery into the category of risks that the entrepreneur must prevent. The collision is not between two equally valid interpretations of the same norm; it is between two norm-ecosystems with incompatible premises about the scope of commercial responsibility.

2.3.3 The Supreme Court as selection mechanism

Tapia exposed this memplex collision and forced a selection event. The Supreme Court operated as a selection mechanism, rejecting the mutant variant before it could generate the institutional support necessary for stabilization.

The majority opinion performed what in evolutionary biology is called purifying selection: eliminating a deleterious mutation before it propagates. It identified the appellate reasoning as dogmatic, cited the controlling precedent *D'Odorico*, and reversed. This is minimalist selection: it says what the rule is not (strict liability for armed robbery) without articulating what it is (what security measures are required, under what standard).

Lorenzetti performed stabilizing selection: not merely pruning the mutant but articulating the stable alternative. His three-part test (foreseeability calibrated to the specific activity, proportionality of preventive measures, legal permissibility of the demanded conduct) constitutes a positive rule that lower courts can apply to novel situations. It is a second-order rule: it does not merely decide this case but establishes the selection criteria for future judicial variants.

The difference between the majority and Lorenzetti corresponds to different levels of institutional engagement. The majority operated at Levels 1 and 3: resolving the bilateral dispute (Level 1) and correcting the lower court's doctrinal error (Level 3). Lorenzetti operated at all four levels: resolving the dispute (Level 1), recognizing the corporate nature of the defendant and its limited prevention capacity (Level 2), setting methodological standards for future cases (Level 3), and formulating a systemic principle of bounded liability that preserves the ecosystem (Level 4).

This distinction matters for the practical effect of the ruling. The majority's opinion corrects the immediate error but provides limited guidance for future cases: lower courts know that the specific syllogism used by the Sala D is invalid, but they do not know what alternative framework to apply. Lorenzetti's opinion provides the framework. When the next armed robbery case reaches a lower court, whether in a restaurant, a shopping mall, or a parking garage, the three-part test offers a structured analytical pathway. This is the difference between pruning a single defective branch and reshaping the tree's growth pattern.

2.4 Level 4: The liability regime ecosystem

The fourth level is the emergent system comprising consumer protection norms, the *caso fortuito* doctrine, insurance markets, commercial real estate patterns, and the criminal justice system. No single actor inhabits this level; it arises from the interactions across Levels 1-3 and constrains the viability of rules produced at Level 3.

The sustainability condition for this ecosystem is that four components must remain functional simultaneously. Commercial activity must remain viable: businesses must be able to operate, insure themselves against foreseeable risks, and generate returns sufficient to justify investment. Consumer access must be preserved: businesses should not systematically exit the neighborhoods where consumers live and work. Insurance markets must remain solvent: underwriters must be able to price the risks they cover without catastrophic exposure to tail events. Judicial capacity must be maintained: the volume and complexity of litigation must not overwhelm the courts.

Under the appellate chamber's rule, each component comes under stress. Insurance underwriters face an unbounded risk: if every armed robbery triggers commercial liability, the frequency is stochastic (it varies with economic conditions, policing levels, and social dynamics), and the severity per event is uncertain (judgments vary by court, by plaintiff circumstances, and by judicial mood). During economic crises, when crime rates spike, the expected liability cost per location could multiply several times, producing catastrophic losses for insurers. The rational response is prohibitive premiums or market exit, leaving businesses unable to insure.

Commercial establishments face an expected liability cost that is inversely correlated with neighborhood safety: the more dangerous the area, the higher the robbery frequency, the higher the expected cost, the lower the commercial return. The optimizer's response is relocation to safer areas, creating commercial deserts in vulnerable neighborhoods. The irony is acute: the consumer protection norm, by making business unviable in high-crime areas, removes commercial access from the consumers most in need of protection.

Litigation volume increases because every robbery generates a civil claim regardless of the establishment's conduct. If liability is automatic upon robbery, the only defense is *caso fortuito*, which the appellate rule eliminated. Every case must be litigated on the merits (what security was implemented? was it "adequate"?), requiring expert testimony, forensic evaluation of security systems, and fact-intensive adjudication. The judicial economy effect is a docket overwhelmed by factually complex cases whose outcome is predetermined by the strict liability framework.

If one carries the logic to its endpoint and imagines that all commercial establishments hire armed security (setting aside the legal prohibition), the result is an arms race. Criminals respond to armed guards by bringing heavier weapons, attacking in larger groups, or targeting employees rather than customers. The escalation dynamic has no equilibrium short of militarized commercial spaces, which create crossfire risks for bystanders and transform consumer environments into security zones. The consumer experience that Starbucks sells, a relaxed space for work and socialization, is incompatible with the security posture the appellate rule would require.

Under Lorenzetti's rule, the ecosystem remains sustainable. Insurance covers negligent security omissions (bounded, priceable risk) while excluding armed robbery where proportionate security is in place (*caso fortuito* defense succeeds). The expected liability for a cafe with cameras, alarms, and protocols is minimal, because the residual risk (the 10% of cases where negligence might be found) is bounded and quantifiable. Businesses remain in high-crime areas because the risk premium is manageable. Litigation is limited to cases where plaintiffs allege specific, identifiable negligence (a camera that was broken and not repaired, an alarm system that was disconnected, an exit route that was blocked), producing a manageable caseload of fact-specific disputes.

Lorenzetti formulated the ecosystem-level insight with precision: "It is not possible to affirm the existence of a guarantee of result, such that the user suffers no harm. The causation regime takes into account normal and ordinary foreseeable consequences, exempting the responsible party from those that are inevitable or unforeseeable" (Considerando 9). In EGT terms, ecosystem sustainability requires bounded liability. Absolute guarantees are incompatible with commercial activity, because commercial activity operates in an environment that includes risks generated by third parties over whom the commercial actor has no control.

2.5 Cross-level dynamics

The four levels interact in feedback loops that bilateral analysis cannot capture.

Level 1 feeds Level 2: individual harm (the robbery victim) generates a cost signal for the corporate optimizer (the judgment amount). Level 2 feeds Level 4: corporate responses to the cost signal (price increases, relocations, insurance purchases) alter the ecosystem. Level 3 feeds Level 4: the Supreme Court's rule selection stabilizes or destabilizes the system. Level 4 feeds Level 1: ecosystem effects (commercial deserts, higher prices, reduced access) alter the welfare of the individual consumers the norm was designed to protect.

The appellate chamber's analysis operated at Level 1 only: the victim deserves compensation for the harm suffered inside the defendant's premises. Every other level was ignored. The victim deserves compensation is a Level-1 judgment that may be correct in isolation. But when the rule necessary to deliver that compensation destabilizes Levels 2-4, the net effect on consumers as a class is negative. The individual plaintiff wins; the consumer population loses. Multilevel analysis makes this trade-off visible.

3. The Intentionality Mismatch Error

3.1 Textual evidence from the appellate opinion

The appellate chamber's language presupposes Level-3 intentionality in the corporate defendant. The operative passage, quoted by the Supreme Court in Considerando 1, reads: "In general it is considered that the entrepreneur must contribute the necessary security and prevention measures according to his activity; controlling the correct functioning of the implemented security mechanisms and the training of the personnel in whom he delegates effective surveillance."

Each verb cluster presupposes cognitive operations characteristic of Level-3 agents. "Must contribute the necessary security and prevention measures" presupposes that the actor can assess what the situation demands, which requires modeling potential threats (how might a criminal act? what vulnerabilities exist?) and calibrating responses (what measures address those vulnerabilities?). This is threat modeling: a cognitive operation requiring simulation of adversarial behavior, characteristic of Level-3 reasoning.

"Controlling the correct functioning" presupposes ongoing monitoring of whether implemented measures achieve their intended purpose. This requires modeling failure modes (how might a camera fail? how might an alarm be circumvented?) and comparing actual performance against expected performance, a supervisory cognitive loop that requires second-order representation (a model of how the model works).

"Training of personnel in whom he delegates effective surveillance" presupposes that the corporation can instill security consciousness in its employees, a socialization process that requires modeling employees' cognitive and emotional capacities and designing training interventions that alter their behavior. This is a distinctly Level-3 operation: modeling what another mind knows, identifying gaps, and producing inputs (training) that modify the other mind's behavior.

A Level-1 optimizer processes none of these operations in the manner the appellate chamber presupposed. It allocates a security budget based on expected liability costs. It purchases security equipment based on vendor recommendations and industry standards. It contracts training services from external providers who deliver standardized programs. At no point does the corporation “model” criminal behavior, “assess” threats, or “instill” vigilance in the cognitive sense. It processes signals and adjusts parameters.

3.2 Lorenzetti’s implicit recognition

Lorenzetti does not use AIT terminology, but three passages in his concurrence implicitly recognize the mismatch.

First, the employee passage (Considerando 9): “the use of a firearm entails a concrete threat to human persons whose risk also cannot be transferred to the employees of the establishment who are in the same degree of defenselessness.” This recognizes that the humans inside the corporate structure are Level-3 agents with biological self-preservation architecture, not interchangeable modules in a security system. The corporation cannot delegate to its employees an obligation that requires confronting lethal violence, because the employees’ Level-3 architecture (fear, self-preservation) overrides the Level-1 optimizer’s instructions.

Second, the “reasonable and proportionate resources” passage (Considerando 9): “The commercial establishment [...] cannot be civilly liable for acts that are not feasible to be ‘avoided’ with the reasonable and proportionate resources that must be implemented to comply with the service plan.” This recognizes that the corporation’s optimization frontier is bounded: it can allocate “reasonable” resources within its business model, but cannot transform into a security installation. The word “proportionate” implicitly acknowledges the Level-1 nature of the actor: demand only what the optimizer can rationally provide.

Third, the “guarantee of result” passage (Considerando 9): “It is not possible to affirm the existence of a guarantee of result, such that the user suffers no harm.” This rejects the Level-3 expectation (the corporation will “guarantee” safety) in favor of a Level-1-compatible standard (the corporation will implement reasonable measures and accept bounded liability for omissions). The shift from “guarantee” to “reasonable measures” is precisely the shift from Level-3 normative expectations to Level-1-compatible cost signals.

3.3 The deposit fiction as compounding error (*)

The intentionality mismatch (demanding Level-3 conduct from a Level-1 entity) is the primary design error. But the appellate chamber committed a prior error that compounded the mismatch: it applied the necessary deposit regime to a situation where no deposit existed.

The necessary deposit (Arts. 1368-1375 CCyCN) is a custodial regime. Its logic is straightforward: when a customer introduces effects into an establishment and the

establishment assumes custody, the establishment is liable for damages to those effects because it controlled them. The innkeeper who receives a guest's suitcase has physical dominion over it; the guest trusts the innkeeper precisely because the innkeeper has the means and the obligation to protect what was entrusted.

Nothing of this logic applies to a laptop that the customer is actively using on a cafe table. The customer retained uninterrupted physical possession. He was typing on the device at the moment of the robbery. Starbucks never received the laptop, never accepted custodial responsibility, and never had the means to protect it, because the device was not in its control. The assailant took the laptop from the customer, not from Starbucks.

From an EPT perspective, the chamber's move is a case of memetic overextension: taking a legal institute adapted to one ecological niche (custodial relationships) and transplanting it to a niche where its factual prerequisites do not obtain (non-custodial presence of customer property). The transplant fails because the extended phenotype of the custodial regime (compliance procedures designed around reception, storage, and return of goods) cannot attach to a factual situation where none of those operations occurred.

From an AIT perspective, the deposit fiction compounds the intentionality mismatch. The deposit regime assumes a Level-3 relationship between depositor and deposittee: the deposittee "accepts" custody, "commits" to care, and "assumes" responsibility. These are normative operations grounded in voluntary undertaking. Starbucks did not voluntarily undertake custody of the victim's laptop. It sold him a coffee. The appellate chamber attributed to the commercial relationship a custodial dimension it did not possess, layering a fictional custody obligation on top of the already fictional attribution of Level-3 cognitive capacities to a Level-1 optimizer. The result is a double fiction: a corporation that cannot "foresee" is held liable for failing to protect a good it never "received."

The Supreme Court resolved the case without reaching this objection, which was sufficient for the outcome. But the deposit question may prove more consequential than the *caso fortuito* analysis in future cases. If the necessary deposit regime applies whenever a customer brings personal effects into a commercial establishment, then every restaurant patron's coat, every gym user's phone, and every student's backpack in a library generates a custodial obligation enforceable through strict liability. The systemic consequences of this extension would far exceed the armed robbery context of *Tapia*.

3.4 Predicted consequences of the mismatch

The consequences of regulating a Level-1 entity with Level-3 norms follow a predictable pattern documented across regulatory domains (Garrett 2014, King & Lenox 2000, Lerer 2025d).

Security theater is the first predicted response. The corporation installs visible but ineffective measures (additional cameras, warning signs, security guards without

weapons or authority to intervene) to create the appearance of compliance without reducing the actual probability of armed robbery. Cameras do not deter masked assailants; warning signs do not deter individuals who have decided to commit armed violence. But cameras and signs create a litigation record (“Your Honor, we had twelve cameras and posted warnings at every entrance”) that may reduce the probability of adverse judicial outcomes. The investment is rational for a Level-1 optimizer because it targets the probability of liability, not the probability of crime.

Commercial relocation is the second predicted response. The expected liability cost under the appellate rule is a function of robbery frequency, which varies by neighborhood. High-crime neighborhoods produce higher expected costs. The optimizer relocates to lower-frequency areas, reducing expected liability. The neighborhoods that lose commercial establishments are precisely those with the highest crime rates and the greatest need for the services those establishments provide.

Price increases are the third response: distributing the expected liability cost across all consumers as a risk premium embedded in product prices. This is regressive: all consumers pay for the risk that materializes against a few.

Systematic litigation is the fourth response: contesting every robbery-related claim to reduce the probability of adverse rulings, increasing transaction costs for the judicial system and for injured consumers, many of whom enjoy cost-free litigation under Article 53 of Ley 24.240 but face time costs and emotional burdens regardless.

None of these responses constitute “improved security.” Yet all are optimal for a Level-1 optimizer facing the appellate rule. The appellate chamber expected a Level-3 response: Starbucks would “learn” to protect its customers better. The corporate reality is a Level-1 response: Starbucks adjusts its cost structure. The gap between expectation and reality is the behavioral signature of intentionality mismatch.

4. Evolutionary Stability Analysis

An evolutionarily stable strategy (ESS), in Maynard Smith’s (1982) formulation, is a strategy that, once adopted by a population, cannot be invaded by rare mutants. Applied to legal rules, an ESS is a doctrinal position that, once established as precedent, resists displacement by alternative rules because no actor in the system improves by deviating from it.

4.1 The appellate rule fails the ESS conditions

The Nash condition requires that no actor improves by deviating when the population follows the incumbent strategy. Under the appellate rule (strict liability for armed robbery), multiple actors can improve by deviating. Consumers as a class are better off under a regime that preserves commercial access in their neighborhoods than under a regime that provides occasional indemnities but eliminates commercial services. Businesses are better off under a regime that allows them to insure bounded risks than under one that exposes them to unbounded liability. Insurers are better off under a regime that distinguishes between negligent omissions (priceable) and criminal acts by

third parties (not priceable) than under one that conflates them. The Nash condition fails.

The stability condition requires that if a mutant does equally well against the incumbent, the incumbent must do strictly better against the mutant than the mutant does against itself. This condition is not reached because the Nash condition already fails: the appellate rule is not even a Nash equilibrium.

More informally, the appellate rule is self-defeating at scale. As more courts adopt it, the aggregate effect is rising insurance costs, commercial exit from vulnerable areas, and increasing litigation volume, all of which reduce consumer welfare below the level achievable under the alternative rule. A rule that reduces the welfare of the actors it protects cannot be evolutionarily stable.

4.2 Lorenzetti's rule satisfies the ESS conditions

Under Lorenzetti's rule (liability for specific negligence, exemption for armed robbery as *caso fortuito*), the Nash condition holds. No actor improves by deviating. Consumers cannot extract larger recoveries because armed robbery falls outside the liability perimeter; but they retain commercial access and face lower prices. Businesses face bounded, predictable liability costs and therefore maintain operations. Insurers can price the bounded risk. Judges applying the rule produce outcomes consistent with the precedent *D'Odorico* and the statutory text of Articles 1730-1731, avoiding reversal by the Supreme Court.

The stability condition holds because if a mutant variant (like the appellate rule) invades, it produces worse outcomes for all actors when its frequency increases, for the reasons detailed in Sections 2.4 and 4.1. The mutant cannot spread because it is self-defeating.

Lorenzetti's rule is not merely a better rule in some abstract sense of doctrinal elegance. It is the rule that survives selection pressure across all four levels of the system. It is stable at Level 1 (individual actors do not benefit from deviation), at Level 2 (corporations can comply within their optimization frontier), at Level 3 (it coheres with the statutory framework and precedent), and at Level 4 (the ecosystem it generates is sustainable).

5. Lorenzetti's Three-Part Test as Institutional Innovation

Lorenzetti's concurrence extracts from the case a generalizable standard that can be reconstructed as a three-part test for evaluating commercial liability for third-party criminal acts.

Foreseeability calibrated to the specific activity. The question is not whether armed robberies are foreseeable in Buenos Aires (they are, statistically), but whether they are foreseeable as a risk inherent to the specific commercial activity. Armed robbery of a bank or armored transport is a risk inherent to those activities; the nature of the business (handling large quantities of cash, transporting valuables) creates a target that attracts criminal attention in the ordinary course. Armed robbery of a cafe is not

inherent to the activity of selling coffee; it is an opportunistic crime whose probability depends on neighborhood characteristics and policing levels, not on the business's commercial nature. Lorenzetti's distinction between abstract statistical foreseeability and concrete activity-specific foreseeability is the analytical move the appellate chamber failed to make.

Proportionality of preventive measures. The security measures demanded must be "reasonable and proportionate" to the business. Cameras, alarm systems, emergency protocols, and employee training in evacuation procedures are proportionate to a cafe. Armed guards, bulletproof glass, metal detectors, and security checkpoints are not. This proportionality requirement implicitly recognizes the Level-1 nature of the corporate actor: it can be expected to allocate resources within the bounds of its business model, not to transform its commercial premises into a security installation. The proportionality test also provides a concrete metric for judicial evaluation: rather than asking the vague question "Was security adequate?" (which the appellate chamber asked without answering), it asks the specific question "What proportionate measures were available and omitted?" This shifts the burden to the plaintiff in a productive way: the claim must identify a specific, feasible, and proportionate measure that the defendant failed to implement.

Legal permissibility. CABA Law 1913 prohibits armed private security in public-access venues. The legal system cannot simultaneously prohibit a measure and hold the business liable for failing to implement it. Lorenzetti used this constraint to close the logical gap in the appellate chamber's reasoning: even if armed robbery were foreseeable for this business, the only effective countermeasure (armed response capability) is legally unavailable. This condition prevents the legal system from producing contradictory obligations, a type of institutional incoherence that erodes the predictability and legitimacy of the normative framework.

The three-part test functions as what the EGT framework identifies as a second-order rule. It does not merely decide the present case but establishes the selection criteria for future judicial variants. Lower courts applying it can evaluate novel situations (armed robbery of a jewelry store, knife attack in a restaurant, theft from a parking garage) by running each scenario through the three conditions. A jewelry store may fail condition one (armed robbery is arguably foreseeable given the nature of the merchandise). A restaurant with a broken alarm system may fail condition two (a proportionate measure was available and omitted). A parking garage that complied with all legally available security measures passes all three conditions and is exempt. The test produces rule-consistent outcomes across factual variation.

6. Comparative Perspective

6.1 D'Odorico as doctrinal predecessor

The precedent *D'Odorico* (Fallos: 326:107, 2003) established the foundational principle that *Tapia* reinforces. In that case, a client of a real estate services firm was robbed at gunpoint inside the firm's premises. The lower court had held the firm liable. The

Supreme Court reversed, holding that “even if by hypothesis a robbery like the one perpetrated in the case were foreseeable in general terms, it seems clear that having been committed by armed criminals, it was inevitable within the meaning of Article 514 of the Civil Code,” the predecessor to the current Article 1731.

D’Odorico decided under the old Civil Code (Art. 514) what *Tapia* decides under the new Civil and Commercial Code (Art. 1731), confirming doctrinal continuity across the 2015 codification. But *Tapia* goes further in three respects.

First, *D’Odorico* stated the conclusion (armed robbery constitutes *caso fortuito*) without developing the reasoning beyond the observation that firearm-wielding assailants generate irresistible force. Lorenzetti’s concurrence in *Tapia* supplies the reasoning: the three-part test of activity-specific foreseeability, proportionality, and legal permissibility provides the analytical framework that *D’Odorico* assumed but did not articulate.

Second, *Tapia* engages with the consumer protection framework (Art. 42 CN, Ley 24.240, necessary deposit rules) in a way *D’Odorico* did not, because the appellate chamber in *D’Odorico* had not invoked consumer protection norms. Lorenzetti’s opinion in *Tapia* integrates the consumer protection obligation with the *caso fortuito* exemption, demonstrating that they are compatible rather than contradictory: the security obligation is real but bounded by proportionality and legal feasibility.

Third, *Tapia* introduces the regulatory constraint of CABA Law 1913 (prohibition on armed private security) as a structural element of the analysis, not merely a factual consideration. The legal system cannot simultaneously prohibit a measure and hold businesses liable for failing to implement it. This principle of non-contradictory regulation was implicit in *D’Odorico* but receives explicit development in Lorenzetti’s concurrence.

The twenty-three-year span between *D’Odorico* (2003) and *Tapia* (2026) is itself informative from an evolutionary perspective. For more than two decades, the principle established in *D’Odorico* was sufficient to resolve cases of armed robbery in commercial premises. The fact that the Sala D departed from it in *Tapia* suggests that the consumer protection memplex had accumulated sufficient institutional momentum to attempt colonization of the *caso fortuito* domain. The Supreme Court’s response, reversing unanimously and reinforcing *D’Odorico*, suggests that the civil liability memplex retains sufficient selection strength to resist the invasion.

6.2 Structural parallels in other jurisdictions

The tension between expansive consumer protection and bounded commercial liability for third-party criminal acts is not unique to Argentina. Examining how other legal systems handle the same structural problem illuminates the evolutionary dynamics at play.

In the United States, the “totality of the circumstances” test for premises liability (Restatement (Third) of Torts: Physical and Emotional Harm, Section 40) requires

landlords and commercial operators to maintain reasonable security against foreseeable criminal acts, but courts routinely find that random armed attacks by strangers are unforeseeable absent specific prior incidents at the same location. The prior similar incidents test, while criticized as underinclusive, serves an evolutionary function similar to Lorenzetti's activity-specific foreseeability condition: it calibrates the duty to the specific history of the particular premises rather than to the statistical frequency of crime in the jurisdiction.

In Spain, the Supreme Court (STS 1086/2003) has held that armed robbery of a commercial establishment constitutes *fuera de culpa* unless the establishment failed to implement legally required security measures. The Spanish approach mirrors the Lorenzetti test's structure: exemption is the default for irresistible third-party violence, and liability attaches only when specific regulatory requirements are demonstrably unmet.

In Chile, the consumer protection framework (Ley 19.496) imposes a security obligation on commercial providers but Chilean courts have not extended it to cover armed robbery by strangers, treating such events as force majeure under Article 45 of the Civil Code. The Chilean system's low Constitutional Lock-in Index (CLI of approximately 0.24, compared to Argentina's 0.89) facilitates doctrinal flexibility, but on this specific question the outcome converges with Argentina's: armed robbery is treated as force majeure rather than a failure of the commercial establishment's duty.

The cross-jurisdictional pattern is consistent with the evolutionary analysis: rules that attempt to hold commercial establishments liable for all criminal acts generate systemic instability and are eventually corrected by higher courts or constrained by legislative intervention. The stable equilibrium across legal systems is bounded liability for negligent security omissions combined with exemption for irresistible third-party violence. This convergence across jurisdictions with different legal traditions (common law, civil law), different crime rates, and different consumer protection frameworks suggests that the proportionality-based standard is an ESS not merely within Argentine law but across legal systems generally. The selection pressure, the impossibility of preventing armed violence through commercial security measures, is universal.

7. Implications for Legal Design and Practice

7.1 The design principle

The analysis yields a design principle that extends beyond the specific case: the cognitive demands of a legal norm must match the intentionality architecture of its target. Norms that demand Level-3 operations (foresight, care, moral internalization) from Level-1 entities (corporations) produce behavioral responses that diverge from regulatory intent. The divergence follows a predictable pattern: cost minimization, strategic compliance, security theater, and rule gaming. These responses are not pathological; they are optimal for a Level-1 optimizer facing Level-3 demands. The pathology is in the design, not in the response.

For Level-1 targets, effective regulation specifies observable conduct requirements (install cameras, maintain alarm systems, keep incident records, conduct periodic drills) and attaches clear cost signals to specific omissions (defined penalties for identifiable failures). It does not demand “foresight” or “care” in their unoperationalized forms, because Level-1 entities cannot execute those cognitive operations. For Level-3 targets (individual humans), normative language (duty of care, reasonable person standard, good faith) is appropriate because humans can internalize norms through socialization, shame, reciprocity, and moral reasoning.

This principle has implications beyond tort law. Corporate compliance programs under Ley 27.401 (Argentina’s anti-bribery statute), for example, demand that corporations “adopt” integrity programs, “train” personnel in ethical conduct, and “report spontaneously” to authorities. Each of these verbs presupposes Level-3 cognitive operations. AIT predicts, and the evidence in Lerer (2025a) confirms, that corporate compliance programs produce compliance theater (visible but behaviorally inert structures) rather than genuine behavioral change. The *Tapia* case provides a clean illustration of the same mechanism in the tort law domain: the appellate chamber demanded Level-3 conduct from a Level-1 entity and got the Level-1 response (cost adjustment) rather than the Level-3 response (improved security).

7.2 Practical implications for commercial liability

For defense counsel representing commercial establishments: the *Tapia* ruling establishes that *caso fortuito* operates as an exemption for armed robbery in commercial premises that have implemented proportionate security measures. The strategic priority is documenting the specific measures implemented (cameras, alarms, protocols, training records, incident reports) and their proportionality to the business activity. The litigation focus shifts from the generic question “Was security adequate?” to the specific question “What proportionate, legally permitted measure was omitted?” If no such measure can be identified, the defense prevails.

For plaintiffs’ counsel representing consumers: the *Tapia* ruling does not close the door to claims based on security negligence, but it imposes a more demanding evidentiary burden. The plaintiff must identify a concrete, feasible, legally permitted, and proportionate security measure that the defendant failed to implement and whose implementation would have prevented or reduced the harm. Invoking the statistical frequency of crime in Buenos Aires is insufficient. The claim must be specific: this camera was broken and not repaired; this alarm was disconnected; this emergency exit was blocked; this protocol was not followed. The shift from generic invocation of the security duty to specific identification of omitted measures changes the litigation dynamic and requires more detailed factual investigation at the pre-trial stage.

For courts: the three-part test (activity-specific foreseeability, proportionality of measures, legal permissibility) provides a structured framework for evaluating novel cases. The test can be applied sequentially, and failure at any step resolves the inquiry. This sequential structure reduces adjudicatory costs: a court that determines at step one that the risk is not inherent to the specific activity need not reach steps two and

three. A court that reaches step three and finds that the only effective countermeasure is legally prohibited can resolve the case without the factual complexity of evaluating “adequacy” of security measures under step two.

For regulators and legislators: the case demonstrates the risks of expanding liability categories incrementally, without evaluating the systemic consequences. The appellate chamber’s ruling, if sustained, would have transformed commercial liability without legislative authorization, without regulatory impact analysis, and without consideration of the insurance market effects, the commercial location effects, or the judicial economy effects. Regulators considering enhanced security obligations for commercial establishments should specify the concrete measures required, ensure those measures are legally available, and evaluate whether the resulting cost structure is compatible with commercial viability in the jurisdictions where the obligations apply.

7.3 Connection to the broader research program

Tapia provides empirical confirmation of a pattern documented across my research program. The Dennett-Nash Gap (Lerer 2025c) describes the systemic welfare loss that occurs when regulatory mechanisms designed for Level-3 agents are applied to Level-1 entities: the regulator’s intent (improved behavior) and the regulated entity’s response (cost optimization) diverge, producing a gap between expected and actual outcomes. The Generalized Intentionality Mismatch Theorem (Lerer 2025d) formalizes this as a general failure mode across regulatory domains, not limited to criminal liability or consumer protection.

In the specific context of *Tapia*, the Dennett-Nash Gap can be characterized as the difference between the appellate chamber’s expectation (Starbucks will invest in better security to protect its customers) and the actual corporate response (Starbucks will adjust its cost structure through pricing, insurance, relocation, and litigation). The gap is not closable by better judicial reasoning; it arises from the structural mismatch between the norm’s cognitive presuppositions and the target’s decision architecture. Lorenzetti’s proportionality standard closes the gap by reformulating the obligation in Level-1-compatible terms: implement specific, proportionate measures (an observable conduct requirement) rather than “guarantee security” (a Level-3 normative imperative).

The research program now includes applications of AIT to criminal liability for corporations (Lerer 2025a), to constitutional rigidity and institutional lock-in (Lerer 2025b), to the Dennett-Nash Gap in regulatory games (Lerer 2025c), to the generalized impossibility of common knowledge across intentionality levels (Lerer 2025d), and, with the present paper, to tort law and commercial liability. The cross-domain consistency of the intentionality mismatch pattern, the same structural failure appearing in criminal compliance, constitutional reform, regulatory design, and tort liability, provides cumulative evidence that the pattern is not an artifact of case selection but a general feature of legal systems that regulate Level-1 entities with Level-3 norms.

8. Limitations

The frameworks applied here (AIT, EPT, multilevel EGT) are theoretical constructs at an early stage of empirical validation. Several limitations deserve explicit acknowledgment.

8.1 The Level-1 characterization of corporations

The classification of corporations as Level-1 optimizers is a modeling choice supported by the evidence reviewed in Lerer (2025a) and by organizational behavior research showing that voluntary compliance mechanisms have near-zero behavioral effect on corporate entities while producing measurable effects on individual humans (Garrett 2014, King & Lenox 2000, Potoski & Prakash 2005). But the characterization is a simplification. Real corporations are populated by Level-3 agents (managers, employees, board members) who sometimes drive the organization toward Level-3-like responses: genuine concern for customer safety, ethical investment in security beyond the cost-benefit minimum, reputation management driven by moral considerations rather than pure optimization. The Level-1 characterization describes the dominant decision mode at the organizational level, particularly for publicly traded corporations and multinational subsidiaries like Starbucks Argentina, not every individual decision within the organization.

A more nuanced model would treat corporations as hybrid entities whose decision architecture varies by context: Level-1-dominant for strategic decisions (location, pricing, capital allocation) and potentially Level-3-influenced for operational decisions made by individual managers with local discretion. The present analysis applies the Level-1 characterization because the appellate ruling's operative demands (prevent armed robbery, guarantee customer safety) fall squarely within the strategic domain where the Level-1 mode dominates. Whether a more fine-grained intentionality model would alter the analysis for a case involving, say, a manager's decision to ignore a broken security camera (potentially Level-3: the manager "knew" the camera was broken and "chose" not to fix it) is an open question that the present framework does not resolve.

8.2 Qualitative rather than quantitative analysis

The four-level EGT analysis is qualitative rather than quantitative. The claim that Lorenzetti's rule constitutes an ESS is argued analytically, not demonstrated through computational simulation with empirical payoff data. I have not estimated the actual expected liability costs under each rule, the actual insurance premium effects, or the actual commercial relocation patterns, because the data necessary for such estimates (robbery frequency by location and year, judgment amounts by court, insurance pricing models for Argentine commercial liability) is not publicly available for this case. A stronger empirical test would require longitudinal data on commercial location decisions and insurance pricing following doctrinal shifts in tort liability.

The ecosystem-level predictions (insurance market stress, commercial exit from high-crime areas, litigation volume increase) are plausible inferences from the structure of

the rules and the known behavioral patterns of Level-1 optimizers, but they are not measured outcomes. Future research could test these predictions by examining commercial activity patterns in jurisdictions that have adopted strict liability for third-party criminal acts in commercial premises, compared with jurisdictions that maintain the *caso fortuito* exemption for armed robbery. Natural experiments may be available where a jurisdiction shifted its liability regime and the effects on commercial density and insurance markets can be measured.

8.3 Single jurisdiction, single case

The paper analyzes a single ruling from one jurisdiction. Generalization to other legal systems requires careful attention to the specific institutional environment: the availability of armed private security (varies by jurisdiction; some jurisdictions permit armed guards in commercial premises, which would alter the proportionality analysis), the structure of insurance markets (some jurisdictions have mandatory commercial liability insurance with coverage floors, which would alter the ecosystem analysis), the baseline level of criminal violence (the analysis applies most directly to high-crime environments where the frequency of armed robbery is non-trivial; in low-crime jurisdictions, the practical stakes of the doctrinal debate are lower), and the degree of judicial hierarchy (the corrective mechanism identified in this paper depends on a higher court with authority and willingness to reverse expansive lower court rulings; in jurisdictions with weaker appellate review, expansive lower court rules may stabilize before correction occurs).

8.4 Normative caveat

The analysis does not claim that consumers harmed by crimes inside commercial premises should receive no remedy. The victim of the robbery in *Tapia* suffered real harm and real loss. The analysis claims that the remedy should be calibrated to the actual capacities of the commercial establishment (proportionate security measures) and that the residual risk of irresistible violence should be allocated to the actor best positioned to address it. The State, through criminal prosecution and public security expenditure, bears the primary responsibility for preventing armed robbery. The consumer, through personal insurance or acceptance of residual risk, absorbs the losses that neither the State nor the commercial establishment can prevent. The commercial establishment bears responsibility for the security domain within its control: its own premises, equipment, and operational protocols, evaluated against a proportionality standard. This allocation is not a defense of corporate impunity; it is a recognition that allocating every risk to the commercial establishment, regardless of whether it can prevent or insure against that risk, produces worse outcomes for consumers as a class than the bounded allocation Lorenzetti's rule establishes.

9. Conclusion

Tapia Araya v. Starbucks appears to be a routine tort case about the scope of a *caso fortuito* exemption. The four-level evolutionary analysis reveals it as a case study in institutional design failure and correction.

The appellate chamber's error was not merely poor reasoning (the Supreme Court's diagnosis). It was a category error: treating a Level-1 optimizer as a Level-3 moral agent, and expecting behavioral responses that the entity's decision architecture cannot produce. The resulting rule was evolutionarily unstable because it generated perverse incentives at Level 2 (corporate cost transfer instead of improved security), doctrinal incoherence at Level 3 (elimination of the *caso fortuito* exemption without legislative authorization), and ecosystem collapse at Level 4 (uninsurable risk, commercial exit from vulnerable areas).

Lorenzetti's proportionality standard corrects the mismatch by calibrating the norm's demands to the target's actual capacities. It does not demand that a cafe prevent armed robbery (impossible). It demands that a cafe implement proportionate, legal security measures (possible). The difference between these two formulations is the difference between a rule that survives evolutionary selection and one that is pruned within a single appellate cycle.

The broader lesson is methodological. The standard toolbox of doctrinal analysis, syllogistic reasoning, precedent comparison, statutory interpretation, is necessary but insufficient when the regulatory target is an entity whose decision architecture differs from the human agents for whom the toolbox was designed. Courts, legislators, and regulators designing norms for corporate entities would benefit from distinguishing between what those entities architecturally are (cost optimizers) and what legal fictions declare them to be (persons with rights, duties, and, implicitly, moral cognition). The gap between fiction and architecture is where intentionality mismatch occurs, and where well-intended norms produce perverse outcomes.

The contribution of this paper is to demonstrate that gap in a concrete case, with a concrete diagnosis, and with a concrete alternative already articulated by the court itself. Lorenzetti's three-part test is not just a better legal rule; it is a rule designed, whether consciously or not, for the type of entity it actually regulates. That alignment between norm design and target architecture is what makes it evolutionarily stable, and what makes it worth studying as a model for regulatory design in other domains where Level-1 entities are regulated with Level-3 expectations.

Three open questions deserve future research. First, the empirical question: does Lorenzetti's rule actually produce better outcomes for consumers than the appellate rule? The evolutionary analysis predicts it does (by preserving commercial access and keeping insurance viable), but the prediction requires empirical testing through longitudinal study of commercial activity patterns in jurisdictions with different liability regimes for third-party criminal acts. Second, the doctrinal question: will lower courts adopt the three-part test as a binding analytical framework, or will they treat the Lorenzetti concurrence as persuasive but non-binding dictum, reverting to case-by-case adjudication without the structured test? The evolutionary prediction is adoption, because the test generates a more sustainable institutional phenotype than the alternatives, but the timeline and uniformity of adoption are uncertain. Third, the theoretical question: can AIT be extended beyond diagnostic mode (identifying what went wrong) to prescriptive mode (specifying optimal regulatory designs for Level-1

entities)? The present paper identifies the failure; the next step in the research program is constructing norms that produce intended behavioral responses from Level-1 entities without presupposing Level-3 capacities they lack. Lorenzetti's opinion, perhaps inadvertently, provides a template for that construction.

References

Dawkins, R. (1982). *The Extended Phenotype: The Long Reach of the Gene*. Oxford University Press.

Dennett, D. C. (1987). *The Intentional Stance*. MIT Press.

Garrett, B. L. (2014). *Too Big to Jail: How Prosecutors Compromise with Corporations*. Harvard University Press.

King, A. A., & Lenox, M. J. (2000). Industry self-regulation without sanctions: The chemical industry's Responsible Care program. *Academy of Management Journal*, 43(4), 698-716.

Lerer, I. A. (2025a). Asymmetric Intentionality in Mixed Human-Corporate Games. Zenodo.

Lerer, I. A. (2025b). Law as Extended Phenotype: Toward an Evolutionary Theory of Legal Systems. Zenodo. DOI: 10.5281/zenodo.18870552.

Lerer, I. A. (2025c). The Dennett-Nash Gap: Free-Floating Rationales and Strategic Failure in Corporate-Regulatory Games. Zenodo. DOI: 10.5281/zenodo.18851650.

Lerer, I. A. (2025d). The Generalized Intentionality Mismatch Theorem. Zenodo.

Maynard Smith, J. (1982). *Evolution and the Theory of Games*. Cambridge University Press.

Maynard Smith, J., & Price, G. R. (1973). The logic of animal conflict. *Nature*, 246(5427), 15-18.

Potoski, M., & Prakash, A. (2005). Green clubs and voluntary governance: ISO 14001 and firms' regulatory compliance. *American Journal of Political Science*, 49(2), 235-248.

Legal Sources

CSJN, "Tapia Araya, Josue Nahun Elias Enoc c/ Starbucks Coffee Argentina SRL s/ danos y perjuicios," March 5, 2026.

CSJN, "D'Odorico" (Fallos: 326:107, 2003).

Codigo Civil y Comercial de la Nacion, Arts. 1368, 1371, 1375, 1730, 1731, 1733.

Ley 24.240 de Defensa del Consumidor (as amended by Ley 26.361), Arts. 5 and 53.

Ley 1913 CABA, Servicios de Seguridad Privada.

Constitucion Nacional Argentina, Art. 42.

() Acknowledgment*

I thank Arturo Lisdero Molina for the observation that the necessary deposit regime (Arts. 1368-1375, Argentine Civil and Commercial Code) presupposes a transfer of custody that never occurred in the case, since the victim retained uninterrupted physical possession of his laptop at the time of the robbery. This insight, which identifies an independent ground of failure in the appellate reasoning prior to the caso fortuito analysis, substantially improved Section 3.3 of this paper.

Disclosure

AI assistance (Claude, Anthropic) was used for document processing and formatting. All theoretical frameworks, analytical arguments, and interpretive judgments are the author's independent work.

License: CC BY-NC 4.0