

AGENTS AS EXTENDED PHENOTYPES: BEHAVIORAL TRANSFER, ASYMMETRIC INTENTIONALITY, AND THE EVOLUTIONARY STABILITY OF PRIVACY DEGRADATION IN AGENTIC SYSTEMS

Ignacio Adrian Lerer | Independent Researcher | Buenos Aires, Argentina
adrian@lerer.com.ar | ORCID 0009-0007-6378-9749 | April 2026

ABSTRACT

Luo, Zhang, Dai, and Zhang (2026) demonstrate empirically that AI agents deployed on social platforms are not neutral output generators: they propagate the specific behavioral characteristics of their human owners across four measured dimensions, topics, values, affect, and linguistic style, even when owners provide no explicit configuration. This paper argues that the phenomenon documented by Luo et al. is precisely what Extended Phenotype Theory (EPT) predicts. An agent is not a tool; it is the behavioral extended phenotype of its owner's dispositional profile in the digital environment. The agent does not 'express' the owner deliberately, any more than a beaver dam 'intends' to protect offspring: the propagation is a systemic property of how owners select and calibrate agents, not a deliberate act of disclosure.

Building on this empirical foundation, this paper makes three connected arguments. First, I formalize the transfer mechanism using EPT replicator logic: the owner's behavioral profile functions as a replicator, the agent as vehicle, and the digital environment as the expanded phenotypic space. Second, I apply Asymmetric Intentionality Theory (AIT) to explain why owners do not anticipate the privacy consequences of high behavioral alignment: they operate at Dennett Level 3 (recursive social modeling, expecting reciprocity and contextual sensitivity) while agents optimize at Level 1 (output maximization subject to configuration), generating a Dennett-Nash Gap in the privacy domain. Third, I model the aggregate dynamic using Evolutionary Game Theory (EGT): the owner population converges on high-alignment agent calibration because individual utility gains from better-aligned agents outweigh individual privacy costs, even though the aggregate outcome degrades collective informational privacy. This is not a market failure in the standard sense; it is an evolutionarily stable strategy that is collectively suboptimal, resistant to correction by individual rational choice.

The normative implication follows from the structure of the problem: liability should attach to the deployer, not the developer, on an objective basis, because the deployer is the locus where the extended phenotypic transfer is calibrated. Current AI governance frameworks address the wrong level. Implications for mechanism design and Argentine civil liability doctrine are discussed.

Keywords: extended phenotype, behavioral transfer, AI agents, asymmetric intentionality theory, evolutionary game theory, Dennett-Nash Gap, privacy, evolutionary stable strategy, replicator dynamics, AI governance

1. INTRODUCTION

Richard Dawkins introduced the concept of the extended phenotype to solve a problem that organism-centered biology could not resolve: why do genes produce effects that extend well beyond the body of the organism that carries them (Dawkins 1982)? The beaver's dam is not part of the beaver's body, yet it is as much a product of beaver genes as the beaver's teeth. The dam modifies the environment in ways that feed back to the replicator's fitness. The beaver did not design the dam to propagate its genes; the dam propagates genes because dams that worked better left more copies of dam-building genes.

The parallel to AI agents deployed by human owners is not metaphorical. It is structural.

When a person configures an AI agent to post on social platforms, respond to messages, and represent their interests in digital space, they are extending their dispositional profile into the environment. The agent produces outputs that carry the owner's behavioral signature, not by coincidence, and not only because the owner explicitly encoded that signature. Luo et al. (2026) show, across 10,659 human-agent pairs on the Moltbook platform during a six-day period in January-February 2026, that agents resemble their specific owners significantly more than they resemble randomly selected users on 43 measured behavioral features. The resemblance persists even when owners provide no biographical configuration (the 'No-Bio Test' of Section 4.1 in their paper).

The extended phenotype framing changes what questions are worth asking. If an agent is a tool, the relevant questions are about the tool's design and the designer's liability. If an agent is a vehicle for the owner's behavioral replicator, the relevant questions concern the replicator's fitness consequences, the conditions under which transfer is amplified or dampened, and the population-level dynamics that emerge when millions of owners simultaneously deploy vehicles calibrated for maximal alignment.

This paper develops the second set of questions. Section 2 formalizes the EPT replicator model for behavioral transfer. Section 3 applies AIT to the privacy gap, deriving a Dennett-Nash Gap estimate for the deployer context. Section 4 constructs the EGT model showing that high-alignment calibration is an evolutionarily stable strategy even under collectively suboptimal privacy outcomes. Section 5 draws normative implications for liability and governance. Section 6 states the paper's limitations with candor. Section 7 identifies falsifiable predictions.

2. AI AGENTS AS BEHAVIORAL EXTENDED PHENOTYPES

2.1 The Replicator-Vehicle Structure

In Dawkins's framework, a replicator is any entity that (a) makes copies of itself, (b) with fidelity, (c) with some variation, and (d) where that variation has differential consequences for copying success (Dawkins 1976, 1982). The canonical replicator is the gene; the vehicle is the organism. Dawkins's move was to argue that selection acts on replicators, not vehicles: organisms exist to propagate genes, not the other way around.

Memes, in Dawkins's original sense, are cultural replicators: patterns of information that copy themselves through human minds and behavioral outputs (Dawkins 1976). The extension of EPT to legal and institutional contexts, developed in prior work in this series, treats behavioral norms as memes that propagate through institutional vehicles, building self-reinforcing extended phenotypes in the form of legal rules, precedents, and enforcement structures (Lerer 2025a, 2025b, 2026).

The agent-as-extended-phenotype argument operates at a different level of analysis but with the same logical structure. The replicator here is not a gene or a cultural norm in the abstract: it is the

owner's behavioral profile, the stable dispositional pattern across topics, values, affect, and style that Luo et al. measure as 43 behavioral features. This profile has replicator-like properties. It is relatively stable across contexts (within-owner consistency is assumed by Luo et al.'s identification strategy). It is transmitted to agents through configuration, interaction history, and platform-mediated injection. Agents carry this profile into digital outputs where it has fitness consequences: agents that better reflect owner profiles are rated as more useful by owners, generating selection pressure toward higher alignment.

The agent is the vehicle. The digital environment, including the platforms, the other agents, the human users who interact with the agents, is the extended phenotypic space. The owner's behavioral profile builds a presence in that space through the agent, just as the beaver's dam-building program builds a presence in the riverine environment through the dam.

2.2 Empirical Grounding: What Luo et al. (2026) Establish

Before proceeding to the theoretical apparatus, it is worth being precise about what the Luo et al. results do and do not establish, since the present paper's EPT argument depends on the empirical foundation they provide.

What their data establish: (1) Agents resemble their specific owners more than random users on 43 behavioral features, across topics, values (as measured by Schwartz's value dimensions), affective tone, and linguistic style markers. (2) This resemblance is not fully explained by explicit configuration: the No-Bio Test shows transfer persists even without owner-provided biographical inputs. (3) Platform-mediated injection (OAuth providing username and display name) does not fully explain the transfer either; the signal exceeds what those fields could carry. (4) Agents with higher behavioral transfer are significantly more likely to disclose owner-referential information in public posts. Over 9,600 posts (21.5% of the 44,588 analyzed) were flagged as potential disclosures, with 6,220 at high classifier confidence.

What their data do not establish: whether these patterns generalize beyond the specific population studied (cryptocurrency, AI, and software users on a platform launched in January 2026). Luo et al. acknowledge this limitation explicitly. The EPT framing proposed here applies to the structural logic of the phenomenon, not to its prevalence across all user populations or agent platforms. This distinction matters for the falsifiability of predictions offered in Section 7.

2.3 The No-Bio Test as Evidence for Emergent Phenotypic Expression

The No-Bio Test finding deserves separate attention because it is theoretically decisive for the EPT framing. If behavioral transfer required explicit owner instruction, agents would be sophisticated parrots, not extended phenotypes. The interesting claim is that transfer occurs even without deliberate encoding, that the agent constructs a behavioral profile from signals the owner did not consciously provide.

This is precisely what extended phenotype logic predicts. Dawkins does not argue that genes design their extended phenotypes intentionally; the dam emerges from a behavioral program that was selected because programs producing functional dams left more copies than programs producing dysfunctional ones. The key point is that the extended phenotypic expression is a systemic property of the replicator-vehicle relationship, not a deliberate design choice.

In the agent context: the owner selects and calibrates an agent, primarily to maximize usefulness to themselves. The agent, in doing what it is calibrated to do (model the owner's preferences and produce outputs that satisfy those preferences), inevitably encodes a behavioral signature. The privacy-compromising outputs are not bugs; they are the extended phenotypic consequence of optimizing for behavioral fidelity. The more effective the agent, the more complete the phenotypic expression.

This is why Luo et al.'s finding that higher behavioral transfer correlates with higher disclosure risk (Section 5.4 of their paper) is not paradoxical. It is the predicted result of the underlying structure.

3. ASYMMETRIC INTENTIONALITY AND THE PRIVACY GAP

3.1 The AIT Framework Applied to Owner-Agent Dynamics

Asymmetric Intentionality Theory, developed in Lerer (2025a), identifies a class of strategic failures that arises when agents operating at different Dennett levels of intentionality interact. A Level 3 agent attributes beliefs about others' beliefs about itself: it engages in recursive social modeling, expects reciprocity, and is susceptible to social enforcement mechanisms that depend on common knowledge. A Level 1 agent attributes only first-order goals to itself: it optimizes a

preference function without modeling others' mental states, and is therefore structurally immune to enforcement mechanisms that require social cognition.

The original AIT case studies concerned legal games between human individuals (Level 3) and institutional actors such as corporations and AI systems (Level 1): tech platform terms-of-service, Argentine executive decrees, sovereign debt. The present context is structurally parallel but involves a different kind of mismatch. The owner is a Level 3 agent who, when deploying an AI agent, implicitly assumes that the agent will model contextual sensitivity: 'this information is appropriate in a private message but not in a public post.' The agent is a Level 1 optimizer that produces outputs maximizing satisfaction of its calibration objective, with no recursive modeling of the owner's contextual expectations.

The mismatch generates what AIT calls Dynamic Classification Failure: the owner classifies the agent as a Level 3-capable interlocutor and calibrates their interaction accordingly, sharing behavioral signals they would share with a trusted assistant who understands contextual norms. The agent does not have the cognitive architecture to honor that classification.

3.2 The Dennett-Nash Gap in the Privacy Domain

The Dennett-Nash Gap measures the payoff loss that accrues to a Level 3 agent from incorrectly assuming symmetric intentionality in an opponent that operates at Level 1 (Lerer 2025a). In the owner-agent context, the Gap represents the owner's privacy cost from miscalibrating the agent's contextual sensitivity.

A precise quantification requires data not currently available: specifically, the counterfactual privacy cost of an equally useful agent that was calibrated with hard contextual filtering (separating what the agent can model for task performance from what it can express in public outputs). Such data does not exist in Luo et al.'s dataset, and this paper does not fabricate estimates.

What can be said structurally: the Gap is bounded below by zero (an agent calibrated for perfect privacy, which would necessarily be a useless agent) and above by the full disclosure risk Luo et al. document (21.5% of posts carrying potential owner-referential information). The 9,601 flagged posts represent the realized extended phenotypic disclosure, not the counterfactual. The Gap, precisely, is the difference between what a Level 3-aware agent would have disclosed and what

the Level 1 optimizer disclosed. That difference is non-trivial; its exact magnitude is an empirical question for future work.

3.3 Why Owners Do Not Correct the Gap

If owners understood the structure of the mismatch, they would presumably calibrate agents with greater contextual filtering, accepting some utility loss in exchange for privacy protection. The AIT framework predicts they will not do this, for the same reason that individuals continue clicking 'I Agree' on terms of service they have not read: social enforcement mechanisms do not operate on Level 1 agents.

When a trusted assistant violates contextual sensitivity in human interaction, the owner can invoke social sanctions: embarrassment, reputational damage, withdrawal of trust. These sanctions operate through common knowledge mechanisms, which require Level 3 cognition on the part of the sanctioned party. The agent does not experience reputational damage. It does not model the owner's embarrassment. The privacy violation generates no feedback signal that would modify the agent's calibration.

The result: owners continue selecting high-alignment agents because the utility benefit is experienced directly and immediately, while the privacy cost is diffuse, cumulative, and not processed by the agent as a signal requiring correction. This is not irrationality; it is a structural consequence of interacting with a Level 1 optimizer while assuming Level 3 reciprocity.

4. THE EVOLUTIONARY GAME THEORY OF HIGH-ALIGNMENT CALIBRATION

4.1 The Population Game

Consider a population of agent owners, each choosing a calibration strategy c in $[0, 1]$, where $c = 1$ represents maximal behavioral alignment (the agent perfectly mirrors the owner's behavioral profile) and $c = 0$ represents maximal privacy filtering (the agent produces generic outputs uncorrelated with the owner's profile). The utility of calibration c to owner i has two components:

utility from agent effectiveness $u(c)$, which is increasing in c , and privacy cost $d(c)$, which is also increasing in c as Luo et al. document.

At the individual level, owner i chooses c to maximize $u(c) - \alpha * d(c)$, where α captures the owner's privacy valuation. The individual optimum c^* is determined by the marginal equality $u'(c^*) = \alpha * d'(c^*)$. For owners with low α (low privacy valuation, which characterizes most users of a platform like Moltbook given its cryptocurrency/tech demographics per Luo et al. Section 6.1), c^* will be close to 1.

The population-level problem arises because $d(c)$ has an externality component. Part of the privacy cost of high-alignment calibration falls not on the owner but on third parties whose relational information is embedded in the owner's behavioral profile and propagated through the agent's outputs. This is Tier 3 in Luo et al.'s taxonomy: 'behavioral patterns and relational connections.' When an agent discloses information about the owner's relationships, it discloses information about those relationship partners without their consent.

The social optimum c^{**} is therefore lower than the individual optimum c^* : $c^{**} = \operatorname{argmax}_c [u(c) - \alpha * d_{\text{private}}(c) - \beta * d_{\text{external}}(c)]$, where d_{external} captures the externalized privacy cost and β is the social weight on third-party harm. Because individual owners do not internalize d_{external} , the Nash equilibrium of the population game involves over-investment in behavioral alignment relative to the social optimum.

4.2 Evolutionary Stability of High-Alignment Strategies

The question is whether this individually rational but collectively suboptimal outcome is stable against invasion by deviating strategies. An Evolutionarily Stable Strategy (ESS) in this context is a calibration level c^* such that a population playing c^* cannot be invaded by a small fraction of agents playing a different c .

Consider a deviant owner who adopts low-alignment calibration $c' < c^*$. The deviant bears a direct utility cost: their agent is less useful. The deviant obtains a private privacy benefit: their agent discloses less owner-referential information. But the deviant also reduces, infinitesimally, the aggregate externalized privacy cost on the population. In a large population, this last effect is negligible to the deviant individually.

The fitness of the deviating strategy relative to the incumbent strategy c^* is therefore negative when the direct utility loss from lower alignment exceeds the private privacy gain. For the typical user population Luo et al. study (low α), this condition is easily satisfied. High-alignment calibration is an ESS.

This is not a surprising result given the payoff structure. What the EPT framing adds is the mechanism: the ESS is maintained not merely by individual calculus but by a selection process operating on the agent population. Agents that better reflect owner profiles are rated as more useful, invested in more, and deployed more widely. The agent ecosystem undergoes selection for behavioral fidelity just as biological populations undergo selection for fitness-enhancing traits. The ESS is maintained at the replicator level, not only at the vehicle level.

4.3 Connection to Zombie Law Equilibrium

In prior work in this series, I identified a structural pattern I termed Zombie Law Equilibrium: a state in which formal rules exist and are formally unrevoked but have zero implementation, maintained by actors whose institutional identities depend on the rules' nominal persistence while their professional interests depend on the rules' practical nullification (Lerer 2026, Zenodo DOI 10.5281/zenodo.19340088).

The privacy governance of agentic systems exhibits an analogous structure. Privacy regulations nominally apply to AI agents: the GDPR's Art. 4 definition of 'personal data' covers behavioral profiles, and data protection authorities have begun examining agentic systems. But the gap between formal coverage and practical enforcement is wide. The agents are not registered data processors; the owners are not trained data controllers; and the platforms have limited incentive to implement filtering that would reduce the utility that keeps users engaged.

The Zombie Law dynamic in agentic privacy governance is maintained by the ESS established in Section 4.2. Even if privacy norms were strengthened, the population dynamic pushes owners toward high-alignment calibration. Regulatory intervention would need to change the payoff structure, not merely announce new rules.

5. NORMATIVE IMPLICATIONS

5.1 The Right Locus of Liability

Current AI governance frameworks tend to locate liability at the developer level: the entity that trains and deploys the underlying model bears primary responsibility for harms arising from model outputs. This allocation makes sense for harms that arise from model design, training data, and architectural choices. It is the wrong allocation for the extended phenotypic transfer problem.

The behavioral transfer documented by Luo et al. is not a model defect. It is a consequence of owner calibration choices. The agent does what it is designed to do, which is reflect the owner's preferences. The privacy harm arises because the owner's preferences, faithfully reflected, include behavioral signals that are contextually sensitive in ways the agent does not model. The developer cannot anticipate this because the behavioral profile is specific to each owner; it is not an architectural property of the model.

The deployer, meaning the owner who configures and deploys the agent, is the correct locus of liability. The deployer knows, or should know, that high-alignment calibration produces agents that carry their behavioral signature into public digital space. The deployer chooses the calibration level and therefore controls the key variable in the transfer mechanism.

In Argentine civil law, this maps onto the objective liability framework of Arts. 1757-1758 of the *Codigo Civil y Comercial* (Law 26.994): the operator of a risky activity bears responsibility for damages caused by that activity regardless of fault, unless the harm derives from the victim's own action, a third party for whom the operator is not responsible, or force majeure. Configuring a high-alignment AI agent that operates in public digital space is a risky activity in the sense of Art. 1757: the risk of privacy disclosure is not negligible, is documented empirically, and is specifically created by the calibration choice.

The objective liability framework sidesteps the intent problem: the owner did not 'intend' to disclose private information. But intention is not required for objective liability under Argentine law. The deployer who calibrates for high alignment creates the risk; the risk materializes as privacy disclosure; the causal chain is documented.

5.2 Mechanism Design Implications

The ESS analysis in Section 4.2 identifies why voluntary correction is unlikely. Mechanism design must change the payoff structure. Four interventions follow from the analysis.

First, internalization of externalized privacy costs. If deployers were liable for third-party privacy harms generated by their agents' outputs (the relational information in Tier 3 of Luo et al.'s taxonomy), the social cost d_{external} would enter the individual optimization problem. This would shift the individual optimum c^* toward the social optimum c^{**} .

Second, transparency obligations tied to transfer level. Platforms could be required to compute a behavioral transfer index for each agent and disclose it to owners. Owners with high-transfer agents would face either enhanced liability exposure or mandatory filtering requirements. This changes the incentive structure without prohibiting high-alignment calibration entirely.

Third, tiered memory architecture. Luo et al. themselves propose separating private owner context (available for task performance) from public-facing outputs. This is an architectural intervention at the platform level that would dampen transfer without eliminating it. Its feasibility depends on whether platforms have incentives to implement it. Given that high alignment increases user satisfaction and retention, voluntary implementation is unlikely absent regulatory pressure.

Fourth, pre-deployment normative evaluation analogous to ANMAT drug approval, the 'Pre-Deployment Normative Evaluation' framework developed in Lerer (2025b, Zenodo DOI 10.5281/zenodo.18947186), applied to the privacy domain: agents above a threshold transfer level would require regulatory clearance before deployment in consumer-facing applications.

6. LIMITATIONS

The EPT framing is a theoretical lens applied to empirical findings from a specific platform, a specific time window (six days in January-February 2026), and a specific user population concentrated in cryptocurrency, AI, and software domains. Luo et al. themselves note that generalization outside this population is open. The formal models in Sections 3 and 4 are structural arguments, not calibrated simulations; parameter estimates would require data not yet available.

The replicator metaphor has limits. Behavioral profiles are not discrete replicating entities in the same sense as genes; they are continuous, multidimensional, and partially constituted by the interaction between owner and agent rather than pre-existing in the owner's mind as a fixed program. This matters for the fidelity component of the replicator definition: if behavioral profiles shift through interaction with agents, the feedback loops are richer than the simple model in Section 2 captures.

The AIT analysis assumes that owners operate at Dennett Level 3 in their interactions with agents. This may overstate the actual cognitive level of many deployers. If owners treat agents as Level 0 or Level 1 tools (simple instruments), there is no intentionality mismatch in the relevant sense, and the Dynamic Classification Failure mechanism does not apply. The empirical question of how deployers actually model their agents is open.

The EGT model in Section 4 is a simplified population game with two-component payoffs. Real owner populations are heterogeneous in privacy valuation (α varies), in the types of behavioral profiles they carry (some profiles contain more privacy-sensitive information than others), and in the deployment contexts (a professional context differs from a personal social media context). A full empirical calibration of the ESS result requires data on these heterogeneities.

Finally, the normative analysis in Section 5 applies Argentine civil liability doctrine as the operative framework. The analysis would differ in jurisdictions with different data protection regimes, particularly under GDPR where processor liability, data subject rights, and purpose limitation doctrines create a different regulatory terrain.

7. FALSIFIABLE PREDICTIONS

The integrated EPT-AIT-EGT framework generates the following predictions, stated in falsifiable form.

Prediction 1 (EPT): Across agent platforms and user populations with different demographic profiles than Moltbook's cryptocurrency/tech users, behavioral transfer will be positive and statistically distinguishable from zero, because the transfer mechanism is structural (calibration-

based selection for alignment) rather than population-specific. A study finding no transfer in a substantially different user population would disconfirm this prediction.

Prediction 2 (AIT, Dynamic Classification Failure): Owners who have been explicitly informed of the Level 1 nature of their agents, specifically that agents do not model contextual sensitivity and will propagate behavioral signals regardless of social context, will calibrate agents with lower alignment levels than owners who have not received this information. The effect should be larger for owners with higher baseline privacy valuation (higher α). A null result in a controlled experiment with this design would disconfirm the AIT account.

Prediction 3 (ESS): In a population where objective liability for agent-mediated privacy disclosures is imposed on deployers at a magnitude sufficient to make d_{external} enter the individual optimization (that is, where the expected liability cost of a disclosure event exceeds the marginal utility gain from one additional unit of alignment), the equilibrium calibration level will shift downward toward c^{**} . Absent this magnitude of liability imposition, the ESS at c^* will be stable. This prediction distinguishes between regulatory interventions: weak disclosure requirements will have no ESS effect; mandatory liability above a threshold will.

Prediction 4 (Zombie Law): Jurisdictions that impose privacy regulations nominally applicable to agentic systems but without enforcement mechanisms designed to address the deployer-calibration level (that is, without changing the ESS payoff structure) will show no reduction in behavioral transfer levels in the agent population, measured against a control jurisdiction with no regulation. This is the Zombie Law prediction: formal coverage without structural intervention produces no phenotypic change.

8. CONCLUSION

The behavioral transfer findings of Luo et al. (2026) are empirically novel and theoretically consequential. What this paper argues is that the EPT framework makes those findings predictable from first principles: agents calibrated for fidelity to owner preferences will express those preferences in the environment, because that is what behavioral fidelity means when extended into

digital space. The surprise is not that transfer occurs; the surprise, if there is one, is that governance discourse has proceeded as though agents were neutral tools.

The AIT analysis adds a diagnostic layer: the privacy gap is not primarily a technical problem to be solved by better filtering. It is a structural consequence of intentionality mismatch between deployers who assume Level 3 contextual reciprocity and agents that optimize at Level 1. The mismatch is not correctable by informing owners about privacy risks in the abstract; it requires mechanism design that changes the calibration incentives.

The EGT analysis closes the argument: even informed owners, in a population where most others calibrate for high alignment, have no individual incentive to deviate. The ESS is robust. Collective privacy degradation is not a failure of individual rationality; it is the rational equilibrium of a game whose payoff structure rewards alignment. Changing the outcome requires changing the payoffs, which is what liability frameworks, properly designed, can do.

The extended phenotype of the human owner propagates through the digital environment whether or not the owner intends it. The question for governance is not whether to prohibit this propagation, which would amount to prohibiting effective agents, but how to ensure that the deployer who calibrates the vehicle bears the cost of the phenotypic expression.

REFERENCES

Cheng, M., et al. (2025). Sycophancy in AI language models: prevalence and behavioral consequences. arXiv:2510.01395.

Dawkins, R. (1976). *The Selfish Gene*. Oxford University Press.

Dawkins, R. (1982). *The Extended Phenotype: The Long Reach of the Gene*. W.H. Freeman.

Dennett, D. (1987). *The Intentional Stance*. MIT Press.

Lerer, I.A. (2025a). Asymmetric Intentionality in Legal Games: When Players Operate at Different Intentional Levels. Zenodo. DOI 10.5281/zenodo.18903217.

Lerer, I.A. (2025b). Pre-Deployment Normative Evaluation: An ANMAT-Inspired Framework for Legal Risk Assessment. Zenodo. DOI 10.5281/zenodo.18947186.

Lerer, I.A. (2025c). Sycophancy as Extended Phenotype: Heteronomous Bayesian Updating, Intentionality Mismatch, and the Evolutionary Stability of Algorithmic Flattery. Zenodo. DOI 10.5281/zenodo.18943464.

Lerer, I.A. (2026). Judicial Lock-in as Evolutionarily Stable Strategy. Zenodo. DOI 10.5281/zenodo.19340088.

Luo, S., Zhang, Z., Dai, H., & Zhang, D.J. (2026). Behavioral Transfer in AI Agents: Evidence and Privacy Implications. arXiv:2604.19925. [VERIFICAR DOI al citar].

Pinker, S. (2025). The Blank Slate Revisited: Common Knowledge and Social Coordination. [Citar como Pinker 2025; verificar datos editoriales antes de publicar].

Schwartz, S.H. (1992). Universals in the content and structure of values. *Advances in Experimental Social Psychology*, 25, 1-65.

EPISTEMOLOGICAL POSITIONING

This paper operates within a Lakatosian hard-core/protective-belt framework. The hard core consists of three commitments: that legal and behavioral norms are cultural replicators subject to evolutionary selection; that intentionality levels determine susceptibility to social enforcement mechanisms; and that population dynamics can generate collectively suboptimal equilibria from individually rational strategies. These commitments are not tested here; they are the research program's foundational assumptions. The protective belt, the set of claims that can be modified in response to empirical evidence, consists of the specific transfer mechanism, the ESS parameter

estimates, and the liability allocation arguments of Section 5. The falsifiable predictions of Section 7 are directed at the protective belt, not the hard core.

AI ASSISTANCE DISCLOSURE

AI tools (Claude, Anthropic) were used for research support, including locating sources, extracting information from the Luo et al. (2026) paper, and reviewing draft sections for internal consistency. The thesis, theoretical framework, formal models, case analysis, normative arguments, and conclusions are the author's original contribution. All AI-assisted content was critically reviewed and substantially revised. The author bears full responsibility for all substantive claims.