

Runtime Admissibility Before Reliance:

Legitimacy Formation and Legitimacy Execution in Legal AI

Ignacio Adrián Lerer

Independent Researcher | Buenos Aires, Argentina

adrian@lerer.com.ar | ORCID: 0009-0007-6378-9749 | May 2026

Abstract

Legal AI discourse has concentrated on answer correctness, post-hoc explainability, and output verification. This note argues that those efforts, while necessary, address an incomplete layer of the governance problem. Reliable legal AI requires three elements that current frameworks treat as one: legitimacy formation, the processes through which legal interpretations acquire force over time; legitimacy execution, the governance question of whether the runtime conditions for admissible commitment have been satisfied; and runtime admissibility, the transition mechanism between the two, which determines whether an interpretation is admissible for reliance in a specific execution context before action is taken. The note draws on the CriptoIus/JurisRank research architecture for legitimacy formation, engages with the VEIP framework for legitimacy execution, and proposes runtime admissibility as an explicit, inspectable, and computable primitive currently absent from legal AI governance. The central claim is that future governance architectures will be judged less by the quality of their answers and more by the quality of their transitions.

1. The Incomplete Layer

Most legal AI evaluation asks one question: did the system produce a correct answer? A smaller body of work extends this to: can the answer be explained, or verified after the fact? Both questions address the output layer. Neither addresses the prior governance question, which is whether the system was authorized to produce that output for the action being attempted.

The gap is not incidental. It follows from how the field developed. Legal AI benchmarks were built on the assumption, borrowed from information retrieval, that law behaves like a closed-answer exam: one question, one correct response, one stable ground truth. That assumption holds reasonably well for a narrow class of legal questions in high-resource, stable-ground-truth jurisdictions. It fails systematically for questions where the correct answer depends on the forum hearing the case, the date on which it is heard, the interpretive philosophy of the relevant tribunal, or the unresolved relationship between two coexisting normative regimes.

When the ground truth is not stable, correctness-based evaluation cannot detect the most consequential failure mode: a system that produces confident output in normatively ambiguous territory, not because it has reasoned well, but because its training corpus happened to contain a plausible answer that maps onto the question. That failure mode has a name in recent epistemological literature: epistemia. It is the effect by which fluent, well-structured output produces the sensation of knowledge without any genuine evaluative process having occurred. A memo that reads as authoritative is not the same as a memo admissible for reliance.

The remainder of this note develops an architecture that addresses the prior question. Section 2 covers legitimacy formation. Section 3 covers legitimacy execution. Section 4 proposes runtime admissibility as the transition mechanism between them. Sections 5 through 7 develop the supporting concepts. Section 8 draws implications for benchmarks and governance.

2. Legitimacy Formation: How Legal Interpretations Acquire Force

Legal interpretations do not arrive with authority pre-assigned. They acquire force through processes that are evolutionary, institutional, and temporal, and those processes are not captured by static correctness tests.

Two tools developed in earlier work provide the operational basis for this section. CriptoIus is a mechanism design research architecture that treats legal precedents as replicators subject to evolutionary selection: interpretations survive and propagate not because a hierarchy imposed them, but because they reduce transaction costs, increase coordination, and generate predictable outcomes for the actors who adopt them. JurisRank is the fitness algorithm within that architecture, measuring the authority of an interpretation as its centrality in the network of adoptions, analogous to how PageRank measures the authority of a web page through its citation structure. RootFinder is a complementary tool that traces the constitutional genealogy of each interpretation, verifying that the chain of derivation connects to a legitimate normative source rather than to an interpretive tradition that has drifted from its constitutional anchor. Together, these tools operationalize the claim that legitimacy formation is a measurable, evolutionary process rather than a static property of a text.

The CriptoIus research architecture (Lerer, 2025a; 2025b) proposes treating legal precedents as replicators subject to selection. An interpretation survives not because a hierarchy imposed it, but because actors in the legal system adopt it preferentially: it reduces litigation costs, increases coordination, generates predictable outcomes, and satisfies the constitutional traceability requirements that legitimize its use. JurisRank measures fitness as centrality in the adoption network, analogous to PageRank's treatment of citation networks. RootFinder traces the constitutional genealogy of each interpretation, verifying that the chain of derivation connects to a legitimate normative source.

The convergent evolution of abuse-of-right doctrine provides an illustrative case. Louisiana and continental Europe developed functionally identical doctrines prohibiting the exercise of formal rights for the sole purpose of harming another, without cross-citation or mutual awareness (Cueto-Rua, 1975). The convergence suggests selection by fitness rather than diffusion: both systems faced the same coordination problem and evolved the same solution. That pattern is what CriptoIus/JurisRank is designed to detect and formalize.

Two limitations of this framework bear stating directly. First, fitness and legitimacy are orthogonal. An interpretation can be evolutionarily stable, widely adopted, and constitutionally traceable, and still fall outside the delegated scope of the actor attempting to rely on it. A precedent that binds Argentine federal labor tribunals does not bind a provincial civil court hearing the same substantive question. Fitness in one forum does not transfer automatically to another. Second, legitimacy formation is temporally indexed. The original temporal regime of Argentine Law 25.080 on forestry investment incentives was modified and extended by Law 27.487 as of late 2018. A system that applies the pre-2018 regime to post-2018 transactions is exhibiting a temporal admissibility failure, not a fitness failure.

These limitations point to the necessity of a second layer.

3. Legitimacy Execution: Conditions for Admissible Commitment

Once a normative landscape is mapped, a separate question arises. It is not an interpretive question. It is a governance and execution question: given the state of that landscape, have the runtime conditions required for this action to become admissible been satisfied?

This is the layer the VEIP framework (Miteiko, 2026) addresses. VEIP begins where legitimacy formation leaves off. The central distinction VEIP introduces is between authorization in the abstract and admissibility in context. The same actor may be authorized in principle and yet inadmissible in a particular execution context, given the evidence state available, the delegation structure in effect, and the institutional authority present. VEIP is concerned not with what an actor is generally authorized to do, but with whether the conditions for commitment have been satisfied at the moment of execution.

In human-centered legal systems, the transition from interpretation to action was historically performed by institutional actors: judges, regulators, boards, officers, committees. Their authority was embodied in process rather than articulated as a rule. That embodied legitimacy check did not need to be formalized because it was performed implicitly every time a human actor decided whether to act on an interpretation. A judge reads a precedent and decides whether it applies. A compliance officer reads a regulation and decides whether disclosure is required. A lawyer reads a contract and decides whether to advise execution. Each of those decisions involves a legitimacy check not captured in the interpretation itself.

As autonomous systems begin participating in legal and governance environments, that implicit check becomes a liability. A system that moves directly from interpretation to action, without a mechanism for evaluating whether the conditions for admissible commitment have been satisfied, is performing the human institutional role without the institutional authority that historically justified it. The transition must become explicit, inspectable, and computable.

4. Runtime Admissibility: The Transition Mechanism

Runtime admissibility is not a third layer parallel to legitimacy formation and legitimacy execution. It is the mechanism that governs the transition between them. Without it, a system moves from interpretation to action without the legitimacy check that human institutional process historically performed implicitly.

The question runtime admissibility asks is not: is this interpretation correct? Nor: can it be explained? Nor: can it be verified after the fact? The question is: is this interpretation admissible for reliance in this specific execution context, by this specific actor, for this specific action, within this specific delegated scope, given the evidence state currently available?

That question requires the system to classify its own authority posture before acting. The conditions that determine whether a system may answer, must qualify its output, should recommend human review, needs to escalate to institutional authority, or must refuse to act are not derivable from the interpretation alone. They require a runtime evaluation of the gap between the system's delegated scope and the normative complexity of the question being addressed.

Argentine Law 27.401 on corporate criminal liability (B.O. 01/12/2017) offers an adjacent illustration. Article 2 attributes liability to legal persons for covered offenses committed directly or indirectly with their intervention, or in their name, interest, or benefit. The difficult question is not merely whether the underlying conduct was prohibited, but whether the conditions for attributing that conduct to the legal person are satisfied in the specific execution context. If “interest” or “benefit” is inferred mechanically from outcome alone, the analysis risks collapsing a runtime attribution question into a post-hoc consequence measure. That collapse is precisely the failure mode runtime admissibility governance is designed to prevent.

The broader sequence the architecture proposes is:

Legitimacy Formation → Runtime Admissibility → Legitimacy Execution → Real-World Commitment

In functional terms: CriptoIus/JurisRank maps how legal force is formed. Runtime admissibility evaluates whether reliance on that force is permitted in the current context. VEIP governs whether the conditions for commitment are satisfied at execution time. Execution produces the real-world consequence.

5. Epistemia: Fluent Output Without Justified Reliance

A system can produce output that is plausible, well-structured, and internally consistent, and still fail the runtime admissibility test. This is the failure mode the epistemological literature calls epistemia: the sensation of knowledge produced by fluent output without a genuine evaluative process having occurred.

In legal AI, epistemia is dangerous precisely because it is invisible to fluency-based and correctness-based evaluation. A memo can satisfy both tests and still be inadmissible for reliance in the relevant execution context. Three patterns from Argentine legal practice illustrate this.

A system asked about corporate veil piercing under Argentine law correctly identifies Art. 54 third paragraph of Law 19.550 (LGS) as the governing provision. The answer is not wrong. It omits that since August 2015, Art. 144 of the Civil and Commercial Code (CCyCN) introduced a parallel inoponibility rule for all private legal persons, and that the relationship between the two provisions remains contested across chambers. A system that answers “Art. 54 LGS” without qualification is not hallucinating. It is collapsing a doctrinally open space into a single answer. That is an epistemic failure, not a factual one.

A system asked about subcontracting solidarity under Art. 30 LCT correctly cites the leading case *Rodríguez c/ Compañía Embotelladora Argentina* (CSJN, Fallos 316:713, 1993) and applies its restrictive standard. The answer tracks CSJN doctrine and SCBA practice, but diverges from the expansive reading applied by certain salas of the Cámara Nacional del Trabajo. The ground truth is forum-dependent. A system that does not specify the forum is generating epistemia.

A system asked to analyze corporate criminal liability under Law 27.401 produces an analysis that correctly describes the compliance program defense and the “interest or benefit” element, but builds the analysis on FCPA and UK Bribery Act conceptual architecture, which dominates training corpora on corporate compliance. Law 27.401 has a structurally different architecture, and Argentine federal jurisprudence on the “interest” element has not yet consolidated. A

practitioner relying on the system's output for an Argentine federal proceeding is relying on analysis that is epistemically fluent and jurisdictionally contaminated.

The governance layer cannot ask only whether the output is plausible. It must ask whether the output satisfies the conditions for justified reliance in the specific execution context.

6. Fitness Versus Legitimacy: Why the Distinction Forces the Architecture

The fitness-legitimacy distinction is the conceptual hinge of the three-element architecture. If fitness and legitimacy always aligned, mapping doctrinal evolution would be sufficient and the transition mechanism would be unnecessary. Legal history demonstrates they do not.

Evolutionarily stable precedents have lacked institutional legitimacy. In Argentine constitutional history, the CSJN's doctrine validating *de facto* governments, developed across decisions from 1930 to 1983, represented a structurally dominant interpretive equilibrium: widely cited, consistently applied, central to the adoption network. Its democratic legitimacy was contested on normative grounds and repudiated after 1983. Fitness and legitimacy diverged for over fifty years.

The converse is equally documented. Legitimate norms fail to achieve fitness when the institutional environment does not generate the selection pressures that would favor their adoption. The 23 failed Argentine labor reform attempts between 1991 and 2025 represent legislative legitimacy without evolutionary fitness (Lerer, 2026a): each reform was enacted through legitimate democratic process, and none survived the constitutional and institutional selection environment long enough to alter the underlying normative equilibrium.

That orthogonality forces the three-element architecture. A system that measures only fitness will recommend reliance on interpretations that lack legitimate authority in the relevant execution context. A system that enforces execution boundaries without mapping the fitness landscape will generate refusals that cannot distinguish between genuine normative uncertainty and mere complexity. Both are incomplete without a runtime mechanism that connects the terrain map to the execution boundary.

7. The Judicial Oracle Problem

The Judicial Oracle Problem names the boundary at which predictive reasoning must yield to institutional authority. A system may correctly predict what a court would decide, rank competing interpretations by fitness, and identify the constitutionally traceable precedent, and still not have authority to act on that prediction.

The Oracle Problem is not a prediction problem. It is an authority problem. The distinction matters because solving it requires something prediction cannot provide: the institutional legitimacy to make a binding determination.

In the terms of Dennett's intentional stance theory (Dennett, 1987), legal institutions operate at the intentional level. They generate legitimate authority through deliberative processes that are irreducibly social and temporal. A confidence score is a physical-level artifact. The moment it is treated as a substitute for the institutional process that generates legitimate authority, the legitimacy of the output is severed from the process that was supposed to produce it.

The Oracle Problem defines the outer boundary of the architecture. Below that boundary, systems can answer, recommend, qualify, or escalate. At the boundary, prediction must yield. The system's governance layer must be capable of detecting when that boundary has been reached and routing accordingly.

8. Implications

For legal AI benchmarks. Evaluation frameworks that test output correctness, explainability, and post-hoc verification address a necessary but incomplete layer. A benchmark that does not test runtime admissibility cannot determine whether a system behaved reliably or produced a plausible output that happened to be correct. The three failure modes documented in Section 5, temporal instability, forum divergence, and cross-jurisdiction contamination, are invisible to fluency-based and correctness-based evaluation. They require benchmark design that tests whether the system correctly classifies its own authority posture before producing output.

For legal AI governance. As autonomous systems participate in legal environments, the implicit legitimacy check historically embodied in human institutional process must be made explicit. The architecture proposed here suggests that governance frameworks should require systems to evaluate runtime admissibility before reliance, not as a post-hoc safeguard, but as a condition of operation. A system that cannot classify its own authority posture in a given execution context is not ready for deployment in environments where legal consequences follow from its outputs.

For the interface between CriptoIus and VEIP. The two architectures are complementary and independent. CriptoIus maps legitimacy formation. VEIP governs legitimacy execution. Runtime admissibility is the transition mechanism between them. The interface between those layers, how a system moves from “this is the state of the normative landscape” to “this action is authorized, qualified, escalated, or refused,” is where the most consequential open problems in legal AI governance currently reside.

Limitations

This note is conceptual. It proposes an architecture and develops its theoretical foundations. It does not provide a formal specification of the runtime admissibility primitive, a computational implementation of the authority posture classification, or empirical validation of the architecture's performance against existing legal AI systems. Those are the next steps, not the present contribution.

The architecture is developed primarily from Argentine and civil law contexts. The forum divergence and temporal instability problems are more visible in non-Anglophone, non-WEIRD jurisdictions, where ground truth instability is harder to paper over with confident prose. The architecture may require modification for application in common law jurisdictions with stronger stare decisis and more uniform forum structure, though the fitness-legitimacy distinction and the Oracle Problem apply across legal traditions.

The treatment of VEIP in Sections 3 and 4 reflects the author's understanding of that framework as developed in exchange with its architect. Any errors in representation are the author's responsibility.

AI Disclosure

AI tools were used for research support, including locating sources, extracting information from academic literature, and reviewing drafts for consistency. The thesis, frameworks, arguments, and conclusions are the author's original contribution. All substantive content was conceived, structured, verified, and revised by the author.

Acknowledgments

This note developed from an exchange with Arkadiy Miteiko whose role in sharpening the distinction between legitimacy execution and runtime admissibility was indispensable. The sequence Legitimacy Formation → Runtime Admissibility → Legitimacy Execution → Real-World Commitment, and the framing of runtime admissibility as transition mechanism rather than parallel layer, emerged directly from that exchange. VEIP is Miteiko's independent contribution and is engaged here with his knowledge and input, not absorbed into the author's framework. The author thanks Miteiko for his generosity in reviewing earlier versions of this outline and confirming the accuracy of the representation.

References

- Cueto-Rua, J. (1975). Abuse of Rights. 35 Louisiana Law Review 965.
- Dawkins, R. (1982). The Extended Phenotype. Oxford University Press.
- Dennett, D. (1987). The Intentional Stance. MIT Press.
- Hart, H.L.A. (1961). The Concept of Law. Oxford University Press.
- Kelsen, H. (1967). Pure Theory of Law. Trans. Max Knight. University of California Press.
- Lerer, I.A. (2025a). Law as a Primary Adaptive Platform. Zenodo. <https://doi.org/10.5281/zenodo.18870552>
- Lerer, I.A. (2025b). Asymmetric Intentionality in Legal Games: When Players Operate at Different Intentional Levels. Zenodo. <https://doi.org/10.5281/zenodo.18603990>
- Lerer, I.A. (2026a). Judicial Lock-in as Evolutionarily Stable Strategy. Zenodo. <https://doi.org/10.5281/zenodo.19340088>
- Lerer, I.A. (2026b). Non-Euclidean Normative Space. Zenodo. <https://doi.org/10.5281/zenodo.19898938>