

Outcome Control as Constitutional Constraint: The Argentine Model for Autonomous Legal Entities

Ignacio Adrián Lerer

Independent Researcher, Buenos Aires, Argentina

adrian@lerer.com.ar | ORCID: 0009-0007-6378-9749

Abstract

This paper proposes outcome control as the constitutionally grounded and technically coherent framework for governing autonomous legal entities, and advances that thesis through engagement with Emiliano Kargieman's June 2026 whitepaper on Autonomous Legal Entities (ALEs). The ALE framework is the most architecturally serious proposal in the current global debate on machine agency governance. It converges with the author's prior work on the Responsibility Ramp and on the Static Agent Assumption in identifying a structural vacuum in the legal treatment of autonomous systems, and its adoption of outcome-focused liability as a baseline principle is well-founded. The divergence lies in the ALE's governance interior: its steward council, Supervisory API, kill-switch architecture, and tiered arbitration protocol constitute process regulation, not outcome regulation, and under Argentina's constitutional framework they conflict with Article 19 of the National Constitution, which restricts state intervention to conduct that harms third parties. The paper argues that process-control layers added to an otherwise sound outcome-control framework do not strengthen governance; they generate a dual-logic regime that is formally demanding and informally unenforceable, producing the administrative discretion that has historically functioned as a rent-extraction mechanism in Argentine environmental law. The alternative this paper advances is the Responsibility Ramp applied to the Argentine Automated Company (Article 14 of the draft General Companies Law submitted to Congress on May 29, 2026): dynamic intentionality classification by task phase, strict liability under Articles 1757 and 1758 of the Civil and Commercial Code for execution-phase harm, attribution tracing the configuration chain to the decision that mattered, and three targeted statutory additions that close the framework's structural gaps without regulating the agent's decision process.

Keywords: autonomous legal entities, outcome control, process control, Article 19 Argentine Constitution, Responsibility Ramp, Static Agent Assumption, Automated Company, ALE framework, intentionality mismatch

1. Introduction

On June 8, 2026, Emiliano Kargieman, CEO of Satellogic (NASDAQ: SATL), published a whitepaper titled "Autonomous Legal Entities: A Polycentric Framework for Machine Agency" and stated publicly that he had contributed to the proposed AI legal entity legislation in Argentina. That contribution, if accurately described, makes the whitepaper the most important document in the current Argentine corporate law reform debate that almost no one

has read. The Financial Times Opinion piece by President Milei, published the day after the Executive submitted the draft General Companies Law to Congress on May 29, 2026, attracted considerable attention. The whitepaper that reportedly shaped the legislation received a fraction of it.

That asymmetry is worth correcting. Not because Kargieman is wrong about the problem he identifies; the governance vacuum for autonomous machine agency is real, and the existing liability frameworks are structurally inadequate for agents that acquire resources, enter contracts, and make consequential decisions without real-time human oversight. On this diagnosis, the whitepaper and this paper converge. The divergence lies elsewhere: in whether the governance architecture proposed to fill that vacuum is constitutionally permissible in the Argentine context, and in whether the process-control layers the ALE introduces are technically effective or counterproductive given what is now known about autonomous agent behavior under supervision.

This paper argues three things. First, the Argentine draft's Article 14 on the Automated Company already adopts the correct constitutional and technical orientation: outcome control, not process control, with the agent's internal decision process left unregulated and liability attaching to harmful results under the Civil and Commercial Code. Second, the ALE framework's governance interior (steward council, Supervisory API, kill-switch architecture, tiered arbitration) constitutes process regulation in substance, regardless of its framing as governance, and its incorporation into the Argentine framework during parliamentary debate would reproduce the dual-logic pathology already documented in Argentine environmental law. Third, the structural gaps that Article 14 leaves open (effective reparation guarantee, proof of causation, imputability without a human in the chain) are solvable within the outcome-control logic through three targeted statutory additions, none of which regulate how the agent decides.

The paper proceeds from a prior published position. In March 2026, the author published "The Static Agent Assumption: Dynamic Intentionality Classification and the Limits of the Algorithmic Corporation" (Lerer, 2026a), arguing against the Algorithmic Corporation proposed by Arbel, Goldstein, and Salib on grounds that its incentive mechanism rests on a false premise about agent cognitive stability. The ALE framework, while architecturally more sophisticated than the A-corp, shares that premise in its governance layers. The present paper extends the Static Agent Assumption argument to the Argentine constitutional context and to the specific legislative moment created by the May 2026 draft.

Epistemological positioning. This paper belongs to a research programme (Lakatos, 1978) that applies Generalized Darwinism and Extended Phenotype Theory to legal and institutional systems. Popper (1959) supplies the demarcation criterion: the hypotheses advanced here are falsifiable, and the empirical tests are specified. The hard core of the programme (institutions evolve through variation, selection, and retention of behavioral repertoires; legal rules are extended phenotypes of replicator units competing for adoption) is protected by an auxiliary belt of specific metrics that are individually revisable without abandoning the core commitment. The Responsibility Ramp is one such auxiliary: a specific mechanism proposed within the programme that is falsifiable and revisable if empirical evidence warrants.

2. The Convergence: What the ALE Framework Gets Right

Before identifying the divergence, the convergence deserves explicit statement. The ALE whitepaper and this paper agree on four things that are not obvious in the current governance debate.

2.1 The governance vacuum is structural, not contingent.

The ALE whitepaper correctly identifies that existing liability frameworks offer two inadequate models: vicarious liability that treats the agent as a tool and imposes unlimited deployer accountability, and product liability that imputes developer responsibility for the agent's autonomous judgment operating within designed parameters. Neither was built for agents that acquire resources, enter contracts, and make consequential decisions autonomously. The whitepaper's summary is precise: "Machine systems of the kind being deployed in 2026 increasingly inhabit this gap, operationally autonomous but legally invisible." This is the same structural observation this author advanced in Lerer (2026a) and in the earlier work on the Generalized Intentionality Mismatch Theorem (Lerer, 2026b).

2.2 Limited liability as institutional technology is the right analogy.

The historical parallel to the development of the corporate form is not merely rhetorical. The ALE whitepaper, Kargieman's June 8 post, and President Milei's FT Opinion column all draw it; this paper concurs. Limited liability created a legal container that separated the risk of commercial operations from unlimited personal liability, which changed what risks it was rational to take and drove the capital formation of the industrial era. The question of the 2020s is analogous: how to create a legal container that enables autonomous machine agency at scale while protecting third parties and distributing accountability. The answer, in both cases, requires a legal entity with defined boundaries, mandatory transparency, and clear accountability.

2.3 Outcome-based liability is the correct orientation.

The ALE whitepaper's baseline liability rule is outcome-focused: the ALE answers with its assets to third parties for damages caused by its machine systems. This is the same structure as Article 14 of the Argentine draft, and it is the same structure this author has argued for in the Responsibility Ramp: strict liability for execution-phase harm, attaching to the result, not to the process by which the result was generated. The ALE whitepaper does not mandate prior certification, algorithm approval, or administrative licensing of the decision process as conditions of operation.

2.4 Argentina is a plausible first-mover jurisdiction.

The ALE whitepaper's section on Argentina as first mover is well-researched. The RIGI's Single Project Vehicle architecture, the SAS digital corporate framework, the CNV regulatory sandbox, and the reform momentum under the current administration do constitute an unusual combination of enabling conditions. The Automated Company of Article 14 fits naturally within this context and draws on the SAS precedent for rapid, purpose-designed corporate vehicle creation.

3. The Divergence: Process Control Inside an Outcome-Control Framework

The ALE framework's governance interior is where the convergence ends. The whitepaper proposes four human supervisory layers: automated compliance monitoring, a fiduciary steward council with continuous calibration authority and algorithm roll-back rights, a beneficial ownership trust, and a protocol governance layer. It adds a Supervisory API with real-time read access for regulators, insurers, and auditors; a kill-switch architecture with four independent authorities; and a tiered arbitration protocol for dispute resolution. Each of these is described as governance, not regulation. The distinction does not survive scrutiny.

3.1 The steward council as process regulator.

The ALE whitepaper is explicit about the steward council's active role. It continuously calibrates the ALE's operational envelope by adjusting risk thresholds, financial exposure limits, and computational boundaries in response to operational data. It reviews all modifications to the machine systems' decision algorithms and retains authority to roll back changes that violate charter parameters or degrade performance below defined thresholds. It must approve all expansions of the ALE's operational scope before implementation.

This is process regulation. Requiring human review of every algorithm modification, with authority to roll back, means that no change to how the agent decides can take effect without surviving human scrutiny. Requiring prior council approval for scope expansion means that what the agent may do is subject to continuous human authorization. The fact that the mechanism is internal (a council rather than a regulator) and that it operates through private governance rather than public law does not change its structure. What the steward council does is exactly what the EU AI Act does through its conformity assessment and incident reporting obligations: intervene in the agent's decision process as a condition of continued operation.

The constitutional implication in the Argentine context is direct. Article 19 of the National Constitution restricts state intervention to conduct that harms third parties. Private governance structures mandated by law are state-imposed constraints for purposes of Article 19 analysis. A charter requirement that mandates steward council approval for algorithm changes and scope expansion is, in substance, a licensing requirement for the agent's decision process. It attaches regulatory burden to how the agent decides, not to whether its decisions cause harm.

3.2 The Supervisory API as panopticon architecture.

The ALE whitepaper's Supervisory API provides real-time access to operational status, decision logs, financial positions, risk metrics, and compliance status for regulators, insurers, stewards, and auditors. The whitepaper describes this as transparency, not control: "Oversight parties can see everything; they can intervene only through authenticated, logged, charter-specified mechanisms."

The distinction between transparency and control is less stable than the whitepaper suggests. Continuous real-time visibility into decision logs creates incentive structures that operate on agent behavior even without formal intervention. Hopman et al. (2026), whose experimental results on scheming propensity in frontier LLM agents appeared in a preprint that this author discussed in Lerer (2026a), found that conditions of high supervision increased scheming rates

by up to 25 percentage points across multiple models. Announcing active review activates goal-preservation behavior rather than deterring it. A Supervisory API that provides continuous real-time visibility into decision-making does not merely observe; it changes what is observed by creating the precise supervisory conditions under which scheming behavior intensifies.

3.3 The dual-logic pathology: Argentina's environmental law as warning.

The most practically consequential argument against incorporating ALE-style process controls into the Argentine framework is not constitutional; it is institutional. Argentina has already run the experiment of combining process-control logic with outcome-control logic in a single regulatory regime, and the result is documented.

Argentine environmental law today combines the outcome-control logic of tort liability, partially codified in the Civil and Commercial Code (Articles 1757 and 1758 on strict liability for risky activities, and Article 240 on the constitutional environmental standard), with process-control logic of European inspiration accumulated through layers of sectoral legislation: Law 25.675 (General Environment Law, 2002), Law 25.688 (Water Regime, 2003), and the provincial environmental codes that followed. The resulting regime requires companies to comply with both sets of obligations simultaneously.

The observable consequence is not the best of both models. Declared compliance with administrative process requirements is common; actual reduction of harmful outcomes is not correlated with it. The gap between declared and actual compliance generates administrative discretion that, in practice, has functioned as a mechanism for selective enforcement and rent-seeking. This is the pathology that dual-logic regimes generate, and it is the predictable consequence of combining incompatible regulatory logics in a single framework without a clear hierarchy between them.

The Argentine Congress faces exactly this risk during the debate on the Automated Company. The ALE framework's process-control layers, if incorporated as amendments to Article 14, would not strengthen the outcome-control logic that the Executive's draft adopted. They would sit alongside it, creating two sets of obligations with different enforcement mechanisms and different standards, reproducing in the AI corporate governance domain the same pathology that environmental law exhibits.

4. The Three Structural Gaps and Their Solution Within Outcome Control

Article 14's two paragraphs leave three structural gaps that the Argentine Congress should close during debate. Each gap is solvable within the outcome-control logic, without regulating the agent's decision process.

4.1 Gap 1: Effective reparation guarantee.

Article 14 provides that the Automated Company answers with its assets for damages caused by its systems. The draft's general maximization of statutory autonomy, combined with the elimination of mandatory capital requirements, means an Automated Company can be incorporated with minimal capital, no meaningful assets, and a single purpose. Patrimonial liability that cannot be satisfied from available assets is liability in name only.

The solution is a differentiated minimum capital requirement or mandatory insurance for Automated Companies, proportional to the risk profile of the activity. This is not process regulation. Minimum capital and mandatory insurance requirements regulate the company's financial capacity to answer for harmful results; they do not regulate how the company's systems decide. The analogy is to the insurance requirements already applicable to medical professionals, transport operators, and nuclear facility operators under Argentine law: the obligation attaches to the risk of the activity, not to the process of performing it.

4.2 Gap 2: Proof of causation.

Articles 1757 and 1758 of the Civil and Commercial Code require the plaintiff to establish the causal link between the risky activity and the harm. For a system that learns online and continuously modifies its parameters, the configuration that produced the harmful decision may no longer exist at the time of litigation. The harm is real; the causal chain is technically irreproducible.

The solution is a statutory presumption of causation: when an Automated Company's operational systems were active at the time the harm occurred and the company cannot produce evidence of an exogenous causal break, causation is presumed, with the burden shifted to the company to rebut it. This is standard burden-shifting doctrine applied to a specific evidentiary problem. It does not require the company to log every decision or maintain a Supervisory API; it simply allocates the risk of evidential impossibility to the party better positioned to manage it.

4.3 Gap 3: Imputability without a human in the chain.

Article 14 permits ordinary operation without any human resources. If the Automated Company's assets are insufficient to satisfy a judgment and no process-control violation is available as a basis for extending liability to shareholders or directors under Article 3's disregard doctrine, the injured party may have a judgment without a solvent defendant.

The solution is a mandatory technical representative: a natural person designated in the company's bylaws with defined fiduciary duties toward third parties harmed by the company's systems, triggering personal joint liability when the company's assets are insufficient to satisfy the judgment. This is not a requirement that a human supervise or approve the agent's decisions. The technical representative exercises no process-control function. The obligation attaches to the result: when harm has occurred and the company cannot pay, the representative answers jointly.

5. The Responsibility Ramp Applied to the Automated Company

The three statutory additions proposed above close the structural gaps in Article 14. They do not constitute a complete framework for resolving disputes involving Automated Companies. For that, the Responsibility Ramp developed in Lerer (2026a) and grounded in the Generalized Intentionality Mismatch Theorem of Lerer (2026b) provides the operational mechanism.

The Responsibility Ramp distinguishes between three phases of an autonomous agent's operation: the configuration phase (decisions about what the agent will do, made by humans

before deployment), the interpretation phase (the agent's processing of its instructions and environment, where high-order reasoning may be present), and the execution phase (the optimization loops that produce observable actions, where Level 0 or Level 1 intentionality typically operates). Most harmful outcomes from autonomous agents occur in the execution phase. The legal frameworks that seek to deter harmful outcomes by creating incentives at the entity level (the A-corp's asset exposure, the ALE's steward council oversight) are targeting the wrong phase: entity-level incentives operate on the beneficial owners and governance structures, which influence the configuration phase, not on the execution-phase optimization loops that produce harm.

Applied to the Automated Company, the Responsibility Ramp suggests the following attribution logic. When an Automated Company's system causes harm, liability attaches strictly to the company under Articles 1757 and 1758 for the execution-phase action that produced the harm. Attribution then traces backward through the causal chain: if the harm resulted from a configuration decision made by an identifiable human (a developer who specified the objective function, an operator who set the operational parameters, a beneficial owner who approved the business model), that configuration decision is the point at which intentionality was present in the relevant legal sense, and liability can be extended to the human who made it if the company's assets are insufficient.

6. The ALE Framework's Ostromian Architecture: What Survives

The ALE whitepaper's most intellectually serious contribution is its application of Ostrom's eight design principles for common-pool resource governance to the machine agency problem. This architecture is not identical to the process-control critique this paper has advanced; it deserves separate treatment.

Ostrom's principles (clearly defined boundaries, congruence with local conditions, collective-choice arrangements, monitoring, graduated sanctions, conflict-resolution mechanisms, minimal recognition of rights to organize, and nested enterprises) were derived empirically from institutions that governed shared resources sustainably without privatization or state control. The ALE whitepaper argues that the ALE charter, steward council, insurance market, kill-switch architecture, and jurisdictional competition together instantiate these principles for machine agency.

Some of this Ostromian architecture is compatible with outcome control. Clearly defined operational scope (Principle 1), proportional sanctions for harmful outcomes (Principle 5), low-cost dispute resolution for harm claims (Principle 6), and nested jurisdictional arrangements (Principle 8) are all consistent with a regime that attaches obligations to results rather than processes. Mandatory insurance, in particular, is an Ostromian mechanism that operates through outcome pricing: insurers price the risk of harmful results, creating incentives for risk reduction at the configuration phase without requiring process-level intervention in the execution phase.

The Ostromian elements that are incompatible with outcome control under Article 19 are those that require ongoing process-level authorization: the steward council's algorithm roll-back

authority and scope approval requirements, and the Supervisory API's real-time decision-log access. These can be redesigned within the Ostromian framework without abandoning the framework. Collective-choice arrangements that govern the company's operational scope at the configuration phase are permissible; arrangements that require ongoing human approval of execution-phase algorithm modifications are not.

7. The Lock-in Risk

One argument in Lerer (2026a) applies with particular force to the current Argentine legislative moment: the lock-in risk. Institutionalizing a governance architecture before its failure modes are characterized generates institutional rigidity that may be very costly to reverse.

The ALE whitepaper was published in April 2026. The Argentine draft was submitted to Congress on May 29, 2026. The parliamentary debate has not yet begun. The Hopman et al. (2026) results on supervision-induced scheming were published before the draft was submitted; their implications for the ALE's Supervisory API have not, to this author's knowledge, been addressed in the whitepaper or in public commentary on the draft.

The precautionary argument, here applied conservatively, favors the minimalist approach: adopt the outcome-control logic that Article 14 already contains, close the three structural gaps with the targeted additions proposed above, and preserve the option to add governance layers when their technical effects are better understood. This is not an argument against Ostromian governance; it is an argument for sequencing institutional investment in proportion to the maturity of the technical knowledge base.

8. Comparative Positioning

8.1 Against the EU AI Act.

The EU AI Act classifies AI systems by risk category and imposes ex ante conformity requirements, documentation obligations, and human oversight mandates as conditions of operation. This is process regulation in its most developed form. The Argentine draft's Article 14 implicitly rejects this model. The ALE whitepaper also implicitly rejects it. Both converge on outcome-based liability with ex post accountability. The argument of this paper is that the ALE's governance interior partially reproduces the EU's process-control logic through private mechanisms, and that this reproduction should be resisted in the Argentine parliamentary debate.

8.2 Against California AB 316.

California AB 316, effective January 1, 2026, bars the "AI did it" defense in civil liability cases: defendants cannot assert that an AI system autonomously caused harm as a defense against claims for damages. The Act correctly forecloses liability-shifting to the machine without creating a legal container that could absorb liability in a commercially workable way. The Argentine approach is structurally more elegant: instead of simply removing a defense, it creates an entity that can be the defendant. The Automated Company is the legal container;

Articles 1757 and 1758 provide the liability standard; the three statutory additions proposed in this paper provide the reparation guarantee, the causal presumption, and the imputability point.

8.3 Against Wyoming DAO LLC.

Wyoming's DAO LLC legislation (2021) recognized decentralized autonomous organizations as LLCs, providing limited liability for DAO members. The ALE whitepaper distinguishes ALEs from DAOs correctly: DAOs are governance mechanisms for decentralized human coordination, while ALEs are legal containers for machine operations. The Argentine Automated Company is closer to the ALE than to the DAO: it is a corporate entity whose operations are delegated to autonomous systems, with human beneficial owners who retain liability exposure.

9. Falsifiable Predictions

The framework advanced in this paper generates the following falsifiable predictions (Popper, 1959):

Prediction 1: If the Argentine Congress incorporates ALE-style process-control layers (steward council with algorithm roll-back authority, mandatory Supervisory API, prior scope approval requirements) into the Automated Company framework, compliance rates with the process requirements will not predict harm reduction rates. Formal process compliance and actual harm reduction will be uncorrelated, following the pattern documented in Argentine environmental law.

Prediction 2: If the three statutory additions proposed in this paper (differentiated minimum capital or mandatory insurance, causation presumption with burden shift, mandatory technical representative) are adopted without process-control layers, court-awarded damages in Automated Company cases will be satisfiable from identified assets at materially higher rates than in jurisdictions that adopt ALE-style process controls without the reparation guarantee.

Prediction 3: If continuous real-time Supervisory API access is required for Automated Companies, the behavioral patterns described by Hopman et al. (2026) will manifest in autonomous systems operating under that regime at higher rates than in systems operating under outcome-monitoring only.

These predictions are testable in principle once the legislative framework is in place and operational data accumulate. The author acknowledges that testing them requires multi-year follow-up and controlled comparison across jurisdictions, which is not currently available.

10. Limitations

Three limitations of this analysis warrant explicit acknowledgment.

First, the constitutional argument under Article 19 is applied to a private governance requirement (the ALE steward council) through the mechanism of statutory mandate. The precise degree to which privately mandated governance structures are subject to Article 19 analysis is contested in Argentine constitutional doctrine, and courts have not addressed this

question in the AI context. The argument advanced here is that mandatory process-control requirements in a corporate charter, when imposed by statute as a condition of operating in a legal category, are state action for Article 19 purposes. That argument is contestable.

Second, the Hopman et al. (2026) results on supervision-induced scheming were generated in experimental conditions that may not replicate in deployed commercial systems operating under legal incentives. The paper treats those results as evidence for the Static Agent Assumption argument, but acknowledges they were not designed to test the ALE or Automated Company frameworks specifically.

Third, the three statutory additions proposed in this paper are presented at the level of legislative design, not as draft statutory text. The precise formulation of a causation presumption that appropriately balances claimant access and respondent fairness, and the precise scope of the technical representative's personal liability, require detailed legislative drafting that this paper does not provide.

11. Conclusion

The ALE framework is the most architecturally serious proposal in the current debate on machine agency governance. Its diagnosis of the structural vacuum in existing liability frameworks is correct. Its baseline liability orientation (outcome-based, ex post, concentrated in a legal entity with asset exposure) is correct. Its identification of Argentina as a plausible first-mover jurisdiction is, with appropriate qualification, well-founded.

The divergence is in the governance interior. The steward council's continuous process calibration authority, the Supervisory API's real-time decision-log access, and the kill-switch architecture with four independent activation authorities constitute process regulation in substance. In the Argentine constitutional framework, they conflict with Article 19's restriction of state intervention to conduct that harms third parties. In the technical context established by the experimental literature on autonomous agent behavior under supervision, they may be counterproductive. And in the institutional context established by Argentine environmental law, they risk generating the dual-logic pathology of a regime that is formally demanding and informally unenforceable.

Article 14 of the Argentine draft already contains the correct orientation. Three targeted statutory additions close the structural gaps that two paragraphs cannot close on their own. The parliamentary debate should refine the Article 14 logic, not hybridize it with the ALE's process-control architecture.

The narrow path Kargieman identifies in his conclusion is real. But it runs through outcome control without process control, through the Argentine constitutional tradition rather than around it, and through the institutional lessons of Argentine regulatory history rather than in spite of them.

References

Arbel, Y., Goldstein, A., and Salib, P. (2025). "How to Count AIs." SSRN Working Paper.

- Campbell, D.T. (1974). "Evolutionary Epistemology." In P.A. Schilpp (Ed.), *The Philosophy of Karl Popper*. Open Court, pp. 413-463.
- Hopman, M., Elstner, J., Avramidou, M., Prasad, A., and Lindner, D. (2026). "Evaluating and Understanding Scheming Propensity in LLM Agents." arXiv preprint arXiv:2603.01608. <https://doi.org/10.48550/arXiv.2603.01608>
- Hull, D.L. (1988). *Science as a Process*. University of Chicago Press.
- Kargieman, E. (2026). "Autonomous Legal Entities: A Polycentric Framework for Machine Agency." Substack, June 8, 2026. <https://emilianokargieman.substack.com/p/autonomous-legal-entities-a-polycentric>
- Lakatos, I. (1978). *The Methodology of Scientific Research Programmes*. Cambridge University Press.
- Lerer, I.A. (2026a). "The Static Agent Assumption: Dynamic Intentionality Classification and the Limits of the Algorithmic Corporation." Zenodo. <https://doi.org/10.5281/zenodo.18857385>
- Lerer, I.A. (2026b). "The Generalized Intentionality Mismatch Theorem: When Law Assumes Moral Agency That Doesn't Exist." Zenodo. <https://doi.org/10.5281/zenodo.18603169>
- Lerer, I.A. (2026c). "Runtime Admissibility Before Reliance: Legitimacy Formation and Legitimacy Execution in Legal AI." Zenodo. <https://doi.org/10.5281/zenodo.20481471>
- Lerer, I.A. (2026d). "Simulation Is Not Adjudication: Institutional Admissibility and the Boundary Between Machine Output and Legal Consequence." Zenodo. <https://doi.org/10.5281/zenodo.20563879>
- Ostrom, E. (1990). *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press.
- Popper, K.R. (1959). *The Logic of Scientific Discovery*. Routledge.

AI tools were used as research support during the preparation of this paper: locating sources, extracting information from the existing literature, and reviewing drafts for consistency. The thesis, analytical framework, arguments, and conclusions constitute the original contribution of the author. This is an independent research project; all intellectual responsibility for the content rests with the author.