

Regulatory Prophecy as Parasitic Spontaneous Order: Epistemological Clergy, Panic Memes, and the Distortion of AI Governance

Ignacio Adrián Lerer

Independent Researcher, Buenos Aires, Argentina

adrian@lerer.com.ar | ORCID: 0009-0007-6378-9749

Abstract

This paper identifies a structural pattern in the governance of emerging technologies: the panic meme as a parasitic spontaneous order (PSO) that generates epistemological clergy administering regulatory orthodoxy, excommunicating empirical dissent, and producing governance frameworks whose compliance costs fall disproportionately on transparent actors while opaque actors structure around them. The pattern is documented in climate policy through Bjorn Lomborg's twenty-year retrospective on Al Gore's *An Inconvenient Truth* (2006), which demonstrates that predictions with high moral certainty produced policies of extraordinary cost and negligible effectiveness. The paper argues this pattern is replicating in AI governance: the prophetic register shared by Bostrom, Thiel, and the EU AI Act's risk-classification architecture shares four structural features with the climate panic memplex: high moral certainty over uncertain risk, deferred apocalyptic horizon, conversion of empirical dissent into moral transgression, and production of clergy gatekeeping regulatory access. The Argentine draft General Companies Law, whose Article 14 on the Automated Company adopts outcome control rather than process control, constitutes a natural experiment: the risk is not that the draft is wrong, but that parliamentary debate will contaminate it with process-control layers under clergy pressure, reproducing the dual-logic pathology documented in Argentine environmental law. Three falsifiable predictions are derived from the PSO/clergy framework and mapped to observable outcomes in AI governance across jurisdictions.

Keywords: parasitic spontaneous order, epistemological clergy, panic meme, AI governance, outcome control, process control, EU AI Act, climate policy, extended phenotype theory, memetic selection

1. Introduction: Two Panics, One Structure

In May 2026, Bjorn Lomborg published a retrospective on Al Gore's documentary *An Inconvenient Truth* twenty years after its release. The data he reports are striking. According to Lomborg (2026), deaths from climate-related disasters have fallen more than 97 percent over the past century despite global population quadrupling; global hurricane frequency and total energy have shown a slight decline since full-coverage satellite data became available in 1980; the area burned annually by wildfires has fallen more than 25 percent over the past quarter century; polar bear populations have more than doubled, from approximately 12,000 in the 1960s to over 26,000 today; the global cost of climate policies since 2006 has exceeded USD 16 trillion; and fossil fuels' share of total global energy supply has moved from 82.6 percent in

2006 to 81.1 percent in 2023 (Lomborg, 2026, citing IEA data). At that rate of transition, Lomborg calculates, zero fossil fuel dependency lies approximately six centuries away.

Lomborg's conclusion: *"The main lesson of An Inconvenient Truth is that panic is a terrible policy adviser."*

That conclusion is not an argument against climate science. Lomborg explicitly accepts that climate change is real and that warming will impose costs. His argument is about the structural relationship between prophetic certainty, the policy frameworks certainty generates, and the empirical outcomes those frameworks produce. High moral certainty over uncertain risk, operationalized through regulatory frameworks whose costs are front-loaded and whose benefits are deferred, tends to produce compliance theater at scale rather than actual harm reduction.

This paper argues that the same structural pattern is replicating in AI governance, and that identifying it now, before the regulatory lock-in that characterized climate policy consolidates, is the work that the governance literature currently requires.

The pattern has a precise theoretical description within the Extended Phenotype Theory (EPT) research programme (Lerer, 2025a; 2026a). It is an instance of Parasitic Spontaneous Order (PSO): an order that arises without a designer, through the competitive replication of memes that enhance their own transmission at the expense of systemic efficiency. The PSO of regulatory prophecy generates epistemological clergy (Lerer, 2025b) that administer the orthodoxy of the panic, excommunicate empirical dissent, and produce regulatory frameworks that transfer rents to the actors best positioned to navigate compliance complexity. The content of the prophecy matters less than its architecture: the same structural features that made Gore's climate predictions effective at generating regulatory momentum also made them poor predictors of actual outcomes.

Epistemological positioning. This paper belongs to a Lakatosian research programme (Lakatos, 1978) applying Generalized Darwinism and Extended Phenotype Theory to legal and institutional systems. The hard core (institutions evolve through variation, selection, and retention of behavioral repertoires; legal rules are extended phenotypes of replicator units competing for adoption) is protected by an auxiliary belt of specific instruments. PSO and the epistemological clergy framework are auxiliary belt components, falsifiable and revisable. Popper (1959) supplies the demarcation criterion: the predictions in Section 6 are falsifiable against observable governance outcomes.

2. The Prophetic Pattern: Four Structural Features

The panic meme in governance contexts shares four structural features across instances that are otherwise ideologically diverse. These features are not accidental; they are adaptive. Each enhances the meme's capacity to generate regulatory momentum, which is the primary mechanism through which the PSO of regulatory prophecy extracts value from institutional environments.

2.1 High moral certainty over uncertain risk.

Prophetic governance discourse operates at a higher confidence level than the underlying empirical uncertainty warrants. Gore's documentary presented its predictions as settled science when the relationship between climate change and specific disaster outcomes remained, and remains, actively debated among climate scientists. The Intergovernmental Panel on Climate Change's own assessments, on which the documentary nominally drew, assigned probability ranges to projections that the documentary converted into certainties.

The AI governance literature exhibits the same feature. Nick Bostrom's *Superintelligence* (2014) assigns high probability to catastrophic outcomes from artificial general intelligence on timescales of decades. The EU AI Act's preamble treats the risks of certain AI system categories as established facts warranting prior restraint rather than contested hypotheses warranting continued research. The empirical basis for these probability assignments is, to put it charitably, thin: there are no confirmed cases of artificial general intelligence, no validated models of capability trajectories that would ground the specific probability estimates that prophetic discourse requires, and substantial expert disagreement about both the timeline and the nature of the relevant risks.

The mismatch between confidence level and evidentiary base is not an error to be corrected by better data. It is a functional feature of the prophetic pattern. High confidence is what generates the urgency that regulatory momentum requires.

2.2 The deferred apocalyptic horizon.

Prophetic governance predictions are characteristically unfalsifiable in the short run because the predicted catastrophe is always sufficiently far in the future to preclude immediate testing, and sufficiently near to generate present urgency. Gore's 2006 documentary implied that the most severe consequences of inaction were one to two decades away. Two decades later, the predicted consequences have not materialized at the predicted scale, but the prophetic register has updated its horizon rather than acknowledging the prediction failure.

The AI governance literature operates on the same temporal structure. Bostrom's original formulations placed transformative AI capability within 'decades.' As decades pass without the predicted developments, the horizon is revised rather than the model revised. The EU AI Act's risk classification system is designed for systems that may not yet exist at the capability levels that justify its most restrictive provisions, but the regulatory framework is built as if they already do. This temporal architecture is self-protecting: it cannot be falsified in the short run because the catastrophe is always ahead, and by the time the prediction window closes, the regulatory framework has generated sufficient institutional interests to resist revision regardless of outcomes.

2.3 Conversion of empirical dissent into moral transgression.

The most diagnostic feature of the prophetic pattern, and the one that most directly links it to the epistemological clergy framework (Lerer, 2025b), is the conversion of empirical disagreement into moral accusation. Lomborg's treatment in the climate debate is the paradigm case. His empirical claims are, on balance, more accurate than Gore's: his predictions about

disaster mortality trends, polar bear populations, and the pace of energy transition have proven more precise than the apocalyptic forecasts they challenged. Yet his reception in mainstream climate discourse has been characterized not by engagement with his empirical claims but by attacks on his moral standing and motivations.

The AI governance context is generating the same conversion. Scholars and practitioners who argue for outcome-control frameworks are characterized as naive about catastrophic risk, as aligned with technology industry interests, or as failing to take the gravity of AI risks seriously. The empirical argument, that prior process regulation may be less effective at reducing harm than outcome-focused liability, and may generate compliance theater similar to what climate policy produced, does not receive sustained engagement. It is converted into a position on the moral register: insufficient seriousness about existential risk.

2.4 Production of clergy that administers regulatory access.

The prophetic pattern generates, as its institutional output, a specialized community that controls access to the regulatory discourse it has helped create. In climate policy, this community consists of the network of consultants, NGOs, standards bodies, and academic departments whose existence and funding depend on the continuation of the regulatory framework the prophecy justified. In AI governance, it is emerging as the network of AI safety researchers, EU AI Act compliance consultants, conformity assessment bodies, and academic centers whose existence the Act's provisions require or incentivize.

This community is not a conspiracy. It is a PSO output: an order that arises without a designer from the competitive dynamics of actors whose fitness is enhanced by the continuation of the regulatory framework. Climate policy generated a USD 16 trillion compliance industry over twenty years; that industry has strong structural incentives to sustain the regulatory apparatus that created it, independent of the apparatus's effectiveness at achieving its stated goals.

3. The Clergy Mechanism in AI Governance

The epistemological clergy framework (Lerer, 2025b) identifies five operational criteria for diagnosing clergy-dominated discourse. Applying them to the current AI governance debate yields a concerning pattern.

Sacralization. The proposition that AI systems above certain capability thresholds require prior process regulation before deployment has acquired the character of a transcendental truth in EU policy discourse. It is not treated as a hypothesis to be tested against evidence of regulatory effectiveness, but as a conclusion from which specific regulatory architectures follow. The EU AI Act does not contain provisions for revising its risk classification system if ex ante process requirements prove ineffective at reducing harm; its amendment procedures require the same legislative process that produced it.

Costly signaling. Endorsing the most expansive interpretations of AI risk, including existential risk framings that most AI researchers consider speculative, has become a costly signal of serious engagement in AI safety discourse. The costliness is real: endorsing existential risk framings makes one vulnerable to empirical challenge and imposes reputational risk if the

predictions do not materialize on the predicted timescale. The willingness to bear that risk signals genuine commitment to the safety agenda.

Clergy gatekeeping. The institutions that control access to AI policy discourse, major AI safety research centers, EU technical advisory bodies, and the journals that specialize in AI governance, exhibit systematic skewing toward process-control frameworks. The published literature on AI governance is dominated by frameworks that take ex ante process regulation as their baseline assumption and treat outcome-control alternatives as proposals requiring justification, rather than the reverse.

Generational transmission. AI safety as a field is young enough that its transmission mechanisms are visible rather than naturalized. The major AI safety research organizations are explicitly training the next generation of researchers within a framework that treats certain risk propositions as foundational. Graduate students who enter AI safety research through these channels are socialized into the risk framing as a precondition for participation, not as a conclusion from evidence they assess independently.

Epistemic closure. The conversion described in Section 2.3 is the operational form of epistemic closure in AI governance discourse. Proposals for outcome-control frameworks are not engaged on their empirical merits; they are converted into positions on the risk-seriousness register and addressed at that level. The governance literature has not produced a serious comparative analysis of ex ante process control versus ex post outcome control in terms of harm reduction effectiveness, because that comparison presupposes that outcome control is a legitimate alternative worth comparing, which the clergy structure does not recognize.

4. The Lomborg Test: Empirical Precision as Heresy

Lomborg's twenty-year retrospective provides what can be called the Lomborg Test: a methodology for evaluating governance frameworks by comparing the precision of their predictions against empirical outcomes, controlling for the institutional treatment of the predictor.

Applied to climate governance, the test yields a striking result. Lomborg's predictions about disaster mortality trends, polar bear population dynamics, the pace of energy transition, and the cost-effectiveness of renewable energy deployment have proven more accurate than Gore's. Yet Lomborg's institutional treatment, characterized by exclusion from mainstream climate policy discourse, academic criticism focused on his motives rather than his methodology, and sustained attempts to delegitimize him as a 'denier,' has been inversely proportional to his predictive accuracy. The clergy mechanism operates precisely as the framework predicts: empirical precision that challenges the prophetic narrative is treated as heresy rather than as evidence.

The Lomborg Test generates a prediction for AI governance: as the predictions of the prophetic register fail to materialize on their predicted timescales, the institutional response will be prediction update (revising the horizon) rather than framework revision (acknowledging that the confidence levels were miscalibrated). This is already observable. The timelines for

transformative AI capability that characterized AI safety discourse in 2014 have been revised repeatedly without the underlying confidence architecture being revised.

A secondary prediction follows: researchers who propose outcome-control frameworks with falsifiable predictions, and whose predictions prove more accurate than those of the prophetic register, will receive treatment analogous to Lomborg's rather than being incorporated into the mainstream governance discourse through normal evidential updating.

5. The Argentine Case as Natural Experiment

The Argentine draft General Companies Law, submitted to Congress on May 29, 2026, provides a natural experiment for the PSO/clergy framework at a moment when the experiment is still in its initial conditions.

Article 14 of the draft creates the Automated Company: a corporate vehicle that operates through autonomous algorithmic systems or AI agents without requiring human employees for ordinary operations, and that answers with its assets to third parties for damages caused by those systems. The article adopts outcome control as its regulatory orientation without using that terminology: it attaches liability to results, not to process, and imposes no prior certification, algorithm approval, or administrative licensing of the decision process.

As argued in Lerer (2026b), this orientation is constitutionally grounded in Article 19 of the Argentine National Constitution, which restricts state intervention to conduct that harms third parties. It is also technically coherent with the experimental evidence on supervision-induced scheming in autonomous systems documented by Hopman et al. (2026): process-control requirements may aggravate rather than mitigate the behavioral risks they seek to address.

The natural experiment consists of observing whether the parliamentary debate on Article 14 is contaminated by clergy pressure toward process-control additions. The pressure is predictable: the EU AI Act is the dominant reference framework in international AI policy discourse, and Argentine legislators seeking to demonstrate seriousness about AI governance will face strong incentives to incorporate EU-style process requirements as amendments. The contamination, if it occurs, will reproduce in AI governance the dual-logic pathology that Argentine environmental law exhibits: a formally demanding, informally unenforceable regime whose compliance gap generates administrative discretion functioning as rent extraction.

Argentine environmental law provides the counterfactual within the same jurisdiction. Law 25.675 (General Environment Law, 2002) installed European process-control logic on top of the Civil and Commercial Code's outcome-control damage liability framework. The resulting dual-logic regime is formally demanding and informally unenforceable; the gap between declared compliance and actual harm reduction generates administrative discretion that has functioned as a rent extraction mechanism across enforcement cycles. The prediction is that AI governance contaminated by the same dual-logic hybridization will exhibit the same pathology.

6. Falsifiable Predictions

The PSO/clergy framework applied to AI governance generates three falsifiable predictions (Popper, 1959) that are testable against observable governance outcomes.

Prediction 1: Compliance rates and harm reduction will be uncorrelated in EU AI Act jurisdictions. If the EU AI Act's ex ante process requirements function as clergy-generated compliance theater analogous to what climate policy produced, then the jurisdictions that implement them most fully will show no systematically better outcomes in AI-related harm incidents than jurisdictions that adopt outcome-control frameworks. The metric is the ratio of compliance rate (process) to harm reduction rate (outcome) across jurisdiction pairs with different regulatory architectures, measured over a five-year implementation window. This prediction is falsified if EU AI Act jurisdictions show systematically better harm outcomes than outcome-control jurisdictions, controlling for the type and density of AI system deployment.

Prediction 2: Compliance costs will fall disproportionately on domestic transparent actors relative to offshore-structured opaque actors. The PSO of regulatory prophecy extracts value specifically from actors who cannot structure around its compliance requirements. This structural asymmetry is the mechanism by which the PSO generates rents for the compliance industry that services domestic transparent actors. This prediction is testable as EU AI Act implementation proceeds: the question is whether compliance cost incidence is systematically higher for domestic transparent actors than for equivalent deployments structured through offshore or lower-regulation jurisdictions.

Prediction 3: The intensity of clergy response to outcome-control proposals will increase as process-control frameworks fail to produce predicted harm reductions. The Lomborg Test predicts that the treatment of outcome-control proposals will become more hostile, not less, as process-control frameworks demonstrate effectiveness gaps. This counterintuitive prediction follows from the PSO/clergy mechanism: clergy response to anomalies that challenge orthodoxy is intensification of the orthodoxy, not revision. The observable test is the bibliometric and policy treatment of outcome-control AI governance proposals over the medium term.

7. Limitations

Three limitations of this analysis warrant explicit acknowledgment.

First, the parallel between climate governance and AI governance is structural, not empirical. The paper argues that both instances share the four architectural features of the prophetic pattern. It does not argue that the empirical content of the claims is equally wrong. Climate change is real; AI risks are real. The argument is about the relationship between the confidence architecture of prophetic governance discourse and the quality of the policy outputs it generates, not about the truth of the underlying risk claims.

Second, the PSO/clergy framework is applied here at a high level of abstraction. Demonstrating that the EU AI Act exhibits the five clergy criteria requires more systematic bibliometric and policy analysis than this paper provides. The criteria application in Section 3 is diagnostic rather than rigorous; a full application of the methodology from Lerer (2025b) to AI governance discourse would require a dedicated empirical study.

Third, the Argentine natural experiment identified in Section 5 has only one observation at the point of writing. The parliamentary debate on Article 14 has not yet occurred. The prediction about dual-logic contamination is testable but not yet tested.

8. Conclusion

Lomborg's twenty-year retrospective on *An Inconvenient Truth* offers a precise empirical test of a theoretical claim: that prophetic governance discourse, characterized by high moral certainty over uncertain risk, converts empirical dissent into moral transgression, generates epistemological clergy administering regulatory access, and produces frameworks whose costs are real and front-loaded while whose benefits are deferred and often unrealized.

The test results are clear for climate governance. The question is whether the same structural pattern is replicating in AI governance before the lock-in that characterized climate policy consolidates.

The Argentine case is valuable precisely because it is early. Article 14 of the draft General Companies Law has adopted an outcome-control orientation that is constitutionally grounded and technically coherent. The clergy pressure toward process-control contamination is predictable and already observable. The parliamentary debate is still open.

Panic may be a poor policy adviser. The more important question is whether the institutional mechanisms that panic generates can be identified early enough to preserve the space for alternatives.

References

- Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- Dennett, D.C. (2006). *Breaking the Spell: Religion as a Natural Phenomenon*. Viking.
- Gore, A. (2006). *An Inconvenient Truth*. Paramount Classics.
- Hopman, M., Elstner, J., Avramidou, M., Prasad, A., and Lindner, D. (2026). "Evaluating and Understanding Scheming Propensity in LLM Agents." arXiv preprint arXiv:2603.01608. <https://doi.org/10.48550/arXiv.2603.01608>
- Hull, D.L. (1988). *Science as a Process*. University of Chicago Press.
- Lakatos, I. (1978). *The Methodology of Scientific Research Programmes*. Cambridge University Press.
- Lerer, I.A. (2025a). "Law as Extended Phenotype: Reframing Legal Theory Through Evolutionary Epistemology." SSRN Working Paper No. 5737383.
- Lerer, I.A. (2025b). "Epistemological Clergies: When Orthodoxy Blocks Rational Debate." SSRN Working Paper No. 5768423. <https://doi.org/10.2139/ssrn.5768423>
- Lerer, I.A. (2026a). "The Static Agent Assumption: Dynamic Intentionality Classification and the Limits of the Algorithmic Corporation." Zenodo. <https://doi.org/10.5281/zenodo.18857385>

Lerer, I.A. (2026b). "Outcome Control as Constitutional Constraint: The Argentine Model for Autonomous Legal Entities." Zenodo. <https://doi.org/10.5281/zenodo.20613712>

Lerer, I.A. (2026c). "The Generalized Intentionality Mismatch Theorem: When Law Assumes Moral Agency That Doesn't Exist." Zenodo. <https://doi.org/10.5281/zenodo.18603169>

Lomborg, B. (2026). "Al Gore's inaccurate untruths distorted policy." Financial Post, May 27, 2026.

Lomborg, B. (2020). *False Alarm: How Climate Change Panic Costs Us Trillions, Hurts the Poor, and Fails to Fix the Planet*. Basic Books.

Ostrom, E. (1990). *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press.

Popper, K.R. (1959). *The Logic of Scientific Discovery*. Routledge.

AI tools were used as research support during the preparation of this paper: locating sources, extracting information from the existing literature, and reviewing drafts for consistency. The thesis, analytical framework, arguments, and conclusions constitute the original contribution of the author. This is an independent research project; all intellectual responsibility for the content rests with the author.