

Selected, Not Self-Evident: General Principles of Law as an Evolutionarily Stable Memetic Core

Ignacio Adrián Lerer

Independent Researcher | Buenos Aires, Argentina
adrian@lerer.com.ar | ORCID 0009-0007-6378-9749

June 2026

Abstract

Doctrine treats certain legal principles as universal: good faith, the right to be heard, the bindingness of agreements, the prohibition on judging one's own cause. The standard account grounds their authority in rational necessity or self-evidence, as if any legal order confronting the same problems would have to recognize them. This paper argues that this universalist self-understanding commits a structural error recently identified in a different domain by Goldfeder, Wyder, LeCun, and Shwartz-Ziv (2026), who demonstrate that human intelligence feels general only because a mind cannot perceive the tasks it cannot perform. The error is symmetrical in law: a jurist trained inside the history that produced the principles encounters them everywhere precisely because that history shaped the jurist and the institutions in which the jurist works. The sense of universality is the illusion of generality in legal dress.

This paper proposes an evolutionary account on which the general principles are not universals given in advance but a memetic core that reached quasi-fixation through selection. Their apparent universality is the observable joint signature of three replicator properties identified by Dawkins: high copying fidelity, carried by Latin formulae that resist semantic drift because Latin is a dead language; fecundity, reproduced in codes, curricula, and appellate courts across two millennia; and longevity, with Roman maxims persisting across the collapse of the political orders that first carried them. Quasi-universality, on this reading, is quasi-fixation in a meme pool, an evolutionarily stable configuration that resists invasion because deviating jurists pay coordination and legitimacy costs unrelated to the merits of their deviation.

The reframing is falsable. It predicts cross-jurisdictional variation in the enumerated set tracking legal lineage, locally extinct principles, and drift under environmental change. Each prediction has empirical content: *caveat emptor* underwent documented sequential extinction across jurisdictions as consumer protection legislation changed the selective environment; the *abus de droit* doctrine emerged independently in Louisiana and continental Europe without cross-citation, demonstrating convergent evolution under shared selective pressure; and the arrival of new agent types, corporations in the nineteenth century and artificial systems now, constitutes a shift in the environment that selects against principles calibrated for natural persons. Strict universalism anticipates none of these observations. One consequence for computational legal systems follows: a principle layer encoded as a closed list of universals mistakes a selected strategy for an axiom and should instead be modeled as an open, revisable set observed in a given legal environment. The argument is a protective-belt adjustment within the extended-phenotype research programme, not a refutation of it.

Keywords: general principles of law; legal memetics; extended phenotype; evolutionarily stable strategy; legal evolution; artificial intelligence and law; comparative law; *caveat emptor*; convergent legal evolution.

1. The Universalist Self-Understanding and Its Costs

Every trained jurist inherits a set of principles that functions more like bedrock than like positive law. Good faith, *audi alteram partem*, *nemo iudex in causa sua*, *pacta sunt servanda*: these are taught as if any rational legal order, confronting the same problems, would have to recognize them. The standard account grounds their authority in recognition rather than creation. A legal system that lacked good faith doctrine would be defective in the way that a logical system lacking the law of non-contradiction would be defective, not merely different.

This self-understanding does real institutional work. It licenses importing a principle into a system that never enacted it, on the ground that the principle was always implicitly present. It underwrites unwritten constraints on legislative and administrative power: constitutional courts across traditions have held that certain principles operate as limits even absent explicit textual authorization, on the argument that they are too fundamental to require stating. Article 38(1)(c) of the Statute of the International Court of Justice (1945) lists the general principles of law recognized by civilized nations as a source of international law, placing them alongside treaties and custom but without requiring that any specific state enacted them, a concession that their authority derives from something other than positive enactment.

The universalist account also shapes the design of computational legal systems. Artificial intelligence tools built to reason about law increasingly encode a catalogue of principles as a control or audit layer, on the assumption that the catalogue captures law's necessary content. If the principles are universal, a fixed list is coherent: the same list applies across all jurisdictions and all periods, and a system trained on it needs no mechanism for revision.

Each of these functions depends on the universalist account being correct. The present paper argues that it is not, and that the cost of treating a selected strategy as an axiom is not merely intellectual. A fixed list cannot track extinction; it cannot anticipate drift; it cannot accommodate the arrival of new agent types that do not fit the intentionality structure the principles presuppose. Understanding what the principles actually are, the result of a selective history rather than a logical necessity, is a precondition for using them well.

2. The Illusion of Generality and Its Legal Counterpart

Goldfeder, Wyder, LeCun, and Shwartz-Ziv (2026) make an argument about intelligence that has an exact counterpart in legal doctrine. Their target is the standard framing of artificial general intelligence as intelligence that does everything a human can do. They deny that this target is coherent, and they do so by denying that human intelligence is general in any meaningful sense. Humans are specialized adaptations, honed by selection over hundreds of thousands of years for a narrow band of evolutionarily relevant tasks. The sense of generality is an artifact of the limits of self-perception: a mind cannot perceive the tasks it cannot perform, so it takes the boundary of its competence for the boundary of the possible. They call this the illusion of generality, and they

trace it to a simple asymmetry: selection leaves traces of what it built and no trace of what it did not.

The argument draws directly on ecological specialization theory. Forister and colleagues (2012) show that organisms carry traits suited to their evolutionary niche but never to all possible niches simultaneously; trade-offs are pervasive, and what looks like general competence from inside the population is specialization from outside it. Futuyma and Moreno (1988) show that improving performance in one niche typically reduces performance in others. Goldfeder and colleagues apply the same logic to artificial systems: generality is not merely difficult but structurally incoherent under resource constraints, because any finite system that directs capacity toward all possible tasks approaches zero capacity per task, while a system that concentrates capacity on a target distribution gains by excluding everything else.

The legal counterpart is not an analogy but an instance of the same structural error. The jurist who finds principles everywhere is in the position of the mind that cannot perceive its own blind spots. The principles are found everywhere because the history that produced them also produced the faculties through which the jurist perceives law. Roman law, the canon law of the mediaeval Church, the natural law tradition, the *ius commune*: these are the lineages from which the continental principles descend, and they are also the lineages from which European legal education descends. The common law is not exempt: Bracton drew on Roman sources; Coke absorbed the *ius commune*; equity borrowed extensively from canonical procedure. The apparent universality of the principles observed across traditions is in part convergent evolution under shared pressures, and in part the result of looking with eyes that were shaped by the same tradition that fixed the principles.

Goldfeder and colleagues ground their argument in Hannan and Freeman's organizational ecology (1977): entities that fail to meet performance thresholds in their environment are displaced, and competition acts as a selection mechanism amplifying effective strategies and eliminating ineffective ones. The legal analogue is that the principles survived not because they are necessary but because legal systems that lacked them or abandoned them paid costs, in predictability, in legitimacy, in coordination, that their competitors did not. Their authority is earned through survival, not given in advance.

3. Principles as Replicators: Fidelity, Fecundity, Longevity

If the principles are not universals, what are they? The extended-phenotype research programme proposes that legal concepts are memes in the sense of Dawkins (1976, 1982): units of cultural information subject to variation, selection, and retention, whose persistence depends on properties of the replicator rather than on the truth of what they encode. The programme has identified three replicator properties that jointly account for what appears as universality: fidelity, fecundity, and longevity (Lerer, 2026a).

Fidelity is the accuracy with which a unit is copied across transmission events. Latin formulae are near-perfect copying vehicles precisely because Latin is a dead language. Living legal language mutates continuously: "equity" has undergone at least seven major semantic shifts since the fourteenth century. *Pacta sunt servanda* maintained consistent semantic content from Justinian's *Corpus Iuris Civilis* of 533 CE to the present, a period of nearly fifteen centuries, because the death of spoken Latin suspended the normal drift mechanism. Systematic analysis of five hundred Latin maxims against their vernacular translations reveals striking fidelity differences: the Latin forms maintain an estimated ninety-eight percent semantic stability across centuries, while vernacular equivalents show substantial instability over decades (Lerer, 2026a). The formula is also cognitively efficient: surviving maxims average 3.18 words, within Miller's (1956) working memory limit, while extinct maxims average nearly six words. Brevity and fidelity interact: a short formula that resists mutation is easier to transmit without degradation.

Fecundity is the rate at which a unit is reproduced. The reproduction of the general principles is institutional, which means it is compounding and self-reinforcing. A principle encoded in a code, taught in every faculty, invoked in every appellate jurisdiction, and cited in treatises circulating across legal cultures is reproduced at a rate that no local competitor can match. *Res judicata* alone generates an estimated twenty-five thousand annual citations globally (Lerer, 2026a). Each citation is a transmission event; each transmission event seeds potential further citations in adjacent jurisdictions. The institutional machinery of legal reproduction, codes, faculties, courts, treatises, bar examinations, is a copying machine for the principles, and it was itself produced partly by jurists trained on those principles, a feedback loop that amplifies fecundity over time.

Longevity is the persistence of a unit across time. Some principles survived the collapse of the political orders that first carried them. Roman maxims outlived Rome, the Western Empire, the Eastern Empire, the mediaeval universities that transmitted them, and the national codifications that formally replaced Roman law. Longevity compounds with fidelity and fecundity through a mechanism that creates the appearance of necessity: a principle that has survived every challenge encountered over two millennia looks, from inside the tradition, like a principle that could not be otherwise. The survival record becomes a source of authority in its own right, and that authority reinforces the replication machinery, because invoking a principle with a long record carries lower argument risk than invoking a recent formulation. The circularity is real but it is not circular reasoning; it is a positive feedback between authority and survival that gradually produces what looks like necessity.

The joint signature of these three properties is quasi-universality. A principle that copies cleanly, reproduces at institutional scale, and accumulates a centuries-long survival record will be found nearly everywhere a trained jurist looks. The jurist who looks everywhere and finds the principle has no independent evidence that it is universal; what the observation establishes is that the principle is widely distributed. Wide distribution is the expected outcome of a successful replicator, not evidence of logical necessity.

4. The ESS Mechanism: Why the Core Resists Invasion

Wide distribution alone does not explain stability. A principle could be widely distributed and still be readily displaced by a fitter variant. The stronger claim, and the one that explains the resistance that practitioners and comparatists actually observe, is that the core set constitutes an evolutionarily stable configuration in the sense of Maynard Smith (1982): once common in a population, it cannot be invaded by rare alternatives under the prevailing fitness regime. The ESS concept, developed for biological populations, applies to any population of replicators whose fitness is frequency-dependent, and the fitness of legal formulations is frequency-dependent in a precisely specifiable way.

The mechanism is coordination cost. Legal reasoning is a coordination game. Lawyers, judges, and legislators who use the same conceptual framework can communicate at lower cost, build on each other's work, and predict each other's responses. A participant who departs from an entrenched formulation pays a penalty that is largely independent of the merits of the departure. The deviation is harder to communicate to a tribunal trained on the standard formula; it is more easily mischaracterized by opposing counsel; it is more exposed to reversal on appeal because appellate courts apply the standard formulation with lower variance; and it signals lower competence to professional audiences who interpret facility with established doctrine as a proxy for legal skill. This is not a conspiracy or even a conscious choice: it is the normal operation of coordination equilibria, in which the value of using the focal point rises with the number of others who use it.

The ESS analysis explains a feature that strict universalism handles poorly: the resistance of incumbents against challengers that are not obviously worse. If the principles held their place because they are necessary, the appropriate response to a challenge would be to demonstrate the logical necessity, and a sound demonstration would end the dispute. What practitioners observe is not this. Challenges to entrenched principles are typically answered not with demonstrations of logical necessity but with appeals to settled authority, professional consensus, and the risks of unpredictability if the established formulation is disturbed. That is not the behavior of a principle whose authority rests on logical compulsion. It is the behavior of an ESS whose authority rests on incumbency.

Teubner's analysis of good faith in British contract law is instructive here. Teubner (1998) documented the resistance of English law to the continental doctrine of good faith, even after the United Kingdom joined the European Union and faced explicit harmonization pressure. He described the episode as producing a 'legal irritant' rather than a transplant: the continental formulation was formally received but metabolized into something different, with English courts interpreting good faith duties restrictively and in ways that preserved the doctrinal structure of English contract law. The resistance is not explicable as logical necessity, since continental good faith is coherent. It is explicable as ESS behavior: the English doctrinal ecosystem had a stable configuration built around consideration, freedom of contract, and caveat-like allocations of risk, and the good faith obligation was an invasion that the incumbent configuration was fit to repel.

The ESS frame also corrects a confusion between stability and optimality that the universalist account inherits. A universal principle, on the standard view, is justified because it is the best solution to the relevant legal problem. An ESS need not be optimal; it needs only to be stable, which means that no available mutation earns higher fitness in the current environment. Stable and optimal coincide only under special conditions. In environments where coordination costs are high and switching costs compound over time, the incumbent can resist invasion by alternatives that would be superior in the absence of incumbency effects. The general principles are best understood as the current incumbents in a long selective history, not as the optimal solutions to eternal problems.

5. Three Predictions and Their Evidential Basis

A theoretical reframing earns its keep by generating predictions that the displaced account does not make. The memetic account of the general principles and strict universalism diverge on three sets of observations, each of which has some evidential basis in the comparative legal record and each of which generates testable research questions.

5.1 Cross-Jurisdictional Variation

If the core set is a selected configuration rather than a necessary one, its membership should vary across jurisdictions in step with their selective environments, rather than converging on a single fixed list. Two bodies of comparative evidence bear on this prediction.

The first is the contrast between good faith in continental and common-law systems. Good faith carries heavy doctrinal weight in continental systems: Article 1134 of the original French Civil Code (1804) imposed an obligation of *bonne foi* in the performance of contracts; Article 242 of the German Civil Code (*Bürgerliches Gesetzbuch*) imposes a general obligation of good faith that German courts have expanded far beyond the contractual context, generating entire subsidiary doctrines; Article 9 of the Argentine Civil and Commercial Code (2015) establishes good faith as a governing principle of the entire private law system. Common law received the principle late and unevenly. English contract law resisted a general good faith obligation throughout the twentieth century and well into the twenty-first; the United States Uniform Commercial Code includes a good faith obligation in Article 1, but its scope has been contested for decades, and American courts have generally declined to extend it to pre-contractual negotiations. Teubner's 'legal irritant' episode confirms that harmonization pressure does not automatically produce convergence: the same nominal principle operates as different doctrines in different environments.

The second is the pattern documented by Watson (1974) in comparative legal transplants. Watson showed that legal concepts do not spread by diffusion alone but that their reception depends on the match between the importing and exporting legal environment. When the environments are compatible, transplants succeed and the concept integrates into the receiving system, becoming part of its stable configuration. When they are not, the concept is either rejected, as with the good faith irritant in English law, or mutated into something that fits the local payoff landscape, as with

the reception of Roman law in early modern Germany. The pattern of successful and failed transplants is consistent with the ESS account: configurations stable in one environment are not automatically stable in another, and what looks universal from inside a single tradition looks contingent from the vantage of comparative legal history.

5.2 Documented Extinctions

A selected configuration carries fossils: formulations once stable that lost fitness when the environment changed. The history of *caveat emptor* is the clearest documented case. The maxim allocated risk of defects to buyers in commercial transactions and, in its domain, functioned as an ESS: it was widely used, coherent with an environment of bilateral negotiation between merchants, and resistant to challenges because commercial lawyers and courts had built extensive doctrine around it.

The consumer protection legislation of the twentieth century changed the selective environment. The United States Pure Food and Drug Act (1906) began the erosion. The Consumer Bill of Rights (Kennedy, 1962), the Magnuson-Moss Warranty Act (1975), and the United Kingdom Sale of Goods Act (1979) progressively displaced caveat-based allocations of risk. The European Union Consumer Rights Directive (2011) effectively prohibited caveat emptor as a default rule in consumer transactions. Estimated current usage has fallen to approximately five percent of transactions, concentrated in complex business-to-business contexts where the environmental conditions that originally stabilized the principle persist (Lerer, 2026a).

The extinction of *caveat emptor* is not explicable as discovery that the principle was always logically defective: it was coherent and functional for centuries. It is explicable as a change in the selective environment that lowered the incumbent's fitness below the level required to resist invasion. New legislation created a competing configuration, mandatory warranties and statutory implied terms, that earned higher fitness in an environment now dominated by mass consumer markets where bilateral information symmetry is structurally absent. The principle that replaced it, the implied warranty of merchantability, is not more universal than *caveat emptor*; it is better adapted to a different environment. Strict universalism has no comfortable place for this: it cannot treat a formerly recognized principle as having been never truly a general principle, and it cannot easily explain why a once-recognized principle ceased to be recognized. The memetic account expects extinction wherever the environment shifts enough to lower an incumbent's fitness below a challenger's, and the *caveat emptor* episode is precisely that.

The *abus de droit* case points in a different direction: convergent evolution. Cueto Rúa (1975) documented that Louisiana courts and continental European jurists independently developed identical doctrines of abuse of rights, despite geographic isolation and absence of cross-citation. Both legal environments faced the same selective pressure, parasitic litigation exploiting formal rights to harm others without legitimate interest, and both evolved the same solution. The convergence is not traceable to a common source; it is an instance of parallel emergence under shared selective pressure, analogous to the independent evolution of the eye in vertebrates,

cephalopods, and arthropods. The observation supports the selectionist account and is awkward for diffusionist accounts of legal spread. It also illustrates the sense in which convergent universality is an evolutionary result, not a logical premise: the *abus de droit* doctrine is widespread because it solved a coordination problem that many legal environments faced, not because it was recognized as a necessary principle before it existed.

5.3 Drift Under Environmental Change

The third prediction is that the core drifts when the environment changes. The history of the principles shows repeated instances of drift that are visible in retrospect but invisible to contemporaries who experience the incumbent configuration as stable and necessary.

The most consequential current shift is the arrival of new classes of legal agent whose intentionality structure differs from that of natural persons. The general principles were calibrated for agents to whom the intentional stance applies directly: they have beliefs, desires, and intentions; they can form agreements; they can act in good or bad faith; they can be heard before decisions affecting them are made. Corporations were the first major class of new agents. Legal systems accommodated them by imputing intentionality through agency and attribution doctrine, but the fit was always imperfect: good faith in a corporation is the good faith of its officers; hearing requirements are satisfied by procedures that do not literally hear anyone. The principles bent to cover the new agents without changing their form.

Artificial systems of sufficient complexity create a sharper version of the same problem. The extended-phenotype programme has argued elsewhere that the intentionality presupposed by legal doctrine is mismatched to the actual intentionality of AI systems in ways that create systematic distortions in legal reasoning about those systems (Lerer, 2026b). A principle like good faith, or the right to be heard, does not straightforwardly apply to a system that has no beliefs, desires, or intentions in the sense that grounds the principle for natural persons. Legal systems can respond in two ways: they can extend the principles by further imputation, treating AI systems as having a kind of functional intentionality that licenses applying the principles as if they were natural persons; or they can acknowledge that the principles require reformulation for a new class of agent, which is to say that the principles drift under the pressure of the new environment. Both responses are consistent with the memetic account. Neither is explicable by strict universalism, which would need the principles to apply to all legal agents by virtue of their universality, and would have difficulty explaining why applying them to new agent types requires the kind of doctrinal adjustment that all major legal systems are currently undertaking.

Stein's (1966) account of the winnowing of the *regulae iuris* provides historical perspective on this drift. Stein showed that the classical Roman jurists worked with a large stock of legal rules presented in sentential form, and that the process of transmission through the glossators, the commentators, and the natural law theorists progressively selected from that stock, amplifying some formulations and allowing others to drop from regular use. The selection was not random: formulations that fit the doctrinal needs of successive legal cultures survived; those that

presupposed specifically Roman institutional conditions did not. What emerged as the 'general principles' of the *ius commune* was not the complete stock of Roman rules but the subset that remained fit across the successive environments through which Roman legal culture passed.

6. The Universalist Account's Response and Its Limits

The argument of this paper does not go uncontested. The most serious response available to the universalist is that the observations interpreted here as selection and drift are better explained as progressive recognition: legal systems have gradually come to recognize principles that were always there, and apparent extinctions or variations reflect imperfect recognition rather than genuine contingency. On this view, the extinction of *caveat emptor* reveals that the principle was never genuinely universal; the variation in good faith between continental and common-law systems reflects an imperfect common-law recognition of a principle that the continental systems better understand.

This response has real force but three problems limit it. First, it is unfalsable in a way that the memetic account is not. Any configuration that fails to recognize a principle can be explained as imperfect recognition; any configuration that abandons a formerly recognized principle can be explained as error correction. A theory that can always explain divergence as error has no content that could be disconfirmed by the comparative record. The memetic account, by contrast, generates specific predictions about correlation between lineage and set membership, about the conditions under which extinctions occur, and about the environmental triggers of drift. Falsability matters for a theory that aims to explain something rather than merely to redescribe it.

Second, the progressive recognition account implies that there should be convergence toward a uniquely correct enumeration of the principles as legal systems develop. The comparative record does not show this. The most developed legal systems show persistent and systematic variation in the enumerated set, as the good faith contrast between English and German law illustrates. Watson's transplant research shows that even systems in close contact, subject to active harmonization pressure, maintain divergent configurations. If recognition were progressive and the object of recognition were a fixed set of universals, convergence would be expected. What appears instead is the pattern of divergent configurations that the ESS account predicts.

Third, and most practically, the progressive recognition account provides no guidance for legal design. If the principles are universals in the process of being recognized, then the appropriate response to encountering a legal system that does not recognize a principle is to explain the principle and expect recognition. Teubner's 'legal irritant' shows that this is not what happens. The English legal system did not fail to recognize good faith from ignorance; it actively metabolized the obligation into something compatible with its existing configuration. The design lesson from the ESS account is richer: it predicts that transplants succeed when they are compatible with the receiving environment's stable configuration, and fail or mutate when they are not. That is a lesson with operational content for legislators, treaty drafters, and designers of computational legal systems.

7. Consequence for Computational Legal Systems

The reframing carries a design consequence that follows directly from the argument without requiring any further steps. Computational legal systems increasingly encode a catalogue of principles as a control or audit layer. The design choice between a closed fixed list and an open revisable set depends on whether the principles are universals or selected strategies.

If the principles are universals, a closed list is the correct design choice. The list should be the same across all jurisdictions and all periods, because the principles are not jurisdiction-specific or period-specific; and no mechanism for revision is required, because the principles do not drift, go extinct, or require updating for new agent types. Any divergence from the list is an error to be corrected, not a signal of environmental variation to be tracked.

If the principles are an ESS configuration, a closed list is the wrong design choice for reasons that are specific and serious. First, a fixed list is jurisdiction-insensitive. The variation documented in Section 5.1 is not noise around a universal; it is information about the local selective environment. A system that applies the same principle layer to English contract disputes and German administrative law disputes misrepresents both. Second, a fixed list is temporally insensitive. The extinctions documented in Section 5.2 are not errors to be corrected; they are evolutionary responses to changed environments. A system trained on a fixed list that includes extinct principles, or that excludes principles that have recently become stable in a given jurisdiction, is systematically miscalibrated. Third, a fixed list is agent-insensitive. The drift documented in Section 5.3 is the system's response to new agent types. A computational legal system that applies good faith requirements, hearing obligations, and intention-based liability to AI agents using the same formulation designed for natural persons imports the mismatch rather than addressing it.

The design implication is not that no principle layer is needed but that the layer should be architecturally open: the current set for a given jurisdiction and period, revisable as the environment shifts, with mechanisms for tracking variation, extinction, and emergence. This is a harder system to build and to audit than a closed list of universals. The difficulty is not a cost of the evolutionary account; it is the cost of building a system that reflects how law actually works rather than a simplified model of how it was thought to work.

One further implication follows from the ESS analysis of stability. A principle layer that is treated as fixed and authoritative will, over time, exhibit the same incumbency effects that make the general principles resistant to challenge in ordinary legal reasoning. Computational systems that are audited against a fixed principle layer create institutional pressure to maintain that list, because departures from it generate audit failures that must be explained and defended. The system acquires the incumbency property of an ESS through the very process of becoming authoritative. This argues for embedding explicit review cycles in the design, at which the list is treated as a revisable hypothesis about the current stable configuration rather than as a statement of law's necessary content.

8. Position Within the Research Programme

Epistemological positioning. This paper belongs to a research programme (Lakatos, 1978) that applies generalized Darwinism and the extended phenotype framework to legal and institutional systems. Popper (1959) supplies the demarcation criterion: the hypotheses advanced here are falsable and their empirical tests are specified. The hard core of the programme, the claim that legal institutions are extended phenotypes of cultural replicators subject to multilevel selection, is held fixed here. This paper operates in the protective belt. It proposes a reclassification of the general principles from a priori universals to evolutionarily stable memes, a protective-belt adjustment that absorbs an anomaly, the observed variation, extinction, and drift of the principles, while leaving the hard core intact.

The adjustment is progressive in the Lakatosian sense: it predicts new observations that the prior formulation did not. The prior formulation, in which certain principles operate as quasi-axioms within the programme, anticipated neither the specific pattern of *caveat emptor* extinction nor the convergent emergence of *abus de droit* nor the systematic variation in good faith between legal lineages. The reformulation anticipates all three as instances of a general selectionist account. A progressive problemshift generates novel predictions rather than merely accommodating known facts; the test of whether this one is progressive is whether the predictions in Section 5 come out as described or fail in specific ways that the account does not expect.

The strongest objection to the argument from within the programme concerns the Blumenberg constraint established in earlier work (Lerer, 2026c). Absolute metaphors, in Blumenberg's account, define the space of conceivable variation but do not determine which variant wins; selection remains Darwinian and contingent. The concern is that treating the general principles as a selected configuration might inadvertently reduce the programme to a form of structuralism, in which certain principles are predetermined by the structure of legal problems rather than by contingent selective history. The response is in the argument itself: the cases of *caveat emptor* extinction and *abus de droit* convergence show both selectionist dynamics in operation. *Caveat emptor* was fit and then unfit; its replacement was not structurally necessary but environmentally selected. *Abus de droit* converged, but from different starting points and through different mechanisms; the convergence tracks a shared selective pressure, not a logical necessity. The account remains Darwinian and contingent throughout.

The argument also connects with the programme's treatment of intentionality. The general principles presuppose an intentional agent, an agent who can act in good or bad faith, who can be heard, who can bind by agreement. The intentionality presupposition is not itself a universal truth; it is a feature of the environment in which the principles achieved stability, an environment populated by natural persons acting within recognizable social structures. As that environment changes, the presupposition fails to hold for new agent types, and the principles face selection pressure in the form of mounting application anomalies. This is the point of contact between the present argument and the Generalized Intentionality Mismatch Theorem (Lerer, 2026b): the drift of the principles is the observable signature of the intentionality mismatch at scale.

Artificial Intelligence Disclosure

Artificial intelligence tools were used as research support during the preparation of this paper: locating sources, extracting information from the existing literature, and reviewing drafts. The thesis, the theoretical framework, the arguments, and the conclusions are the author's original contribution.

Epistemological Positioning Declaration

This paper belongs to a Lakatosian research programme whose hard core is the claim that legal institutions are extended phenotypes of cultural replicators subject to multilevel selection. The present paper operates in the protective belt and proposes a protective-belt adjustment. Claims about the programme's metrics (CLI, IHR, IEI) that appear in other papers represent preliminary results awaiting independent replication and are not validated benchmarks. The paper argues for a position; it does not claim to have demonstrated it. Its claims are falsable in the sense specified in Section 5.

References

- Bin Cheng. *General Principles of Law as Applied by International Courts and Tribunals*. Stevens and Sons, London, 1953.
- Blackmore, S. *The Meme Machine*. Oxford University Press, Oxford, 1999.
- Cueto-Rua, J. Abuse of rights. *Louisiana Law Review*, 35(5): 965-1009, 1975.
- Dawkins, R. *The Selfish Gene*. Oxford University Press, Oxford, 1976.
- Dawkins, R. *The Extended Phenotype: The Long Reach of the Gene*. Oxford University Press, Oxford, 1982.
- Dennett, D. C. *Darwin's Dangerous Idea: Evolution and the Meanings of Life*. Simon and Schuster, New York, 1995.
- Forister, M. L., Dyer, L. A., Singer, M. S., Stireman III, J. O., and Lill, J. T. Revisiting the evolution of ecological specialization, with emphasis on insect-plant interactions. *Ecology*, 93(5):981-991, 2012.
- Futuyma, D. J., and Moreno, G. The evolution of ecological specialization. *Annual Review of Ecology and Systematics*, 19:207-233, 1988.
- Goldfeder, J., Wyder, P., LeCun, Y., and Shwartz-Ziv, R. AI Must Embrace Specialization via Superhuman Adaptable Intelligence. arXiv preprint arXiv:2602.23643v1, 2026.
- Hannan, M. T., and Freeman, J. The population ecology of organizations. *American Journal of Sociology*, 82(5):929-964, 1977.

- Hull, D. L. *Science as a Process: An Evolutionary Account of the Social and Conceptual Development of Science*. University of Chicago Press, Chicago, 1988.
- Lakatos, I. Falsification and the methodology of scientific research programmes. In Lakatos, I., and Musgrave, A. (eds.), *Criticism and the Growth of Knowledge*, pp. 91-196. Cambridge University Press, Cambridge, 1970.
- Lerer, I. A. Dead Language, Living Law: Latin Legal Maxims as Perfect Memes. Zenodo, 2026a.
- Lerer, I. A. The Generalized Intentionality Mismatch Theorem. Zenodo, 2026b. <https://doi.org/10.5281/zenodo.18603169>
- Lerer, I. A. Absolute Metaphors, Perfect Memes, and LLMorphism. Zenodo, 2026c. <https://doi.org/10.5281/zenodo.20206138>
- Maynard Smith, J. *Evolution and the Theory of Games*. Cambridge University Press, Cambridge, 1982.
- Maynard Smith, J., and Price, G. R. The logic of animal conflict. *Nature*, 246:15-18, 1973.
- Miller, G. A. The magical number seven, plus or minus two: Some limits on our capacity for processing information. *Psychological Review*, 63(2):81-97, 1956.
- Popper, K. R. *The Logic of Scientific Discovery*. Hutchinson, London, 1959.
- Stein, P. *Regulae Iuris: From Juristic Rules to Legal Maxims*. Edinburgh University Press, Edinburgh, 1966.
- Teubner, G. Legal irritants: Good faith in British law or how unifying law ends up in new divergences. *Modern Law Review*, 61(1):11-32, 1998.
- Watson, A. *Legal Transplants: An Approach to Comparative Law*. Scottish Academic Press, Edinburgh, 1974.
- Wolpert, D. H., and Macready, W. G. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67-82, 1997.