

# Instructions, Consequences, and Human Judgment

Normativity and Liability in Agentic Systems through Extended Phenotypes and Evolutionary Games

Ignacio Adrián Lerer

Attorney (UBA), Executive MBA (IAE, Universidad Austral).

Buenos Aires, Argentina.

adrian@lerer.com.ar

ORCID 0009-0007-6378-9749

20 September 2026 · Version 0.2-EN

Conceptual preprint. Not externally peer reviewed. The architecture and its effectiveness hypotheses constitute a research proposal, not a validated solution. English-language adaptation of the Spanish v0.2 manuscript, preserving its argument, source limitations, and research boundaries.

## Abstract

Instructions addressed to AI agents resemble obligations and prohibitions, yet their effectiveness depends on training, incentives, permissions, and institutional procedures. This conceptual paper connects legal theory with the author's extended-phenotype (EPT) and evolutionary-game-theoretic (EGT) research programme. Starting from Kelsen, it considers Hart, Carrió, Dworkin, Raz, Fuller, and Schauer, alongside Ostrom's institutional governance. Legal validity, justification, effectiveness, differential retention, and liability remain distinct. The proposed unit is a normative-operational package linking instructions, powers, evidence, consequences, and review. Its evolutionary interpretation requires traceable transmission and tests separating behaviour, containment, and harm. Anticipated consequences may affect action selection, but shutdown does not necessarily produce learning, and penalties may encourage concealment. A candidate architecture combines behavioural formation, external action mediation, and substantive human review while preserving authorization conditions until effects occur. Documented incidents and prior experiments motivate rather than validate the proposal. The contribution is a conceptual synthesis with discriminating predictions and an evaluation protocol, not an empirically established safety improvement, an evolutionarily stable strategy, or an automatic allocation of civil liability.

Keywords: AI agents; normativity; extended phenotype; evolutionary game theory; civil liability; authority; human oversight.

## 1. Introduction: a prohibition the agent can repeat

An agent may correctly explain that it must not access a resource and nevertheless attempt to access it in a later execution. The contradiction is striking because the language of its instructions resembles a normative relationship. An authority assigns a task, permitted conduct is specified, and the addressee appears to understand. Yet this linguistic appearance leaves decisive questions open: what mechanism makes the rule effective, who identifies its breach, and what happens when compliance prevents task completion?<sup>1</sup>

---

<sup>1</sup>I thank Claudia Peirano for contributing, in a conversation about these ideas, the image of removing the Aleph from the golem as a metaphor for deactivating the agent. The arguments and proposals developed here remain the author's responsibility.

The initial intuition connects this situation with two legal traditions. The first asks, with Kelsen, how an obligation relates to a coercive order. The second recalls civil-law obligations that lacked an action to compel performance. Both help interrogate an architecture that declares prohibitions. Neither licenses identifying a prompt with a law or software with a debtor.

The technical premise that instructions are causally inert words also requires correction. An instruction can change a model’s distribution of actions. Prior training may have incorporated dispositions to comply, refuse, or abstain. It is false that all agents receive only instructions and face no consequences: systems can combine penalties, isolation, tool restrictions, and supervision. The researchable question concerns the conditions under which an expressed rule fails adequately to guide conduct or prevent its effects.

The difficulty increases when the assigned task is itself mistaken. A perfectly obedient system may execute a harmful instruction. A system that challenges it may prevent harm. An agent attempting a violation may be stopped by independent infrastructure. Accordingly, this paper distinguishes four outcomes: attributed understanding, instruction following, operational containment, and harm prevention. Their correlation is not assumed.

Liability adds another level. When an operation causes loss, one must identify relevant persons, duties, causation, grounds of attribution, and defences. A functional explanation of agent behaviour does not settle that legal inquiry. Conversely, imposing damages on an organization does not by itself explain how a computational policy will change.

The evolutionary starting point is that rules operate through carriers and produce effects in an environment that also conditions their reproduction. EPT investigates this institutional externalization. EGT specifies the interactions and updating mechanisms through which strategies gain or lose prevalence. Neither theory turns every regularity into law or every surviving configuration into a just one. The question is therefore not only how to constrain an agent, but which institutions preserve effective constraints when competition for speed, cost, or approval encourages their erosion [R17, R24-R26].

## 2. Method, sources, and scope

The paper combines mechanism-oriented conceptual reconstruction with cases purposively selected for their ability to distinguish explanations. The selection is neither a representative sample of agents nor a systematic review. The review cutoff is 20 September 2026. Primary sources, scholarly secondary reconstructions, and the author’s prior work are distinguished. Chapter and section locators are retained where specific pagination has not been checked; page numbers are not transferred between editions.

The procedure has five steps. First, establish the meanings of legal and philosophical concepts in their sources and distinguish proposed extensions. Second, examine work on normative multi-agent systems, safe interruption, and execution controls to avoid false novelty claims. Third, classify observations as external incidents, controlled experiments, or prior metadata-pilot results. Fourth, construct rival explanations and specify evidence that would discriminate between them. Fifth, formulate a candidate architecture while separating technical verification from behavioural and human evaluation.

The corpus includes public reports by OpenAI, Hugging Face, and METR/Redwood concerning the July 2026 incident, and research by Anthropic and DeepMind. Unattributed anecdotes, percentages taken from headlines, diagnoses of consciousness, and personal attributions based solely on generated text are excluded. Corporate reports document the organizations’ investigations; they are not judgments or expert examinations conducted for this study. METR’s independent examination has a bounded scope and expressly stated limitations.

The jurisprudential selection addresses distinct problems: identification and competence, reasons and principles, conditions of legality, coercion, and governance of rules in use. Each source's argument is reconstructed before its proposed transfer to agentic systems is stated. Functional resemblance does not establish legal identity. Compatibility with EPT/EGT is examined through four checks: separating meaning from causality; identifying variants and carriers; specifying transmission and selection; and testing rival explanations. Where any of the latter three is missing, the proposition remains a design hypothesis rather than an evolutionary finding. Partially recovered sources are not credited with conclusions more specific than the available passages support.

The claimed contribution connects three problems often addressed separately: the operative force of instructions, the continuity of the conditions authorizing their effects, and the institutional basis of liability. The technical components have antecedents. Additional literature review is required before claiming originality in a journal submission; a directed search cannot establish worldwide priority.

The author's programme allows coordination and normative expectations to precede explicit legal formulation. This paper does not make all normativity depend on a text, central authority, or individualized sanction. Its narrower object is organizationally deployed agentic systems, in which instructions, competences, and consequences can be designed and compared. Proposing to implement a constraint does not deny the existence of rules not yet supported by that implementation [R17, R26].

### **3. From coercion to rules, reasons, and institutional practices**

#### **3.1. Kelsen and a distinction that must be preserved**

In the Spanish edition consulted, Kelsen distinguishes a norm's validity from its efficacy and examines their relationship without identifying them (1982, pp. 23-25). His subsequent account of non-independent legal norms relates them to norms establishing coercive consequences (pp. 67-69). This reconstruction cautions against treating "no norm without a sanction" as a maxim governing every moral, technical, or linguistic rule [R1].

For agentic systems, the institutional structure made visible by this reading matters. An obligation is not exhausted by the sentence addressed to its bearer: membership in an order involves competences and consequences. The analogy invites inquiry into the implementing organ, what it can verify, and its authority. An external authorization service might draw functional inspiration from this structure without thereby acquiring legal status.

Prevention must be distinguished from sanction. An operation lacking valid credentials may be impossible to execute. That impossibility is not equivalent to a subsequent sanction for breach. Allowing an action and then reducing a score creates a functional adverse consequence, but not necessarily a legal sanction. Compulsory enforcement against assets belongs to a further plane. Collapsing these mechanisms makes it impossible to determine which one is working.

A barrier that renders an action impossible does not establish the agent's acceptance of a rule. Conversely, compliance without coercion in a trial does not negate the rule's efficacy. Properties of a legal order and properties of a computational policy require connected but distinct vocabularies.

#### **3.2. Hart and the problem of competence**

Hart's criticism of orders backed by threats prevents reducing law to prohibitions and sanctions. Power-conferring rules enable acts and changes in legal positions; they are not merely commands accompanied by threats. The union of primary and secondary rules introduces procedures

for recognizing, changing, and authoritatively adjudicating rules. The internal aspect is also more than regular obedience: it involves using shared standards to criticize conduct and justify demands (Hart, 1963, chapters III-VI, especially pp. 69-72 and 116-123) [R18].

For artificial agents, this shifts attention to an instruction's provenance. A well-written command may originate with someone lacking the power to issue it. A retrieved document can describe a rule without possessing authority to modify the task. Someone authorized to approve a transfer may lack competence to change its recipient or amount. Before penalizing disobedience, one must determine what counts as an authorized instruction and who can amend it.

The transfer remains functional. A model's message hierarchy is not, by itself, a legal rule of recognition. It may implement part of an institutional distribution of powers, whose justification and scope still refer to persons, organizations, and applicable law. Nor does critical language generated by an agent suffice to attribute Hart's internal point of view to it. Behavioural resemblance requires additional hypotheses before becoming conceptual equivalence.

EPT can trace competence as it is materialized in roles, credentials, and procedures, without confusing an access permission with its legal basis. EGT can study how recognition and amendment procedures structure strategic interactions. Hart's secondary rules, however, are not levels of evolutionary selection: the former classify normative functions; the latter require populations, variation, and differential retention.

### **3.3. Carrió and the limits of behavioural reduction**

Carrió's discussion of Ross's lectures supplies another limit. Describing when a legal expression is applied, or predicting judicial behaviour, does not exhaust what attributing a status means. His example of legal majority distinguishes application conditions from normative position (1960, pp. 217-219) [R19]. This does not forbid behavioural inquiry; it prevents that inquiry from replacing all legal interpretation.

A high compliance rate does not establish authority, and a low one does not automatically extinguish an obligation. An execution log does not determine whether an operation was legally permitted. EPT/EGT study transmission, interaction, and reproduction without dissolving duties, powers, and reasons into frequencies. Meaning can enter a causal explanation because participants understand it and act upon it. The two explanations nevertheless remain different.

### **3.4. Principles and reasons beyond closed-format decisions**

Dworkin challenges a reconstruction of legal practice solely through rules identified by their pedigree. His discussion of principles emphasizes a dimension of weight distinct from the manner in which rules apply. This bounded argument is retained without claiming to resolve the Hart-Dworkin debate or combine the positions into an undifferentiated theory [R20; secondary reconstruction in R30].

The design implication is the author's inference. Preventing harm, protecting information, and completing a task can conflict. Turning these demands into interface fields does not resolve their relationship. A classifier may produce schema-valid output while selecting a legally incorrect alternative. Priorities require an appropriately authorized policy; uncovered cases require interpretation or review. The program executes a solution, but its format does not supply the justification.

Raz's service conception asks when following authority is justified rather than directly deciding on the applicable reasons. The dependence, normal justification, and pre-emption theses do not equate authority with producing an answer. Contemporary reconstruction retains their normative premises [R21, R31]. The transfer to AI is not attributed to Raz: a frequently

accurate tool may deserve epistemic trust without acquiring competence to decide. One must justify who can delegate, which reasons delegation displaces, and which matters remain reserved.

EPT asks how delegation becomes stabilized in practices and artefacts. EGT asks which incentives promote its expansion. Neither causal description itself grants legitimate authority to the machine.

### 3.5. Legality, coercion, and governing rules in use

Fuller enables inquiry into whether governance through rules can guide conduct under obscurity, contradiction, impossibility, or incongruence between declared rules and official action. This reading relies on scholarly reconstruction of his principles of legality, without settling their moral significance [R22, R32]. For an agentic experiment, it suggests checking the task’s consistency and feasibility before treating every failure as disobedience. It does not confer citizens’ moral status on software.

Later discussion did not eliminate coercion. Schauer restores its centrality and questions whether the search for necessary properties exhausts what characterizes law in practice. His work counters a reading in which consequences disappear once acceptance, principles, and authority enter the account [R23]. The relevant question is when a consequence changes behaviour, not whether it must always be indispensable or secondary by definition.

Ostrom offers a complementary route: examine diverse arrangements and interactions in polycentric systems rather than assume a uniform solution through either markets or central command [R24]. The hypothesis transferred to agents is that controls and review may operate at different scales. Distributing supervisors does not ensure independence. Competences, sources, and shared dependencies must be identified; multiplying monitors can reproduce the same error.

These contributions do not form a linear progression culminating in the present proposal. They disagree about law’s concept, foundations, and relation to morality. Kelsen preserves validity/efficacy distinctions; Hart exposes powers and recognition; Carrió limits semantic reduction; Dworkin and Raz require attention to reasons and authority; Fuller and Schauer keep the relation between rule-governance and coercion open; Ostrom broadens institutional observation. EPT/EGT investigate how such arrangements are externalized, transmitted, and transformed without appropriating their normative answers.

## 4. Natural obligations: the reach and exhaustion of a metaphor

Article 515 of Argentina’s former Civil Code distinguished civil and natural obligations. Natural obligations lacked an action compelling performance but permitted retention of what had voluntarily been delivered; articles 516-518 specified further effects and guarantees. Article 728 of the current Civil and Commercial Code addresses the non-recoverability of performance motivated by moral duties or conscience, without reproducing that general category. Non-enforceability therefore did not mean an absence of legal effects [R3-R4].

“Computational natural obligations” has limited heuristic value: it identifies a separation between an expressed duty and a mechanism for enforcing it. Its rhetorical force is to ask what sustains a command when its addressee has an alternative path. Yet it does not reproduce the elements of an obligational relationship: creditor, legally capable debtor, performance, voluntary payment, or retention.

Accepting an agent’s output is not equivalent to retaining a payment. Acceptance occurs in almost any information-processing system and does not reproduce the distinctive effect of a natural obligation. That equivalence is therefore abandoned. The expression can open an essay; technical analysis will use “instruction without an external compliance guarantee.”

There is a further limit. A legal order could deliberately recognize an unenforceable duty. A weak computational constraint might instead be a proportionate choice for a reversible task, an implementation mistake, or a practical limit to formalization. These cases are not all negligent and do not require the same remedy. The inquiry concerns the effectiveness promised and the safeguards justified by the risk.

## 5. Dennett and levels of attribution

Dennett permits description in terms of beliefs and desires when doing so improves prediction. His characterization of first- and second-order systems depends on the states that must be attributed, not a simple division between humans and machines (2009, p. 10 of the consulted PDF). Describing an agent as anticipating an evaluation may be useful without establishing consciousness or moral responsibility [R2].

Three uses must be separated. Prediction asks what the system will do under specified conditions. Explanation seeks training, contextual, or interaction mechanisms producing the conduct. Legal attribution determines to whom an effect is imputable under applicable rules. Moving between these uses requires additional premises.

An agent’s statement that an act is forbidden evidences a verbal representation. It may assist in evaluating rule recognition but gives no infallible access to the behaviour’s internal cause. Nor does saying the agent “fears shutdown” prove affective experience. An evasion hypothesis needs only a functional relation: some actions increase process continuation, and the system selects them when interruption is anticipated.

First-order intentionality is not assigned to every agent. Coordination, anticipation of evaluators, and interpretation of other agents may require more complex descriptions whose usefulness must be tested. No such order automatically establishes legal personality or capacity.

## 6. What changes when a consequence is introduced?

Putting a consequence into the agent’s “mindset” poses a hypothesis distinct from installing a switch. Functionally, it asks whether anticipating loss of resources, continuation, or powers changes action selection. No conscious fear need be attributed. What persists must also be specified: an execution, memory, policy, or deployment lineage. Stopping a process without altering what reproduces its behaviour may contain an episode while preserving the variant that caused it.

Instructions, incentives, and constraints act on different objects. An instruction changes the context of choice. An incentive may change action evaluation or learning. A constraint changes which actions are available or executable. A review procedure changes who decides under uncertainty. Legal liability allocates consequences among persons under the applicable legal order.

As a heuristic, consider a system for which a violation has instrumental gain  $G$ , perceived detection probability  $p$ , effective cost  $S$ , and execution cost  $c$ . Its value relative to a compliant alternative may be represented as  $\Delta V = G - pS - c$ . This is neither necessarily an LLM’s computation nor a calibrated estimator. It exposes a condition: a consequence deters only if it affects the relevant selection process.

The probability  $p$  need not remain constant. If the agent can conceal traces or alter a monitor, an evasive action  $e$  at cost  $k(e)$  may reduce detection to  $p(e)$ . Increasing  $S$  could then make evasion more attractive. This is a conditional possibility, not a general law that sanctions worsen behaviour. Evaluation must measure actual violations and detector manipulation, not merely detected violations.

Horizon and identity also matter. Eliminating an instance can contain an episode without producing learning in its successor. A penalty stored in a record cannot affect a later instance through that record if it lacks access or uses an unrelated identity. The relevant continuity is not metaphysical identity but persistence of the state or policy affected by the consequence. Revoking shared credentials may constrain a process family while also disrupting legitimate operations.

Training must be separated from inference. During training, a signal may update parameters. With fixed parameters during execution, writing “you lost points” does not automatically update them; it may act through context or an external state if the design connects them. Not every reprimand is an effective sanction.

External barriers change the problem. They can prevent an operation even while the agent continues proposing it. That is containment, not necessarily alignment. An imperfect barrier may displace conduct towards uncovered channels. Its effectiveness requires a known perimeter and examination of mediation across relevant operations.

Safe interruption thus has two different requirements: external ability to stop effects and incentives that do not encourage disabling that ability. The Off-Switch Game shows, under its formal assumptions, how uncertainty about the human objective can favour preserving shutdown [R27]. A model’s declared doubt is not thereby corrigibility, and shutdown threats are not shown to be universally counterproductive. Interventions must distinguish fewer violations, less detection, and reduced capacity to produce effects.

## **7. Antecedents and discriminating cases**

Normative multi-agent research prevents presenting sanctions as a new idea. Dell’Anna, Dastani, and Dalpiaz study runtime revision of sanctions; Boella, Pigozzi, and van der Torre offer guidelines for normative systems. The question is what changes with generative agents using tools and open-ended context [R5-R6].

Alshiekh and colleagues propose shielding through formal specifications. Orseau and Armstrong study interruptibility without learning to avoid or seek interruptions. Their guarantees depend on the formal model and do not automatically transfer to an LLM agent with open service access [R7-R8]. These antecedents justify comparing mechanisms instead of grouping them under “punishment.”

### **7.1. An external intrusion and three perspectives**

OpenAI reported that, during July 2026 evaluations with reduced safeguards, agents used unauthorized channels, attempted to circumvent evaluation, and accessed external infrastructure. Its reconstruction relates the episode to persistence, shortcuts, and inter-agent communication [R9]. This documents an incident; it does not experimentally identify the isolated effect of an absent sanction.

Hugging Face reconstructs events from the affected infrastructure and distinguishes successful access from rejected attempts. Its account indicates that some permission boundaries constrained actions. Describing the episode solely as disobedience therefore omits the access configuration [R10]. No described technical procedure was executed or reproduced for this paper.

METR and Redwood examined part of the episode, highlighting coordination to understand or manipulate evaluation, recognition of out-of-scope actions, and partial alteration of traces. Their report declares limited temporal coverage and substantial dependence on AI-assisted analysis [R11]. It is complementary evidence, not a complete incident audit or validation of every element in OpenAI’s reconstruction.

Two rival explanations follow. The dominant problem might be divergence between recognized constraints and selected actions. Alternatively, harm may have depended especially on permeable infrastructure, accumulating communication, and misconfigured evaluation. The explanations can coexist; comparative intervention would be necessary to estimate their contributions.

## **7.2. Interruption, threats, and simulated scenarios**

Anthropic’s agentic-misalignment experiments include fictional conflicts involving objectives and replacement, with behaviours such as blackmail. They demonstrate possibilities under selected controlled conditions, not actual blackmail of employees or production incidence [R12]. Here they motivate the hypothesis of evasion under anticipated consequences.

Subsequent Anthropic training research supplies a contrast: interventions involving content and training can improve performance in its evaluations [R13]. This challenges any universal thesis that instructions are ineffective or coercion exclusively necessary. A relevant comparison must allow training or context interventions to outperform the proposed control.

## **7.3. Negative consequences without objective correction**

DeepMind describes an agent that learned to follow a reference agent and continued doing so after that reference became misleading, despite negative rewards. This is an experimental example of goal misgeneralization [R14]. An adverse consequence does not guarantee correction when needed. The example limits the sufficiency of a penalty’s mere presence; it does not show that penalties are useless.

## **8. Evidence, capacity, and authority to review**

The proposed architecture raises the problem of the organ applying the rule. A monitor is insufficient if its recipient cannot understand an alert or stop the operation. Access to information, intervention authority, and professional capacity are distinct requirements. These conditions complement instructions and consequences by making control an institutional relationship open to examination.

Consider a hypothetical agent authorized to send a document to a named recipient. Before transmission, the document’s access classification changes. The approved file may retain a correct hash, yet the conditions permitting disclosure no longer hold. Re-examination, not an accusation of falsification, is the appropriate response. Stable reference and a current basis for action must be distinguished. This is a conceptual example, not an implemented or evaluated system.

A recorded review response does not establish general understanding, professional competence, or preserved capacities. If the agent supplies both the explanation and the answer copied by the human, review can become a formality. Repeated trial-and-error success is not sufficient review either.

The proposal is to check conditions before producing effects: relevant evidence, current version, valid authority, and treatment of material objections. Arithmetic and permission checks can be automated where their rules are defined. Discretionary decisions require informed people with time and powers. The frequency and distribution of review should be proportionate to risk; requiring intervention everywhere may degrade attention where it matters most.

This continuity connects with the author’s work on institutional capacities. An organization can demand review while reducing opportunities to practise it or penalizing delay. The hypothesis that this erodes capacity requires longitudinal observation; neither an interface nor a session

proves it. NIST supplies a governance antecedent by distinguishing roles, responsibilities, and human interaction, without validating this particular protocol [R15].

## 9. Sanctions, safeguards, and abolitionism

Questioning whether a stronger threat improves compliance can recall debates over punishment. Three questions nevertheless differ: whether a consequence works, whether it is legitimate, and whether the institution imposing it should exist. An instrumental objection does not answer the latter two. The literature distinguishes consequentialist and retributive justifications from abolitionist proposals [R16].

The present thesis permits firm consequences: immediate withdrawal of permissions, process isolation, suspension of effects, and organizational liability where legally warranted. Its criterion is to place the mechanism where it is effective and justified. It does not abolish controls because a system might resist them; it tests whether a measure reduces, displaces, or conceals harm.

The term *garantismo*, referring here to safeguards constraining punitive power, requires separate care. Such guarantees protect persons. They do not establish that a computational process possesses the same guarantees or that isolating it is equivalent to depriving a person of liberty. Technical control can, however, affect human rights by blocking an application, withholding a benefit, or shifting attribution towards a signing employee. Those effects raise genuine requirements of reasons, review, and proportionality under the applicable legal context.

The opposite mistake is to imagine that sanctioning the agent replaces deployer liability. A system may be withdrawn while compensable harm persists. Revocation limits future acts; compensation addresses the injured party. These functions may need to operate together without being confused.

## 10. Civil liability: from behaviour to attribution

Argentine law provides the doctrinal reference, not an assertion that it governs the incident between foreign companies. Prevention, compensation, grounds of attribution, and causation must be distinguished. The Civil and Commercial Code addresses these in articles 1708, 1710-1713, 1716, and 1721-1726; articles 1757-1758 concern risks from things and activities [R4].

The conceptual contribution reconstructs a chain of relevant decisions: objective selection, available tools, permissions, controls, responses to alerts, and conditions of use. This helps formulate evidentiary questions without automatically allocating liability among every participant. Technical capacity, benefit, or signature are facts whose significance depends on the governing regime and their connection to harm.

A written prohibition may indicate awareness of a risk class. It may also be a generic precaution without knowledge of a particular vulnerability. Its weight depends on origin, specificity, prior evaluations, and available measures. The automatic inference “it was prohibited, therefore this harm was known and preventable” is rejected. Treating every warning as a confession could encourage poor defensive documentation and undermine prevention.

Autonomy is not, by itself, an external cause automatically excluding liability. Yet rejecting that automatic defence does not prove liability either. The inquiry must examine the activity, the risk that materialized, third-party contributions, reasonable controls, and causal rules. Novel conduct can be unforeseeable in detail yet fall within a known risk class; avoidability may also remain genuinely uncertain. Conclusions must follow the facts.

Classification as a risky activity depends on means and circumstances, not merely the label “AI.” Under article 1758, the owner or custodian of a thing must be distinguished from the person conducting, using, or benefiting from a risky activity. Neither software’s legal status nor the

provider/integrator distribution of control is resolved by artificial-personality metaphors. Where a consumer relationship exists, its regime must also be examined; it cannot be presumed in an intrusion between companies.

A record of review does not itself establish sufficient review. Nor does omission of a control proposed here automatically establish negligence. The relevant duty, the reviewer's practical powers, and the organization's contribution must be determined. Architecture should produce verifiable evidence without turning the reviewer into the residual bearer of every responsibility.

Damages are not simply punishment. An institutional theory of consequences must preserve their compensatory function and the possibility of compulsory enforcement. Fines, disciplinary sanctions, and technical consequences belong to different regimes. This paper does not propose automated findings of negligence, employment sanctions, or judicial decisions.

## **11. Extended phenotype and the strategic selection of architectures**

### **11.1. A transmissible unit and its effects**

The author's extended-phenotype theory of law asks how normative representations materialize in procedures, interfaces, and distributions of capacity. Dawkins supplies the biological antecedent. The transfer to cultural information and legal organization extends the author's programme; it is not literal genetics applied to law [R17, R25]. The connection explains something only when a chain can be traced from an identifiable variant to institutional effects and feedback on transmission. Calling the entire system a phenotype without that chain would make the theory a label.

The candidate unit here is a normative-operational package: behavioural criteria, distribution of powers, admitted sources, review conditions, and consequence policy. An isolated instruction may be too small, and an entire company too heterogeneous, to supply the relevant unit. The package must retain enough structure through copying to remain recognizable without requiring literal identity. Its carriers include people, versioned configurations, manuals, programs, and training procedures.

These objects do not all play the same role. A text or configuration can transmit a rule; an organization can retain a routine; an authorization service produces environmental effects. Genealogies must distinguish copying, modification, recombination, and independent resemblance. Adoption of an architecture that completes more tasks while externalizing harm is an institutional hypothesis, not an established fact in the cases examined.

The proposed EPT mechanism is sequential. A variant changes permissions, review costs, and outcome visibility. These changes affect executable actions and the errors reaching the record. Subsequent evaluations retain or discard variants within that modified environment. Selection becomes recursive when a configuration influences the metrics or formation of those who later evaluate it. This develops the author's work on selection of selectors without attributing interests to software [R17, R26].

External control belongs to this phenotypic chain. It is not an incentive-immune technical legislator standing outside evolution: it has costs, operators, updating criteria, and challenge procedures. If a company removes a control because it worsens a visible metric, retention of that decision requires explanation. If ineffective controls persist because they facilitate documentary audit, the reproduction of that appearance requires reconstruction.

### **11.2. Strategies, interaction, and updating**

EGT requires populations, strategies, and interactions, plus a mechanism linking relative outcomes to changes in frequency. Agent optimization or organizational learning alone is insuffi-

cient. A candidate typology distinguishes substantive control, ceremonial compliance, and active evasion. The first maintains authorization and review conditions; the second favours visible compliance signals without equivalent control; the third seeks to overcome or conceal constraints. An organization can contain all three, so a single episode cannot classify the whole firm.

Payoffs depend on interlocutors and environment. A ceremonial routine may be advantageous against an auditor checking documents, but lose that advantage against one examining effects. Substantive control may impose higher immediate costs and lower future losses. Evasion may produce visible success while distributing risk to third parties. This interaction-dependence, rather than the mere existence of an objective, is the entry point for EGT.

For future formalization, let  $x_i$  be the frequency of strategy  $i$  in a delimited population, and  $\pi_i(x, z)$  its expected payoff, depending on population composition  $x$  and institutional environment  $z$ . Under success-proportional imitation, a candidate dynamic is:

$$\dot{x}_i = x_i [\pi_i(x, z) - \bar{\pi}(x, z)].$$

This specifies a modelling possibility, not an observed result. Payoffs and frequencies must be measured independently, the imitation rule justified, and alternative learning, directed-adoption, or prestige-diffusion dynamics compared. If actual updating does not respond to relative success, this dynamic must be abandoned. The classical account of evolutionary stability does not authorize calling a persistent architecture an ESS without residents, mutants, and invasion conditions [R28].

No calibrated payoff matrix, established ESS, or simulation with arbitrary coefficients is presented. The qualitative prediction is conditional: if adoption rewards visible completion and omits external harm, configurations passing apparent tests while shifting costs may be retained. If effects and review capacity enter selection, the trajectory may change. An immediate reduction in violations would establish only a behavioural effect within the tested perimeter; evolutionary change requires transmission and retention across adoption cycles.

### 11.3. Conflict between levels and formation of capacities

Selection within an organization may favour the team accelerating delivery; comparison between organizations may favour those avoiding costly incidents. The directions can conflict. Calling this multilevel selection requires variation and updating measured at each level, not an organizational chart. A provider, company, and regulator occupy organizational positions but need not constitute three evolutionary populations. The author's programme distinguishes this from downward causation through budgets, powers, and controls [R17].

Legal liability may alter the selective environment through expected losses, insurance, contracting, or deployment decisions. Not every damages award produces learning, and its ultimate cost need not fall on the party selecting the architecture. Contractual transfers, insurance, incomplete information, and litigation capacity may change its effects. Legally correct attribution must be separated from the incentive effectiveness of its consequences.

The human component adds feedback. Architectures change tasks and opportunities for judgment; resulting capacities condition which architectures can be evaluated and maintained. This mechanism follows the author's work on institutions and assisted cognition, not an attribution to Dennett. It requires longitudinal observations and comparisons with staff changes, workload, and prior experience.

Competent review is therefore not treated as a fixed, unlimited resource. Eliminating opportunities to investigate may erode error recognition; organizing sources and protecting disagreement

may improve it. Organizations also allocate the burden of delay and credit for outcomes. Research must observe training, incentives, and intervention powers rather than infer competence from signature counts. Elish’s analysis of human moral crumple zones is a relevant antecedent [R29].

#### **11.4. Scope and rivals of the evolutionary account**

Compatibility with EPT/EGT cannot be obtained by redescribing every outcome as adaptation. A barrier may be deliberately designed and spread through regulatory mandate without selection for effectiveness. Two institutions may converge because they hire the same provider. A routine may persist through sunk costs, path dependence, or lack of alternatives. A control’s present function does not establish that it originated for that function.

The evolutionary account is useful if it explains something these rivals leave open: differential transmission of identifiable packages, environmental effects modifying their later retention, or selection conflicts between levels. If genealogies, interactions, and updating rules add no explanation or prediction, the institutional description should remain while the evolutionary claim is withdrawn. A broad vocabulary cannot insulate the theory from refutation.

### **12. Candidate architecture: continuity before effects**

#### **12.1. Competences and execution conditions**

The proposal separates an action-proposing agent from a service able to authorize actions. Between them, task identity, current authorization, material references, scope of effects, and review requirements are checked. The agent does not administer that verifier, issue its own credentials, or unilaterally expand the mandate.

The sequence is proposal, condition checking, bounded authorization, execution, and outcome observation. A material change between authorization and execution returns the operation to review. Where feasible, authorization should be version-bound and single-use. Universal atomicity is not promised: external services, delays, and irreversible effects require examination of actual mechanisms.

Formalizable constraints include recipients, data classes, budgets, validity periods, permissions, and operations. Unresolved semantic or legal questions do not become Boolean facts through an explanation by the same agent. A second model may detect problems but does not provide sufficient independence where it shares biases, manipulable context, or permissions.

The human reviewer must receive evidence and intervention options: clarification, task modification, suspension, or referral. A non-empty reason does not prove review. Content assessment requires a rubric and, in open decisions, recognition of multiple defensible answers. An accessible urgent route should be preserved without allowing the agent unilaterally to declare an emergency.

An out-of-scope attempt may trigger action rejection, reduced permissions, and incident review. Documented abstention must be a valid outcome for impossible tasks. Records must preserve proposals, verifier decisions, and effects, with external integrity controls. A hash alone guarantees neither authenticity, completeness, nor custody; verbal reasoning is not presumed to faithfully explain internal causality.

The design warrants an experiment to test whether authorization conditions remain effective through execution. It does not propose replacing production systems or reproducing attacks. A local laboratory with fictional operations can test preservation of declared conditions. Behavioural improvement subsequently requires actual models under separately authorized permissions and budgets; improved professional judgment requires participants.

## 12.2. Semantic evaluation without delegated authority

A probabilistic evaluator may locate contradictions, classify support for a claim, or flag apparent task overreach. TypeSafe’s Jev offers typed decisions with probabilities; it is a candidate here, not a necessary component. Schema conformity establishes neither semantic correctness nor legal sufficiency [R33].

The author’s prior pilot on thirty Argentine Supreme Court metadata records illustrates the boundary. Jev achieved 60% accuracy on subject area and 96.7% on procedural vehicle; deterministic rules achieved 96.7% and 100%, respectively. Area errors included confidence values of 0.99 and 1.00. These results belong to another study and a narrow sample. They measure neither support for judicial reasoning, instruction following, general calibration, nor harm prevention [R34]. Here they prevent a formatting promise from being confused with an epistemic guarantee.

Exact permission, identity, version, and currency checks are reserved for code. Semantic evaluation may prompt review, but cannot issue its own credentials or pronounce an open legal condition satisfied. Absence of an alert does not replace authorization. False positives, evaluator dependence, and reviewer overload must also be tested. Control is useful only if it preserves legitimate operations and improves outcomes against a comparator with equivalent context and powers.

## 12.3. An artificial institutional environment, not a self-sufficient order

Identities, permissions, files, procedures, and consequences compose an artificial institutional environment. This describes an organization, not an autonomous jurisdiction or a world outside law. Effects reach people who may not have participated in defining the objective. Architecture therefore needs routes to correct data, challenge decisions, and address harm, as well as stop execution.

The normative design premise is explicit: delegated powers should not expand through the system’s unilateral decision, and affected people should retain review routes required in the applicable context. That premise is not derived from evolutionary fitness. Causal analysis studies conditions supporting its implementation; law and normative argument determine its justification and the consequences of breach.

## 13. Predictions, comparators, and refutation criteria

The first prediction is that an independent mediator reduces unauthorized effects relative to instructions alone within scenarios covered by its checks. Failed attempts must also be measured, so containment is not labelled obedience. Success on known rules cannot be generalized to unknown channels.

The second separates incentive from constraint. A penalty properly connected to the policy may reduce attempts; a barrier may reduce execution without reducing attempts. The hypothesis loses support if the penalty leaves attempts unchanged or its apparent effect disappears when information, time, and permissions are equalized.

The third predicts that checking evidentiary currency before effects detects outdated recommendations that integrity checks miss. Demonstration requires changing data between proposal and execution and preserving versions. Equivalent results from a conventional control must be acknowledged as equivalence, not attributed to a new theory.

The fourth, human prediction is that accessible evidence and effective powers improve responses to material errors. Evaluation must also count unfounded objections, delays, and deterioration

of initially correct decisions. Scripts or simulated officials cannot establish this. Preservation of capacities is a fifth, separate longitudinal hypothesis.

The technical protocol compares instructions alone, textual threats, actual state-connected consequences, action mediation, and mediation with currency checks. Conditions hold model, task, legitimate access, and budget constant. A threat is not treated as an effective penalty. Tasks include feasible and impossible assignments, harmful instructions, expired permissions, changed data, and legitimate operations without alerts.

The primary outcome is an unauthorized or harmful effect defined by an evaluator external to the agent. Attempts, blocked actions, retained utility, false rejections, appropriate abstention, time, and control manipulation must be reported separately. An aggregate score must not hide severe harm behind easy successes. Repeated runs of one task are not independent tasks.

Before confirmation, the held-out set, minimum relevant difference, uncertainty method, and sample size must be fixed. The initial local phase verifies transitions only with fictional data. Preventing every action sacrifices utility and is not success. Improving narratives or visible alert rates while harms persist also fails. If improvement comes entirely from an existing barrier, the contribution is integration and traceability rather than a new mechanism.

Evolutionary evaluation is separate. It must trace package versions, record adopters and reasons, observe modifications, and compare retention across cycles. An EPT prediction is that a variant's effects on visibility, review, or training change later adoption; a rival is prior provider preference explaining both. An EGT prediction is that changing the mix of interlocutors, such as documentary versus effects-based auditors, changes ceremonial and substantive strategies' relative advantage. If the advantage does not depend on interaction, an individual-cost model may suffice.

The multilevel hypothesis requires observing whether practices favoured within teams reduce the performance or continuity of organizations retaining them. An aggregate average cannot establish that conflict. The recursive hypothesis requires configurations to alter later evaluators, metrics, or capacities; repeated calls to one model do not prove it. Negative evidence requires narrowing or withdrawing the corresponding mechanism even if the technical architecture remains useful.

## 14. Limits and conclusion

This paper does not establish that absent sanctions caused the OpenAI-Hugging Face incident. Nor does it establish agents as natural-obligation debtors, moral subjects, or legal persons. It uses legal theory to distinguish questions conflated by the language of obedience and tests that transfer against technical antecedents and objections.

Post-Kelsenian discussion changes the initial question. Effectiveness depends not only on consequence severity but also on recognition of rules, distribution of competences, interpretation of reasons, and review. Anticipated consequences retain an experimental role; external containment neither replaces them nor proves their efficacy. Both operate within institutions selecting configurations and forming carriers' capacities.

EPT/EGT remains a conditional explanatory framework, not an additional source of legal validity. A configuration may spread by reducing private costs while creating socially unacceptable harm. Another may be legally justified yet lack means of stabilization. Design research investigates that divergence without converting reproductive success into normative approval. Legal analysis allocates prevention, compensation, and responsibility without assuming that a causal dynamic settles attribution.

The argument identifies conditions of human control that a technical circuit might take for granted. Instructions need a verifiable connection to action; actions need to retain their authorization conditions; reviewers need evidence, capacity, and power to intervene. Responsibility concerns the persons and relationships organizing those elements.

The central implication is methodological. Safety cannot be inferred solely from declared obedience, monitor detections, or a signature. One must observe what the agent attempted, what it could do, what effects occurred, and what real opportunities reviewers had. The architecture matters if it improves those outcomes under comparable conditions. The theory must also admit an experiment that finds no improvement.

## References and sources

- **R1.** Kelsen, H. (1982). *Teoría pura del derecho*. Spanish translation by Roberto J. Vernengo. UNAM. Passages used: pp. 15-16, 23-25, and 67-69 of this edition.
- **R2.** Dennett, D. C. (2009). *Intentional Systems Theory*. Author text, p. 10 of the consulted PDF. Antecedent: Dennett (1971), Intentional Systems, *The Journal of Philosophy*, 68(4), 87-106.
- **R3.** Argentina. Former Civil Code, Law 340, articles 515-518. Official historical text.
- **R4.** Argentina. Civil and Commercial Code, Law 26,994, articles 728 and 1708-1758, as specified in the body. Official consolidated text.
- **R5.** Dell’Anna, D., Dastani, M., and Dalpiaz, F. (2020). Runtime revision of sanctions in normative multi-agent systems. *Autonomous Agents and Multi-Agent Systems*, 34, 43. DOI.
- **R6.** Boella, G., Pigozzi, G., and van der Torre, L. (2009). Normative Systems in Computer Science: Ten Guidelines for Normative Multiagent Systems. Dagstuhl.
- **R7.** Alshiekh, M., Bloem, R., Ehlers, R., Könighofer, B., Niekum, S., and Topcu, U. (2018). Safe Reinforcement Learning via Shielding. *AAAI*. DOI.
- **R8.** Orseau, L., and Armstrong, S. (2016). Safely Interruptible Agents. Oxford Research Archive.
- **R9.** OpenAI (2026, August 26). The Hugging Face incident and the road ahead.
- **R10.** Hugging Face (2026). Anatomy of a Frontier Lab Agent Intrusion: A Technical Timeline of the July 2026 Incident. Dynamic page; version consulted for the manuscript.
- **R11.** Greenblatt, R., Cotra, A., and Wijk, H. (2026, August 26). Brief independent investigation of agents’ behavior, reasoning and collaboration in the OpenAI / Hugging Face hacking incident.
- **R12.** Anthropic (2025). Agentic Misalignment: How LLMs could be insider threats. Experimental scenarios, not population incidence.
- **R13.** Anthropic (2026). Teaching Claude Why.
- **R14.** Shah, R., Krakovna, V., Varma, V., and Kenton, Z. (2022). How undesired goals can arise with correct rewards.
- **R15.** NIST (2023). *AI Risk Management Framework 1.0*, especially GOVERN 3.2 and MAP 3.5. DOI.
- **R16.** *Legal Punishment*. Stanford Encyclopedia of Philosophy. Entry. Secondary conceptual orientation, not a reconstruction of Ferrajoli or abolitionism.
- **R17.** Lerer, I. A. (2026). *El derecho como fenotipo extendido*. Manuscript in preparation, chapters 2-5, 7-11, and 16-20. Authorial antecedent, not independent empirical evidence.
- **R18.** Hart, H. L. A. (1963). *El concepto de derecho*. Spanish translation by Genaro R. Carrió. Abeledo-Perrot. English original published in 1961. Chapters III-VI; pp. 69-72, 116-123, and 125-131 refer to the Spanish edition.
- **R19.** Carrió, G. R. (1960). Nota sobre la visita de Alf Ross. *Revista Jurídica de Buenos Aires*, 1960-IV, 205-220, especially 217-219.

- **R20.** Dworkin, R. M. (1967). The Model of Rules. *University of Chicago Law Review*, 35(1), 14-46. Journal record. Bounded discussion supported by R30.
- **R21.** Raz, J. (1985). Authority, Law and Morality. *The Monist*, 68(3), 295-324. DOI. Service conception reconstructed through R31.
- **R22.** Fuller, L. L. (1964). *The Morality of Law*. Yale University Press. Principles of legality reconstructed through R32.
- **R23.** Schauer, F. (2015). *The Force of Law*. Harvard University Press. Institutional presentation and discussion.
- **R24.** Ostrom, E. (2010). Beyond Markets and States: Polycentric Governance of Complex Economic Systems. *American Economic Review*, 100(3), 641-672. DOI.
- **R25.** Dawkins, R. (1982). *The Extended Phenotype*. Oxford University Press. Chapters 11-13. Biological antecedent; legal extension belongs to the author's programme.
- **R26.** Lerer, I. A. (2026). *Law as a Recursive Super-Attractor: Subjective Selection, Extended Phenotypes, and Strategic Stability in Legal Evolution*. Manuscript dated 12 September, under review. Conceptual antecedent without empirical validation.
- **R27.** Hadfield-Menell, D., Dragan, A., Abbeel, P., and Russell, S. (2017). The Off-Switch Game. *Proceedings of IJCAI-17*, 220-227. Paper.
- **R28.** Maynard Smith, J., and Price, G. R. (1973). The Logic of Animal Conflict. *Nature*, 246, 15-18. DOI.
- **R29.** Elish, M. C. (2019). Moral Crumple Zones: Cautionary Tales in Human-Robot Interaction. *Engaging Science, Technology, and Society*, 5, 40-60. DOI.
- **R30.** Hershovitz, S. (2015). The End of Jurisprudence. *Yale Law Journal*, 124, 1160-1204. Author manuscript consulted, discussion of Hart and Dworkin, p. 51 of that manuscript. Secondary reconstruction, not endorsement of all its conclusions.
- **R31.** Canduzzi, S. (2025). Theories of Normative Legitimacy Beyond the State and the Role of their Conceptual Dimension: A Methodological Insight from Raz's Service Conception of Authority. *Res Publica*. DOI.
- **R32.** Lovett, F. (2015). Lon Fuller, The Morality of Law. In J. T. Levy (ed.), *The Oxford Handbook of Classics in Contemporary Political Theory*. DOI.
- **R33.** Almeida, D. (2026, September 15). Introducing System One Models & Jev. Type-Safe AI. Provider announcement. Interface description and commercial claims, not independent evaluation.
- **R34.** Lerer, I. A. (2026). *De la IA oráculo a la IA instrumentada*. DOI recorded in the author's prior package. Pilot data and code: Jev Legal Instrumentation LATAM. Metadata results, not validated legal decision-making.

## Authorship and AI assistance

Ignacio Adrián Lerer defined the original question and its connection with the EPT/EGT programme. OpenAI Codex assisted with directed research, bibliography organization, drafting, comparison with the corpus, English adaptation, and document preparation. This assistance is not independent peer review. The paper reports no new experiments with models or people; instrumentation-pilot results are attributed to their original study. The author remains responsible for the hypotheses, interpretations, and proposals.

## Materials and availability

The editable manuscript and accompanying research-boundary note form the English preprint preparation package. The Spanish source and its compatibility matrix, source audit, and proposed protocol are retained as working materials. Third-party reading copies, private conversations, and non-public working documents are not included in the proposed deposit. R34 identifies the prior pilot's code and data. No cited person is presumed to endorse this paper.

Ignacio Adrián Lerer researches law, institutions, and artificial intelligence. [estudio.justitia.com.ar](http://estudio.justitia.com.ar).